

EE 590 A

Applied High-Performance GPU Computing

LECTURE 5 : Thursday, October 25, 2018

Instructor: Dr. Colin Reinhardt

University of Washington - Professional Masters Program
Autumn 2018



Lecture 5 : Outline

- ▶ HW3 Discussion and Examples
- ▶ Intel VTune Amplifier 2018 Tour
 - ▶ with In-class Exercise 5a
- ▶ SIMD/SPMD/SIMT, OpenCL, and Intel
- ▶ MMUL OPT5: blocked/tiled/partitioned
- ▶ In-class Exercise 5b

Reminders

- ▶ Office hours this Sunday 3-5 PM
- ▶ HW-4: Due Thursday 11/8 by 5:59 PM
 - ▶ will be posted by Fri 11/26 midnight

HW3 Discussion and Examples

- ▶ Using OpenCL C printf() in kernels
- ▶ User-defined types (UDTs)
 - ▶ Alignment of built-in Types (C Spec, 6.1.5)
 - ▶ Unions and Structs (C Spec, 6.11.1)
 - ▶ `__attribute__((aligned (x)))`
 - ▶ `__attribute__((packed))`

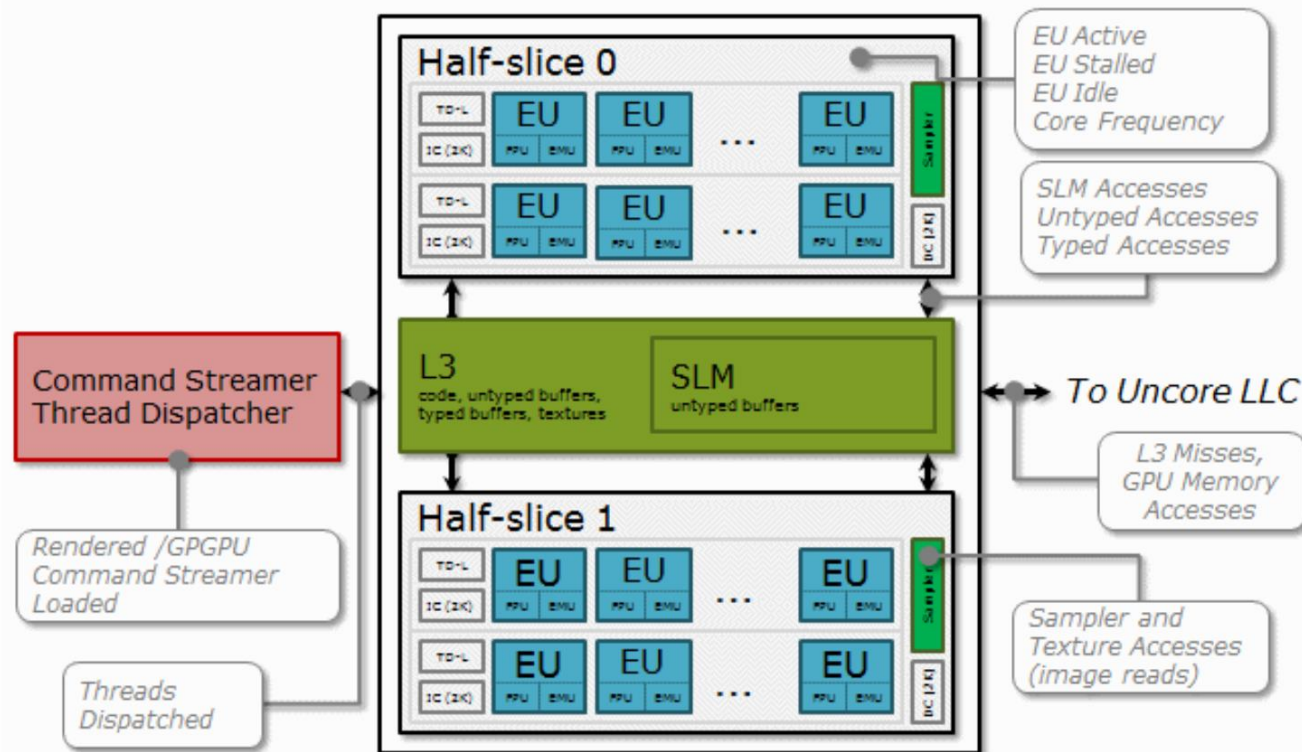


In-class Exercise 5a

Introduction to Intel VTune Amplifier 2018

VTune GPU Metrics

- Requires Administrator privileges



Metric	Formula
EU Array Active	$\frac{\sum_{\text{across all EUs}} \text{cycles when EU executes instructions}}{\sum_{\text{across all EUs}} \text{all cycles}}$
EU Array Stalled	$\frac{\sum_{\text{across all EUs}} \text{cycles when EU does not execute instructions and at least one thread is scheduled on EU}}{\sum_{\text{across all EUs}} \text{all cycles}}$
EU Array Idle	$\frac{\sum_{\text{across all EUs}} \text{cycles when no threads scheduled on EU}}{\sum_{\text{across all EUs}} \text{all cycles}}$

VTune Amplifier provides platform-specific presets of the hardware metrics. All presets collect data about execution units (EUs) activity: EU Array Active, EU Array Stalled, EU Array Idle, Computing Threads Started, and Core Frequency.

- Overview** event set also includes metrics that track general GPU memory accesses such as Memory Read/Write Bandwidth, GPU L3 Misses, Sampler Busy, Sampler Is Bottleneck, and GPU Memory Texture Read Bandwidth. These metrics can be useful for both graphics and compute-intensive applications.
- Compute Basic (with global/local memory accesses)** event group also includes metrics that distinguish accessing different types of data on a GPU: Untyped Memory Read/Write Bandwidth, Typed Memory Read/Write Transactions, SLM Read/Write Bandwidth, Render/GPGPU Command Streamer Loaded, and GPU EU Array Usage. These metrics are useful for compute-intensive workloads on the GPU.
- Compute Extended** event group includes metrics targeted only for GPU analysis on the Intel processor code name Broadwell and higher. For other systems, this preset is not available.
- Full Compute** event group is a combination of the **Overview** and **Compute Basic** event sets.



SIMD, SPMD, SIMT

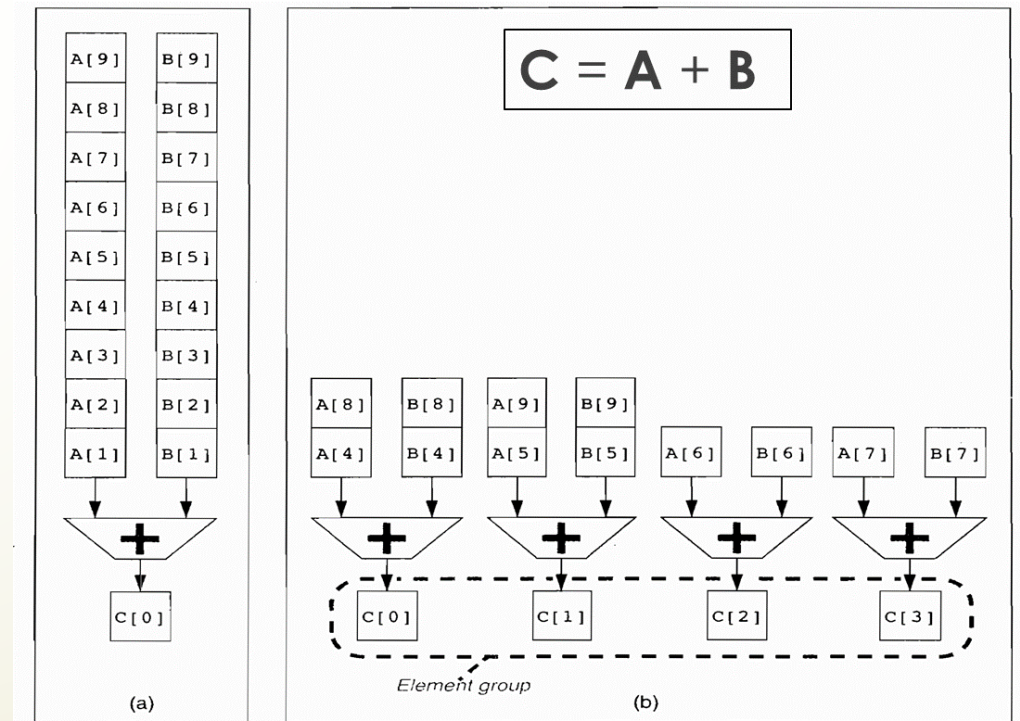
Start with definitions of the acronyms:

- SIMD: Single-Instruction Multiple-Data
- SPMD: Single-Program Multiple-Data
- SIMT: Single-Instructions Multiple Threads

To understand these concepts requires we take a look at the hardware...

First there was SIMD

- Flynn's Taxonomy (1966)
 - Originally basis for vector supercomputers in early 1970-1980s (Cray)
 - older SIMD architectures: "long vector" 64- 64K vector lengths
 - newer SIMD architectures: "short vector" 2-16 Words
-
- **Data-level parallelism** model
-
- **vector** versus **array** processors
 - pipelining limited functional units, or
 - multiple general functional units



Array vs. Vector Processors

Instruction Stream

LD VR \leftarrow A[3:0]
ADD VR \leftarrow VR, 1
MUL VR \leftarrow VR, 2
ST A[3:0] \leftarrow VR

Definitions:

VR: vector register

LD: LOAD instruction (from mem)

ADD: ADD instruction


MUL: Multiply instruction

ST: STORE instruction (to mem)

ARRAY PROCESSOR



Same op @ same time



Different ops @ same space

Time

Space

VECTOR PROCESSOR



Different ops @ time



Same op @ space

Space

GPUs are SIMD Engines Underneath

- ▶ The instruction pipeline operates like a SIMD pipeline (e.g., an array processor)
- ▶ However, the programming is done using **threads**, NOT SIMD instructions
- ▶ We will consider a simple parallelizable loop code example

```
for (i=0; i < N; i++)  
    C[i] = A[i] + B[i];
```

- ▶ Need to distinguish between
 - ▶ Programming Model (Software)
 - vs.
 - ▶ Execution Model (Hardware)


Programming Model vs. Hardware Execution Model

- ▶ Programming Model refers to **how the programmer expresses the code**
 - ▶ E.g., Sequential (von Neumann), Data Parallel (SIMD), Multi-threaded (MIMD), Single-program multiple-data (SPMD), ...
- ▶ Execution Model refers to **how the hardware executes the code underneath**
 - ▶ E.g., Dynamic Out-of-order execution, Vector processor, Array processor, Multiprocessor, Multithreaded processor, ...
- ▶ **Execution Model can be very different from the Programming Model**
 - ▶ E.g., von Neumann model implemented by an OoO processor
 - ▶ E.g., SPMD model implemented by a SIMD processor (**GPUs**)

How Can You Exploit Parallelism Here?

```
for (i=0; i < N; i++)  
    C[i] = A[i] + B[i];
```

Scalar Sequential Code




Consider three programming options to exploit the instruction-level parallelism present in this sequential code:

1. Sequential (SISD)
2. Data-Parallel (SIMD)
3. Multithreaded (MIMD/SPMD)

Prog. Model 1: Sequential (SISD)

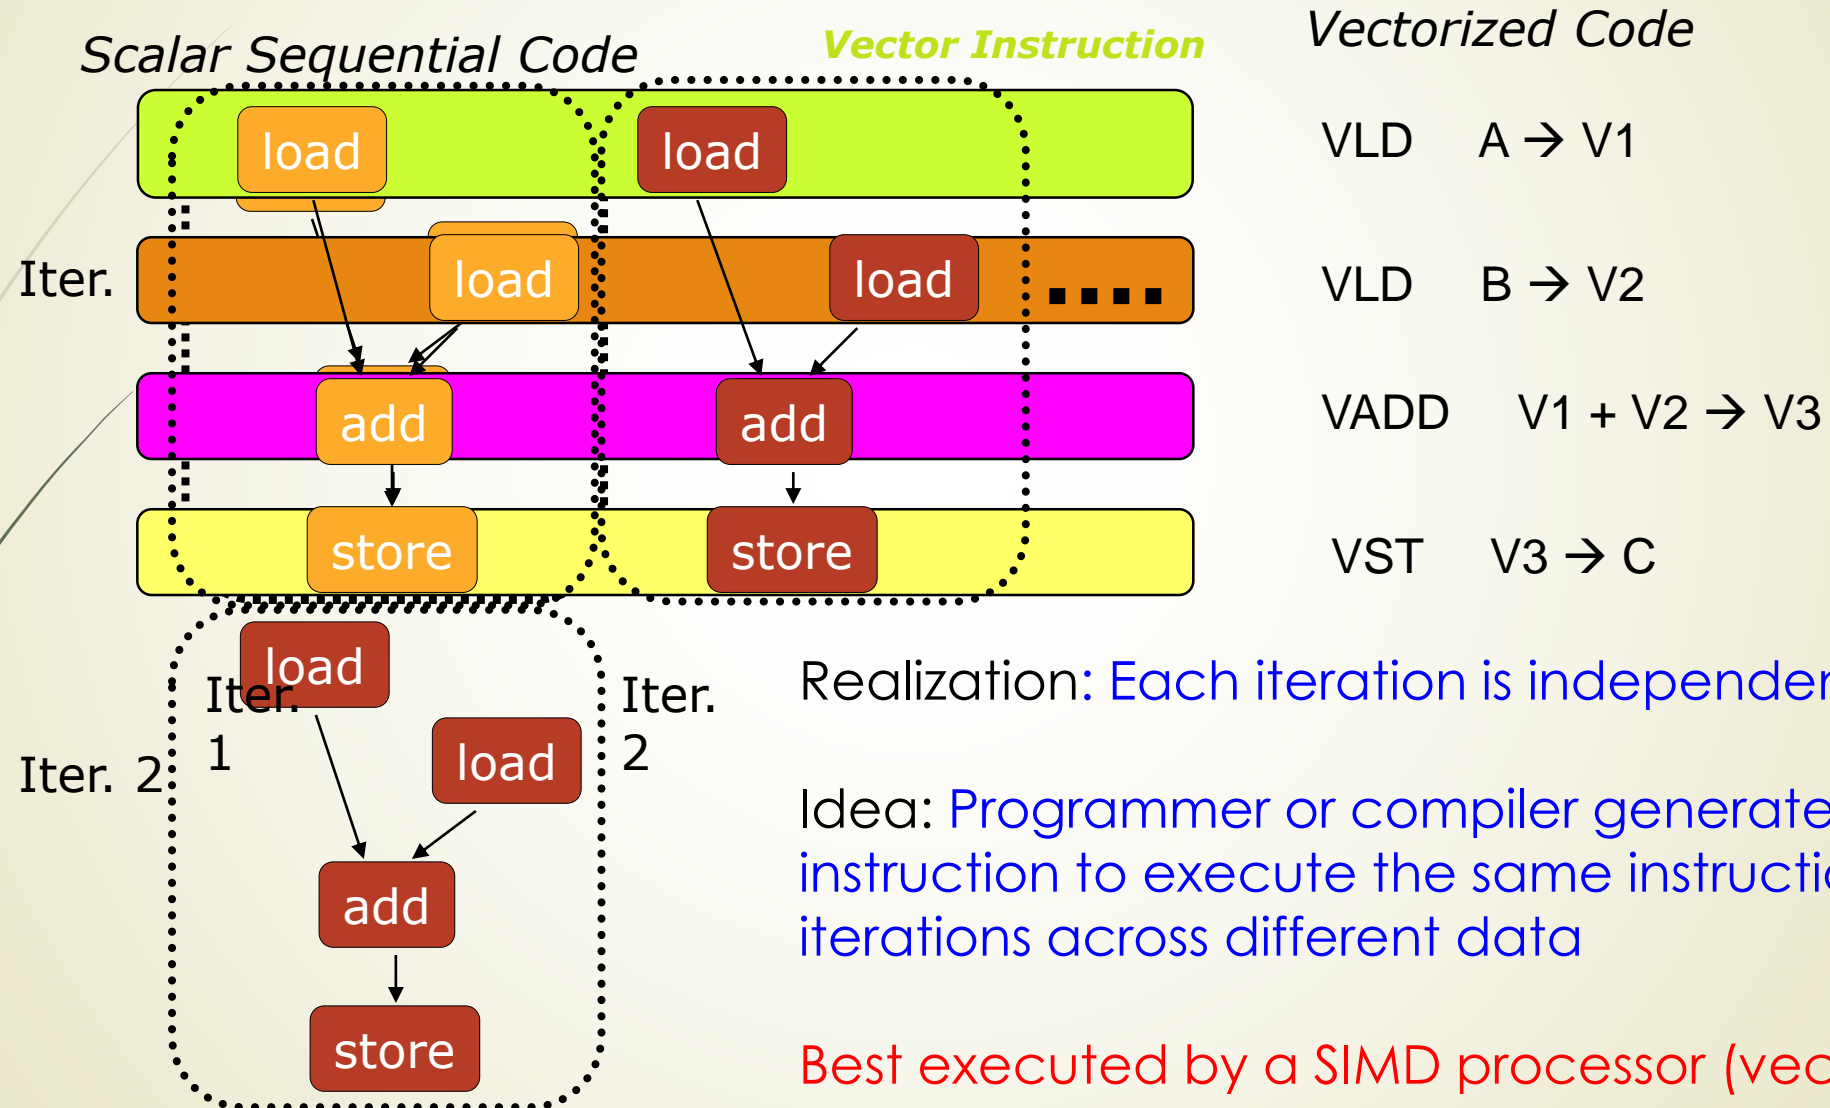
```
for (i=0; i < N; i++)  
    C[i] = A[i] + B[i];
```



- Can be executed on
 - Pipelined processor
 - Out-of-order execution processor
 - ❑ Independent instructions executed when ready
 - ❑ Different iterations are present in the instruction window and can execute in parallel in multiple functional units
 - ❑ In other words, the loop is dynamically unrolled by the hardware
 - Superscalar or VLIW processor
 - ❑ Can fetch and execute multiple instructions per cycle

Prog. Model 2: Data Parallel (SIMD)

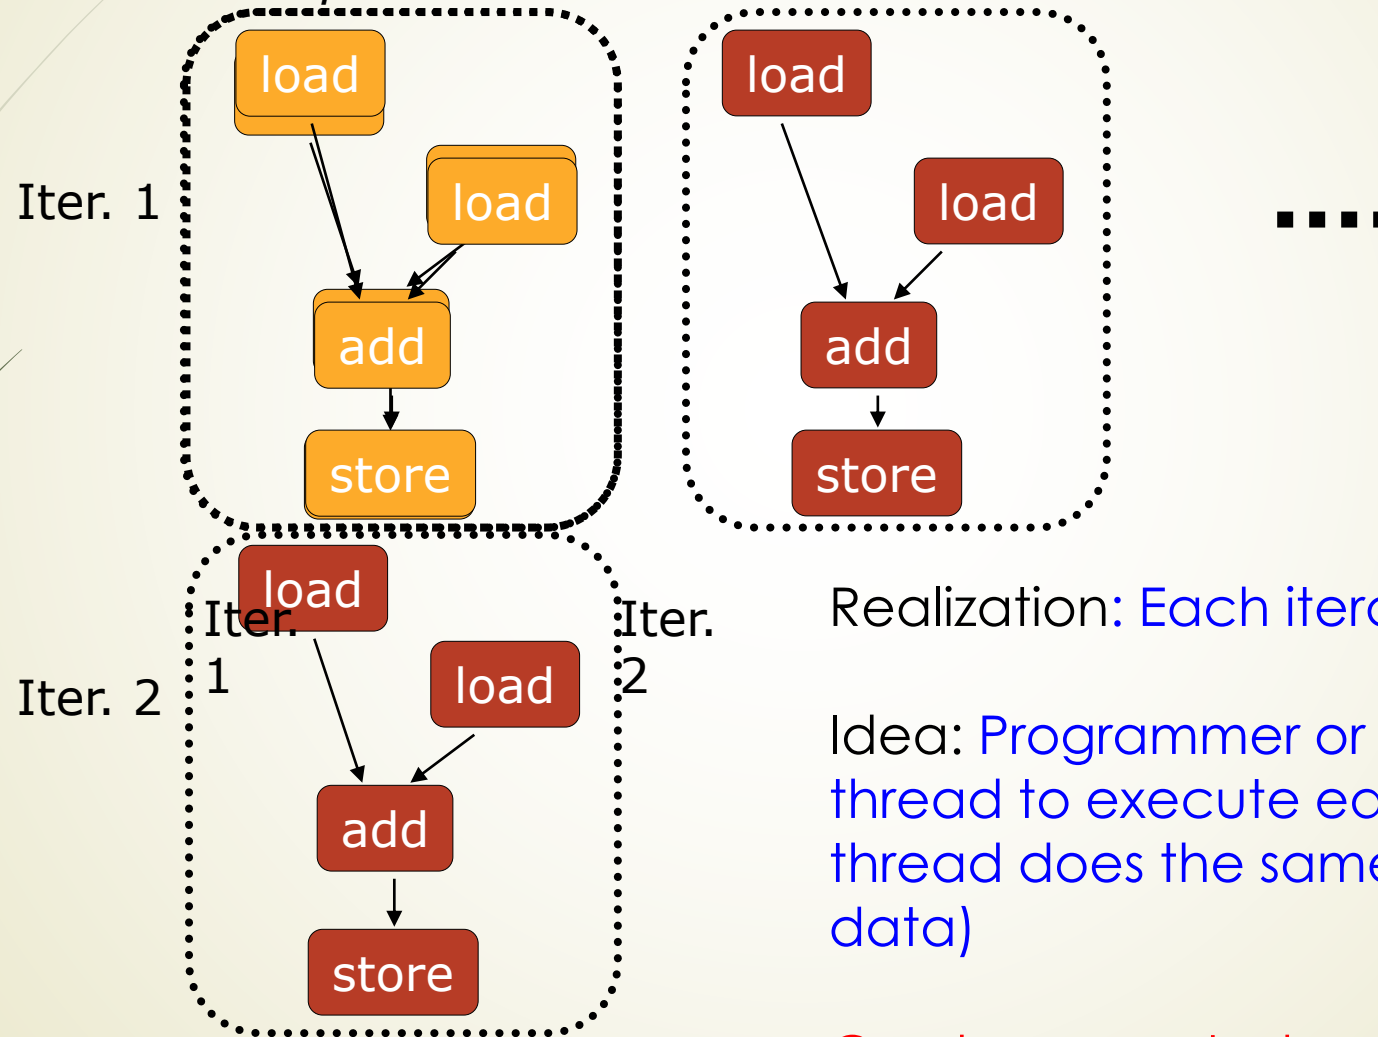
```
for (i=0; i < N; i++)  
    C[i] = A[i] + B[i];
```



Prog. Model 3: Multithreaded

```
for (i=0; i < N; i++)  
    C[i] = A[i] + B[i];
```

Scalar Sequential Code




Realization: Each iteration is independent

Idea: Programmer or compiler generates a thread to execute each iteration. Each thread does the same thing (but on different data)

Can be executed on a MIMD machine

Prog. Model 3: Multithreaded

```
for (i=0; i < N; i++)  
    C[i] = A[i] + B[i];
```



Realization: Each iteration is independent

This particular model is also called:

SPMD: Single Program Multiple Data

Can be executed on a SIMT machine

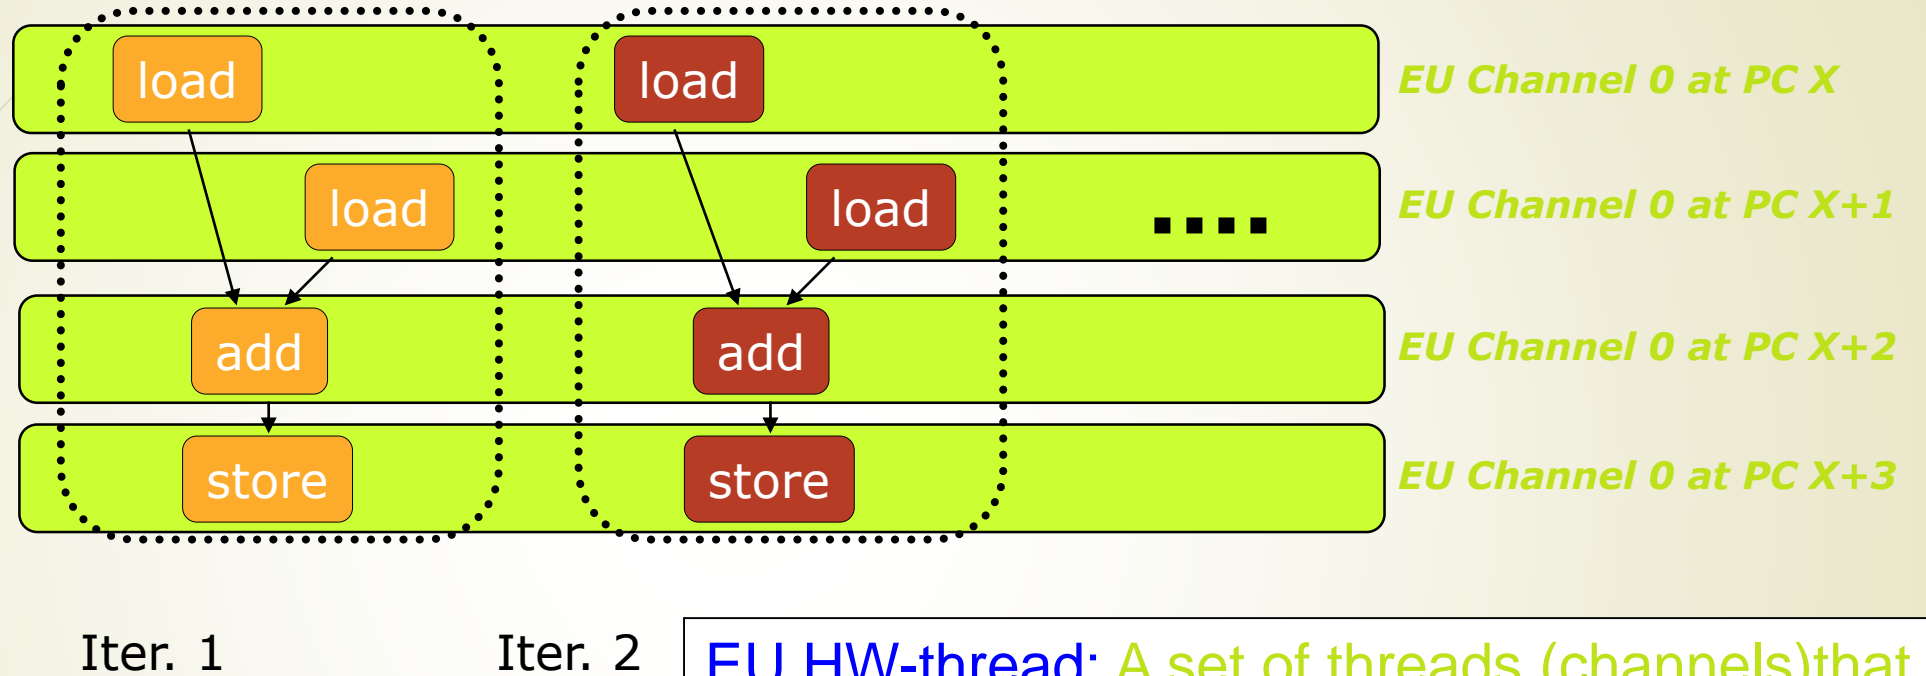
Single Instruction Multiple Thread

A GPU is a SIMD (SIMT) Machine

- ▶ Except it is not programmed using SIMD instructions
- ▶ It is programmed using threads (**SPMD** programming model)
 - ▶ Each thread executes the same code but operates a different piece of data
 - ▶ Each thread has its own context (state registers) (i.e., can be treated/restarted/executed independently)
- ▶ A set of threads executing the same instruction are dynamically grouped into an **EU HW-thread** (A.K.A. warp (NVIDIA), wavefront (AMD) by the hardware)
 - ▶ essentially a SIMD operation formed by hardware!

SPMD on SIMT Machine

```
for (i=0; i < N; i++)  
    C[i] = A[i] + B[i];
```



EU HW-thread: A set of threads (channels) that execute the same instruction (i.e., same PC)

This particular model is also called:

SPMD: Single Program Multiple Data

A GPU executes it using the SIMT model:
Single Instruction Multiple Thread

Intel SIMD

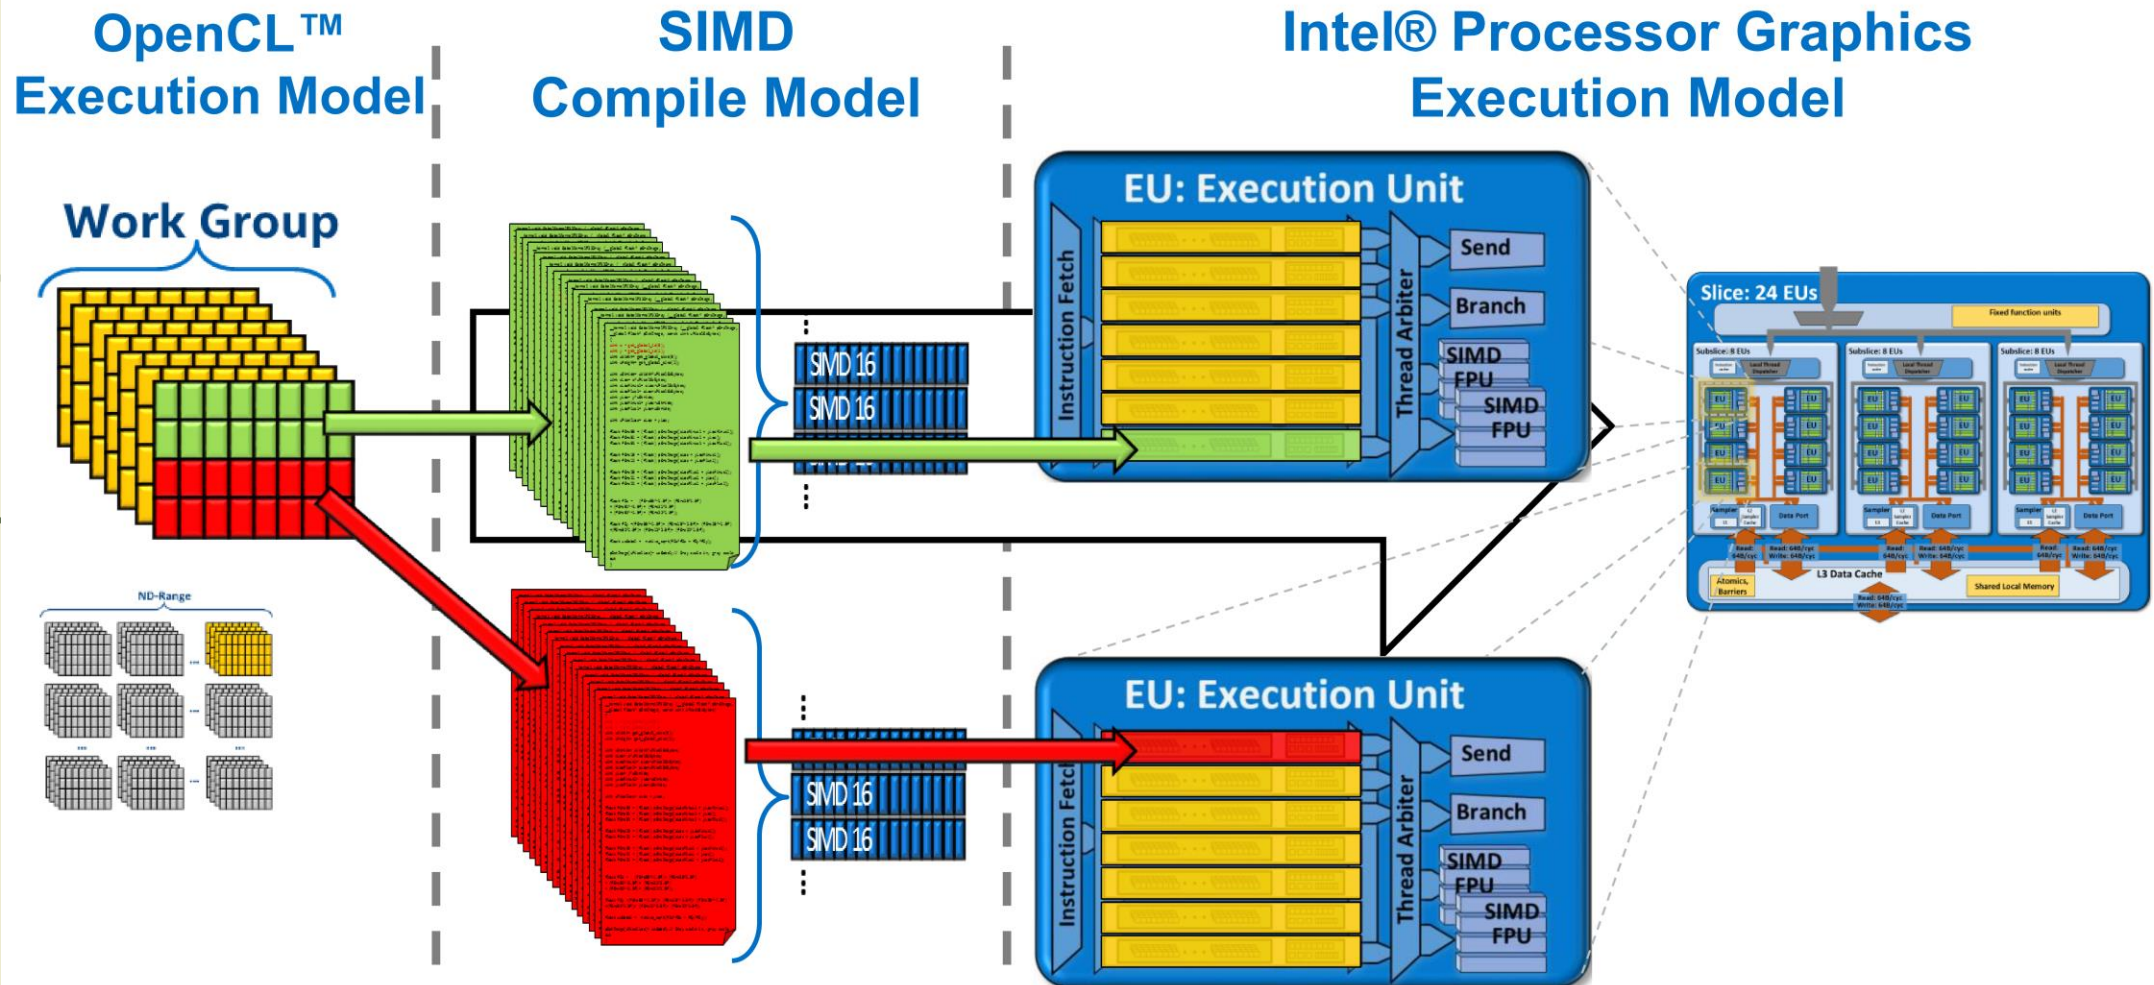
- ▶ CPU ISAs (x86, x64): SIMD-extensions: MMX, SSE, AVX, ...
 - ▶ SIMD intrinsics (AVX-512): `__m512i _mm512_add_epi32(__m512i a, __m512i b)`
- ▶ Processor Graphics (GENx GPUs)
 - ▶ 2 x SIMD FPUs in each EU
 - ▶ support both floating-point and integer computation
 - ▶ parallel execute 4 x 32-bit (Word) floating-point/integer operations
 - ▶ or 8 x 16-bit (Half/Short) floating-point/integer operations
 - ▶ each SIMD FPU can complete simultaneous ADD and MULT (MAD) operations each cycle
 - ▶ Thus each EU capable of: $(\text{ADD} + \text{MUL}) \times 2 \text{ FPUs} \times \text{SIMD-4} = 16 \text{ 32-bit FLOP/cycle}$
 - ▶ EU General-purpose Register Files (GRFs)
 - ▶ 128 GRFs per EU HW thread
 - ▶ Each register stores 32 Bytes, accessible as SIMD 8-element vector of 32-bit (Word) data
 - ▶ $128 \times 32 = 4096 \text{ bytes} = 4 \text{ KB} / \text{EU HW thread}$



Intel: Logical versus Physical SIMD

- ▶ EU ISA and GRF are designed to support a flexible SIMD width
 - ▶ **Physical = HW** (Execution model)
 - ▶ For 32-bit data types, Gen9 FPU are *physically* SIMD 4-wide, GRFs 8-wide.
 - ▶ **Logical = SW** (Programming-Compilation model)
 - ▶ The FPUs can be targeted with SIMD instructions and registers that are *logically* 1-wide, 2-wide, 4-wide, 8-wide, 16-wide, or 32-wide.
- ▶ GEN Compiler: scalarization and auto-vectorization passes
- ▶ If the kernel is compiled SIMD-8, then a subgroup is made up of 8 work items
 - ▶ share 4 KB of register space of an EU HW thread and execute together

Intel OpenCL SIMD Model



OpenCL™ Work Groups Assigned to One or More EU Threads, Across Multiple EUs

Intel OpenCL Extension: subgroups


What is a Subgroup?

A Subgroup is a Collection of Work Items

- Another Level in the Execution Hierarchy
- Between Work Groups and Work Items

Key Takeaways:



- On Intel® Processor Graphics, work items in a subgroup execute on the same EU Thread!
- Subgroups can use specialized SIMD instructions for “block operations”



Subgroup Functions bring “Explicit SIMD” to OpenCL kernels!

SIMD/SIMT Execution Degradation

- Work-group divergence:
 - Inefficient utilization
- Analogous to pipeline structural hazards and stalls
- Strive to avoid this!**
 - Avoid branching in kernels**





SIMD/SPMD/SIMT Summary



- ▶ Don't worry if you feel unsure about exactly what and how these work!
- ▶ We will continue to dig into the concepts during the course.
 - ▶ Intel OpenCL extensions: subgroups
- ▶ Still lots of unknowns due to lack of documentation and vendor IP
- ▶ Opportunities for Final Project topics:
 - ▶ Microbenchmarking
 - ▶ Intel NEO Open-source driver and compiler
 - ▶ auto-scalarization and vectorization
 - ▶ Intermediate representations (IR) and native ISA

Block Matrices and Algorithms

Block Matrix Terminology

- ▶ In general we can partition both the rows and columns of an m -by- n matrix A to obtain

$$A = \begin{bmatrix} A_{11} & \dots & A_{1r} \\ \vdots & & \vdots \\ A_{q1} & \dots & A_{qr} \end{bmatrix} \begin{matrix} m_1 \\ \\ m_q \end{matrix}$$

$n_1 \qquad n_r$

where $m_1 + \dots + m_q = m$, $n_1 + \dots + n_r = n$

and $A_{\alpha\beta}$ designates the (α, β) block (submatrix, partition, tile).

With this notation block $A_{\alpha\beta}$ has dimension m_α -by- n_β and we say that $A = (A_{\alpha\beta})$ is a q -by- r block matrix.

Block Matrices and Algorithms

- ▶ Block matrix addition

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \\ A_{31} & A_{32} \end{bmatrix} + \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \\ B_{31} & B_{32} \end{bmatrix} = \begin{bmatrix} A_{11} + B_{11} & A_{12} + B_{12} \\ A_{21} + B_{21} & A_{22} + B_{22} \\ A_{31} + B_{31} & A_{32} + B_{32} \end{bmatrix}$$

- ▶ Block matrix multiplication

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \\ A_{31} & A_{32} \end{bmatrix} \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} A_{11}B_{11} + A_{12}B_{21} & A_{11}B_{12} + A_{12}B_{22} \\ A_{21}B_{11} + A_{22}B_{21} & A_{21}B_{12} + A_{22}B_{22} \\ A_{31}B_{11} + A_{32}B_{21} & A_{31}B_{12} + A_{32}B_{22} \end{bmatrix}$$

- ▶ Requires that row and column dimensions of the blocks satisfy the necessary constraints
 - ▶ Operands are **partitioned conformably**.

Partitioned (blocked, tiled) Matrix Multiplication

► Golub & Van Loan, *Matrix Computations*, 4th, 2013

Theorem 1.3.1. *If*

$$A = \begin{bmatrix} A_{11} & \dots & A_{1s} \\ \vdots & & \vdots \\ A_{q1} & \dots & A_{qs} \end{bmatrix} \begin{matrix} m_1 \\ \\ m_q \end{matrix}, \quad B = \begin{bmatrix} B_{11} & \dots & B_{1r} \\ \vdots & & \vdots \\ B_{s1} & \dots & B_{sr} \end{bmatrix} \begin{matrix} p_1 \\ \\ p_s \end{matrix},$$

$p_1 \qquad p_s \qquad n_1 \qquad n_r$

and we partition the product $C = AB$ as follows,

$$C = \begin{bmatrix} C_{11} & \dots & C_{1r} \\ \vdots & & \vdots \\ C_{q1} & \dots & C_{qr} \end{bmatrix} \begin{matrix} m_1 \\ \\ m_q \end{matrix},$$

$n_1 \qquad n_r$

then for $\alpha = 1:q$ and $\beta = 1:r$ we have $C_{\alpha\beta} = \sum_{\gamma=1}^s A_{\alpha\gamma} B_{\gamma\beta}$.

Block Matrix Multiplication: Design

- ▶ So we know theoretically/analytically, we can decompose large matrices into a large number of smaller submatrices which can be computed independently.
- ▶ This is important, but
 - ▶ How big should the submatrix blocks be?
 - ▶ How will the data communication and memory hierarchy be used?
 - ▶ Remember
 - ▶ Each EU HW thread has limited **private** memory general-purpose register (GRF) space: 4KB
 - ▶ the amount available to each work-item depends on SW-SIMD width (subgroup size)
 - ▶ All work-items in a work-group are confined to a subslice, share 64KB **local** memory
 - ▶ The L3 cache holds data from **global** memory within the slice: 512KB for application data.
 - ▶ How will we choose the best
 - ▶ SW-SIMD width
 - ▶ NDRange dimensions, global work size, local work size?
 - ▶ How best to layout the data, based on above??

Block Matrix Mult: Intuition



OPT5 Block Matrix Mult: Algorithm

- ▶ Require all matrices (A,B,C) to be NxN (to keep it simple)
- ▶ Partition matrices into PxP blocks
- ▶ Each work-group (WG) calculates an output block
- ▶ Each work-item (WI) calculates an output block row

Kernel Pseudocode:

- determine current output block: WG ID [0]=row, [1]=column
- FOR P block-products of $A_{(\text{row}, 0 \dots P-1)} \times B_{(0 \dots P-1, \text{col})}$
 - Each WI put a block row $A_{(\text{row}, i)}(\text{WI}\#, :)$ into private memory (GRF) array
 - All WIs put block $B_{(i, \text{col})}(\text{WI}\#, :)$ into local memory (SLM)
 - ---WG barrier---
 - each WI compute block product for i^{th} row: $a_{ik}b_{kj}$
 - accumulate partial results for output row Ci in GRF array
 - ---WG barrier---
- Write final results back to global memory

```

// Version 0: (V0) - Initial Implementation
__attribute__((reqd_work_group_size(4,4,1)))
__kernel void mmul_opt5_tiled_v0(const int N, const int P, __global float* A, __global float* B, __global float* C)
{
    int i,k;
    int iloc = get_local_linear_id();
    int wgRow = get_group_id(0); // using dim-0 as row dim
    int wgCol = get_group_id(1); // using dim-1 as column dim
    //int glid = get_global_linear_id();

    int blkSize = N/P;
    int blkRowOffset = blkSize*wgRow;
    int blkColOffset = blkSize*wgCol;

    //printf("<mmul_opt5_tiled> glid = %d: iloc: %d, wgRow: %d, wgCol: %d, blocksize: %d, blkoffset: (row: %d, col: %d)\n",glid,iloc,wgRow,wgCol,blkSize,blkRowOffset,blkColOffset);

    float outtmp[16] = {0.0f}; // this value needs to be updated based on N and P
    float Apriv[16] = {0.0f}; // this value needs to be updated based on N and P
    local float Bloc[256];

    // each WI computes a row i of partition sub-matrix C(wgRow,wgCol)
    for(int nP=0; nP<P; nP++) // for each partition in INPUT matrices row,col
    {
        // copy global row of partition submatrix A to WI private registers
        for(k=0; k<blkSize; k++)
            Apriv[k] = A[(blkRowOffset+iloc)*N + nP*blkSize + k];

        // copy global col of partition submatrix B to WG shared local memory (SLM)
        // WIs copy all columns of partition submatrix in parallel
        for(k=0; k<blkSize; k++)
            Bloc[k*blkSize+iloc] = B[(nP*blkSize+k)*N + blkColOffset + iloc];

        work_group_barrier(CLK_LOCAL_MEM_FENCE); // WG sync

        // each WI compute partition row sub-product
        for(i=0; i < blkSize; i++) {
            for(k=0; k < blkSize; k++) {
                outtmp[i] += Apriv[k] * Bloc[k*blkSize+i];
            }
        }
        work_group_barrier(CLK_LOCAL_MEM_FENCE); // WG sync
    }
    // write final results back to global
    for(i=0; i< blkSize; i++)
        C[(blkRowOffset+iloc)*N + blkColOffset + i] = outtmp[i];
}

```



OPT5 Block Matrix Mult: Next Steps

- ▶ We have a basic algorithm and implementation that works.
- ▶ But we're not done yet!
- ▶ Next steps:
 - ▶ profile it (get a performance baseline)
 - ▶ analyze its memory and compute performance
 - ▶ VTune!
 - ▶ Additional Optimizations



Exercise 5b: VTune Analysis of MMUL

Refer to course website: [Files/InclassExercises/EX5](#)

Lecture 5 Summary

- ▶ We looked at
 - ▶ Using OpenCL C `printf()`
 - ▶ OpenCL data alignment between Host and Device
 - ▶ particularly with UDTs
 - ▶ Attribute qualifiers
 - ▶ SIMD/SPMD/SIMT and Intel HW and SW (OpenCL)
 - ▶ OPT5: Blocked Matrix Multiplication kernel (several variants)
 - ▶ Intel VTune Amplifier 2018
 - ▶ Setting up and running GPU/OpenCL Hotspot analysis
 - ▶ beginning to interpret the results for kernel performance tuning.