



EE 524 P

# Applied High-Performance GPU Computing

**LECTURE 8 : Thursday, November 15, 2018**

Instructor: Dr. Colin Reinhardt

University of Washington - Professional Masters Program

Autumn 2018



# Lecture 8 : Outline

- ▶ HW4 Extension: Due 11/18 by 11:59 PM (**SUNDAY**)
- ▶ Final Project Reminders
  - ▶ Project Proposals DRAFT: due TODAY
  - ▶ Project Proposals FINAL: due WED 11/21 by 11:59 PM
  - ▶ Proposal Designs: due 11/29 at 6:00 PM
  - ▶ Final Project: due 12/14 by 6:00 PM (no late submissions)
- ▶ Final Project Details
- ▶ Spatial Domain Image Processing
  - ▶ OpenCL Image kernel corrections
  - ▶ EX 8a
- ▶ Frequency Domain Image Processing: Image Enhancement
  - ▶ cIFFT
- ▶ Parallelizing Partial Differential Equations (PDEs)
- ▶ EX 8b



# Final Project: Proposals

Project Proposal: 1-2 pages

MUST include:

- Summary of topic/problem to be studied
- List of primary references to be used
- Role of OpenCL in problem solution

If you still don't have a final project idea or are unsure

- email me
- come to Sunday office hour
- look at the course website page:  
<https://canvas.uw.edu/courses/1260593/pages/final-project-ideas>

# Final Project Details

## ► Desired Elements to Include

1. Theoretical analysis of expected performance
  - a. Use the theory and metrics we've studied:
    - a. Operational Intensity
    - b. Speedup, Efficiency (theoretical/asymptotic, measured/empirical), serial fraction
2. Optimization of NDRange configuration (workgroup, workitem sizes)
3. Performance profiling
  - a. host side timing : Windows Performance Counters (WPC)
  - b. device side timing: various events QUEUED-COMPLETE, START-END, etc...
4. VTune analysis of performance
  - a. Execution Unit (EU) utilization, occupancy
  - b. Memory subsystem utilization
  - c. Bottleneck analysis



# Spatial-Domain Image Processing

Image Enhancement with OpenCL

# Corrections to the Conv Kernel

```
__kernel void img_conv_filter(__read_only image2d_t inImg, __write_only image2d_t outImg, __constant float* convfilter, uint filtWidth)
{
    // use global IDs for output coords
    int x = get_global_id(0); // columns
    int y = get_global_id(1); // rows

    int halfWidth = (int)(filtWidth/2); // auto-round nearest int

    float sum = 0.0f;

    int filtIdx = 0; // filter kernel passed in as linearized buffer array
    int2 coords;

    for(int i = -halfWidth; i <= halfWidth; i++) // iterate filter rows
    {
        coords.y = y + i;
        for(int j = -halfWidth; j <= halfWidth; j++) // iterate filter cols
        {
            coords.x = x + j;
            float pixel = convert_float(read_imageui(inImg, sampler, coords).x); // operate on single component (x = r)
            sum += pixel * convfilter[filtIdx];
            filtIdx++;
        }
    }

    //write resultant filtered pixel to output image
    coords = (int2)(x,y);
    write_imageui(outImg, coords, convert_uint4((float4)(sum,sum,sum,1.0f))); // leave a-channel unchanged
}
```



# Spatial-domain image processing

Original



- ▶ Fix for non-32bit images
  - ▶ use `stbi_load( )` : force last argument = 4
  - ▶ also ensure you allocate your host-side output-image buffer size accordingly

## Example Results

2D Sobel



Basic Laplacian



Composite Laplacian



5x5 Gaussian Blur



Full Chain





# In-class Exercise 8a

- ▶ Implement the corrections your stencil kernels
- ▶ Test the Gauss Blur, Sobel, and Laplacian kernels
- ▶ Further experiments
  - ▶ Sobel – one dimension only
  - ▶ Laplacian – basic versus composite
- ▶ Get the full Multiqueue Device-side enqueue processing chain working!
  - ▶ Between consecutive kernel device-enqueues...
    - ▶ be careful with input-output image object re-ordering
    - ▶ be careful with input-output event synchronization
  - ▶ Also on host when enqueueing ReadImage... which image is real output??





# Frequency-Domain Image Processing

Image Enhancement with OpenCL





# cIFFT

- Go here: <http://clmathlibraries.github.io/cIFFT/index.html#Outline>
- We'll walk through a code example...




# Image Filtering in Frequency Domain

- ▶ Image smoothing with Low-pass filters
  - ▶ Image sharpening with High-pass filters
- 



# Parallelizing PDEs



- ▶ A VAST topic area
  - ▶ Numerical methods for PDEs is a VAST subset
    - ▶ Parallel numerical methods for PDEs is a smaller but increasingly important sub-subset
- ▶ We won't even begin to scratch the surface in this class
- ▶ Will give a brief taste of some approaches, considerations, and techniques
  - ▶ some basic techniques are very general and common
- ▶ Will gloss over a lot of important details in order to focus on parallelization
  - ▶ consistency, stability, convergence
  - ▶ underpinnings of various finite methods

# PDEs

## ► Classification:

### ► Elliptical

► Laplace's Equation  $\nabla^2 u(x, y, z) = 0$

► Poisson Equations  $\nabla^2 u(x, y, z) = f(x, y, z)$

### ► Hyperbolic

► 2<sup>nd</sup>-order Wave Equation,

$$\frac{\partial^2 u}{\partial t^2} = c^2 \frac{\partial^2 u}{\partial x^2}$$

► 1D Convection (Advection) Equation

$$\frac{\partial u}{\partial t} + c \frac{\partial u}{\partial x} = 0$$

### ► Parabolic

► Heat (Diffusion) Equation

$$\frac{\partial u}{\partial t} = \alpha \frac{\partial^2 u}{\partial x^2}$$



# 1D Linear Convection Equation

## Finite Different Method (FDM)

- Step 1: Discretize the continuous equation

$$\frac{\partial u}{\partial t} + c \frac{\partial u}{\partial x} = 0 \quad \longrightarrow \quad \frac{u_j^{n+1} - u_j^n}{\Delta t} + c \frac{u_{j+1}^{n+1} - u_{j-1}^{n+1}}{2\Delta x} = 0$$

- Truncation Error (T.E.)  $O[\Delta t, (\Delta x)^2]$
- Step 2 : rewrite with unknowns at time (n+1) on LHS, known  $u_j^n$  on RHS


$$a_{j,j+1}^{n+1} + d_j u_j^{n+1} + b_j u_{j-1}^{n+1} = C_j$$

- where  $a_j = \frac{v}{2}$  and the Courant (CFL) number is  $v = c \frac{\Delta t}{\Delta x}$   
 $d_j = 1$   
 $b_j = \frac{v}{2}$   
 $C_j = u_j^n$

- This is an implicit scheme which is *unconditionally stable* for all time steps

# 1D Linear Convection Equation : Implicit FD scheme

- Consider computational mesh with  $M+2$  grid points in  $x$  direction
  - known initial conditions at  $n=0$ , boundary conditions  $u_0^{n+1}$  and  $u_{M+1}^{n+1}$



- A **tridiagonal** matrix results from solving the system of  $M$  linear algebraic equations at each  $(n+1)$  time level

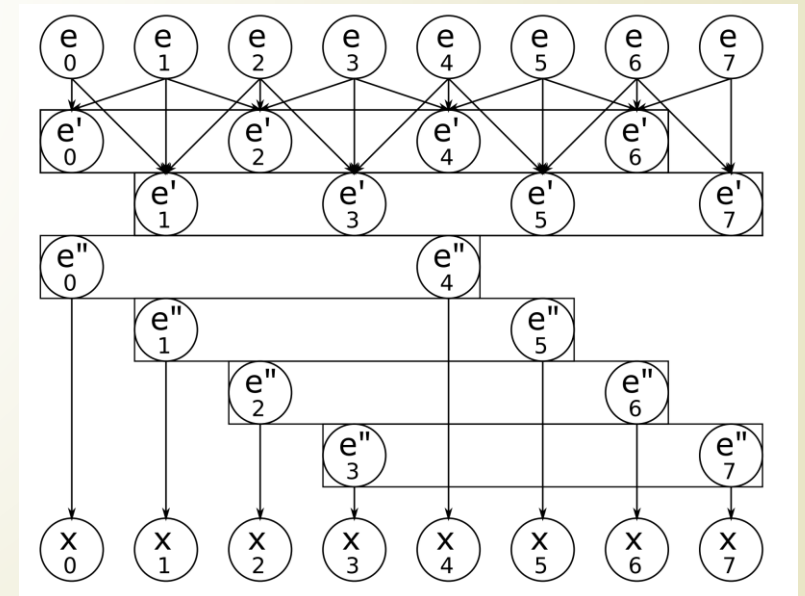
$$\begin{array}{c} [A] \end{array}
 \begin{bmatrix} d_1 & a_1 & 0 & \cdot & \cdot & \cdot & \cdot & 0 \\ b_2 & d_2 & a_2 & & & & & \cdot \\ 0 & b_3 & d_3 & a_3 & & & & \cdot \\ \cdot & & & & & & & \cdot \\ \cdot & & & & & & & \cdot \\ \cdot & & & & & & & 0 \\ \cdot & & & & & b_{M-1} & d_{M-1} & a_{M-1} \\ 0 & \cdot & \cdot & \cdot & \cdot & 0 & b_M & d_M \end{bmatrix}
 \begin{array}{c} [u] \end{array}
 \begin{bmatrix} u_1^{n+1} \\ u_2^{n+1} \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ u_{M-1}^{n+1} \\ u_M^{n+1} \end{bmatrix}
 =
 \begin{array}{c} [C] \end{array}
 \begin{bmatrix} C_1 \\ C_2 \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ C_{M-1} \\ C_M \end{bmatrix}$$

# Parallel Tridiagonal Solver

- Parallel Cyclic Reduction (PCR) (*Hockney, 1981*)
  - only used forward-reduction phase
  - performed on both odd and even rows

$$\begin{matrix} e_1 \\ e_2 \\ e_3 \\ e_4 \end{matrix} \begin{bmatrix} b_1 & c_1 & & \\ a_2 & b_2 & c_2 & \\ & a_3 & b_3 & c_3 \\ & & a_4 & b_4 \end{bmatrix} \rightarrow \begin{matrix} e'_1 \\ e'_2 \\ e'_3 \\ e'_4 \end{matrix} \begin{bmatrix} b'_1 & 0 & c'_1 & \\ 0 & b'_2 & 0 & c'_2 \\ a'_3 & 0 & b'_3 & 0 \\ & a'_4 & 0 & b'_4 \end{bmatrix} \rightarrow \frac{\begin{bmatrix} b'_1 & c'_1 \\ a'_3 & b'_3 \end{bmatrix}}{\begin{bmatrix} b'_2 & c'_2 \\ a'_4 & b'_4 \end{bmatrix}}$$

- Data access pattern of PCR algorithm
  - multi-level tree decomposition
  - obtain independent tasks



- Computational complexity:  $O(n \log n)$
- Required number of elimination steps:  $\log n + 1$

# 2D Shallow Water Equations (SWEs)

$$\begin{bmatrix} h \\ hu \\ hv \end{bmatrix}_t + \begin{bmatrix} hu \\ hu^2 + \frac{1}{2}gh^2 \\ huv \end{bmatrix}_x + \begin{bmatrix} hv \\ huv \\ hv^2 + \frac{1}{2}gh^2 \end{bmatrix}_y = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

Matrix form:

$$Q_t + F(Q)_x + G(Q)_y = 0$$

$Q$  is vector of conserved variables

$F, G$  are flux functions

$h$  is water depth

$hu, hv$  are momenta components along  $x, y$  directions

$g$  is gravitational acceleration

## Finite Volume Method (FVM)

➤ Step 1: Discretize the continuous equations : *Lax-Friedrichs explicit scheme*

$$Q_{ij}^{n+1} = \frac{1}{4} (Q_{i,j+1}^n + Q_{i,j-1}^n + Q_{i+1,j}^n + Q_{i-1,j}^n) - \frac{\Delta t}{2\Delta x} [F(Q_{i+1,j}^n) - F(Q_{i-1,j}^n)] - \frac{\Delta t}{2\Delta y} [G(Q_{i,j+1}^n) - G(Q_{i,j-1}^n)]$$

➤ This is a stencil form!

➤ four nearest spatial neighbors

➤ Time-marching method

# Parallel 2D SWEs

$$Q_{ij}^{n+1} = \frac{1}{4}(Q_{i,j+1}^n + Q_{i,j-1}^n + Q_{i+1,j}^n + Q_{i-1,j}^n) - \frac{\Delta t}{2\Delta x} [F(Q_{i+1,j}^n) - F(Q_{i-1,j}^n)] - \frac{\Delta t}{2\Delta y} [G(Q_{i,j+1}^n) - G(Q_{i,j-1}^n)]$$

## Serial Algorithm (for computing $h^{n+1}$ )

```
for (int j=1; j<ny-1; ++j) {
    for (int i=1; i<nx-1; ++i) {
        int n=(j+1)*nx+i, s=(j-1)*nx+i, e=j*n+nx+i+1, w=j*n+nx+i-1;
        h_new[j*n+nx+i] = 0.25*(h[n]+h[s]+h[e]+h[w])
            - 0.5*dt/dx*(hu[e]-hu[w]) - 0.5*dt/dy*(hv[n]-hv[s]);
    }
}
```

## Simple OpenCL Kernel

```
__kernel void SWE_LaxFriedrichs(global float* h_new, global float* h, global float* hu,
                                global float* hv, int nx, int ny)
{
    int i = get_global_id(0);
    int j = get_global_id(1);
    if (i > 0 && i < nx-1 && j > 0 && j < ny-1) {
        int n=(j+1)*nx+i, s=(j-1)*nx+i, e=j*n+nx+i+1, w=j*n+nx+i-1;
        h_new[j*n+nx+i] = 0.25*(h[n]+h[s]+h[e]+h[w])
            - 0.5*dt/dx*(hu[e]-hu[w]) - 0.5*dt/dy*(hv[n]-hv[s]);
    }
}
```





# In-class Exercise 8b

- ▶ cIFFT
  - ▶ See procedures on class website </InclassExercises/Ex7>
- 