

# The **OpenCL** Specification

*Version: 2.0*

*Document Revision: 29*

Khronos OpenCL Working Group

*Editors:* Lee Howes and Aaftab Munshi

<b>1.</b>	<b>INTRODUCTION.....</b>	<b>11</b>
<b>2.</b>	<b>GLOSSARY.....</b>	<b>13</b>
<b>3.</b>	<b>THE OPENCL ARCHITECTURE.....</b>	<b>24</b>
<b>3.1</b>	<b>Platform Model.....</b>	<b>24</b>
<b>3.2</b>	<b>Execution Model.....</b>	<b>26</b>
3.2.1	Execution Model: Mapping work-items onto an NDRange .....	29
3.2.2	Execution Model: Execution of kernel-instances .....	31
3.2.3	Execution Model: Device-side enqueue .....	32
3.2.4	Execution Model: Synchronization .....	33
3.2.5	Execution Model: Categories of Kernels.....	34
<b>3.3</b>	<b>Memory Model .....</b>	<b>35</b>
3.3.1	Memory Model: Fundamental Memory Regions .....	35
3.3.2	Memory Model: Memory Objects.....	38
3.3.3	Memory Model: Shared Virtual Memory.....	40
3.3.4	Memory Model: Memory Consistency Model .....	41
3.3.5	Memory Model: Overview of atomic and fence operations .....	44
3.3.6	Memory Model: Memory Ordering Rules.....	45
3.3.6.1	Memory Ordering Rules: Atomic Operations.....	48
3.3.6.2	Memory Ordering Rules: Fence Operations.....	51
3.3.6.3	Memory Ordering Rules: Work-group Functions.....	53
3.3.6.4	Memory Ordering Rules: Host-side and Device-side Commands .....	53
<b>3.4</b>	<b>The OpenCL Framework .....</b>	<b>56</b>
3.4.1	OpenCL Framework: Mixed Version Support .....	57
<b>4.</b>	<b>THE OPENCL PLATFORM LAYER.....</b>	<b>58</b>
<b>4.1</b>	<b>Querying Platform Info .....</b>	<b>58</b>
<b>4.2</b>	<b>Querying Devices.....</b>	<b>60</b>
<b>4.3</b>	<b>Partitioning a Device.....</b>	<b>79</b>
<b>4.4</b>	<b>Contexts.....</b>	<b>84</b>
<b>5.</b>	<b>THE OPENCL RUNTIME .....</b>	<b>91</b>
<b>5.1</b>	<b>Command Queues .....</b>	<b>91</b>
<b>5.2</b>	<b>Buffer Objects.....</b>	<b>97</b>
5.2.1	Creating Buffer Objects.....	97
5.2.2	Reading, Writing and Copying Buffer Objects .....	103
5.2.3	Filling Buffer Objects.....	115
5.2.4	Mapping Buffer Objects .....	117

<b>5.3</b>	<b>Image Objects .....</b>	<b>121</b>
5.3.1	Creating Image Objects .....	121
5.3.1.1	Image Format Descriptor .....	124
5.3.1.2	Image Descriptor .....	127
5.3.2	Querying List of Supported Image Formats .....	129
5.3.2.1	Minimum List of Supported Image Formats .....	130
5.3.2.2	Image format mapping to OpenCL C image access qualifiers .....	132
5.3.3	Reading, Writing and Copying Image Objects .....	133
5.3.4	Filling Image Objects .....	140
5.3.5	Copying between Image and Buffer Objects .....	142
5.3.6	Mapping Image Objects .....	147
5.3.7	Image Object Queries .....	150
<b>5.4</b>	<b>Pipes .....</b>	<b>153</b>
5.4.1	Creating Pipe Objects .....	153
5.4.2	Pipe Object Queries .....	154
<b>5.5</b>	<b>Querying, Unmapping, Migrating, Retaining and Releasing Memory Objects .....</b>	<b>156</b>
5.5.1	Retaining and Releasing Memory Objects .....	156
5.5.2	Unmapping Mapped Memory Objects .....	158
5.5.3	Accessing mapped regions of a memory object .....	160
5.5.4	Migrating Memory Objects .....	161
5.5.5	Memory Object Queries .....	163
<b>5.6</b>	<b>Shared Virtual Memory .....</b>	<b>167</b>
5.6.1	SVM sharing granularity: coarse- and fine- grained sharing .....	167
5.6.2	Memory consistency for SVM allocations .....	180
<b>5.7</b>	<b>Sampler Objects .....</b>	<b>182</b>
5.7.1	Creating Sampler Objects .....	182
5.7.2	Sampler Object Queries .....	184
<b>5.8</b>	<b>Program Objects .....</b>	<b>186</b>
5.8.1	Creating Program Objects .....	186
5.8.2	Building Program Executables .....	192
5.8.3	Separate Compilation and Linking of Programs .....	194
5.8.4	Compiler Options .....	200
5.8.4.1	Preprocessor options .....	200
5.8.4.2	Math Intrinsics Options .....	200
5.8.4.3	Optimization Options .....	201
5.8.4.4	Options to Request or Suppress Warnings .....	202
5.8.4.5	Options Controlling the OpenCL C version .....	203
5.8.4.6	Options for Querying Kernel Argument Information .....	203
5.8.4.7	Options for debugging your program .....	204
5.8.5	Linker Options .....	204
5.8.5.1	Library Linking Options .....	204
5.8.5.2	Program Linking Options .....	204
5.8.6	Unloading the OpenCL Compiler .....	205
5.8.7	Program Object Queries .....	206
<b>5.9</b>	<b>Kernel Objects .....</b>	<b>213</b>
5.9.1	Creating Kernel Objects .....	213
5.9.2	Setting Kernel Arguments .....	216
5.9.3	Kernel Object Queries .....	222
<b>5.10</b>	<b>Executing Kernels .....</b>	<b>230</b>

5.11	Event Objects .....	237
5.12	Markers, Barriers and Waiting for Events .....	247
5.13	Out-of-order Execution of Kernels and Memory Object Commands .....	250
5.14	Profiling Operations on Memory Objects and Kernels .....	251
5.15	Flush and Finish .....	254
6.	THE OPENCL C PROGRAMMING LANGUAGE .....	256
7.	OPENCL NUMERICAL COMPLIANCE .....	256
8.	IMAGE ADDRESSING AND FILTERING .....	256
9.	OPTIONAL EXTENSIONS .....	256
10.	OPENCL EMBEDDED PROFILE .....	257
	APPENDIX A .....	266
A.1	Shared OpenCL Objects .....	266
A.2	Multiple Host Threads .....	267
	APPENDIX B — PORTABILITY .....	268
	APPENDIX C — APPLICATION DATA TYPES .....	273
C.1	Shared Application Scalar Data Types .....	273
C.2	Supported Application Vector Data Types .....	273
C.3	Alignment of Application Data Types .....	274
C.4	Vector Literals .....	274
C.5	Vector Components .....	274
	C.5.1 Named vector components notation .....	275
	C.5.2 High/Low vector component notation .....	275
	C.5.3 Native vector type notation .....	276
C.6	Implicit Conversions .....	276
C.7	Explicit Casts .....	276
C.8	Other operators and functions .....	277

<b>C.9</b>	<b>Application constant definitions.....</b>	<b>277</b>
	<b>APPENDIX D — OPENCL C++ WRAPPER API.....</b>	<b>279</b>
	<b>APPENDIX E — CL_MEM_COPY_OVERLAP.....</b>	<b>280</b>
	<b>APPENDIX F – CHANGES.....</b>	<b>282</b>
<b>F.1</b>	<b>Summary of changes from OpenCL 1.0 .....</b>	<b>282</b>
<b>F.2</b>	<b>Summary of changes from OpenCL 1.1 .....</b>	<b>284</b>
<b>F.3</b>	<b>Summary of changes from OpenCL 1.2 .....</b>	<b>285</b>
	<b>INDEX - APIS .....</b>	<b>287</b>

Copyright (c) 2008-2014 The Khronos Group Inc. All Rights Reserved.

This specification is protected by copyright laws and contains material proprietary to the Khronos Group, Inc. It or any components may not be reproduced, republished, distributed, transmitted, displayed, broadcast or otherwise exploited in any manner without the express prior written permission of Khronos Group. You may use this specification for implementing the functionality therein, without altering or removing any trademark, copyright or other notice from the specification, but the receipt or possession of this specification does not convey any rights to reproduce, disclose, or distribute its contents, or to manufacture, use, or sell anything that it may describe, in whole or in part.

Khronos Group grants express permission to any current Promoter, Contributor or Adopter member of Khronos to copy and redistribute UNMODIFIED versions of this specification in any fashion, provided that NO CHARGE is made for the specification and the latest available update of the specification for any version of the API is used whenever possible. Such distributed specification may be re-formatted AS LONG AS the contents of the specification are not changed in any way. The specification may be incorporated into a product that is sold as long as such product includes significant independent work developed by the seller. A link to the current version of this specification on the Khronos Group web-site should be included whenever possible with specification distributions.

Khronos Group makes no, and expressly disclaims any, representations or warranties, express or implied, regarding this specification, including, without limitation, any implied warranties of merchantability or fitness for a particular purpose or non-infringement of any intellectual property. Khronos Group makes no, and expressly disclaims any, warranties, express or implied, regarding the correctness, accuracy, completeness, timeliness, and reliability of the specification. Under no circumstances will the Khronos Group, or any of its Promoters, Contributors or Members or their respective partners, officers, directors, employees, agents or representatives be liable for any damages, whether direct, indirect, special or consequential damages for lost revenues, lost profits, or otherwise, arising from or in connection with these materials.

Khronos, StreamInput, WebGL, COLLADA, OpenKODE, OpenVG, OpenWF, OpenSL ES, OpenMAX, OpenMAX AL, OpenMAX IL and OpenMAX DL are trademarks and WebCL is a certification mark of the Khronos Group Inc. OpenCL is a trademark of Apple Inc. and OpenGL and OpenML are registered trademarks and the OpenGL ES and OpenGL SC logos are trademarks of Silicon Graphics International used under license by Khronos. All other product names, trademarks, and/or company names are used solely for identification and belong to their respective owners.

# Acknowledgements

The OpenCL specification is the result of the contributions of many people, representing a cross section of the desktop, hand-held, and embedded computer industry. Following is a partial list of the contributors, including the company that they represented at the time of their contribution:

Chuck Rose, Adobe  
Eric Berdahl, Adobe  
Shivani Gupta, Adobe  
David Neto, Altera  
Bill Licea Kane, AMD  
Ed Buckingham, AMD  
Jan Civlin, AMD  
Laurent Morichetti, AMD  
Lee Howes, AMD  
Mark Fowler, AMD  
Michael Mantor, AMD  
Norm Rubin, AMD  
Ofer Rosenberg, AMD  
Victor Odintsov, AMD  
Aaftab Munshi, Apple  
Abe Stephens, Apple  
Alexandre Namaan, Apple  
Anna Tikhonova, Apple  
Chendi Zhang, Apple  
Eric Bainville, Apple  
David Hayward, Apple  
Giridhar Murthy, Apple  
Ian Ollmann, Apple  
Inam Rahman, Apple  
James Shearer, Apple  
MonPing Wang, Apple  
Tanya Lattner, Apple  
Mikael Bourges-Sevenier, Aptina  
Anton Lokhmotov, ARM  
Dave Shreiner, ARM  
Hedley Francis, ARM  
Robert Elliott, ARM  
Scott Moyers, ARM  
Tom Olson, ARM  
Christopher Thompson-Walsh, Broadcom  
Holger Waechtler, Broadcom  
Norman Rink, Broadcom

Andrew Richards, Codeplay  
Maria Rovatsou, Codeplay  
Stephen Frye, Electronic Arts  
Eric Schenk, Electronic Arts  
Brian Horton, IBM  
Brian Watt, IBM  
Gordon Fossum, IBM  
Greg Bellows, IBM  
Joaquin Madruga, IBM  
Mark Nutter, IBM  
Mike Perks, IBM  
Jon Parr, Imagination Technologies  
Robert Quill, Imagination Technologies  
James McCarthy, Imagination Technologies  
Aaron Kunze, Intel  
Aaron Lefohn, Intel  
Adam Lake, Intel  
Allen Hux, Intel  
Andrew Brownsword, Intel  
Andrew Lauritzen, Intel  
Brian Lewis, Intel  
Geoff Berry, Intel  
Jayanth Rao, Intel  
Josh Fryman, Intel  
Hong Jiang, Intel  
Larry Seiler, Intel  
Mike MacPherson, Intel  
Murali Sundaresan, Intel  
Paul Lalonde, Intel  
Stephen Junkins, Intel  
Tim Foley, Intel  
Timothy Mattson, Intel  
Yariv Aridor, Intel  
Jon Leech, Khronos  
Benjamin Bergen, Los Alamos National Laboratory  
Amit Rao, NVIDIA  
Ashish Srivastava, NVIDIA  
Bastiaan Aarts, NVIDIA  
Chris Cameron, NVIDIA  
Christopher Lamb, NVIDIA  
Dibyapran Sanyal, NVIDIA  
Guatam Chakrabarti, NVIDIA  
Ian Buck, NVIDIA



Jaydeep Marathe, NVIDIA  
Jian-Zhong Wang, NVIDIA  
Karthik Raghavan Ravi, NVIDIA  
Kedar Patil, NVIDIA  
Manjunath Kudlur, NVIDIA  
Mark Harris, NVIDIA  
Michael Gold, NVIDIA  
Neil Trevett, NVIDIA  
Richard Johnson, NVIDIA  
Sean Lee, NVIDIA  
Tushar Kashalikar, NVIDIA  
Vinod Grover, NVIDIA  
Xiangyun Kong, NVIDIA  
Yogesh Kini, NVIDIA  
Yuan Lin, NVIDIA  
Allan Tzeng, QUALCOMM  
Alex Bourd, QUALCOMM  
Anirudh Acharya, QUALCOMM  
Andrew Gruber, QUALCOMM  
Andrzej Mamona, QUALCOMM  
Benedict Gaster, QUALCOMM  
Bill Torzewski, QUALCOMM  
Bob Rychlik, QUALCOMM  
Chihong Zhang, QUALCOMM  
Chris Mei, QUALCOMM  
Colin Sharp, QUALCOMM  
David Garcia, QUALCOMM  
David Ligon, QUALCOMM  
Jay Yun, QUALCOMM  
Robert Simpson, QUALCOMM  
Sumesh Udayakumaran, QUALCOMM  
Vineet Goel, QUALCOMM  
Tasneem Brutch, Samsung  
Thierry Lepley, STMicroelectronics  
Alan Ward, Texas Instruments  
Simon McIntosh-Smith, University of Bristol  
Brian Hutsell Vivante  
Mike Cai, Vivante  
Sumeet Kumar, Vivante  
Jeff Fifield, Xilinx  
Henry Styles, Xilinx  
Ralph Wittig, Xilinx



# 1. Introduction






Modern processor architectures have embraced parallelism as an important pathway to increased performance. Facing technical challenges with higher clock speeds in a fixed power envelope, Central Processing Units (CPUs) now improve performance by adding multiple cores. Graphics Processing Units (GPUs) have also evolved from fixed function rendering devices into programmable parallel processors. As today's computer systems often include highly parallel CPUs, GPUs and other types of processors, it is important to enable software developers to take full advantage of these heterogeneous processing platforms.

Creating applications for heterogeneous parallel processing platforms is challenging as traditional programming approaches for multi-core CPUs and GPUs are very different. CPU-based parallel programming models are typically based on standards but usually assume a shared address space and do not encompass vector operations. General purpose GPU programming models address complex memory hierarchies and vector operations but are traditionally platform-, vendor- or hardware-specific. These limitations make it difficult for a developer to access the compute power of heterogeneous CPUs, GPUs and other types of processors from a single, multi-platform source code base. More than ever, there is a need to enable software developers to effectively take full advantage of heterogeneous processing platforms – from high performance compute servers, through desktop computer systems to handheld devices - that include a diverse mix of parallel CPUs, GPUs and other processors such as DSPs and the Cell/B.E. processor.

**OpenCL** (Open Computing Language) is an open royalty-free standard for general purpose parallel programming across CPUs, GPUs and other processors, giving software developers portable and efficient access to the power of these heterogeneous processing platforms.

OpenCL supports a wide range of applications, ranging from embedded and consumer software to HPC solutions, through a low-level, high-performance, portable abstraction. By creating an efficient, close-to-the-metal programming interface, OpenCL will form the foundation layer of a parallel computing ecosystem of platform-independent tools, middleware and applications. OpenCL is particularly suited to play an increasingly significant role in emerging interactive graphics applications that combine general parallel compute algorithms with graphics rendering pipelines.

OpenCL consists of an API for coordinating parallel computation across heterogeneous processors; and a cross-platform programming language with a well-specified computation environment. The OpenCL standard:

-  Supports both data- and task-based parallel programming models
-  Utilizes a subset of ISO C99 with extensions for parallelism
-  Defines consistent numerical requirements based on IEEE 754
-  Defines a configuration profile for handheld and embedded devices
-  Efficiently interoperates with OpenGL, OpenGL ES and other graphics APIs

This document begins with an overview of basic concepts and the architecture of OpenCL, followed by a detailed description of its execution model, memory model and synchronization support. It then discusses the OpenCL platform and runtime API and is followed by a detailed description of the OpenCL C programming language. Some examples are given that describe sample compute use-cases and how they would be written in OpenCL. The specification is divided into a core specification that any OpenCL compliant implementation must support; a handheld/embedded profile which relaxes the OpenCL compliance requirements for handheld and embedded devices; and a set of optional extensions that are likely to move into the core specification in later revisions of the OpenCL specification.

## 2. Glossary

**Application:** The combination of the program running on the *host* and *OpenCL devices*.



**Acquire semantics:** One of the memory order semantics defined for synchronization operations. Acquire semantics apply to atomic operations that load from memory. Given two units of execution, **A** and **B**, acting on a shared atomic object **M**, if **A** uses an atomic load of **M** with acquire semantics to synchronize-with an atomic store to **M** by **B** that used release semantics, then **A**'s atomic load will occur before any subsequent operations by **A**. Note that the memory orders *release*, *sequentially consistent*, and *acquire\_release* all include *release semantics* and effectively pair with a load using acquire semantics.

**Acquire release semantics:** A memory order semantics for synchronization operations (such as atomic operations) that has the properties of both acquire and release memory orders. It is used with read-modify-write operations.


**Atomic operations:** Operations that at any point, and from any perspective, have either occurred completely, or not at all. Memory orders associated with atomic operations may constrain the visibility of loads and stores with respect to the atomic operations (see *relaxed semantics*, *acquire semantics*, *release semantics* or *acquire release semantics*).

**Blocking and Non-Blocking Enqueue API calls:** A *non-blocking enqueue API call* places a *command* on a *command-queue* and returns immediately to the host. The *blocking-mode enqueue API calls* do not return to the host until the command has completed.

**Barrier:** There are two types of *barriers* – a *command-queue barrier* and a *work-group barrier*.

-  The OpenCL API provides a function to enqueue a *command-queue barrier* command. This *barrier* command ensures that all previously enqueued commands to a *command-queue* have finished execution before any following *commands* enqueued in the *command-queue* can begin execution.
-  The OpenCL C programming language provides a built-in *work-group barrier* function. This *barrier* built-in function can be used by a *kernel* executing on a *device* to perform synchronization between *work-items* in a *work-group* executing the *kernel*. All the *work-items* of a *work-group* must execute the *barrier* construct before any are allowed to continue execution beyond the *barrier*.

**Buffer Object:** A memory object that stores a linear collection of bytes. Buffer objects are accessible using a pointer in a *kernel* executing on a *device*. Buffer objects can be manipulated by the host using OpenCL API calls. A *buffer object* encapsulates the following information:

-  Size in bytes.

- ✚ Properties that describe usage information and which region to allocate from.
- ✚ Buffer data.

**Built-in Kernel:** A *built-in kernel* is a *kernel* that is executed on an OpenCL *device* or *custom device* by fixed-function hardware or in firmware. *Applications* can query the *built-in kernels* supported by a *device* or *custom device*. A *program object* can only contain *kernels* written in OpenCL C or *built-in kernels* but not both. See also *Kernel* and *Program*.

**Child kernel:** see *device-side enqueue*.

**Command:** The OpenCL operations that are submitted to a *command-queue* for execution. For example, OpenCL commands issue kernels for execution on a compute device, manipulate memory objects, etc.

**Command-queue:** An object that holds *commands* that will be executed on a specific *device*. The *command-queue* is created on a specific *device* in a *context*. *Commands* to a *command-queue* are queued in-order but may be executed in-order or out-of-order. Refer to *In-order Execution* and *Out-of-order Execution*.

**Command-queue Barrier.** See *Barrier*.

**Command synchronization:** Constraints on the order that commands are launched for execution on a device defined in terms of the synchronization points that occur between commands in host command-queues and between commands in device-side command-queues. See *synchronization points*.

**Complete:** The final state in the six state model for the execution of a command. The transition into this state occurs is signaled through event objects or callback functions associated with a command.

**Compute Device Memory:** This refers to one or more memories attached to the compute device.

**Compute Unit:** An OpenCL *device* has one or more *compute units*. A *work-group* executes on a single *compute unit*. A *compute unit* is composed of one or more *processing elements* and *local memory*. A *compute unit* may also include dedicated texture filter units that can be accessed by its processing elements.

**Concurrency:** A property of a system in which a set of tasks in a system can remain active and make progress at the same time. To utilize concurrent execution when running a program, a programmer must identify the concurrency in their problem, expose it within the source code, and then exploit it using a notation that supports concurrency.

**Constant Memory:** A region of *global memory* that remains constant during the execution of a *kernel*. The *host* allocates and initializes memory objects placed into *constant memory*.

**Context:** The environment within which the *kernels* execute and the domain in which synchronization and memory management is defined. The *context* includes a set of *devices*, the memory accessible to those *devices*, the corresponding memory properties and one or more *command-queues* used to schedule execution of a *kernel(s)* or operations on *memory objects*.

**Control flow:** The flow of instructions executed by a work-item. Multiple logically related work items may or may not execute the same control flow. The control flow is said to be *converged* if all the work-items in the set execution the same stream of instructions. In a *diverged* control flow, the work-items in the set execute different instructions. At a later point, if a diverged control flow becomes converged, it is said to be a re-converged control flow.

**Converged control flow:** see **control flow**.

**Custom Device:** An OpenCL *device* that fully implements the OpenCL Runtime but does not support *programs* written in OpenCL C. A custom device may be specialized non-programmable hardware that is very power efficient and performant for directed tasks or hardware with limited programmable capabilities such as specialized DSPs. Custom devices are not OpenCL conformant. Custom devices may support an online compiler. Programs for custom devices can be created using the OpenCL runtime APIs that allow OpenCL programs to be created from source (if an online compiler is supported) and/or binary, or from *built-in kernels* supported by the *device*. See also *Device*.

**Data Parallel Programming Model:** Traditionally, this term refers to a programming model where concurrency is expressed as instructions from a single program applied to multiple elements within a set of data structures. The term has been generalized in OpenCL to refer to a model wherein a set of instructions from a single program are applied concurrently to each point within an abstract domain of indices.

**Data race:** The execution of a program contains a data race if it contains two actions in different work items or host threads where (1) one action modifies a memory location and the other action reads or modifies the same memory location, and (2) at least one of these actions is not atomic, or the corresponding memory scopes are not inclusive, and (3) the actions are global actions unordered by the global-happens-before relation or are local actions unordered by the local-happens before relation.

**Device:** A *device* is a collection of *compute units*. A *command-queue* is used to queue *commands* to a *device*. Examples of *commands* include executing *kernels*, or reading and writing *memory objects*. OpenCL devices typically correspond to a GPU, a multi-core CPU, and other processors such as DSPs and the Cell/B.E. processor.

**Device-side enqueue:** A mechanism whereby a kernel-instance is enqueued by a kernel-instance running on a device without direct involvement by the host program. This produces *nested parallelism*; i.e. additional levels of concurrency are nested inside a running kernel-instance. The kernel-instance executing on a device (the *parent kernel*) enqueues a kernel-instance (the *child kernel*) to a device-side command queue. Child and parent kernels execute asynchronously though a parent kernel does not complete until all of its child-kernels have completed.

**Diverged control flow:** see *control flow*.

**Ended:** The fifth state in the six state model for the execution of a command. The transition into this state occurs when execution of a command has ended. When a Kernel-enqueue command ends, all of the work-groups associated with that command have finished their execution.

**Event Object:** An *event object* encapsulates the status of an operation such as a *command*. It can be used to synchronize operations in a context.

**Event Wait List:** An *event wait list* is a list of *event objects* that can be used to control when a particular *command* begins execution.

**Fence:** A memory ordering operation without an associated atomic object. A fence can use the *acquire semantics*, *release semantics*, or *acquire release semantics*.

**Framework:** A software system that contains the set of components to support software development and execution. A *framework* typically includes libraries, APIs, runtime systems, compilers, etc.

**Generic address space:** An address space that include the *private*, *local*, and *global* address spaces available to a device. The generic address space supports conversion of pointers to and from private, local and global address spaces, and hence lets a programmer write a single function that at compile time can take arguments from any of the three named address spaces.

**Global Happens before:** see *happens before*.

**Global ID:** A *global ID* is used to uniquely identify a *work-item* and is derived from the number of *global work-items* specified when executing a *kernel*. The *global ID* is a N-dimensional value that starts at (0, 0, ... 0). See also *Local ID*.

**Global Memory:** A memory region accessible to all *work-items* executing in a *context*. It is accessible to the *host* using *commands* such as read, write and map. *Global memory* is included within the *generic address space* that includes the private and local address spaces.

**GL share group:** A *GL share group* object manages shared OpenGL or OpenGL ES resources such as textures, buffers, framebuffers, and renderbuffers and is associated with one or more GL context objects. The *GL share group* is typically an opaque object and not directly accessible.

**Handle:** An opaque type that references an *object* allocated by OpenCL. Any operation on an *object* occurs by reference to that object's handle.

**Happens before:** An ordering relationship between operations that execute on multiple units of execution. If an operation A happens-before operation B then A must occur before B; in particular, any value written by A will be visible to B. We define two separate happens before relations: *global-happens-before* and *local-happens-before*. These are defined in section 3.3.6.



**Host:** The *host* interacts with the *context* using the OpenCL API.

**Host-thread:** the unit of execution that executes the statements in the Host program.

**Host pointer:** A pointer to memory that is in the virtual address space on the *host*.

**Illegal:** Behavior of a system that is explicitly not allowed and will be reported as an error when encountered by OpenCL.

**Image Object:** A *memory object* that stores a two- or three- dimensional structured array. Image data can only be accessed with read and write functions. The read functions use a *sampler*.

The *image object* encapsulates the following information:

- ✚ Dimensions of the image.
- ✚ Description of each element in the image.
- ✚ Properties that describe usage information and which region to allocate from.
- ✚ Image data.

The elements of an image are selected from a list of predefined image formats.

**Implementation Defined:** Behavior that is explicitly allowed to vary between conforming implementations of OpenCL. An OpenCL implementor is required to document the implementation-defined behavior.

**In-order Execution:** A model of execution in OpenCL where the *commands* in a *command-queue* are executed in order of submission with each *command* running to completion before the next one begins. See *Out-of-order Execution*.

**Kernel:** A *kernel* is a function declared in a *program* and executed on an OpenCL *device*. A *kernel* is identified by the `__kernel` or `kernel` qualifier applied to any function defined in a *program*.

**Kernel-instance:** The work carried out by an OpenCL program occurs through the execution of kernel-instances on devices. The kernel instance is the *kernel object*, the values associated with the arguments to the kernel, and the parameters that define the *NDRange* index space.

**Kernel Object:** A *kernel object* encapsulates a specific `__kernel` function declared in a *program* and the argument values to be used when executing this `__kernel` function.

**Launch:** The transition of a command from the *submitted* state to the *ready* state. See *Ready*.

**Local ID:** A *local ID* specifies a unique *work-item ID* within a given *work-group* that is executing a *kernel*. The *local ID* is a N-dimensional value that starts at (0, 0, ... 0). See also *Global ID*.

**Local Memory:** A memory region associated with a *work-group* and accessible only by *work-items* in that *work-group*. *Local memory* is included within the *generic address space* that includes the private and global address spaces.

**Marker:** A *command* queued in a *command-queue* that can be used to tag all *commands* queued before the *marker* in the *command-queue*. The *marker* command returns an *event* which can be used by the *application* to queue a wait on the marker event i.e. wait for all commands queued before the *marker* command to complete.

**Memory Consistency Model:** Rules that define which values are observed when multiple units of execution load data from any shared memory plus the synchronization operations that constrain the order of memory operations and define synchronization relationships. The memory consistency model in OpenCL is based on the memory model from the ISO C11 programming language.

**Memory Objects:** A *memory object* is a handle to a reference counted region of *global memory*. Also see *Buffer Object* and *Image Object*.

**Memory Regions (or Pools):** A distinct address space in OpenCL. *Memory regions* may overlap in physical memory though OpenCL will treat them as logically distinct. The *memory regions* are denoted as *private*, *local*, *constant*, and *global*.

**Memory Scopes:** These memory scopes define a hierarchy of visibilities when analyzing the ordering constraints of memory operations. They are defined by the values of the *memory\_scope* enumeration constant. Current values are **memory\_scope\_work\_item** (memory constraints only apply to a single work-item and in practice apply only to image operations), **memory\_scope\_work\_group** (memory-ordering constraints only apply to work-items executing in a work-group), **memory\_scope\_device** (memory-ordering constraints only apply to work-items executing on a single device) and **memory\_scope\_all\_svm\_devices** (memory-ordering constraints only apply to work-items executing across multiple devices and when using shared virtual memory).

**Modification Order:** All modifications to a particular atomic object M occur in some particular **total order**, called the **modification order** of M. If A and B are modifications of an atomic object M, and A happens-before B, then A shall precede B in the modification order of M. Note that the modification order of an atomic object M is independent of whether M is in local or global memory.

**Nested Parallelism:** See *device-side enqueue*.

**Object:** Objects are abstract representation of the resources that can be manipulated by the OpenCL API. Examples include *program objects*, *kernel objects*, and *memory objects*.

**Out-of-Order Execution:** A model of execution in which *commands* placed in the *work queue* may begin and complete execution in any order consistent with constraints imposed by *event wait lists* and *command-queue barrier*. See *In-order Execution*.

**Parent device:** The OpenCL *device* which is partitioned to create *sub-devices*. Not all *parent devices* are *root devices*. A *root device* might be partitioned and the *sub-devices* partitioned again. In this case, the first set of *sub-devices* would be *parent devices* of the second set, but not the *root devices*. Also see *device*, *parent device* and *root device*.

**Parent kernel:** see *device-side enqueue*.

**Pipe:** The *pipe* memory object conceptually is an ordered sequence of data items. A pipe has two endpoints: a write endpoint into which data items are inserted, and a read endpoint from which data items are removed. At any one time, only one kernel instance may write into a pipe, and only one kernel instance may read from a pipe. To support the producer consumer design pattern, one kernel instance connects to the write endpoint (the producer) while another kernel instance connects to the reading endpoint (the consumer).





**Platform:** The *host* plus a collection of *devices* managed by the OpenCL *framework* that allow an application to share *resources* and execute *kernels* on *devices* in the *platform*.

**Private Memory:** A region of memory private to a *work-item*. Variables defined in one *work-item*'s *private memory* are not visible to another *work-item*.

**Processing Element:** A virtual scalar processor. A work-item may execute on one or more processing elements.

**Program:** An OpenCL *program* consists of a set of *kernels*. *Programs* may also contain auxiliary functions called by the `__kernel` functions and constant data.

**Program Object:** A *program object* encapsulates the following information:

-  A reference to an associated *context*.
-  A *program* source or binary.
-  The latest successfully built program executable, the list of *devices* for which the program executable is built, the build options used and a build log.
-  The number of *kernel objects* currently attached.

**Queued:** The first state in the six state model for the execution of a command. The transition into this state occurs when the command is enqueued into a command-queue.

**Ready:** The third state in the six state model for the execution of a command. The transition into this state occurs when pre-requisites constraining execution of a command have been met; i.e. the command has been launched. When a Kernel-enqueue command is launched, work-groups associated with the command are placed in a device's work-pool from which they are scheduled for execution.

**Re-converged Control Flow:** see *control flow*.

**Reference Count:** The life span of an OpenCL object is determined by its *reference count*—an internal count of the number of references to the object. When you create an object in OpenCL, its *reference count* is set to one. Subsequent calls to the appropriate *retain* API (such as `clRetainContext`, `clRetainCommandQueue`) increment the *reference count*. Calls to the appropriate *release* API (such as `clReleaseContext`, `clReleaseCommandQueue`) decrement the *reference count*. Implementations may also modify the *reference count*, e.g. to track attached objects or to ensure correct operation of in-progress or scheduled activities. After the reference count reaches zero, the object is deleted.

**Relaxed Consistency:** A memory consistency model in which the contents of memory visible to different *work-items* or *commands* may be different except at a *barrier* or other explicit synchronization points.

**Resource:** A class of *objects* defined by OpenCL. An instance of a *resource* is an *object*. The most common *resources* are the *context*, *command-queue*, *program objects*, *kernel objects*, and *memory objects*. Computational resources are hardware elements that participate in the action of advancing a program counter. Examples include the *host*, *devices*, *compute units* and *processing elements*.

**Retain, Release:** The action of incrementing (retain) and decrementing (release) the reference count using an OpenCL *object*. This is a book keeping functionality to make sure the system doesn't remove an *object* before all instances that use this *object* have finished. Refer to *Reference Count*.

**Relaxed Semantics:** A memory order semantics for atomic operations that implies no order constraints. The operation is *atomic* but it has no impact on the order of memory operations.

**Release Semantics:** One of the memory order semantics defined for synchronization operations. Release semantics apply to atomic operations that store to memory. Given two units of execution, **A** and **B**, acting on a shared atomic object **M**, if **A** uses an atomic store of **M** with release semantics to synchronize-with an atomic load to **M** by **B** that used acquire semantics, then **A**'s atomic store will occur *after* any prior operations by **A**. Note that the memory orders *acquire*, *sequentially consistent*, and *acquire\_release* all include *acquire semantics* and effectively pair with a store using release semantics.

**Remainder work-groups:** When the work-groups associated with a kernel-instance are defined, the sizes of a work-group in each dimension may not evenly divide the size of the NDRange in the corresponding dimensions. The result is a collection of work-groups on the boundaries of the NDRange that are smaller than the base work-group size. These are known as *remainder work-groups*.

**Running:** The fourth state in the six state model for the execution of a command. The transition into this state occurs when the execution of the command starts. When a Kernel-enqueue command starts, one or more work-groups associated with the command start to execute.

**Root device:** A *root device* is an OpenCL *device* that has not been partitioned. Also see *device*, *parent device* and *root device*.

**Sampler:** An *object* that describes how to sample an image when the image is read in the *kernel*. The image read functions take a *sampler* as an argument. The *sampler* specifies the image addressing-mode i.e. how out-of-range image coordinates are handled, the filter mode, and whether the input image coordinate is a normalized or unnormalized value.

**Scope inclusion:** Two actions **A** and **B** are defined to have an inclusive scope if they have the same scope **P** such that: (1) if **P** is `memory_scope_work_group`, and **A** and **B** are executed by work-items within the same workgroup, or (2) if **P** is `memory_scope_device`, and **A** and **B** are executed by work-items on the same device, or (3) if **P** is `memory_scope_all_svm_devices`, if **A** and **B** are executed by host threads or by work-items on one or more devices that can share SVM memory with each other and the host process.

**Sequenced before:** A relation between evaluations executed by a single unit of execution. Sequenced-before is an asymmetric, transitive, pair-wise relation that induces a partial order between evaluations. Given any two evaluations A and B, if A is sequenced-before B, then the execution of A shall precede the execution of B.

**Sequential consistency:** Sequential consistency interleaves the steps executed by each unit of execution. Each access to a memory location sees the last assignment to that location in that interleaving.

**Sequentially consistent semantics:** One of the memory order semantics defined for synchronization operations. When using sequentially-consistent synchronization operations, the loads and stores within one unit of execution appear to execute in program order (i.e., the sequenced-before order), and loads and stores from different units of execution appear to be simply interleaved.

**Shared Virtual Memory (SVM):** An address space exposed to both the host and the devices within a context. SVM causes addresses to be meaningful between the host and all of the devices within a context and therefore supports the use of pointer based data structures in OpenCL kernels. It logically extends a portion of the global memory into the host address space therefore giving work-items access to the host address space. There are three types of SVM in OpenCL  
**Coarse-Grained buffer SVM:** Sharing occurs at the granularity of regions of OpenCL buffer memory objects. **Fine-Grained buffer SVM:** Sharing occurs at the granularity of individual loads/stores into bytes within OpenCL buffer memory objects. **Fine-Grained system SVM:** Sharing occurs at the granularity of individual loads/stores into bytes occurring anywhere within the host memory.

**SIMD:** Single Instruction Multiple Data. A programming model where a *kernel* is executed concurrently on multiple *processing elements* each with its own data and a shared program counter. All *processing elements* execute a strictly identical set of instructions.

**SPMD:** Single Program Multiple Data. A programming model where a *kernel* is executed concurrently on multiple *processing elements* each with its own data and its own program counter. Hence, while all computational resources run the same *kernel* they maintain their own instruction counter and due to branches in a *kernel*, the actual sequence of instructions can be quite different across the set of *processing elements*.

**Sub-device:** An OpenCL *device* can be partitioned into multiple *sub-devices*. The new *sub-devices* alias specific collections of compute units within the parent *device*, according to a partition scheme. The *sub-devices* may be used in any situation that their parent *device* may be used. Partitioning a *device* does not destroy the parent *device*, which may continue to be used along side and intermingled with its child *sub-devices*. Also see *device*, *parent device* and *root device*.

**Subgroup:** Subgroups are an implementation-dependent grouping of work-items within a workgroup. The size and number of subgroups is implementation-defined and not exposed in the core OpenCL 2.0 feature set. Implementations of OpenCL that expose the decomposition of a workgroup into multiple subgroups can depend on these subgroups to make *independent forward progress*. Hence, concurrent control structures such as spin locks can be used when they apply between two subgroups of a single work-group.

**Submitted:** The second state in the six state model for the execution of a command. The transition into this state occurs when the command is flushed from the command-queue and submitted for execution on the device. Once submitted, a programmer can assume a command will execute once its prerequisites have been met.

**SVM Buffer:** A memory allocation enabled to work with Shared Virtual Memory (SVM). Depending on how the SVM buffer is created, it can be a coarse-grained or fine-grained SVM buffer. Optionally it may be wrapped by a Buffer Object. See *Shared Virtual Memory (SVM)*.

**Synchronization:** Synchronization refers to mechanisms that constrain the order of execution and the visibility of memory operations between two or more units of execution.

**Synchronization operations:** Operations that define memory order constraints in a program. They play a special role in controlling how memory operations in one unit of execution (such as work-items or, when using SVM a host thread) are made visible to another. Synchronization operations in OpenCL include *atomic operations* and *fences*.

**Synchronization point:** A synchronization point between a pair of commands (A and B) assures that results of command A happens-before command B is launched (i.e. enters the ready state) .

**Synchronizes with:** A relation between operations in two different units of execution that defines a memory order constraint in global memory (*global-synchronizes-with*) or local memory (*local-synchronizes-with*).

**Task Parallel Programming Model:** A programming model in which computations are

expressed in terms of multiple concurrent tasks executing in one or more *command-queues*. The concurrent tasks can be running different *kernels*.

**Thread-safe:** An OpenCL API call is considered to be *thread-safe* if the internal state as managed by OpenCL remains consistent when called simultaneously by multiple *host* threads. OpenCL API calls that are *thread-safe* allow an application to call these functions in multiple *host* threads without having to implement mutual exclusion across these *host* threads i.e. they are also re-entrant-safe.

**Undefined:** The behavior of an OpenCL API call, built-in function used inside a *kernel* or execution of a *kernel* that is explicitly not defined by OpenCL. A conforming implementation is not required to specify what occurs when an undefined construct is encountered in OpenCL.

**Unit of execution:** a generic term for a process, OS managed thread running on the host (a host-thread), kernel-instance, host program, work-item or any other executable agent that advances the work associated with a program.

**Work-group:** A collection of related *work-items* that execute on a single *compute unit*. The *work-items* in the group execute the same *kernel-instance* and share *local memory* and *work-group functions*.

**Work-group Barrier.** See *Barrier*.

**Work-group Function:** A function that carries out collective operations across all the work-items in a work-group. Available collective operations are a barrier, reduction, broadcast, prefix sum, and evaluation of a predicate. A work-group function must occur within a *converged control flow*; i.e. all work-items in the work-group must encounter precisely the same work-group function.

**Work-group Synchronization:** Constraints on the order of execution for work-items in a single work-group.

**Work-pool:** A logical pool associated with a device that holds commands and work-groups from kernel-instances that are ready to execute. OpenCL does not constrain the order that commands and work-groups are scheduled for execution from the work-pool; i.e. a programmer must assume that they could be interleaved. There is one work-pool per device used by all command-queues associated with that device. The work-pool may be implemented in any manner as long as it assures that work-groups placed in the pool will eventually execute.





**Work-item:** One of a collection of parallel executions of a *kernel* invoked on a *device* by a *command*. A *work-item* is executed by one or more *processing elements* as part of a *work-group* executing on a *compute unit*. A *work-item* is distinguished from other work-items by its *global ID* or the combination of its *work-group ID* and its *local ID* within a *work-group*.

### 3. The OpenCL Architecture

**OpenCL** is an open industry standard for programming a heterogeneous collection of CPUs, GPUs and other discrete computing devices organized into a single platform. It is more than a language. OpenCL is a framework for parallel programming and includes a language, API, libraries and a runtime system to support software development. Using OpenCL, for example, a programmer can write general purpose programs that execute on GPUs without the need to map their algorithms onto a 3D graphics API such as OpenGL or DirectX.

The target of OpenCL is expert programmers wanting to write portable yet efficient code. This includes library writers, middleware vendors, and performance oriented application programmers. Therefore OpenCL provides a low-level hardware abstraction plus a framework to support programming and many details of the underlying hardware are exposed.

To describe the core ideas behind OpenCL, we will use a hierarchy of models:

-  Platform Model
-  Memory Model
-  Execution Model
-  Programming Model


#### 3.1 Platform Model

The Platform model for OpenCL is defined in *figure 3.1*. The model consists of a **host** connected to one or more **OpenCL devices**. An OpenCL device is divided into one or more **compute units** (CUs) which are further divided into one or more **processing elements** (PEs). Computations on a device occur within the processing elements.

An OpenCL application is implemented as both host code and device kernel code. The host code portion of an OpenCL application runs on a host processor according to the models native to the host platform. The OpenCL application host code submits the kernel code as commands from the host to OpenCL devices. An OpenCL device executes the command's computation on the processing elements within the device.

An OpenCL device has considerable latitude on how computations are mapped onto the device's processing elements. When processing elements within a compute unit execute the same sequence of statements across the processing elements, the control flow is said to be *converged*. Hardware optimized for executing a single stream of instructions over multiple processing elements is well suited to converged control flows. When the control flow varies from one processing element to another, it is said to be *diverged*. While a kernel always begins execution with a converged control flow, due to branching statements within a kernel, converged and diverged control flows may occur within a single kernel. This provides a great deal of flexibility in the algorithms that can be implemented with OpenCL.





**Figure 3.1:** Platform model ... one host plus one or more compute devices each with one or more compute units composed of one or more processing elements.

Programmers write device code in OpenCL C, a variant of C99. A platform provides a compiler to translate an OpenCL C source program into an executable program object. The device code compiler may be *online* or *offline*. An *online compiler* is available during host program execution using standard APIs. An *offline compiler* is invoked outside of host program control, using platform-specific methods. The OpenCL runtime allows developers to get a previously compiled device program executable and be able to load and execute a previously compiled device program executable.

OpenCL defines two kinds of platform profiles: a *full profile* and a *reduced-functionality embedded profile*. A full profile platform must provide an online compiler for all its devices. An embedded platform may provide an online compiler, but is not required to do so.

A device may expose special purpose functionality as a *built-in function*. The platform provides APIs for enumerating and invoking the built-in functions offered by a device, but otherwise does not define their construction or semantics. A *custom device* supports only built-in functions, and cannot be programmed via OpenCL C.

All device types support the OpenCL execution model, the OpenCL memory model, and the APIs used in OpenCL to manage devices.

The platform model is an abstraction describing how OpenCL views the hardware. The relationship between the elements of the platform model and the hardware in a system may be a fixed property of a device or it may be a dynamic feature of a program dependent on how a compiler optimizes code to best utilize physical hardware.

## 3.2 Execution Model

The OpenCL execution model is defined in terms of two distinct units of execution: **kernels** that execute on one or more OpenCL devices and a **host program** that executes on the host. With regard to OpenCL, the kernels are where the "work" associated with a computation occurs. This work occurs through **work-items** that execute in groups (**work-groups**).

A kernel executes within a well-defined context managed by the host. The context defines the environment within which kernels execute. It includes the following resources:

- + **Devices:** One or more devices exposed by the OpenCL platform.
- + **Kernel Objects:** The OpenCL functions with their associated argument values that run on OpenCL devices.
- + **Program Objects:** The program source and executable that implement the kernels.
- + **Memory Objects:** Variables visible to the host and the OpenCL devices. Instances of kernels operate on these objects as they execute.

The host program uses the OpenCL API to create and manage the context. Functions from the OpenCL API enable the host to interact with a device through a *command-queue*. **Each command-queue is associated with a single device.** The commands placed into the command-queue fall into one of three types:


- + **Kernel-enqueue commands:** Enqueue a kernel for execution on a device.
- + **Memory commands:** Transfer data between the host and device memory, between memory objects, or map and unmap memory objects from the host address space.
- + **Synchronization commands:** Explicit synchronization points that define order constraints between commands.

In addition to commands submitted from the host command-queue, a kernel running on a device can enqueue commands to a device-side command queue. This results in *child kernels* enqueued by a kernel executing on a device (the *parent kernel*). Regardless of whether the command-queue resides on the host or a device, each command passes through six states.

1. **Queued:** The command is enqueued to a command-queue. A command may reside in the queue until it is flushed either explicitly (a call to `clFlush`) or implicitly by some other command.
2. **Submitted:** The command is flushed from the command-queue and submitted for execution on the device. Once flushed from the command-queue, a command will execute after any prerequisites for execution are met.
3. **Ready:** All prerequisites constraining execution of a command have been met. The command, or for a kernel-enqueue command the collection of work groups associated with a command, is placed in a device work-pool from which it is scheduled for execution.
4. **Running:** Execution of the command starts. For the case of a kernel-enqueue command, one or more work-groups associated with the command start to execute.

5. **Ended:** Execution of a command ends. When a Kernel-enqueue command ends, all of the work-groups associated with that command have finished their execution.  
*Immediate side effects*, i.e. those associated with the kernel but not necessarily with its child kernels, are visible to other units of execution. These side effects include updates to values in global memory.
6. **Complete:** The command and its child commands have finished execution and the status of the event object, if any, associated with the command is set to CL\_COMPLETE.

The execution states and the transitions between them are summarized in Figure 3-2. These states and the concept of a device work-pool are conceptual elements of the execution model. An implementation of OpenCL has considerable freedom in how these are exposed to a program. Five of the transitions, however, are directly observable through a profiling interface. These profiled states are shown in Figure 3-2.



**Figure 3-2: The states and transitions between states defined in the OpenCL execution model. A subset of these transitions is exposed through the profiling interface (see section 5.14).**

**Commands communicate their status through *Event objects*.** Successful completion is indicated by setting the event status associated with a command to CL\_COMPLETE. Unsuccessful completion results in abnormal termination of the command which is indicated by setting the event status to a negative value. In this case, the command-queue associated with the abnormally terminated command and all other command-queues in the same context may no longer be available and their behavior is implementation defined.

A command submitted to a device will not launch until prerequisites that constrain the order of commands have been resolved. These prerequisites have two sources. First, they may arise from commands submitted to a command-queue that constrain the order that commands are launched. For example, commands that follow a command queue barrier will not launch until all commands prior to the barrier are complete. The second source of prerequisites is dependencies between commands expressed through events. A command may include an optional list of events. The command will wait and not launch until all the events in the list are in the state CL\_COMPLETE. By this mechanism, event objects define order constraints between commands and coordinate execution between the host and one or more devices.

Command execution can be blocking or non-blocking. Consider a sequence of OpenCL commands. For blocking commands, the OpenCL API functions that enqueue commands don't return until the command has completed. Alternatively, OpenCL functions that enqueue non-blocking commands return immediately and require that a programmer defines dependencies between enqueued commands to ensure that enqueued commands are not launched before needed resources are available. In both cases, the actual execution of the command may occur asynchronously with execution of the host program.

Commands within a single command-queue execute relative to each other in one of two modes:

- ✚ **In-order Execution:** Commands and any side effects associated with commands appear to the OpenCL application as if they execute in the same order they are enqueued to a command-queue.
- ✚ **Out-of-order Execution:** Commands execute in any order constrained only by explicit synchronization points (e.g. through command queue barriers) or explicit dependencies on events.

Multiple command-queues can be present within a single context. Multiple command-queues execute commands independently. Event objects visible to the host program can be used to define synchronization points between commands in multiple command queues. If such synchronization points are established between commands in multiple command-queues, an implementation must assure that the command-queues progress concurrently and correctly account for the dependencies established by the synchronization points. For a detailed explanation of synchronization points, see section 3.2.4.

The core of the OpenCL execution model is defined by how the kernels execute. When a kernel-enqueue command submits a kernel for execution, an index space is defined. The kernel, the argument values associated with the arguments to the kernel, and the parameters that define the index space define a *kernel-instance*. When a kernel-instance executes on a device, the kernel function executes for each point in the defined index space. Each of these executing kernel functions is called a *work-item*. The work-items associated with a given kernel-instance are managed by the device in groups called *work-groups*. These work-groups define a coarse grained decomposition of the Index space.

Work-items have a global ID based on their coordinates within the Index space. They can also be defined in terms of their work-group and the local ID within a work-group. The details of this mapping are described in the following section.

### 3.2.1 Execution Model: Mapping work-items onto an NDRange

The index space supported by OpenCL is called an NDRange. An NDRange is an N-dimensional index space, where N is one, two or three. The NDRange is decomposed into work-groups forming blocks that cover the Index space. An NDRange is defined by three integer arrays of length N:

- ✚ The extent of the index space (or global size) in each dimension.
- ✚ An offset index F indicating the initial value of the indices in each dimension (zero by default).
- ✚ The size of a work-group (local size) in each dimension.

Each work-item's global ID is an N-dimensional tuple. The global ID components are values in the range from F, to F plus the number of elements in that dimension minus one.

If a kernel is compiled as an OpenCL 2.0 kernel, the size of work-groups in an NDRange (the local size) need not be the same for all work-groups. In this case, any single dimension for which the global size is not divisible by the local size will be partitioned into two regions. One region will have workgroups that have the same number of work items as was specified for that dimension by the programmer (the local size). The other region will have workgroups with less than the number of work items specified by the local size parameter in that dimension (the *remainder work-groups*). Workgroup sizes could be non-uniform in multiple dimensions, potentially producing workgroups of up to 4 different sizes in a 2D range and 8 different sizes in a 3D range.

Each work-item is assigned to a work-group and given a local ID to represent its position within the workgroup. A work-item's local ID is an N-dimensional tuple with components in the range from zero to the size of the work-group in that dimension minus one.

Work-groups are assigned IDs similarly. The number of work-groups in each dimension is not directly defined but is inferred from the local and global NDRanges provided when a kernel-instance is enqueued. A work-group's ID is an N-dimensional tuple with components in the range 0 to the ceiling of the global size in that dimension divided by the local size in the same dimension. As a result, the combination of a work-group ID and the local-ID within a work-group uniquely defines a work-item. Each work-item is identifiable in two ways; in terms of a global index, and in terms of a work-group index plus a local index within a work group.

For example, consider the 2-dimensional index space in figure 3-3. We input the index space for the work-items ( $G_x$ ,  $G_y$ ), the size of each work-group ( $S_x$ ,  $S_y$ ) and the global ID offset ( $F_x$ ,  $F_y$ ). The global indices define an  $G_x$  by  $G_y$  index space where the total number of work-items is the product of  $G_x$  and  $G_y$ . The local indices define an  $S_x$  by  $S_y$  index space where the number of

work-items in a single work-group is the product of  $S_x$  and  $S_y$ . Given the size of each workgroup and the total number of work-items we can compute the number of work-groups. A 2-dimensional index space is used to uniquely identify a work-group. Each work-item is identified by its global ID  $(g_x, g_y)$  or by the combination of the work-group ID  $(w_x, w_y)$ , the size of each work-group  $(S_x, S_y)$  and the local ID  $(s_x, s_y)$  inside the work-group such that

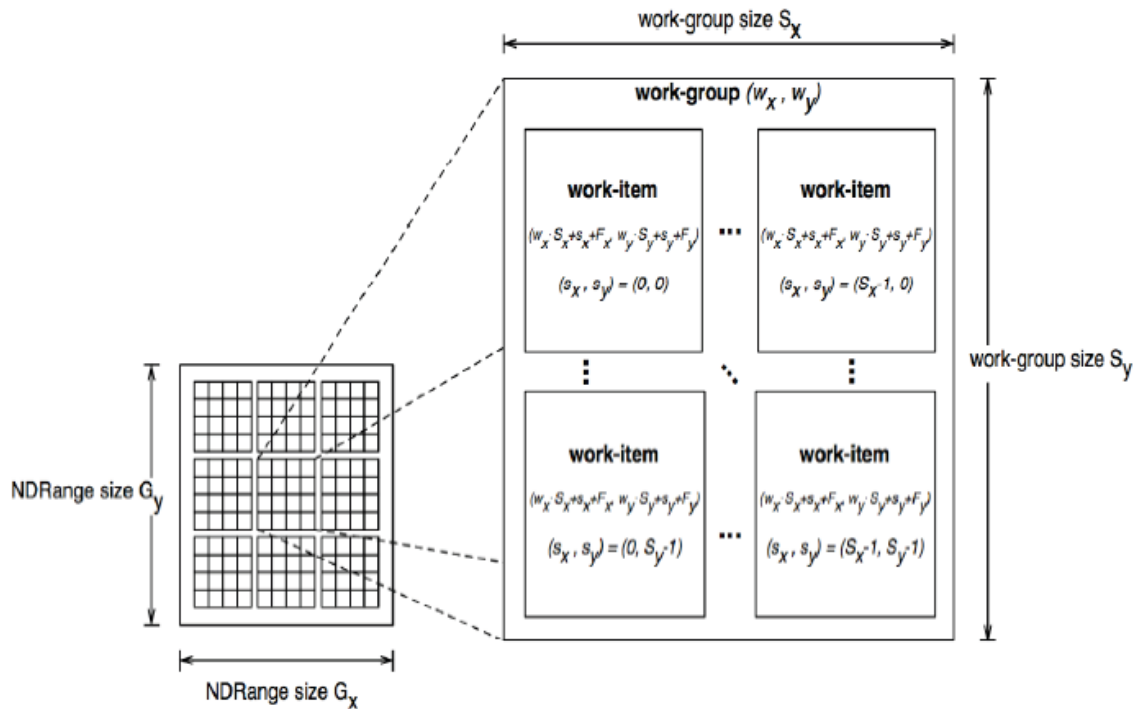
$$(g_x, g_y) = (w_x * S_x + s_x + F_x, w_y * S_y + s_y + F_y)$$

The number of work-groups can be computed as:

$$(W_x, W_y) = (\text{ceil}(G_x / S_x), \text{ceil}(G_y / S_y))$$

Given a global ID and the work-group size, the work-group ID for a work-item is computed as:

$$(w_x, w_y) = ((g_x - s_x - F_x) / S_x, (g_y - s_y - F_y) / S_y)$$



**Figure 3-3: An example of an NDRange index space showing work-items, their global IDs and their mapping onto the pair of work-group and local IDs. In this case, we assume that in each dimension, the size of the work-group evenly divides the global NDRange size (i.e. all work-groups have the same size) and that the offset is equal to zero.**

An implementation of OpenCL may divide each work-group into one or more subgroups. Subgroups make independent forward progress with respect to each other. There is no guarantee, however, that work-items within a subgroup make independent forward progress with respect to each other. The size and number of subgroups is implementation-defined and not exposed in the core OpenCL 2.0 feature set.

### 3.2.2 Execution Model: Execution of kernel-instances

The work carried out by an OpenCL program occurs through the execution of kernel-instances on compute devices. To understand the details of OpenCL's execution model, we need to consider how a kernel object moves from the kernel-enqueue command, into a command-queue, executes on a device, and completes.

A kernel-object is defined from a function within the program object and a collection of arguments connecting the kernel to a set of argument values. The host program enqueues a kernel-object to the command queue along with the NDRange, and the work-group decomposition. **These define a *kernel-instance***. In addition, an optional set of events may be defined when the kernel is enqueued. The events associated with a particular kernel-instance are used to constrain when the kernel-instance is launched with respect to other commands in the queue or to commands in other queues within the same context.

A kernel-instance is submitted to a device. For an in-order command queue, the kernel instances appear to launch and then execute in that same order; where we use the term “appear” to emphasize that when there are no dependencies between commands and hence differences in the order that commands execute cannot be observed in a program, an implementation can reorder commands even in an in-order command queue. For an out of order command-queue, kernel-instances wait to be launched until:

- ✚ Synchronization commands enqueued prior to the kernel-instance are satisfied.
- ✚ Each of the events in an optional event list defined when the kernel-instance was enqueued are set to CL\_COMPLETE.

Once these conditions are met, the kernel-instance is launched and the work-groups associated with the kernel-instance are placed into a pool of “ready to execute” work-groups. This pool is called a ***work-pool***. The work-pool may be implemented in any manner as long as it assures that work-groups placed in the pool will eventually execute. The device schedules work-groups from the work-pool for execution on the compute units of the device. **The kernel-enqueue command is complete when all work-groups associated with the kernel-instance end their execution, updates to global memory associated with a command are visible globally, and the device signals successful completion by setting the event associated with the kernel-enqueue command to CL\_COMPLETE.**

**While a command-queue is associated with only one device, a single device may be associated with multiple command-queues all feeding into the single work-pool. A device may also be associated with command queues associated with different contexts within the same platform, again all feeding into the single work-pool.** The device will pull work-groups from the work-pool and execute them on one or several compute units in any order; possibly interleaving execution of work-groups from multiple commands. A conforming implementation may choose to serialize the work-groups so a correct algorithm cannot assume that work-groups will execute in parallel.



There is no safe and portable way to synchronize across the independent execution of work-groups since once in the work-pool, they can execute in any order.

The work-items within a single work-group execute concurrently but not necessarily in parallel (i.e. they are not guaranteed to make independent forward progress). Therefore, only high-level synchronization constructs (e.g. work-group functions such as barriers) that apply to all the work-items in a work-group are well defined and included in OpenCL.

In the absence of work-group functions (e.g. a barrier), work-items within a workgroup may be serialized. In the presence of work-group functions, work-items within a workgroup may be serialized before any given work-group function, between dynamically encountered pairs of work-group functions and between a work-group function and the end of the kernel.

The potential of the work-items within a workgroup to be serialized means that independent forward progress of the work-items cannot be assumed; therefore, synchronization between subsets of work-items within a work-group (e.g. using spin-locks) cannot portably be supported in OpenCL.

Implementations of OpenCL that support the decomposition of a workgroup into multiple subgroups can depend on these subgroups to make independent forward progress. Hence, concurrent control structures such as spin locks can be used when they apply between two subgroups of a single work-group. The existence of sub-groups and their sizes are highly implementation dependent. Extreme care should be exercised when writing code that uses subgroups if the goal is to write portable OpenCL applications.

### 3.2.3 Execution Model: Device-side enqueue

Algorithms may need to generate additional work as they execute. In many cases, this additional work cannot be determined statically; so the work associated with a kernel only emerges at runtime as the kernel-instance executes. This capability could be implemented in logic running within the host program, but involvement of the host may add significant overhead and/or complexity to the application control flow. A more efficient approach would be to nest kernel-enqueue commands from inside other kernels. This **nested parallelism** can be realized by supporting the enqueueing of kernels on a device without direct involvement by the host program; so-called **device-side enqueue**.

Device-side kernel-enqueue commands are similar to host-side kernel-enqueue commands. The kernel executing on a device (the **parent kernel**) enqueues a kernel-instance (the **child kernel**) to a device-side command queue. This is an out-of-order command-queue and follows the same behavior as the out-of-order command-queues exposed to the host program. Commands enqueued to a device side command-queue generate and use events to enforce order constraints just as for the command-queue on the host. These events, however, are only visible to the parent kernel running on the device. When these prerequisite events take on the value CL\_COMPLETE, the work-groups associated with the child kernel are launched into the device's work pool. The device then schedules them for execution on the compute units of the device. Child and parent



kernels execute asynchronously. However, a parent will not indicate that it is complete by setting its event to CL\_COMPLETE until all child kernels have ended execution and have signaled completion by setting any associated events to the value CL\_COMPLETE. Should any child kernel complete with an event status set to a negative value (i.e. abnormally terminate), the parent kernel will abnormally terminate and propagate the child's negative event value as the value of the parent's event. If there are multiple children that have an event status set to a negative value, the selection of which child's negative event value is propagated is implementation-defined.

### 3.2.4 Execution Model: Synchronization

Synchronization refers to mechanisms that constrain the order of execution between two or more units of execution. Consider the following two domains of synchronization in OpenCL:

- ✚ Work-group synchronization: Constraints on the order of execution for work-items in a single work-group
- ✚ Command synchronization: Constraints on the order of commands launched for execution

Synchronization across all work-items within a single work-group is carried out using a *work-group function*. These functions carry out collective operations across all the work-items in a work-group. Available collective operations are: barrier, reduction, broadcast, prefix sum, and evaluation of a predicate. A work-group function must occur within a converged control flow; i.e. all work-items in the work-group must encounter precisely the same work-group function. For example, if a work-group function occurs within a loop, the work-items must encounter the same work-group function in the same loop iterations. All the work-items of a work-group must execute the work-group function and complete reads and writes to memory before any are allowed to continue execution beyond the work-group function. Work-group functions that apply between work-groups are not provided in OpenCL since OpenCL does not define forward-progress or ordering relations between work-groups, hence collective synchronization operations are not well defined.

Command synchronization is defined in terms of distinct **synchronization points**. The synchronization points occur between commands in host command-queues and between commands in device-side command-queues. The synchronization points defined in OpenCL include:

- ✚ **Launching a command:** A kernel-instance is launched onto a device after all events that kernel is waiting-on have been set to CL\_COMPLETE.
- ✚ **Ending a command:** Child kernels may be enqueued such that they wait for the parent kernel to reach the *end* state before they can be launched. In this case, the ending of the parent command defines a synchronization point.
- ✚ **Completion of a command:** A kernel-instance is complete after all of the work-groups in the kernel and all of its child kernels have completed. This is signaled to the host, a parent

kernel or other kernels within command queues by setting the value of the event associated with a kernel to CL\_COMPLETE.

- ✚ **Blocking Commands:** A blocking command defines a synchronization point between the unit of execution that calls the blocking API function and the enqueued command reaching the complete state.
- ✚ **Command-queue barrier:** The command-queue barrier ensures that all previously enqueued commands have completed before subsequently enqueued commands can be launched.
- ✚ **clFinish:** This function blocks until all previously enqueued commands in the command queue have completed after which clFinish defines a synchronization point and the clFinish function returns.

A synchronization point between a pair of commands (A and B) assures that results of command A happens-before command B is launched. This requires that any updates to memory from command A complete and are made available to other commands before the synchronization point completes. Likewise, this requires that command B waits until after the synchronization point before loading values from global memory. The concept of a synchronization point works in a similar fashion for commands such as a barrier that apply to two sets of commands. All the commands prior to the barrier must complete and make their results available to following commands. Furthermore, any commands following the barrier must wait for the commands prior to the barrier before loading values and continuing their execution.

These *happens-before* relationships are a fundamental part of the OpenCL memory model. When applied at the level of commands, they are straightforward to define at a language level in terms of ordering relationships between different commands. Ordering memory operations inside different commands, however, requires rules more complex than can be captured by the high level concept of a synchronization point. These rules are described in detail in section 3.3.6.

### 3.2.5 Execution Model: Categories of Kernels

The OpenCL execution model supports three types of kernels:

- ✚ **OpenCL kernels** are managed by the OpenCL API as kernel-objects associated with kernel functions within program-objects. They are typically written with the OpenCL C programming language and compiled with the OpenCL C compiler. All OpenCL implementations must support OpenCL kernels and the OpenCL C programming language. Implementations may also support OpenCL programs through an external intermediate representation. For example, SPIR is an optional extension to OpenCL that provides an intermediate representation for OpenCL program objects. Through SPIR, any kernel programming language that targets SPIR can be used to define an OpenCL kernel.
- ✚ **Native kernels** are accessed through a host function pointer. Native kernels are queued for execution along with OpenCL kernels on a device and share memory objects with OpenCL kernels. For example, these native kernels could be functions defined in application code or

exported from a library. The ability to execute native kernels is optional within OpenCL and the semantics of native kernels are implementation-defined. The OpenCL API includes functions to query capabilities of a device(s) and determine if this capability is supported.

- ✚ **Built-in kernels** are tied to particular device and are not built at runtime from source code in a program object. The common use of built in kernels is to expose fixed-function hardware or firmware associated with a particular OpenCL device or custom device. The semantics of a built-in kernel may be defined outside of OpenCL and hence are implementation defined.

All three types of kernels are manipulated through the OpenCL command queues and must conform to the synchronization points defined in the OpenCL execution model.

## 3.3 Memory Model

The OpenCL memory model describes the structure, contents, and behavior of the memory exposed by an OpenCL platform as an OpenCL program runs. The model allows a programmer to reason about values in memory as the host program and multiple kernel-instances execute.

An OpenCL program defines a context that includes a host, one or more devices, command-queues, and memory exposed within the context. Consider the units of execution involved with such a program. The host program runs as one or more host threads managed by the operating system running on the host (the details of which are defined outside of OpenCL). There may be multiple devices in a single context which all have access to memory objects defined by OpenCL. On a single device, multiple work-groups may execute in parallel with potentially overlapping updates to memory. Finally, within a single work-group, multiple work-items concurrently execute, once again with potentially overlapping updates to memory.

The memory model must precisely define how the values in memory as seen from each of these units of execution interact so a programmer can reason about the correctness of OpenCL programs. We define the memory model in four parts.

- ✚ **Memory regions**: The distinct memories visible to the host and the devices that share a context.
- ✚ **Memory objects**: The objects defined by the OpenCL API and their management by the host and devices.
- ✚ **Shared Virtual Memory**: A virtual address space exposed to both the host and the devices within a context.
- ✚ **Consistency Model**: Rules that define which values are observed when multiple units of execution load data from memory plus the atomic/fence operations that constrain the order of memory operations and define synchronization relationships.

### 3.3.1 Memory Model: Fundamental Memory Regions

Memory in OpenCL is divided into two parts.

- ✚ **Host Memory:** The memory directly available to the host. The detailed behavior of host memory is defined outside of OpenCL. Memory objects move between the Host and the devices through functions within the OpenCL API or through a shared virtual memory interface.
- ✚ **Device Memory:** Memory directly available to kernels executing on OpenCL devices.

Device memory consists of four named address spaces or *memory regions*:

- ✚ **Global Memory:** This memory region permits read/write access to all work-items in all work-groups running on any device within a context. Work-items can read from or write to any element of a memory object. Reads and writes to global memory may be cached depending on the capabilities of the device.
- ✚ **Constant Memory:** A region of global memory that remains constant during the execution of a kernel-instance. The host allocates and initializes memory objects placed into constant memory.
- ✚ **Local Memory:** A memory region local to a work-group. This memory region can be used to allocate variables that are shared by all work-items in that work-group.
- ✚ **Private Memory:** A region of memory private to a work-item. Variables defined in one work-item's private memory are not visible to another work-item.

The memory regions and their relationship to the OpenCL Platform model are summarized in figure 3-4. Local and private memories are always associated with a particular device. The global and constant memories, however, are shared between all devices within a given context. An OpenCL device may include a cache to support efficient access to these shared memories

To understand memory in OpenCL, it is important to appreciate the relationships between these named address spaces. The four named address spaces available to a device are disjoint meaning they do not overlap. This is a logical relationship, however, and an implementation may choose to let these disjoint named address spaces share physical memory.

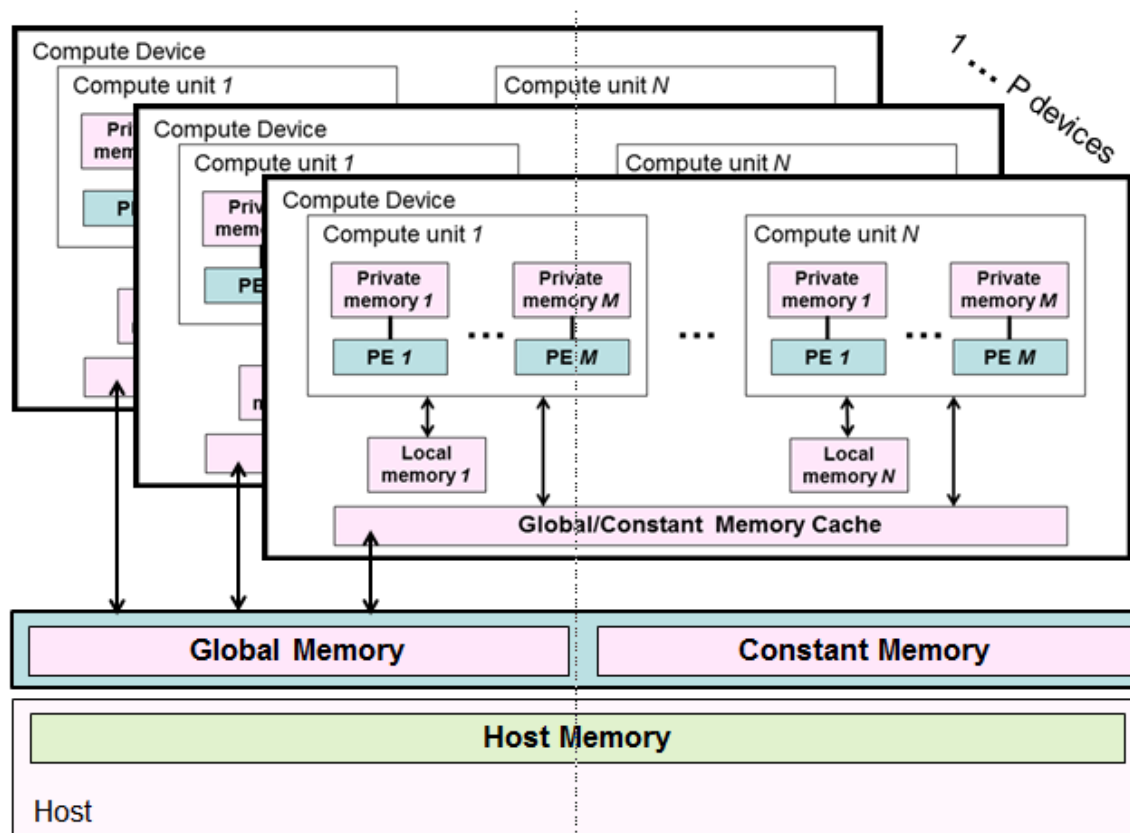
Programmers often need functions callable from kernels where the pointers manipulated by those functions can point to multiple named address spaces. This saves a programmer from the error-prone and wasteful practice of creating multiple copies of functions; one for each named address space. Therefore the global, local and private address spaces belong to a single *generic address space*. This is closely modeled after the concept of a generic address space used in the embedded C standard (ISO/IEC 9899:1999). Since they all belong to a single generic address space, the following properties are supported for pointers to named address spaces in device memory:

- ✚ A pointer to the generic address space can be cast to a pointer to a global, local or private address space
- ✚ A pointer to a global, local or private address space can be cast to a pointer to the generic address space.
- ✚ A pointer to a global, local or private address space can be implicitly converted to a pointer to the generic address space, but the converse is not allowed.

The constant address space is disjoint from the generic address space.

The addresses of memory associated with memory objects in Global memory are not preserved between kernel instances, between a device and the host, and between devices. In this regard global memory acts as a global pool of memory objects rather than an address space. This restriction is relaxed when shared virtual memory (SVM) is used.

SVM causes addresses to be meaningful between the host and all of the devices within a context hence supporting the use of pointer based data structures in OpenCL kernels. It logically extends a portion of the global memory into the host address space giving work-items access to the host address space. **On platforms with hardware support for a shared address space between the host and one or more devices, SVM may also provide a more efficient way to share data between devices and the host.** Details about SVM are presented in section 3.3.3.



**Figure 3-4: The named address spaces exposed in an OpenCL Platform.** Global and Constant memories are shared between the one or more devices within a context, while local and private memories are associated with a single device. Each device may include an optional cache to support efficient access to their view of the global and constant address spaces.

A programmer may use the features of the memory consistency model (section 3.3.4) to manage safe access to global memory from multiple work-items potentially running on one or more devices. In addition, when using shared virtual memory (SVM), the memory consistency model

may also be used to ensure that host threads safely access memory locations in the shared memory region.

### 3.3.2 Memory Model: Memory Objects

The contents of global memory are *memory objects*. A memory object is a handle to a reference counted region of global memory. Memory objects use the OpenCL type *cl\_mem* and fall into three distinct classes.

- ✚ **Buffer:** A memory object stored as a block of contiguous memory and used as a general purpose object to hold data used in an OpenCL program. The types of the values within a buffer may be any of the built in types (such as int, float), vector types, or user-defined structures. The buffer can be manipulated through pointers much as one would with any block of memory in C.
- ✚ **Image:** An image memory object holds one, two or three dimensional images. The formats are based on the standard image formats used in graphics applications. An image is an opaque data structure managed by functions defined in the OpenCL API. To optimize the manipulation of images stored in the texture memories found in many GPUs, OpenCL kernels have traditionally been disallowed from both reading and writing a single image. In OpenCL 2.0, however, we have relaxed this restriction by providing synchronization and fence operations that let programmers properly synchronize their code to safely allow a kernel to read and write a single image.
- ✚ **Pipe:** The *pipe* memory object conceptually is an ordered sequence of data items. A pipe has two endpoints: a write endpoint into which data items are inserted, and a read endpoint from which data items are removed. At any one time, only one kernel instance may write into a pipe, and only one kernel instance may read from a pipe. To support the producer consumer design pattern, one kernel instance connects to the write endpoint (the producer) while another kernel instance connects to the reading endpoint (the consumer).




Memory objects are allocated by host APIs. The host program can provide the runtime with a pointer to a block of continuous memory to hold the memory object when the object is created (`CL_MEM_USE_HOST_PTR`). Alternatively, the physical memory can be managed by the OpenCL runtime and not be directly accessible to the host program.

Allocation and access to memory objects within the different memory regions varies between the host and work-items running on a device. This is summarized in table 3.1 which describes whether the kernel or the host can allocate from a memory region, the type of allocation (static at compile time vs. dynamic at runtime) and the type of access allowed (i.e. whether the kernel or the host can read and/or write to a memory region).

	Global	Constant	Local	Private
Host	Dynamic Allocation	Dynamic Allocation	Dynamic Allocation	No allocation
	Read/Write access to buffers and images but not pipes	Read/Write access	No access	No access
Kernel	Static Allocation for program scope variables	Static Allocation	Static allocation.  Dynamic allocation for child kernel	Static Allocation
	Read/Write access	Read-only access	Read/write access.  No access to child's local memory.	Read/Write access

**Table 3-1: The different memory regions in OpenCL and how memory objects are allocated and accessed by the host and by an executing instance of a kernel.** For the case of kernels, we distinguish between the behavior of local memory with respect to a kernel (self) and it's child kernels.

Once allocated, a memory object is made available to kernel-instances running on one or more devices. In addition to shared virtual memory (section 3.3.3) there are three basic ways to manage the contents of buffers between the host and devices.

-  **Read/Write/Fill commands:** The data associated with a memory object is explicitly read and written between the host and global memory regions using commands enqueued to an OpenCL command queue.
-  **Map/Unmap commands:** Data from the memory object is mapped into a contiguous block of memory accessed through a host accessible pointer. The host program enqueues a *map* command on block of a memory object before it can be safely manipulated by the host program. When the host program is finished working with the block of memory, the host program enqueues an *unmap* command to allow a kernel-instance to safely read and/or write the buffer.
-  **Copy commands:** The data associated with a memory object is copied between two buffers, each of which may reside either on the host or on the device.

In both cases, the commands to transfer data between devices and the host can be blocking or non-blocking operations. The OpenCL function call for a blocking memory transfer returns once the associated memory resources on the host can be safely reused. For a non-blocking memory transfer, the OpenCL function call returns as soon as the command is enqueued.






Memory objects are bound to a context and hence can appear in multiple kernel-instances running on more than one physical device. The OpenCL platform must support a large range of hardware platforms including systems that do not support a single shared address space in hardware; hence the ways memory objects can be shared between kernel-instances is restricted. The basic principle is that multiple read operations on memory objects from multiple kernel-instances that overlap in time are allowed, but mixing overlapping reads and writes into the same memory objects from different kernel instances is only allowed when fine grained synchronization is used with shared virtual memory (see section 3.3.3).

When global memory is manipulated by multiple kernel-instances running on multiple devices, the OpenCL runtime system must manage the association of memory objects with a given device. In most cases the OpenCL runtime will implicitly associate a memory object with a device. A kernel instance is naturally associated with the command queue to which the kernel was submitted. Since a command-queue can only access a single device, the queue uniquely defines which device is involved with any given kernel-instance; hence defining a clear association between memory objects, kernel-instances and devices. Programmers may anticipate these associations in their programs and explicitly manage association of memory objects with devices in order to improve performance.

### 3.3.3 Memory Model: Shared Virtual Memory

OpenCL extends the global memory region into the host memory region through a shared virtual memory (SVM) mechanism. There are three types of SVM in OpenCL

-  **Coarse-Grained buffer SVM:** Sharing occurs at the granularity of regions of OpenCL buffer memory objects. Consistency is enforced at synchronization points and with map/unmap commands to drive updates between the host and the device. This form of SVM is similar to the use of `cl_mem` buffers, with two differences. First, it lets kernel-instances share pointer-based data structures (such as linked-lists) with the host program. Second, concurrent access by multiple kernels on the same device is valid as long as the set of concurrently executing kernels is bounded by synchronization points. This form of SVM is similar to non-SVM use of memory; however, it lets kernel-instances share pointer-based data structures (such as linked-lists) with the host program.
-  **Fine-Grained buffer SVM:** Sharing occurs at the granularity of individual loads/stores into bytes within OpenCL buffer memory objects. Loads and stores may be cached. This means consistency is guaranteed at synchronization points. If the optional OpenCL atomics are supported, they can be used to provide fine-grained control of memory consistency.
-  **Fine-Grained system SVM:** Sharing occurs at the granularity of individual loads/stores into bytes occurring anywhere within the host memory. Loads and stores may be cached so consistency is guaranteed at synchronization points. If the optional OpenCL atomics are supported, they can be used to provide fine-grained control of memory consistency.



	Granularity of sharing	Memory Allocation	Mechanisms to enforce Consistency	Explicit updates between host and device?
Non-SVM buffers	OpenCL Memory objects (buffer)	clCreateBuffer	Host synchronization points on the same or between devices.	yes, through Map and Unmap commands.
Coarse-Grained buffer SVM	OpenCL Memory objects (buffer)	clSVMAlloc	Host synchronization points between devices	yes, through Map and Unmap commands.
Fine Grained buffer SVM	Bytes within OpenCL Memory objects (buffer)	clSVMAlloc	Synchronization points plus atomics (if supported)	No
Fine-Grained system SVM	Bytes within Host memory (system)	Host memory allocation mechanisms (e.g. malloc)	Synchronization points plus atomics (if supported)	No
<b>Table 3-2: A summary of shared virtual memory (SVM) options in OpenCL.</b>				

Coarse-Grained buffer SVM is required in the core OpenCL specification. The two finer grained approaches are optional features in OpenCL. The various SVM mechanisms to access host memory from the work-items associated with a kernel instance are summarized in table 3-2.

### 3.3.4 Memory Model: Memory Consistency Model




The OpenCL memory model tells programmers what they can expect from an OpenCL implementation; which memory operations are guaranteed to happen in which order and which memory values each read operation will return. The memory model tells compiler writers which restrictions they must follow when implementing compiler optimizations; which variables they can cache in registers and when they can move reads or writes around a barrier or atomic operation. The memory model also tells hardware designers about limitations on hardware optimizations; for example, when they must flush or invalidate hardware caches.

The memory consistency model in OpenCL is based on the memory model from the ISO C11 programming language. To help make the presentation more precise and self-contained, we


include modified paragraphs taken verbatim from the ISO C11 international standard. When a paragraph is taken or modified from the C11 standard, it is identified as such along with its original location in the C11 standard.

For programmers, the most intuitive model is the **sequential consistency** memory model. **Sequential consistency interleaves the steps executed by each of the units of execution. Each access to a memory location sees the last assignment to that location in that interleaving.** While sequential consistency is relatively straightforward for a programmer to reason about, implementing sequential consistency is expensive. Therefore, OpenCL implements a relaxed memory consistency model; i.e. it is possible to write programs where the loads from memory violate sequential consistency. Fortunately, if a program does not contain any races and if the program only uses atomic operations that utilize the sequentially consistent memory order (the default memory ordering for OpenCL), OpenCL programs appear to execute with sequential consistency.

Programmers can to some degree control how the memory model is relaxed by choosing the *memory order* for synchronization operations. The precise semantics of synchronization and the memory orders are formally defined in section 3.3.6. Here, we give a high level description of how these memory orders apply to atomic operations on atomic objects shared between units of execution. **OpenCL memory\_order choices are based on those from the ANSI C11 standard memory model.** They are specified in certain OpenCL functions through the following enumeration constants:

-  **memory\_order\_relaxed**: implies no order constraints. This memory order can be used safely to increment counters that are concurrently incremented, but it doesn't guarantee anything about the ordering with respect to operations to other memory locations. It can also be used, for example, to do ticket allocation and by expert programmers implementing lock-free algorithms.
-  **memory\_order\_acquire**: A synchronization operation (fence or atomic) that has acquire semantics "acquires" side-effects from a release operation that synchronises with it: if an acquire synchronises with a release, the acquiring unit of execution will see all side-effects preceding that release (and possibly subsequent side-effects.) As part of carefully-designed protocols, programmers can use an "acquire" to safely observe the work of another unit of execution.
-  **memory\_order\_release**: A synchronization operation (fence or atomic operation) that has release semantics "releases" side effects to an acquire operation that synchronises with it. All side effects that precede the release are included in the release. As part of carefully-designed protocols, programmers can use a "release" to make changes made in one unit of execution visible to other units of execution.

NOTE: In general, no acquire must *always* synchronise with any particular release. However, synchronisation can be forced by certain executions. See 3.3.6.2 for detailed rules for when synchronisation must occur.

-  **memory\_order\_acq\_rel**: A synchronization operation with acquire-release semantics has the properties of both the acquire and release memory orders. It is typically used to order read-modify-write operations.

- ✚ **memory\_order\_seq\_cst**: The loads and stores of each unit of execution appear to execute in program (i.e., sequenced-before) order, and the loads and stores from different units of execution appear to be simply interleaved.

Regardless of which `memory_order` is specified, resolving constraints on memory operations across a heterogeneous platform adds considerable overhead to the execution of a program. An OpenCL platform may be able to optimize certain operations that depend on the features of the memory consistency model by restricting the scope of the memory operations. Distinct memory scopes are defined by the values of the `memory_scope` enumeration constant:

- ✚ **memory\_scope\_work\_item**: memory-ordering constraints only apply within the work-item<sup>1</sup>.
- ✚ **memory\_scope\_work\_group**: memory-ordering constraints only apply to work-items executing within a single work-group.
- ✚ **memory\_scope\_device**: memory-ordering constraints only apply to work-items executing on a single device
- ✚ **memory\_scope\_all\_svm\_devices**: memory-ordering constraints apply to work-items executing across multiple devices and (when using SVM) the host. A release performed with **memory\_scope\_all\_svm\_devices** to a buffer that does not have the `CL_MEM_SVM_ATOMICS` flag set will commit to at least **memory\_scope\_device** visibility, with full synchronization of the buffer at a queue synchronization point (e.g. an OpenCL event).

These memory scopes define a hierarchy of visibilities when analyzing the ordering constraints of memory operations. For example if a programmer knows that a sequence of memory operations will only be associated with a collection of work-items from a single work-group (and hence will run on a single device), the implementation is spared the overhead of managing the memory orders across other devices within the same context. This can substantially reduce overhead in a program. All memory scopes are valid when used on global memory or local memory. For local memory, all visibility is constrained to within a given work-group and scopes wider than **memory\_scope\_work\_group** carry no additional meaning.

In the following subsections (leading up to section 3.4), we will explain the synchronization constructs and detailed rules needed to use OpenCL's relaxed memory models. It is important to appreciate, however, that many programs do not benefit from relaxed memory models. Even expert programmers have a difficult time using atomics and fences to write correct programs with relaxed memory models. A large number of OpenCL programs can be written using a simplified memory model. This is accomplished by following these guidelines.

- ✚ Write programs that manage safe sharing of global memory objects through the synchronization points defined by the command queues.
- ✚ Restrict low level synchronization inside work-groups to the work-group functions such as `barrier`.

---

<sup>1</sup> This value for `memory_scope` can only be used with `atomic_work_item_fence` with flags set to `CLK_IMAGE_MEM_FENCE`.

- ✚ If you want sequential consistency behavior with system allocations or fine-grain SVM buffers with atomics support, use only `memory_order_seq_cst` operations with the scope `memory_scope_all_svm_devices`.
- ✚ If you want sequential consistency behavior when not using system allocations or fine-grain SVM buffers with atomics support, use only `memory_order_seq_cst` operations with the scope `memory_scope_device`.
- ✚ Ensure your program has no races.

If these guidelines are followed in your OpenCL programs, you can skip the detailed rules behind the relaxed memory models and go directly to section 3.4.

### 3.3.5 Memory Model: Overview of atomic and fence operations

The OpenCL 2.0 specification defines a number of *synchronization operations* that are used to define memory order constraints in a program. They play a special role in controlling how memory operations in one unit of execution (such as work-items or, when using SVM a host thread) are made visible to another. There are two types of synchronization operations in OpenCL; *atomic operations* and *fences*.

Atomic operations are indivisible. They either occur completely or not at all. These operations are used to order memory operations between units of execution and hence they are parameterized with the `memory_order` and `memory_scope` parameters defined by the OpenCL memory consistency model. The atomic operations for the OpenCL C kernel programming language are similar to the corresponding operations defined by the C11 standard.

The OpenCL 2.0 atomic operations apply to variables of an atomic type (a subset of those in the C11 standard) including atomic versions of the `int`, `uint`, `long`, `ulong`, `float`, `double`, `half`, `intptr_t`, `uintptr_t`, `size_t`, and `ptrdiff_t` types. However, support for some of these atomic types depends on support for the corresponding regular types.

An atomic operation on one or more memory locations is either an acquire operation, a release operation, or both an acquire and release operation. An atomic operation without an associated memory location is a fence and can be either an acquire fence, a release fence, or both an acquire and release fence. In addition, there are relaxed atomic operations, which do not have synchronization properties, and atomic read-modify-write operations, which have special characteristics. [C11 standard, Section 5.1.2.4, paragraph 5, modified]

The orders `memory_order_acquire` (used for reads), `memory_order_release` (used for writes), and `memory_order_acq_rel` (used for read-modify-write operations) are used for simple communication between units of execution using shared variables. Informally, executing a `memory_order_release` on an atomic object A makes all previous side effects visible to any unit of execution that later executes a `memory_order_acquire` on A. The orders `memory_order_acquire`, `memory_order_release`, and `memory_order_acq_rel` do not provide

sequential consistency for race-free programs because they will not ensure that atomic stores followed by atomic loads become visible to other threads in that order.

The fence operation is `atomic_work_item_fence`, which includes a `memory_order` argument as well as the `memory_scope` and `cl_mem_fence_flags` arguments. Depending on the `memory_order` argument, this operation:

- ✚ has no effects, if `memory_order_relaxed`;
- ✚ is an acquire fence, if `memory_order_acquire`;
- ✚ is a release fence, if `memory_order_release`;
- ✚ is both an acquire fence and a release fence, if `memory_order_acq_rel`;
- ✚ is a sequentially-consistent fence with both acquire and release semantics, if `memory_order_seq_cst`.

If specified, the `cl_mem_fence_flags` argument must be `CLK_IMAGE_MEM_FENCE`, `CLK_GLOBAL_MEM_FENCE`, `CLK_LOCAL_MEM_FENCE`, or `CLK_GLOBAL_MEM_FENCE | CLK_LOCAL_MEM_FENCE`.

The synchronization operations in OpenCL can be parameterized by a `memory_scope`. Memory scopes control the extent that an atomic operation or fence is visible with respect to the memory model. These memory scopes may be used when performing atomic operations and fences on global memory and local memory. When used on global memory visibility is bounded by the capabilities of that memory. When used on a fine-grained non-atomic SVM buffer, a coarse-grained SVM buffer, or a non-SVM buffer, operations parameterized with `memory_scope_all_svm_devices` will behave as if they were parameterized with `memory_scope_device`. When used on local memory, visibility is bounded by the work-group and, as a result, `memory_scope` with wider visibility than `memory_scope_work_group` will be reduced to `memory_scope_work_group`.

Two actions **A** and **B** are defined to have an inclusive scope if they have the same scope **P** such that:

- **P** is `memory_scope_work_group` and **A** and **B** are executed by work-items within the same workgroup.
- **P** is `memory_scope_device` and **A** and **B** are executed by work-items on the same device when **A** and **B** apply to an SVM allocation or **A** and **B** are executed by work-items in the same kernel or one of its children when **A** and **B** apply to a `cl_mem` buffer.
- **P** is `memory_scope_all_svm_devices` if **A** and **B** are executed by host threads or by work-items on one or more devices that can share SVM memory with each other and the host process.

### 3.3.6 Memory Model: Memory Ordering Rules

Fundamentally, the issue in a memory model is to understand the orderings in time of modifications to objects in memory. Modifying an object or calling a function that modifies an object are side effects, i.e. changes in the state of the execution environment. Evaluation of an

expression in general includes both value computations and initiation of side effects. Value computation for an lvalue expression includes determining the identity of the designated object. [C11 standard, Section 5.1.2.3, paragraph 2, modified]

We assume that the OpenCL C kernel programming language and host programming languages have a sequenced-before relation between the evaluations executed by a single unit of execution. This sequenced-before relation is an asymmetric, transitive, pair-wise relation between those evaluations, which induces a partial order among them. Given any two evaluations **A** and **B**, if **A** is sequenced-before **B**, then the execution of **A** shall precede the execution of **B**. (Conversely, if **A** is sequenced-before **B**, then **B** is sequenced-after **A**.) If **A** is not sequenced-before or sequenced-after **B**, then **A** and **B** are unsequenced. Evaluations **A** and **B** are indeterminately sequenced when **A** is either sequenced-before or sequenced-after **B**, but it is unspecified which. [C11 standard, Section 5.1.2.3, paragraph 3, modified]

NOTE: sequenced-before is a partial order of the operations executed by a single unit of execution (e.g. a host thread or work-item). It generally corresponds to the source program order of those operations, and is partial because of the undefined argument evaluation order of OpenCL's kernel C language.

In the OpenCL C kernel programming language, the value of an object visible to a work-item **W** at a particular point is the initial value of the object, a value stored in the object by **W**, or a value stored in the object by another work-item or host thread, according to the rules below. Depending on details of the host programming language, the value of an object visible to a host thread may also be the value stored in that object by another work-item or host thread. [C11 standard, Section 5.1.2.4, paragraph 2, modified]

Two expression evaluations conflict if one of them modifies a memory location and the other one reads or modifies the same memory location. [C11 standard, Section 5.1.2.4, paragraph 4]

All modifications to a particular atomic object **M** occur in some particular total order, called the modification order of **M**. If **A** and **B** are modifications of an atomic object **M**, and **A** happens-before **B**, then **A** shall precede **B** in the modification order of **M**, which is defined below. Note that the modification order of an atomic object **M** is independent of whether **M** is in local or global memory. [C11 standard, Section 5.1.2.4, paragraph 7, modified]

A release sequence begins with a release operation **A** on an atomic object **M** and is the maximal contiguous sub-sequence of side effects in the modification order of **M**, where the first operation is **A** and every subsequent operation either is performed by the same work-item or host thread that performed the release or is an atomic read-modify-write operation. [C11 standard, Section 5.1.2.4, paragraph 10, modified]

OpenCL's local and global memories are disjoint. Kernels may access both kinds of memory while host threads may only access global memory. Furthermore, the *flags* argument of OpenCL's `work_group_barrier` function specifies which memory operations the function will make visible: these memory operations can be, for example, just the ones to local memory, or the

ones to global memory, or both. Since the visibility of memory operations can be specified for local memory separately from global memory, we define two related but independent relations, *global-synchronizes-with* and *local-synchronizes-with*. Certain operations on global memory may global-synchronize-with other operations performed by another work-item or host thread. An example is a release atomic operation in one work-item that global-synchronizes-with an acquire atomic operation in a second work-item. Similarly, certain atomic operations on local objects in kernels can local-synchronize-with other atomic operations on those local objects. [C11 standard, Section 5.1.2.4, paragraph 11, modified]

We define two separate happens-before relations: global-happens-before and local-happens-before.

A global memory action **A** global-happens-before a global memory action **B** if

- ✚ **A** is sequenced before **B**, or
- ✚ **A** global-synchronizes-with **B**, or
- ✚ For some global memory action **C**, **A** global-happens-before **C** and **C** global-happens-before **B**.

A local memory action **A** local-happens-before a local memory action **B** if

- ✚ **A** is sequenced before **B**, or
- ✚ **A** local-synchronizes-with **B**, or
- ✚ For some local memory action **C**, **A** local-happens-before **C** and **C** local-happens-before **B**.

An OpenCL implementation shall ensure that no program execution demonstrates a cycle in either the “local-happens-before” relation or the “global-happens-before” relation.

NOTE: The global- and local-happens-before relations are critical to defining what values are read and when data races occur. The global-happens-before relation, for example, defines what global memory operations definitely happen before what other global memory operations. If an operation **A** global-happens-before operation **B** then **A** must occur before **B**; in particular, any write done by **A** will be visible to **B**. The local-happens-before relation has similar properties for local memory. Programmers can use the local- and global-happens-before relations to reason about the order of program actions.

A visible side effect **A** on a global object **M** with respect to a value computation **B** of **M** satisfies the conditions:

- ✚ **A** global-happens-before **B**, and
- ✚ there is no other side effect **X** to **M** such that **A** global-happens-before **X** and **X** global-happens-before **B**.

We define visible side effects for local objects **M** similarly. The value of a non-atomic scalar object **M**, as determined by evaluation **B**, shall be the value stored by the visible side effect **A**. [C11 standard, Section 5.1.2.4, paragraph 19, modified]

The execution of a program contains a data race if it contains two conflicting actions **A** and **B** in different units of execution, and

- ✚ (1) at least one of **A** or **B** is not atomic, or **A** and **B** do not have inclusive memory scope, and
- ✚ (2) the actions are global actions unordered by the global-happens-before relation or are local actions unordered by the local-happens-before relation.

Any such data race results in undefined behavior. [C11 standard, Section 5.1.2.4, paragraph 25, modified]

We also define the visible sequence of side effects on local and global atomic objects. The remaining paragraphs of this subsection define this sequence for a global atomic object **M**; the visible sequence of side effects for a local atomic object is defined similarly by using the local-happens-before relation.

The visible sequence of side effects on a global atomic object **M**, with respect to a value computation **B** of **M**, is a maximal contiguous sub-sequence of side effects in the modification order of **M**, where the first side effect is visible with respect to **B**, and for every side effect, it is not the case that **B** global-happens-before it. The value of **M**, as determined by evaluation **B**, shall be the value stored by some operation in the visible sequence of **M** with respect to **B**. [C11 standard, Section 5.1.2.4, paragraph 22, modified]

If an operation **A** that modifies an atomic object **M** global-happens before an operation **B** that modifies **M**, then **A** shall be earlier than **B** in the modification order of **M**. This requirement is known as write-write coherence.

If a value computation **A** of an atomic object **M** global-happens-before a value computation **B** of **M**, and **A** takes its value from a side effect **X** on **M**, then the value computed by **B** shall either equal the value stored by **X**, or be the value stored by a side effect **Y** on **M**, where **Y** follows **X** in the modification order of **M**. This requirement is known as read-read coherence. [C11 standard, Section 5.1.2.4, paragraph 22, modified]

If a value computation **A** of an atomic object **M** global-happens-before an operation **B** on **M**, then **A** shall take its value from a side effect **X** on **M**, where **X** precedes **B** in the modification order of **M**. This requirement is known as read-write coherence.

If a side effect **X** on an atomic object **M** global-happens-before a value computation **B** of **M**, then the evaluation **B** shall take its value from **X** or from a side effect **Y** that follows **X** in the modification order of **M**. This requirement is known as write-read coherence.

### 3.3.6.1 Memory Ordering Rules: Atomic Operations

This and following sections describe how different program actions in kernel C code and the host program contribute to the local- and global-happens-before relations. This section discusses ordering rules for OpenCL 2.0's atomic operations.



Section 3.2.3 defined the enumerated type `memory_order`.

- ✚ For `memory_order_relaxed`, no operation orders memory.
- ✚ For `memory_order_release`, `memory_order_acq_rel`, and `memory_order_seq_cst`, a store operation performs a release operation on the affected memory location.
- ✚ For `memory_order_acquire`, `memory_order_acq_rel`, and `memory_order_seq_cst`, a load operation performs an acquire operation on the affected memory location. [C11 standard, Section 7.17.3, paragraphs 2-4, modified]

Certain built-in functions synchronize with other built-in functions performed by another unit of execution. This is true for pairs of release and acquire operations under specific circumstances. An atomic operation **A** that performs a release operation on a global object **M** global-synchronizes-with an atomic operation **B** that performs an acquire operation on **M** and reads a value written by any side effect in the release sequence headed by **A**. A similar rule holds for atomic operations on objects in local memory: an atomic operation **A** that performs a release operation on a local object **M** local-synchronizes-with an atomic operation **B** that performs an acquire operation on **M** and reads a value written by any side effect in the release sequence headed by **A**. [C11 standard, Section 5.1.2.4, paragraph 11, modified]

NOTE: Atomic operations specifying `memory_order_relaxed` are relaxed only with respect to memory ordering. Implementations must still guarantee that any given atomic access to a particular atomic object be indivisible with respect to all other atomic accesses to that object. [C11 standard, Section 7.17.3, paragraph 8]

If one of the following two conditions holds:

- ✚ All executed `memory_order_seq_cst` operations have the scope `memory_scope_all_svm_devices` and all affected memory locations are contained in system allocations or fine grain SVM buffers with atomics support
- ✚ All executed `memory_order_seq_cst` operations have the scope `memory_scope_device` and all affected memory locations are not located in system allocated regions or fine-grain SVM buffers with atomics support

then there shall exist a single total order **S** for all `memory_order_seq_cst` operations that is consistent with the modification orders for all affected locations, as well as the appropriate global-happens-before and local-happens-before orders for those locations, such that each `memory_order_seq` operation **B** that loads a value from an atomic object **M** in global or local memory observes one of the following values:

- ✚ the result of the last modification **A** of **M** that precedes **B** in **S**, if it exists, or
- ✚ if **A** exists, the result of some modification of **M** in the visible sequence of side effects with respect to **B** that is not `memory_order_seq_cst` and that does not happen before **A**, or
- ✚ if **A** does not exist, the result of some modification of **M** in the visible sequence of side effects with respect to **B** that is not `memory_order_seq_cst`. [C11 standard, Section 7.17.3, paragraph 6, modified]

Let  $X$  and  $Y$  be two `memory_order_seq_cst` operations. If  $X$  local-synchronizes-with or global-synchronizes-with  $Y$  then  $X$  both local-synchronizes-with  $Y$  and global-synchronizes-with  $Y$ .

If the total order  $S$  exists, the following rules hold:

- ✚ For an atomic operation  $B$  that reads the value of an atomic object  $M$ , if there is a `memory_order_seq_cst` fence  $X$  sequenced-before  $B$ , then  $B$  observes either the last `memory_order_seq_cst` modification of  $M$  preceding  $X$  in the total order  $S$  or a later modification of  $M$  in its modification order. [C11 standard, Section 7.17.3, paragraph 9]
- ✚ For atomic operations  $A$  and  $B$  on an atomic object  $M$ , where  $A$  modifies  $M$  and  $B$  takes its value, if there is a `memory_order_seq_cst` fence  $X$  such that  $A$  is sequenced-before  $X$  and  $B$  follows  $X$  in  $S$ , then  $B$  observes either the effects of  $A$  or a later modification of  $M$  in its modification order. [C11 standard, Section 7.17.3, paragraph 10]
- ✚ For atomic operations  $A$  and  $B$  on an atomic object  $M$ , where  $A$  modifies  $M$  and  $B$  takes its value, if there are `memory_order_seq_cst` fences  $X$  and  $Y$  such that  $A$  is sequenced-before  $X$ ,  $Y$  is sequenced-before  $B$ , and  $X$  precedes  $Y$  in  $S$ , then  $B$  observes either the effects of  $A$  or a later modification of  $M$  in its modification order. [C11 standard, Section 7.17.3, paragraph 11]
- ✚ For atomic operations  $A$  and  $B$  on an atomic object  $M$ , if there are `memory_order_seq_cst` fences  $X$  and  $Y$  such that  $A$  is sequenced-before  $X$ ,  $Y$  is sequenced-before  $B$ , and  $X$  precedes  $Y$  in  $S$ , then  $B$  occurs later than  $A$  in the modification order of  $M$ .

NOTE: `memory_order_seq_cst` ensures sequential consistency only for a program that is (1) free of data races, (2) exclusively uses `memory_order_seq_cst` synchronization operations, and (3) the `memory_order_seq_cst` operations satisfy one of the two conditions for the existence of a single total order  $S$  mentioned earlier in this section. Any use of weaker ordering will invalidate this guarantee unless extreme care is used. In particular, `memory_order_seq_cst` fences ensure a total order only for the fences themselves. Fences cannot, in general, be used to restore sequential consistency for atomic operations with weaker ordering specifications.

Atomic read-modify-write operations should always read the last value (in the modification order) stored before the write associated with the read-modify-write operation. [C11 standard, Section 7.17.3, paragraph 12]

Implementations should ensure that no "out-of-thin-air" values are computed that circularly depend on their own computation.

Note: Under the rules described above, and independent to the previously footnoted C++ issue, it is known that `x == y == 42` is a valid final state in the following problematic example:

```
global atomic_int x = ATOMIC_VAR_INIT(0);
```

```

local atomic_int y = ATOMIC_VAR_INIT(0);

unit_of_execution_1:
... [execution not reading or writing x or y, leading up to:]
int t = atomic_load(&y, memory_order_acquire);
atomic_store(&x, t, memory_order_release);

unit_of_execution_2:
... [execution not reading or writing x or y, leading up to:]
int t = atomic_load(&x, memory_order_acquire);
atomic_store(&y, t, memory_order_release);

```

This outcome is justified by a cycle that is "split" between the *local-happens-before* and *global-happens-before* relations.

This is not useful behavior and implementations should not exploit this phenomenon. It should be expected that in the future this may be disallowed by appropriate updates to the memory model description by the OpenCL committee.

Implementations should make atomic stores visible to atomic loads within a reasonable amount of time. [C11 standard, Section 7.17.3, paragraph 16]

As long as the following conditions are met, a host program sharing SVM memory with a kernel executing on one or more OpenCL devices may use atomic and synchronization operations to ensure that its assignments, and those of the kernel, are visible to each other:

- 1) Either fine-grained buffer or fine-grained system SVM must be used to share memory. While coarse-grained buffer SVM allocations may support atomic operations, visibility on these allocations is not guaranteed except at map and unmap operations.
- 2) The optional OpenCL 2.0 SVM atomic-controlled visibility specified by provision of the CL\_MEM\_SVM\_ATOMICS flag must be supported by the device and the flag provided to the SVM buffer on allocation.
- 3) The host atomic and synchronization operations must be compatible with those of OpenCL C as described in Section 6.13.11. This requires that the size and representation of the data types that the host atomic operations act on be consistent with the OpenCL C atomic types.

If these conditions are met, the host operations will apply at all\_svm\_devices scope.

### 3.3.6.2 Memory Ordering Rules: Fence Operations

This section describes how the OpenCL 2.0 fence operations contribute to the local- and global-happens-before relations.

Earlier, we introduced synchronization primitives called fences. Fences can utilize the acquire memory\_order, release memory\_order, or both. A fence with acquire semantics is called an acquire fence; a fence with release semantics is called a release fence.

A global release fence **A** global-synchronizes-with a global acquire fence **B** if there exist atomic operations **X** and **Y**, both operating on some global atomic object **M**, such that **A** is sequenced-before **X**, **X** modifies **M**, **Y** is sequenced-before **B**, **Y** reads the value written by **X** or a value written by any side effect in the hypothetical release sequence **X** would head if it were a release operation, and that the scopes of **A**, **B** are inclusive. [C11 standard, Section 7.17.4, paragraph 2, modified.]

A global release fence **A** global-synchronizes-with an atomic operation **B** that performs an acquire operation on a global atomic object **M** if there exists an atomic operation **X** such that **A** is sequenced-before **X**, **X** modifies **M**, **B** reads the value written by **X** or a value written by any side effect in the hypothetical release sequence **X** would head if it were a release operation, and the scopes of **A** and **B** are inclusive. [C11 standard, Section 7.17.4, paragraph 3, modified.]

An atomic operation **A** that is a release operation on a global atomic object **M** global-synchronizes-with a global acquire fence **B** if there exists some atomic operation **X** on **M** such that **X** is sequenced-before **B** and reads the value written by **A** or a value written by any side effect in the release sequence headed by **A**, and the scopes of **A** and **B** are inclusive. [C11 standard, Section 7.17.4, paragraph 4, modified.]

A local release fence **A** local-synchronizes-with a local acquire fence **B** if there exist atomic operations **X** and **Y**, both operating on some local atomic object **M**, such that **A** is sequenced-before **X**, **X** modifies **M**, **Y** is sequenced-before **B**, and **Y** reads the value written by **X** or a value written by any side effect in the hypothetical release sequence **X** would head if it were a release operation, and the scopes of **A** and **B** are inclusive. [C11 standard, Section 7.17.4, paragraph 2, modified.]

A local release fence **A** local-synchronizes-with an atomic operation **B** that performs an acquire operation on a local atomic object **M** if there exists an atomic operation **X** such that **A** is sequenced-before **X**, **X** modifies **M**, and **B** reads the value written by **X** or a value written by any side effect in the hypothetical release sequence **X** would head if it were a release operation, and the scopes of **A** and **B** are inclusive. [C11 standard, Section 7.17.4, paragraph 3, modified.]

An atomic operation **A** that is a release operation on a local atomic object **M** local-synchronizes-with a local acquire fence **B** if there exists some atomic operation **X** on **M** such that **X** is sequenced-before **B** and reads the value written by **A** or a value written by any side effect in the release sequence headed by **A**, and the scopes of **A** and **B** are inclusive. [C11 standard, Section 7.17.4, paragraph 4, modified.]

Let **X** and **Y** be two work item fences that each have both the CLK\_GLOBAL\_MEM\_FENCE and CLK\_LOCAL\_MEM\_FENCE flags set. If **X** either local-synchronizes-with or global-synchronizes-with **Y**, then **X** both local-synchronizes-with **Y** and global-synchronizes-with **Y**.

### 3.3.6.3 Memory Ordering Rules: Work-group Functions

The OpenCL C programming language includes collective operations across the work-items within a single work-group. These are called work-group functions. Besides the work-group barrier function, they include the scan, reduction and pipe work-group functions described in Section 6.13.15. We will first discuss the work-group barrier. The other work group functions are discussed afterwards.

The barrier function provides a mechanism for a kernel to synchronize the work-items within a single work-group: informally, each work-item of the work-group must execute the barrier before any are allowed to proceed. It also orders memory operations to a specified combination of one or more address spaces such as local memory or global memory, in a similar manner to a fence.

To precisely specify the memory ordering semantics for barrier, we need to distinguish between a dynamic and a static instance of the call to a barrier. A call to a barrier can appear in a loop, for example, and each execution of the same static barrier call results in a new dynamic instance of the barrier that will independently synchronize a work-group's work-items.

A work-item executing a dynamic instance of a barrier results in two operations, both fences, that are called the entry and exit fences. These fences obey all the rules for fences specified elsewhere in this chapter as well as the following:

- ✚ The entry fence is a release fence with the same flags and scope as requested for the barrier.
- ✚ The exit fence is an acquire fence with the same flags and scope as requested for the barrier.
- ✚ For each work-item the entry fence is sequenced before the exit fence.
- ✚ If the flags have `CLK_GLOBAL_MEM_FENCE` set then for each work-item the entry fence global-synchronizes-with the exit fence of all other work-items in the same work-group.
- ✚ If the flags have `CLK_LOCAL_MEM_FENCE` set then for each work-item the entry fence local-synchronizes-with the exit fence of all other work-items in the same work-group.

The other work-group functions include such functions as `work_group_all()` and `work_group_broadcast()` and are described in Section 6.13.15. The use of these workgroup functions implies sequenced-before relationships between statements within the execution of a single work-item in order to satisfy data dependencies. For example, a work item that provides a value to a workgroup function must behave as if it generates that value before beginning execution of that workgroup function. Furthermore, the programmer must ensure that all work items in a work group must execute the same workgroup function call site, or dynamic workgroup function instance.

### 3.3.6.4 Memory Ordering Rules: Host-side and Device-side Commands

This section describes how the OpenCL API functions associated with command-queues contribute to happens-before relations. There are two types of command queues and associated API functions in OpenCL 2.0; *host command-queues* and *device command-queues*. The interaction of these command queues with the memory model are for the most part equivalent. In a few cases, the rules only applies to the host command-queue. We will indicate these special cases by specifically denoting the host command-queue in the memory ordering rule. SVM memory consistency in such instances is implied only with respect to synchronizing host commands.

Memory ordering rules in this section apply to all memory objects (buffers, images and pipes) as well as to SVM allocations where no earlier, and more fine-grained, rules apply.

In the remainder of this section, we assume that each command **C** enqueued onto a command-queue has an associated event object **E** that signals its execution status, regardless of whether **E** was returned to the unit of execution that enqueued **C**. We also distinguish between the API function call that enqueues a command **C** and creates an event **E**, the execution of **C**, and the completion of **C** (which marks the event **E** as complete).

The ordering and synchronization rules for API commands are defined as following:

1. If an API function call **X** enqueues a command **C**, then **X** global-synchronizes-with **C**. For example, a host API function to enqueue a kernel global-synchronizes-with the start of that kernel-instance's execution, so that memory updates sequenced-before the enqueue kernel function call will global-happen-before any kernel reads or writes to those same memory locations. For a device-side enqueue, global memory updates sequenced before **X** happens-before **C** reads or writes to those memory locations only in the case of fine-grained SVM.
2. If **E** is an event upon which a command **C** waits, then **E** global-synchronizes-with **C**. In particular, if **C** waits on an event **E** that is tracking the execution status of the command **C1**, then memory operations done by **C1** will global-happen-before memory operations done by **C**. As an example, assume we have an OpenCL program using coarse-grain SVM sharing that enqueues a kernel to a host command-queue to manipulate the contents of a region of a buffer that the host thread then accesses after the kernel completes. To do this, the host thread can call `clEnqueueMapBuffer` to enqueue a blocking-mode map command to map that buffer region, specifying that the map command must wait on an event signaling the kernel's completion. When `clEnqueueMapBuffer` returns, any memory operations performed by the kernel to that buffer region will global- happen-before subsequent memory operations made by the host thread.
3. If a command **C** has an event **E** that signals its completion, then **C** global-synchronizes-with **E**.
4. For a command **C** enqueued to a host-side command queue, if **C** has an event **E** that signals its completion, then **E** global- synchronizes-with an API call **X** that waits on **E**. For example, if a host thread or kernel-instance calls the wait-for-events function on **E** (e.g. the `clWaitForEvents` function called from a host thread), then **E** global-synchronizes-with that wait-for-events function call.

5. If commands **C** and **C1** are enqueued in that sequence onto an in-order command-queue, then the event (including the event implied between **C** and **C1** due to the in-order queue) signaling **C**'s completion global-synchronizes-with **C1**. Note that in OpenCL 2.0, only a host command-queue can be configured as an in-order queue.
6. If an API call enqueues a marker command **C** with an empty list of events upon which **C** should wait, then the events of all commands enqueued prior to **C** in the command-queue global-synchronize-with **C**.
7. If a host API call enqueues a command-queue barrier command **C** with an empty list of events on which **C** should wait, then the events of all commands enqueued prior to **C** in the command-queue global-synchronize-with **C**. In addition, the event signaling the completion of **C** global-synchronizes-with all commands enqueued after **C** in the command-queue.
8. If a host thread executes a `clFinish` call **X**, then the events of all commands enqueued prior to **X** in the command-queue global-synchronizes-with **X**.
9. The start of a kernel-instance **K** global-synchronizes-with all operations in the work items of **K**. Note that this includes the execution of any atomic operations by the work items in a program using fine-grain SVM.
10. All operations of all work items of a kernel-instance **K** global-synchronizes-with the event signaling the completion of **K**. Note that this also includes the execution of any atomic operations by the work items in a program using fine-grain SVM.
11. If a callback procedure **P** is registered on an event **E**, then **E** global-synchronizes-with all operations of **P**. Note that callback procedures are only defined for commands within host command-queues.
12. If **C** is a command that waits for an event **E**'s completion, and API function call **X** sets the status of a user event **E**'s status to `CL_COMPLETE` (for example, from a host thread using a `clSetUserEventStatus` function), then **X** global-synchronizes-with **C**.
13. If a device enqueues a command **C** with the `CLK_ENQUEUE_FLAGS_WAIT_KERNEL` flag, then the end state of the parent kernel instance global-synchronizes with **C**.
14. If a work-group enqueues a command **C** with the `CLK_ENQUEUE_FLAGS_WAIT_WORK_GROUP` flag, then the end state of the work-group global-synchronizes with **C**.

When using an out-of-order command queue, a wait on an event or a marker or command-queue barrier command can be used to ensure the correct ordering of dependent commands. In those cases, the wait for the event or the marker or barrier command will provide the necessary global-synchronizes-with relation.

In the remainder of this section, we discuss a few points regarding the ordering rules for commands with a host command queue.

The OpenCL 1.2 standard describes a synchronization point as a kernel-instance or host program location where the contents of memory visible to different work-items or command-queue

commands are the same. It also says that waiting on an event and a command-queue barrier are synchronization points between commands in command- queues. Four of the rules listed above (2, 4, 7, and 8) cover these OpenCL synchronization points.

A map operation (`clEnqueueMapBuffer` or `clEnqueueMapImage`) performed on a non-SVM buffer or a coarse-grained SVM buffer is allowed to overwrite the entire target region with the latest runtime view of the data as seen by the command with which the map operation synchronizes, whether the values were written by the executing kernels or not. Any values that were changed within this region by another kernel or host thread while the kernel synchronizing with the map operation was executing may be overwritten by the map operation.

Access to non-SVM `cl_mem` buffers and coarse-grained SVM allocations is ordered at synchronization points between host commands. In the presence of an out-of-order command queue or a set of command queues mapped to the same device, multiple kernel instances may execute concurrently on the same device.

In this situation:

- ✚ access to shared locations or disjoint locations in a single `cl_mem` object when using atomic operations from different kernel instances enqueued from the host such that one or more of the atomic operations is a write is implementation-defined and correct behavior is not guaranteed except at synchronization points.
- ✚ access to shared locations or disjoint locations in a single `cl_mem` object when using atomic operations from different kernel instances consisting of a parent kernel and any number of child kernels enqueued by that kernel is guaranteed under the memory ordering rules described earlier in this section.
- ✚ access to shared locations or disjoint locations in a single coarse-grained or fine-grained SVM allocation when using atomic operations from different kernel instances enqueued from the host to a single device is guaranteed under the memory ordering rules described earlier in this section.




If fine-grain SVM is used but without support for the OpenCL 2.0 atomic operations, then the host and devices can concurrently read the same memory locations and can concurrently update non-overlapping memory regions, but attempts to update the same memory locations are undefined. Memory consistency is guaranteed at the OpenCL synchronization points without the need for calls to `clEnqueueMapBuffer` and `clEnqueueUnmapMemObject`. For fine-grained SVM buffers it is guaranteed that at synchronization points only values written by the kernel will be updated. No writes to fine-grained SVM buffers can be introduced that were not in the original program.

## 3.4 The OpenCL Framework

The OpenCL framework allows applications to use a host and one or more OpenCL devices as a single heterogeneous parallel computer system. The framework contains the following



components:

-  **OpenCL Platform layer:** The platform layer allows the host program to discover OpenCL devices and their capabilities and to create contexts.
-  **OpenCL Runtime:** The runtime allows the host program to manipulate contexts once they have been created.
-  **OpenCL Compiler:** The OpenCL compiler creates program executables that contain OpenCL kernels. The OpenCL C programming language implemented by the compiler supports a subset of the ISO C99 language with extensions for parallelism.

### 3.4.1 OpenCL Framework: Mixed Version Support

OpenCL supports devices with different capabilities under a single platform. This includes devices which conform to different versions of the OpenCL specification. There are three version identifiers to consider for an OpenCL system: the platform version, the version of a device, and the version(s) of the OpenCL C programming language supported on a device.

The platform version indicates the version of the OpenCL runtime that is supported. This includes all of the APIs that the host can use to interact with resources exposed by the OpenCL runtime; including contexts, memory objects, devices, and command queues.

The device version is an indication of the device's capabilities separate from the runtime and compiler as represented by the device info returned by **clGetDeviceInfo**. Examples of attributes associated with the device version are resource limits (e.g., minimum size of local memory per compute unit) and extended functionality (e.g., list of supported KHR extensions). The version returned corresponds to the highest version of the OpenCL specification for which the device is conformant, but is not higher than the platform version.

The language version for a device represents the OpenCL programming language features a developer can assume are supported on a given device. The version reported is the highest version of the language supported.

Backwards compatibility is an important goal for the OpenCL C programming language. An OpenCL 1.x device must support at least one 1.x version of the OpenCL C programming language. An OpenCL 2.0 device must support at least one 1.x version of the OpenCL C programming language and the OpenCL C 2.0 programming language. If multiple language versions are supported, the compiler defaults to using the highest OpenCL 1.x language version supported for the device (typically OpenCL 1.2). To utilize the OpenCL 2.0 Kernel programming language, a programmer must specifically set the appropriate compiler flag (-cl-std=CL2.0). The language version is not higher than the platform version, but may exceed the device version (see section 5.8.4.5).

## 4. The OpenCL Platform Layer

This section describes the OpenCL platform layer which implements platform-specific features that allow applications to query OpenCL devices, device configuration information, and to create OpenCL contexts using one or more devices.

### 4.1 Querying Platform Info

The list of platforms available can be obtained using the following function.



```
cl_int  clGetPlatformIDs (cl_uint num_entries,  
                          cl_platform_id *platforms,  
                          cl_uint *num_platforms)
```

*num\_entries* is the number of *cl\_platform\_id* entries that can be added to *platforms*. If *platforms* is not NULL, the *num\_entries* must be greater than zero.

*platforms* returns a list of OpenCL platforms found. The *cl\_platform\_id* values returned in *platforms* can be used to identify a specific OpenCL platform. If *platforms* argument is NULL, this argument is ignored. The number of OpenCL platforms returned is the minimum of the value specified by *num\_entries* or the number of OpenCL platforms available.

*num\_platforms* returns the number of OpenCL platforms available. If *num\_platforms* is NULL, this argument is ignored.

**clGetPlatformIDs** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_VALUE if *num\_entries* is equal to zero and *platforms* is not NULL or if both *num\_platforms* and *platforms* are NULL.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clGetPlatformInfo (cl_platform_id platform,  
                           cl_platform_info param_name,  
                           size_t param_value_size,  
                           void *param_value,  
                           size_t *param_value_size_ret)
```

gets specific information about the OpenCL platform. The information that can be queried using **clGetPlatformInfo** is specified in *table 4.1*.

*platform* refers to the platform ID returned by **clGetPlatformIDs** or can be NULL. If *platform* is NULL, the behavior is implementation-defined.

*param\_name* is an enumeration constant that identifies the platform information being queried. It can be one of the following values as specified in *table 4.1*.

*param\_value* is a pointer to memory location where appropriate values for a given *param\_name* as specified in *table 4.1* will be returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* specifies the size in bytes of memory pointed to by *param\_value*. This size in bytes must be  $\geq$  size of return type specified in *table 4.1*.

*param\_value\_size\_ret* returns the actual size in bytes of data being queried by *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

cl_platform_info	Return Type	Description
CL_PLATFORM_PROFILE	char <sup>2</sup> []	OpenCL profile string. Returns the profile name supported by the implementation. The profile name returned can be one of the following strings:  FULL_PROFILE – if the implementation supports the OpenCL specification (functionality defined as part of the core specification and does not require any extensions to be supported).  EMBEDDED_PROFILE - if the implementation supports the OpenCL embedded profile. The embedded profile is defined to be a subset for each version of OpenCL. The embedded profile for OpenCL 2.0 is described in <i>section 10</i> .
CL_PLATFORM_VERSION	char[]	OpenCL version string. Returns the OpenCL version supported by the implementation. This version string has the following format:

<sup>2</sup> A null terminated string is returned by OpenCL query function calls if the return type of the information being queried is a char[].

		<i>OpenCL</i> <space><major_version.minor_version><space><platform-specific information>  The <i>major_version.minor_version</i> value returned will be 2.0.
<b>CL_PLATFORM_NAME</b>	char[]	Platform name string.
<b>CL_PLATFORM_VENDOR</b>	char[]	Platform vendor string.
<b>CL_PLATFORM_EXTENSIONS</b>	char[]	Returns a space separated list of extension names (the extension names themselves do not contain any spaces) supported by the platform. Extensions defined here must be supported by all devices associated with this platform.

**Table 4.1.**     *OpenCL Platform Queries*

**clGetPlatformInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors<sup>3</sup>:

- ✚ CL\_INVALID\_PLATFORM if *platform* is not a valid platform.
- ✚ CL\_INVALID\_VALUE if *param\_name* is not one of the supported values or if size in bytes specified by *param\_value\_size* is < size of return type as specified in *table 4.1* and *param\_value* is not a NULL value.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 4.2 Querying Devices

The list of devices available on a platform can be obtained using the following function.

```
cl_int      clGetDeviceIDs4(cl_platform_id platform,
                             cl_device_type device_type,
                             cl_uint num_entries,
                             cl_device_id *devices,
                             cl_uint *num_devices)
```

<sup>3</sup> The OpenCL specification does not describe the order of precedence for error codes returned by API calls.

<sup>4</sup> **clGetDeviceIDs** may return all or a subset of the actual physical devices present in the platform and that match *device\_type*.

*platform* refers to the platform ID returned by **clGetPlatformIDs** or can be NULL. If *platform* is NULL, the behavior is implementation-defined.

*device\_type* is a bitfield that identifies the type of OpenCL device. The *device\_type* can be used to query specific OpenCL devices or all OpenCL devices available. The valid values for *device\_type* are specified in *table 4.2*.

<b>cl_device_type</b>	<b>Description</b>
<b>CL_DEVICE_TYPE_CPU</b>	An OpenCL device that is the host processor. The host processor runs the OpenCL implementations and is a single or multi-core CPU.
<b>CL_DEVICE_TYPE_GPU</b>	An OpenCL device that is a GPU. By this we mean that the device can also be used to accelerate a 3D API such as OpenGL or DirectX.
<b>CL_DEVICE_TYPE_ACCELERATOR</b>	Dedicated OpenCL accelerators (for example the IBM CELL Blade). These devices communicate with the host processor using a peripheral interconnect such as PCIe.
<b>CL_DEVICE_TYPE_CUSTOM</b>	Dedicated accelerators that do not support programs written in OpenCL C.
<b>CL_DEVICE_TYPE_DEFAULT</b>	The default OpenCL device in the system. The default device cannot be a <b>CL_DEVICE_TYPE_CUSTOM</b> device.
<b>CL_DEVICE_TYPE_ALL</b>	All OpenCL devices available in the system except <b>CL_DEVICE_TYPE_CUSTOM</b> devices..


**Table 4.2.**     *List of OpenCL Device Categories*

*num\_entries* is the number of *cl\_device\_id* entries that can be added to *devices*. If *devices* is not NULL, the *num\_entries* must be greater than zero.

*devices* returns a list of OpenCL devices found. The *cl\_device\_id* values returned in *devices* can be used to identify a specific OpenCL device. If *devices* argument is NULL, this argument is ignored. The number of OpenCL devices returned is the minimum of the value specified by *num\_entries* or the number of OpenCL devices whose type matches *device\_type*.

*num\_devices* returns the number of OpenCL devices available that match *device\_type*. If *num\_devices* is NULL, this argument is ignored.

**clGetDeviceIDs** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

 **CL\_INVALID\_PLATFORM** if *platform* is not a valid platform.

- ✚ CL\_INVALID\_DEVICE\_TYPE if *device\_type* is not a valid value.
- ✚ CL\_INVALID\_VALUE if *num\_entries* is equal to zero and *devices* is not NULL or if both *num\_devices* and *devices* are NULL.
- ✚ CL\_DEVICE\_NOT\_FOUND if no OpenCL devices that matched *device\_type* were found.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The application can query specific capabilities of the OpenCL device(s) returned by **clGetDeviceIDs**. This can be used by the application to determine which device(s) to use.

The function

```
cl_int          clGetDeviceInfo (cl_device_id device,
                                   cl_device_info param_name,
                                   size_t param_value_size,
                                   void *param_value,
                                   size_t *param_value_size_ret)
```

gets specific information about an OpenCL device. *device* may be a device returned by **clGetDeviceIDs** or a sub-device created by **clCreateSubDevices**. If *device* is a sub-device, the specific information for the sub-device will be returned. The information that can be queried using **clGetDeviceInfo** is specified in *table 4.3*.

*device* is a device returned by **clGetDeviceIDs**.

*param\_name* is an enumeration constant that identifies the device information being queried. It can be one of the following values as specified in *table 4.3*.

*param\_value* is a pointer to memory location where appropriate values for a given *param\_name* as specified in *table 4.3* will be returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* specifies the size in bytes of memory pointed to by *param\_value*. This size in bytes must be  $\geq$  size of return type specified in *table 4.3*.

*param\_value\_size\_ret* returns the actual size in bytes of data being queried by *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

<b>cl_device_info</b>	<b>Return Type</b>	<b>Description</b>
<b>CL_DEVICE_TYPE</b>	cl_device_type	The OpenCL device type. Currently supported values are:  CL_DEVICE_TYPE_CPU, CL_DEVICE_TYPE_GPU, CL_DEVICE_TYPE_ACCELERATOR, CL_DEVICE_TYPE_DEFAULT, a combination of the above types or CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_VENDOR_ID</b>	cl_uint	A unique device vendor identifier. An example of a unique device identifier could be the PCIe ID.
<b>CL_DEVICE_MAX_COMPUTE_UNITS</b>	cl_uint	The number of parallel compute units on the OpenCL device. A work-group executes on a single compute unit. The minimum value is 1.
<b>CL_DEVICE_MAX_WORK_ITEM_DIMENSIONS</b>	cl_uint	Maximum dimensions that specify the global and local work-item IDs used by the data parallel execution model. (Refer to <b>clEnqueueNDRangeKernel</b> ). The minimum value is 3 for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_MAX_WORK_ITEM_SIZES</b>	size_t []	Maximum number of work-items that can be specified in each dimension of the work-group to <b>clEnqueueNDRangeKernel</b> .  Returns <i>n</i> size_t entries, where <i>n</i> is the value returned by the query for CL_DEVICE_MAX_WORK_ITEM_DIMENSIONS.  The minimum value is (1, 1, 1) for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_MAX_WORK_GROUP_SIZE</b>	size_t	Maximum number of work-items in a work-group that a device is capable of executing on a single compute unit, for any given kernel-instance running on the device. (Refer also to <b>clEnqueueNDRangeKernel</b> and <b>CL_KERNEL_WORK_GROUP_SIZE</b> ). The minimum value is 1.
<b>CL_DEVICE_PREFERRED_VECTOR_WIDTH_CHAR</b>	cl_uint	Preferred native vector width size for built-in scalar types that can be put into vectors.

<b>CL_DEVICE_PREFERRED_VECTOR_WIDTH_SHORT</b>  <b>CL_DEVICE_PREFERRED_VECTOR_WIDTH_INT</b>  <b>CL_DEVICE_PREFERRED_VECTOR_WIDTH_LONG</b>  <b>CL_DEVICE_PREFERRED_VECTOR_WIDTH_FLOAT</b>  <b>CL_DEVICE_PREFERRED_VECTOR_WIDTH_DOUBLE</b>  <b>CL_DEVICE_PREFERRED_VECTOR_WIDTH_HALF</b>		<p>The vector width is defined as the number of scalar elements that can be stored in the vector.</p> <p>If double precision is not supported, CL_DEVICE_PREFERRED_VECTOR_WIDTH_DOUBLE must return 0.</p> <p>If the <b>cl_khr_fp16</b> extension is not supported, CL_DEVICE_PREFERRED_VECTOR_WIDTH_HALF must return 0.</p>
<b>CL_DEVICE_NATIVE_VECTOR_WIDTH_CHAR</b>  <b>CL_DEVICE_NATIVE_VECTOR_WIDTH_SHORT</b>  <b>CL_DEVICE_NATIVE_VECTOR_WIDTH_INT</b>  <b>CL_DEVICE_NATIVE_VECTOR_WIDTH_LONG</b>  <b>CL_DEVICE_NATIVE_VECTOR_WIDTH_FLOAT</b>  <b>CL_DEVICE_NATIVE_VECTOR_WIDTH_DOUBLE</b>  <b>CL_DEVICE_NATIVE_VECTOR_WIDTH_HALF</b>	cl_uint	<p>Returns the native ISA vector width. The vector width is defined as the number of scalar elements that can be stored in the vector.</p> <p>If double precision is not supported, CL_DEVICE_NATIVE_VECTOR_WIDTH_DOUBLE must return 0.</p> <p>If the <b>cl_khr_fp16</b> extension is not supported, CL_DEVICE_NATIVE_VECTOR_WIDTH_HALF must return 0.</p>
<b>CL_DEVICE_MAX_CLOCK_FREQUENCY</b>	cl_uint	Maximum configured clock frequency of the device in MHz.
<b>CL_DEVICE_ADDRESS_BITS</b>	cl_uint	<p>The default compute device address space size of the global address space specified as an unsigned integer value in bits. Currently supported values are 32 or 64 bits.</p>
<b>CL_DEVICE_MAX_MEM_ALLOC_SIZE</b>	cl_ulong	<p>Max size of memory object allocation in bytes. The minimum value is max (min(1024*1024*1024, 1/4<sup>th</sup> of <b>CL_DEVICE_GLOBAL_MEM_SIZE</b>), 32*1024*1024) for devices that are not of type CL_DEVICE_TYPE_CUSTOM.</p>
<b>CL_DEVICE_IMAGE_SUPPORT</b>	cl_bool	Is CL_TRUE if images are supported by the OpenCL device and CL_FALSE otherwise.
<b>CL_DEVICE_MAX_READ_IMAGE_ARGS</b>	cl_uint	Max number of image objects arguments



		of a kernel declared with the <code>read_only</code> qualifier. The minimum value is 128 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_MAX_WRITE_IMAGE_ARGS</b>	<code>cl_uint</code>	Max number of image objects arguments of a kernel declared with the <code>write_only</code> qualifier. The minimum value is 64 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_MAX_READ_WRITE_IMAGE_ARGS<sup>5</sup></b>	<code>cl_uint</code>	Max number of image objects arguments of a kernel declared with the <code>write_only</code> or <code>read_write</code> qualifier. The minimum value is 64 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE2D_MAX_WIDTH</b>	<code>size_t</code>	Max width of 2D image or 1D image not created from a buffer object in pixels.  The minimum value is 16384 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE2D_MAX_HEIGHT</b>	<code>size_t</code>	Max height of 2D image in pixels.  The minimum value is 16384 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE3D_MAX_WIDTH</b>	<code>size_t</code>	Max width of 3D image in pixels.  The minimum value is 2048 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE3D_MAX_HEIGHT</b>	<code>size_t</code>	Max height of 3D image in pixels.  The minimum value is 2048 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE3D_MAX_DEPTH</b>	<code>size_t</code>	Max depth of 3D image in pixels.  The minimum value is 2048 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE_MAX_BUFFER_SIZE</b>	<code>size_t</code>	Max number of pixels for a 1D image

<sup>5</sup> NOTE: **CL\_DEVICE\_MAX\_WRITE\_IMAGE\_ARGS** is only there for backward compatibility. **CL\_DEVICE\_MAX\_READ\_WRITE\_IMAGE\_ARGS** should be used instead.

		<p>created from a buffer object.</p> <p>The minimum value is 65536 if CL_DEVICE_IMAGE_SUPPORT is CL_TRUE.</p>
CL_DEVICE_IMAGE_MAX_ARRAY_SIZE	size_t	<p>Max number of images in a 1D or 2D image array.</p> <p>The minimum value is 2048 if CL_DEVICE_IMAGE_SUPPORT is CL_TRUE.</p>
CL_DEVICE_MAX_SAMPLERS	cl_uint	<p>Maximum number of samplers that can be used in a kernel.</p> <p>The minimum value is 16 if CL_DEVICE_IMAGE_SUPPORT is CL_TRUE.</p>
CL_DEVICE_IMAGE_PITCH_ALIGNMENT	cl_uint	<p>The row pitch alignment size in pixels for 2D images created from a buffer. The value returned must be a power of 2.</p> <p>If the device does not support images, this value must be 0.</p>
CL_DEVICE_IMAGE_BASE_ADDRESS_ALIGNMENT	cl_uint	<p>This query should be used when a 2D image is created from a buffer which was created using CL_MEM_USE_HOST_PTR. The value returned must be a power of 2.</p> <p>This query specifies the minimum alignment in pixels of the <i>host_ptr</i> specified to <b>clCreateBuffer</b>.</p> <p>If the device does not support images, this value must be 0.</p>
CL_DEVICE_MAX_PIPE_ARGS	cl_uint	<p>The maximum number of pipe objects that can be passed as arguments to a kernel. The minimum value is 16.</p>
CL_DEVICE_PIPE_MAX_ACTIVE_RESERVATIONS	cl_uint	<p>The maximum number of reservations that can be active for a pipe per work-item in a kernel. A work-group reservation is counted as one reservation per work-item. The minimum value is 1.</p>
CL_DEVICE_PIPE_MAX_PACKET_SIZE	cl_uint	<p>The maximum size of pipe packet in bytes. The minimum value is 1024 bytes.</p>

<b>CL_DEVICE_MAX_PARAMETER_SIZE</b>	size_t	<p>Max size in bytes of all arguments that can be passed to a kernel.</p> <p>The minimum value is 1024 for devices that are not of type CL_DEVICE_TYPE_CUSTOM. For this minimum value, only a maximum of 128 arguments can be passed to a kernel.</p>
<b>CL_DEVICE_MEM_BASE_ADDR_ALIGN</b>	cl_uint	<p>Alignment requirement (in bits) for sub-buffer offsets. The minimum value is the size (in bits) of the largest OpenCL built-in data type supported by the device (long16 in FULL profile, long16 or int16 in EMBEDDED profile) for devices that are not of type CL_DEVICE_TYPE_CUSTOM.</p>
<b>CL_DEVICE_SINGLE_FP_CONFIG<sup>6</sup></b>	cl_device_fp_config	<p>Describes single precision floating-point capability of the device. This is a bit-field that describes one or more of the following values:</p> <p>CL_FP_DENORM – denorms are supported</p> <p>CL_FP_INF_NAN – INF and quiet NaNs are supported.</p> <p>CL_FP_ROUND_TO_NEAREST – round to nearest even rounding mode supported</p> <p>CL_FP_ROUND_TO_ZERO – round to zero rounding mode supported</p> <p>CL_FP_ROUND_TO_INF – round to positive and negative infinity rounding modes supported</p> <p>CL_FP_FMA – IEEE754-2008 fused multiply-add is supported.</p> <p>CL_FP_CORRECTLY_ROUNDED_DIVIDE_SQRT – divide and sqrt are correctly rounded as defined by the IEEE754 specification.</p>

<sup>6</sup> The optional rounding modes should be included as a device capability only if it is supported natively. All explicit conversion functions with specific rounding modes must still operate correctly.

		<p>CL_FP_SOFT_FLOAT – Basic floating-point operations (such as addition, subtraction, multiplication) are implemented in software.</p> <p>For the full profile, the mandated minimum floating-point capability for devices that are not of type CL_DEVICE_TYPE_CUSTOM is: CL_FP_ROUND_TO_NEAREST   CL_FP_INF_NAN.</p> <p>For the embedded profile, see section 10.</p>
CL_DEVICE_DOUBLE_FP_CONFIG <sup>7</sup>	cl_device_fp_config	<p>Describes double precision floating-point capability of the OpenCL device. This is a bit-field that describes one or more of the following values:</p> <p>CL_FP_DENORM – denorms are supported</p> <p>CL_FP_INF_NAN – INF and NaNs are supported.</p> <p>CL_FP_ROUND_TO_NEAREST – round to nearest even rounding mode supported.</p> <p>CL_FP_ROUND_TO_ZERO – round to zero rounding mode supported.</p> <p>CL_FP_ROUND_TO_INF – round to positive and negative infinity rounding modes supported.</p> <p>CP_FP_FMA – IEEE754-2008 fused multiply-add is supported.</p> <p>CL_FP_SOFT_FLOAT – Basic floating-point operations (such as addition, subtraction, multiplication) are implemented in software.</p> <p>Double precision is an optional feature so the mandated minimum double precision floating-point capability is 0.</p>

<sup>7</sup> The optional rounding modes should be included as a device capability only if it is supported natively. All explicit conversion functions with specific rounding modes must still operate correctly.

		If double precision is supported by the device, then the minimum double precision floating-point capability must be: CL_FP_FMA   CL_FP_ROUND_TO_NEAREST   CL_FP_INF_NAN   CL_FP_DENORM.
<b>CL_DEVICE_GLOBAL_MEM_CACHE_TYPE</b>	cl_device_mem_cache_type	Type of global memory cache supported. Valid values are: CL_NONE, CL_READ_ONLY_CACHE and CL_READ_WRITE_CACHE.
<b>CL_DEVICE_GLOBAL_MEM_CACHELINE_SIZE</b>	cl_uint	Size of global memory cache line in bytes.
<b>CL_DEVICE_GLOBAL_MEM_CACHE_SIZE</b>	cl_ulong	Size of global memory cache in bytes.
<b>CL_DEVICE_GLOBAL_MEM_SIZE</b>	cl_ulong	Size of global device memory in bytes.
<b>CL_DEVICE_MAX_CONSTANT_BUFFER_SIZE</b>	cl_ulong	Max size in bytes of a constant buffer allocation. The minimum value is 64 KB for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_MAX_CONSTANT_ARGS</b>	cl_uint	Max number of arguments declared with the __constant qualifier in a kernel. The minimum value is 8 for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_MAX_GLOBAL_VARIABLE_SIZE</b>	size_t	The maximum number of bytes of storage that may be allocated for any single variable in program scope or inside a function in OpenCL C declared in the global address space.  The minimum value is 64 KB.
<b>CL_DEVICE_GLOBAL_VARIABLE_PREFERRED_TOTAL_SIZE</b>	size_t	Maximum preferred total size, in bytes, of all program variables in the global address space. This is a performance hint. An implementation may place such variables in storage with optimized device access. This query returns the capacity of such storage. The minimum value is 0.
<b>CL_DEVICE_LOCAL_MEM_TYPE</b>	cl_device_local_mem_type	Type of local memory supported. This can be set to CL_LOCAL implying dedicated local memory storage such as SRAM, or CL_GLOBAL.

		For custom devices, CL_NONE can also be returned indicating no local memory support.
<b>CL_DEVICE_LOCAL_MEM_SIZE</b>	cl_ulong	Size of local memory region in bytes. The minimum value is 32 KB for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_ERROR_CORRECTION_SUPPORT</b>	cl_bool	Is CL_TRUE if the device implements error correction for all accesses to compute device memory (global and constant). Is CL_FALSE if the device does not implement such error correction.
<b>CL_DEVICE_PROFILING_TIMER_RESOLUTION</b>	size_t	Describes the resolution of device timer. This is measured in nanoseconds. Refer to <i>section 5.14</i> for details.
<b>CL_DEVICE_ENDIAN_LITTLE</b>	cl_bool	Is CL_TRUE if the OpenCL device is a little endian device and CL_FALSE otherwise.
<b>CL_DEVICE_AVAILABLE</b>	cl_bool	Is CL_TRUE if the device is available and CL_FALSE otherwise. A device is considered to be available if the device can be expected to successfully execute commands enqueued to the device.
<b>CL_DEVICE_COMPILER_AVAILABLE</b>	cl_bool	Is CL_FALSE if the implementation does not have a compiler available to compile the program source. Is CL_TRUE if the compiler is available.  This can be CL_FALSE for the embedded platform profile only.
<b>CL_DEVICE_LINKER_AVAILABLE</b>	cl_bool	Is CL_FALSE if the implementation does not have a linker available. Is CL_TRUE if the linker is available.  This can be CL_FALSE for the embedded platform profile only.  This must be CL_TRUE if <b>CL_DEVICE_COMPILER_AVAILABLE</b> is CL_TRUE.

<b>CL_DEVICE_EXECUTION_CAPABILITIES</b>	cl_device_exec_capabilities	<p>Describes the execution capabilities of the device. This is a bit-field that describes one or more of the following values:</p> <p>CL_EXEC_KERNEL – The OpenCL device can execute OpenCL kernels.</p> <p>CL_EXEC_NATIVE_KERNEL – The OpenCL device can execute native kernels.</p> <p>The mandated minimum capability is: CL_EXEC_KERNEL.</p>
<b>CL_DEVICE_QUEUE_ON_HOST_PROPERTIES<sup>8</sup></b>	cl_command_queue_properties	<p>Describes the on host command-queue properties supported by the device. This is a bit-field that describes one or more of the following values:</p> <p>CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE</p> <p>CL_QUEUE_PROFILING_ENABLE</p> <p>These properties are described in <i>table 5.1</i>.</p> <p>The mandated minimum capability is: CL_QUEUE_PROFILING_ENABLE.</p>
<b>CL_DEVICE_QUEUE_ON_DEVICE_PROPERTIES</b>	cl_command_queue_properties	<p>Describes the on device command-queue properties supported by the device. This is a bit-field that describes one or more of the following values:</p> <p>CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE</p> <p>CL_QUEUE_PROFILING_ENABLE</p> <p>These properties are described in <i>table 5.1</i>.</p> <p>The mandated minimum capability is: CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE   CL_QUEUE_PROFILING_ENABLE.</p>
<b>CL_DEVICE_QUEUE_ON_DEVICE_PREFERRED_SIZE</b>	cl_uint	The size of the device queue in bytes preferred by the implementation.

<sup>8</sup> CL\_DEVICE\_QUEUE\_PROPERTIES is deprecated and replaced by CL\_DEVICE\_QUEUE\_ON\_HOST\_PROPERTIES.

		Applications should use this size for the device queue to ensure good performance.  The minimum value is 16 KB
<b>CL_DEVICE_QUEUE_ON_DEVICE_MAX_SIZE</b>	cl_uint	The max. size of the device queue in bytes.  The minimum value is 256 KB for the full profile and 64 KB for the embedded profile
<b>CL_DEVICE_MAX_ON_DEVICE_QUEUES</b>	cl_uint	The maximum number of device queues that can be created per context.  The minimum value is 1.
<b>CL_DEVICE_MAX_ON_DEVICE_EVENTS</b>	cl_uint	The maximum number of events in use by a device queue. These refer to events returned by the enqueue_ built-in functions to a device queue or user events returned by the create_user_event built-in function that have not been released.  The minimum value is 1024.
<b>CL_DEVICE_BUILT_IN_KERNELS</b>	char[]	A semi-colon separated list of built-in kernels supported by the device. An empty string is returned if no built-in kernels are supported by the device.
<b>CL_DEVICE_PLATFORM</b>	cl_platform_id	The platform associated with this device.
<b>CL_DEVICE_NAME</b>	char[]	Device name string.
<b>CL_DEVICE_VENDOR</b>	char[]	Vendor name string.
<b>CL_DRIVER_VERSION</b>	char[]	OpenCL software driver version string. Follows a vendor-specific format.
<b>CL_DEVICE_PROFILE<sup>9</sup></b>	char[]	OpenCL profile string. Returns the profile name supported by the device. The profile name returned can be one of the following strings:  FULL_PROFILE – if the device supports the OpenCL specification (functionality

<sup>9</sup> The platform profile returns the profile that is implemented by the OpenCL framework. If the platform profile returned is FULL\_PROFILE, the OpenCL framework will support devices that are FULL\_PROFILE and may also support devices that are EMBEDDED\_PROFILE. The compiler must be available for all devices i.e. CL\_DEVICE\_COMPILER\_AVAILABLE is CL\_TRUE. If the platform profile returned is EMBEDDED\_PROFILE, then devices that are only EMBEDDED\_PROFILE are supported.



		<p>defined as part of the core specification and does not require any extensions to be supported).</p> <p>EMBEDDED_PROFILE - if the device supports the OpenCL embedded profile.</p>
CL_DEVICE_VERSION	char[]	<p>OpenCL version string. Returns the OpenCL version supported by the device. This version string has the following format:</p> <p><i>OpenCL&lt;space&gt;&lt;major_version.minor_version&gt;&lt;space&gt;&lt;vendor-specific information&gt;</i></p> <p>The <i>major_version.minor_version</i> value returned will be 2.0.</p>
CL_DEVICE_OPENCL_C_VERSION	char[]	<p>OpenCL C version string. Returns the highest OpenCL C version supported by the compiler for this device that is not of type CL_DEVICE_TYPE_CUSTOM. This version string has the following format:</p> <p><i>OpenCL&lt;space&gt;C&lt;space&gt;&lt;major_version.minor_version&gt;&lt;space&gt;&lt;vendor-specific information&gt;</i></p> <p>The <i>major_version.minor_version</i> value returned must be 2.0 if CL_DEVICE_VERSION is OpenCL 2.0.</p> <p>The <i>major_version.minor_version</i> value returned must be 1.2 if CL_DEVICE_VERSION is OpenCL 1.2.</p> <p>The <i>major_version.minor_version</i> value returned must be 1.1 if CL_DEVICE_VERSION is OpenCL 1.1.</p> <p>The <i>major_version.minor_version</i> value returned can be 1.0 or 1.1 if CL_DEVICE_VERSION is OpenCL 1.0.</p>
CL_DEVICE_EXTENSIONS	char[]	<p>Returns a space separated list of extension names (the extension names themselves do not contain any spaces) supported by the device. The list of extension names</p>

		<p>returned can be vendor supported extension names and one or more of the following Khronos approved extension names:</p> <p><b>cl_khr_int64_base_atomics</b>  <b>cl_khr_int64_extended_atomics</b>  <b>cl_khr_fp16</b>  <b>cl_khr_gl_sharing</b>  <b>cl_khr_gl_event</b>  <b>cl_khr_d3d10_sharing</b>  <b>cl_khr_dx9_media_sharing</b>  <b>cl_khr_d3d11_sharing</b>  <b>cl_khr_gl_depth_images</b>  <b>cl_khr_gl_msaa_sharing</b>  <b>cl_khr_initialize_memory</b>  <b>cl_khr_context_abort</b>  <b>cl_khr_spir</b>  <b>cl_khr_srgb_image_writes</b></p> <p><b>The following approved Khronos extension names must be returned by all devices that support OpenCL C 2.0:</b></p> <p><b>cl_khr_byte_addressable_store</b>  <b>cl_khr_fp64</b> (for backward compatibility if double precision is supported)  <b>cl_khr_3d_image_writes</b>  <b>cl_khr_image2d_from_buffer</b>  <b>cl_khr_depth_images</b></p> <p>Please refer to the OpenCL 2.0 Extension Specification for a detailed description of these extensions.</p>
<b>CL_DEVICE_PRINTF_BUFFER_SIZE</b>	size_t	Maximum size in bytes of the internal buffer that holds the output of printf calls from a kernel. The minimum value for the FULL profile is 1 MB.
<b>CL_DEVICE_PREFERRED_INTEROP_USER_SYNC</b>	cl_bool	Is CL_TRUE if the device's preference is for the user to be responsible for synchronization, when sharing memory objects between OpenCL and other APIs such as DirectX, CL_FALSE if the device / implementation has a performant path for performing synchronization of memory object shared between OpenCL and other APIs such as DirectX.

<b>CL_DEVICE_PARENT_DEVICE</b>	cl_device_id	Returns the cl_device_id of the parent device to which this sub-device belongs. If <i>device</i> is a root-level device, a NULL value is returned.
<b>CL_DEVICE_PARTITION_MAX_SUB_DEVICES</b>	cl_uint	Returns the maximum number of sub-devices that can be created when a device is partitioned.  The value returned cannot exceed CL_DEVICE_MAX_COMPUTE_UNITS.
<b>CL_DEVICE_PARTITION_PROPERTIES</b>	cl_device_partition_property[]	Returns the list of partition types supported by <i>device</i> . The is an array of cl_device_partition_property values drawn from the following list:  CL_DEVICE_PARTITION_EQUALLY CL_DEVICE_PARTITION_BY_COUNTS CL_DEVICE_PARTITION_BY_AFFINITY_DOMAIN  If the device cannot be partitioned (i.e. there is no partitioning scheme supported by the device that will return at least two subdevices), a value of 0 will be returned.
<b>CL_DEVICE_PARTITION_AFFINITY_DOMAIN</b>	cl_device_affinity_domain	Returns the list of supported affinity domains for partitioning the <i>device</i> using CL_DEVICE_PARTITION_BY_AFFINITY_DOMAIN. This is a bit-field that describes one or more of the following values:  CL_DEVICE_AFFINITY_DOMAIN_NUMA CL_DEVICE_AFFINITY_DOMAIN_L4_CACHE CL_DEVICE_AFFINITY_DOMAIN_L3_CACHE CL_DEVICE_AFFINITY_DOMAIN_L2_CACHE CL_DEVICE_AFFINITY_DOMAIN_L1_CACHE CL_DEVICE_AFFINITY_DOMAIN_NEXT_PARTITIONABLE  If the device does not support any affinity domains, a value of 0 will be returned.
<b>CL_DEVICE_PARTITION_TYPE</b>	cl_device_partition_property[]	Returns the <i>properties</i> argument specified in <b>clCreateSubDevices</b> if <i>device</i> is a sub-device. In the case where the <i>properties</i> argument to <b>clCreateSubDevices</b> is CL_DEVICE_PARTITION_BY_AFFINITY_DOMAIN, CL_DEVICE_AFFINITY_DOMAIN_NEXT_PARTITIONABLE, the affinity domain used to perform the partition will be returned. This can be one of the following values:

		<p>CL_DEVICE_AFFINITY_DOMAIN_NUMA  CL_DEVICE_AFFINITY_DOMAIN_L4_CACHE  CL_DEVICE_AFFINITY_DOMAIN_L3_CACHE  CL_DEVICE_AFFINITY_DOMAIN_L2_CACHE  CL_DEVICE_AFFINITY_DOMAIN_L1_CACHE</p> <p>Otherwise the implementation may either return a <i>param_value_size_ret</i> of 0 i.e. there is no partition type associated with device or can return a property value of 0 (where 0 is used to terminate the partition property list) in the memory that <i>param_value</i> points to.</p>
CL_DEVICE_REFERENCE_COUNT	cl_uint	Returns the <i>device</i> reference count. If the device is a root-level device, a reference count of one is returned.
CL_DEVICE_SVM_CAPABILITIES	cl_device_svm_capabilities	<p>Describes the various shared virtual memory (a.k.a. SVM) memory allocation types the device supports. Coarse-grain SVM allocations are required to be supported by all OpenCL 2.0 devices. This is a bit-field that describes a combination of the following values:</p> <p>CL_DEVICE_SVM_COARSE_GRAIN_BUFFER – Support for coarse-grain buffer sharing using <b>clSVMAlloc</b>. Memory consistency is guaranteed at synchronization points and the host must use calls to <b>clEnqueueMapBuffer</b> and <b>clEnqueueUnmapMemObject</b>.</p> <p>CL_DEVICE_SVM_FINE_GRAIN_BUFFER – Support for fine-grain buffer sharing using <b>clSVMAlloc</b>. Memory consistency is guaranteed at synchronization points without need for <b>clEnqueueMapBuffer</b> and <b>clEnqueueUnmapMemObject</b>.</p> <p>CL_DEVICE_SVM_FINE_GRAIN_SYSTEM – Support for sharing the host’s entire virtual memory including memory allocated using <b>malloc</b>. Memory consistency is guaranteed at synchronization points.</p>

		<p>CL_DEVICE_SVM_ATOMICS – Support for the OpenCL 2.0 atomic operations that provide memory consistency across the host and all OpenCL devices supporting fine-grain SVM allocations.</p> <p>The mandated minimum capability is CL_DEVICE_SVM_COARSE_GRAIN_BUFFER.</p>
CL_DEVICE_PREFERRED_PLATFORM_ATOMIC_ALIGNMENT	cl_uint	Returns the value representing the preferred alignment in bytes for OpenCL 2.0 fine-grained SVM atomic types. This query can return 0 which indicates that the preferred alignment is aligned to the natural size of the type.
CL_DEVICE_PREFERRED_GLOBAL_ATOMIC_ALIGNMENT	cl_uint	Returns the value representing the preferred alignment in bytes for OpenCL 2.0 atomic types to global memory. This query can return 0 which indicates that the preferred alignment is aligned to the natural size of the type.
CL_DEVICE_PREFERRED_LOCAL_ATOMIC_ALIGNMENT	cl_uint	Returns the value representing the preferred alignment in bytes for OpenCL 2.0 atomic types to local memory. This query can return 0 which indicates that the preferred alignment is aligned to the natural size of the type.

**Table 4.3.**     *OpenCL Device Queries*

The device queries described in *table 4.3* should return the same information for a root-level device i.e. a device returned by **clGetDeviceIDs** and any sub-devices created from this device except for the following queries:

CL\_DEVICE\_GLOBAL\_MEM\_CACHE\_SIZE  
CL\_DEVICE\_BUILT\_IN\_KERNELS  
CL\_DEVICE\_PARENT\_DEVICE  
CL\_DEVICE\_PARTITION\_TYPE  
CL\_DEVICE\_REFERENCE\_COUNT

**clGetDeviceInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

 CL\_INVALID\_DEVICE if *device* is not valid.

- ✚ CL\_INVALID\_VALUE if *param\_name* is not one of the supported values or if size in bytes specified by *param\_value\_size* is < size of return type as specified in *table 4.3* and *param\_value* is not a NULL value or if *param\_name* is a value that is available as an extension and the corresponding extension is not supported by the device.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 4.3 Partitioning a Device

The function

```
cl_int  clCreateSubDevices (cl_device_id in_device,
                           const cl_device_partition_property *properties,
                           cl_uint num_devices,
                           cl_device_id *out_devices,
                           cl_uint *num_devices_ret)
```

creates an array of sub-devices that each reference a non-intersecting set of compute units within *in\_device*, according to a partition scheme given by *properties*. The output sub-devices may be used in every way that the root (or parent) device can be used, including creating contexts, building programs, further calls to **clCreateSubDevices** and creating command-queues. When a command-queue is created against a sub-device, the commands enqueued on the queue are executed only on the sub-device.

*in\_device* is the device to be partitioned.

*properties* specifies how *in\_device* is to be partition described by a partition name and its corresponding value. Each partition name is immediately followed by the corresponding desired value. The list is terminated with 0. The list of supported partitioning schemes is described in *table 4.4*. Only one of the listed partitioning schemes can be specified in *properties*.

cl_device_partition_property enum	Partition value	Description
CL_DEVICE_PARTITION_EQUALLY	unsigned int	Split the aggregate device into as many smaller aggregate devices as can be created, each containing <i>n</i> compute units. The value <i>n</i> is passed as the value accompanying this property. If <i>n</i> does not divide evenly into CL_DEVICE_PARTITION_MAX_COMPUTE_UNITS, then the remaining compute units are not used.
CL_DEVICE_PARTITION_BY_COUNTS	unsigned int	<p>This property is followed by a CL_DEVICE_PARTITION_BY_COUNTS_LIST_END terminated list of compute unit counts. For each non-zero count <i>m</i> in the list, a sub-device is created with <i>m</i> compute units in it. CL_DEVICE_PARTITION_BY_COUNTS_LIST_END is defined to be 0.</p> <p>The number of non-zero count entries in the list may not exceed CL_DEVICE_PARTITION_MAX_SUB_DEVICES.</p>

		<p>The total number of compute units specified may not exceed CL_DEVICE_PARTITION_MAX_COMPUTE_UNITS.</p>
<b>CL_DEVICE_PARTITION_BY_AFFINITY_DOMAIN</b>	cl_device_affinity_domain	<p>Split the device into smaller aggregate devices containing one or more compute units that all share part of a cache hierarchy. The value accompanying this property may be drawn from the following list:</p> <p>CL_DEVICE_AFFINITY_DOMAIN_NUMA – Split the device into sub-devices comprised of compute units that share a NUMA node.</p> <p>CL_DEVICE_AFFINITY_DOMAIN_L4_CACHE – Split the device into sub-devices comprised of compute units that share a level 4 data cache.</p> <p>CL_DEVICE_AFFINITY_DOMAIN_L3_CACHE – Split the device into sub-devices comprised of compute units that share a level 3 data cache.</p> <p>CL_DEVICE_AFFINITY_DOMAIN_L2_CACHE – Split the device into sub-devices comprised of compute units that share a level 2 data cache.</p> <p>CL_DEVICE_AFFINITY_DOMAIN_L1_CACHE – Split the device into sub-devices comprised of compute units that share a level 1 data cache.</p> <p>CL_DEVICE_AFFINITY_DOMAIN_NEXT_PARTITIONABLE – Split the device along the next partitionable affinity domain. The implementation shall find the first level along which the device or sub-device may be further subdivided in the order NUMA, L4, L3, L2, L1, and partition the device into sub-devices comprised of compute units that share memory subsystems at this level.</p> <p>The user may determine what happened by calling <b>clGetDeviceInfo(CL_DEVICE_PARTITION_TYPE)</b> on the sub-devices.</p>

**Table 4.4**      *List of supported partition schemes by **clCreateSubDevices***



*num\_devices* is the size of memory pointed to by *out\_devices* specified as the number of *cl\_device\_id* entries.

*out\_devices* is the buffer where the OpenCL sub-devices will be returned. If *out\_devices* is NULL, this argument is ignored. If *out\_devices* is not NULL, *num\_devices* must be greater than or equal to the number of sub-devices that *device* may be partitioned into according to the partitioning scheme specified in *properties*.

*num\_devices\_ret* returns the number of sub-devices that *device* may be partitioned into according to the partitioning scheme specified in *properties*. If *num\_devices\_ret* is NULL, it is ignored.

**clCreateSubDevices** returns CL\_SUCCESS if the partition is created successfully. Otherwise, it returns a NULL value with the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_DEVICE if *in\_device* is not valid.
- ✚ CL\_INVALID\_VALUE if values specified in *properties* are not valid or if values specified in *properties* are valid but not supported by the device.
- ✚ CL\_INVALID\_VALUE if *out\_devices* is not NULL and *num\_devices* is less than the number of sub-devices created by the partition scheme.
- ✚ CL\_DEVICE\_PARTITION\_FAILED if the partition name is supported by the implementation but *in\_device* could not be further partitioned.
- ✚ CL\_INVALID\_DEVICE\_PARTITION\_COUNT if the partition name specified in *properties* is CL\_DEVICE\_PARTITION\_BY\_COUNTS and the number of sub-devices requested exceeds CL\_DEVICE\_PARTITION\_MAX\_SUB\_DEVICES or the total number of compute units requested exceeds CL\_DEVICE\_PARTITION\_MAX\_COMPUTE\_UNITS for *in\_device*, or the number of compute units requested for one or more sub-devices is less than zero or the number of sub-devices requested exceeds CL\_DEVICE\_PARTITION\_MAX\_COMPUTE\_UNITS for *in\_device*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

A few examples that describe how to specify partition properties in *properties* argument to **clCreateSubDevices** are given below:

To partition a device containing 16 compute units into two sub-devices, each containing 8 compute units, pass the following in *properties*:

```
{ CL_DEVICE_PARTITION_EQUALLY, 8, 0 }
```

To partition a device with four compute units into two sub-devices with one sub-device containing 3 compute units and the other sub-device 1 compute unit, pass the following in properties argument:

```
{ CL_DEVICE_PARTITION_BY_COUNTS,  
  3, 1, CL_DEVICE_PARTITION_BY_COUNTS_LIST_END, 0 }
```




To split a device along the outermost cache line (if any), pass the following in properties argument:

```
{ CL_DEVICE_PARTITION_BY_AFFINITY_DOMAIN,  
  CL_DEVICE_AFFINITY_DOMAIN_NEXT_PARTITIONABLE,  
  0 }
```

The function

```
cl_int      clRetainDevice (cl_device_id device)
```

increments the *device* reference count if *device* is a valid sub-device created by a call to **clCreateSubDevices**. If *device* is a root level device i.e. a *cl\_device\_id* returned by **clGetDeviceIDs**, the *device* reference count remains unchanged. **clRetainDevice** returns CL\_SUCCESS if the function is executed successfully or the device is a root-level device. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_DEVICE if *device* is not a valid sub-device created by a call to **clCreateSubDevices**.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int      clReleaseDevice (cl_device_id device)
```

decrements the *device* reference count if *device* is a valid sub-device created by a call to **clCreateSubDevices**. If *device* is a root level device i.e. a *cl\_device\_id* returned by **clGetDeviceIDs**, the *device* reference count remains unchanged. **clReleaseDevice** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_DEVICE if *device* is not a valid sub-device created by a call to **clCreateSubDevices**.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

After the *device* reference count becomes zero and all the objects attached to *device* (such as command-queues) are released, the *device* object is deleted. Using this function to release a reference that was not obtained by creating the object or by calling **clRetainDevice** causes undefined behavior.

## 4.4 Contexts

The function

```
cl_context    clCreateContext (const cl_context_properties *properties,
                                cl_uint num_devices,
                                const cl_device_id *devices,
                                void (CL_CALLBACK *pfn_notify)(const char *errinfo,
                                                                const void *private_info, size_t cb,
                                                                void *user_data),
                                void *user_data,
                                cl_int *errcode_ret)
```

creates an OpenCL context. An OpenCL context is created with one or more devices. Contexts are used by the OpenCL runtime for managing objects such as command-queues, memory, program and kernel objects and for executing kernels on one or more devices specified in the context.

*properties* specifies a list of context property names and their corresponding values. Each property name is immediately followed by the corresponding desired value. The list is terminated with 0. The list of supported properties is described in table 4.5. *properties* can be NULL in which case the platform that is selected is implementation-defined.

cl_context_properties enum	Property value	Description
CL_CONTEXT_PLATFORM	cl_platform_id	Specifies the platform to use.
CL_CONTEXT_INTEROP_USER_SYNC	cl_bool	Specifies whether the user is responsible for synchronization between OpenCL and other APIs. Please refer to the specific sections in the OpenCL 2.0 extension specification that describe sharing with other APIs for restrictions on using this flag.  If CL_CONTEXT_INTEROP_USER_SYNC is not specified, a default of CL_FALSE is assumed.

**Table 4.5** List of supported properties by **clCreateContext**

*num\_devices* is the number of devices specified in the *devices* argument.

*devices* is a pointer to a list of unique devices<sup>10</sup> returned by **clGetDeviceIDs** or sub-devices created by **clCreateSubDevices** for a platform.

*pfn\_notify* is a callback function that can be registered by the application. This callback function will be used by the OpenCL implementation to report information on errors during context creation as well as errors that occur at runtime in this context. This callback function may be called asynchronously by the OpenCL implementation. **It is the application's responsibility to ensure that the callback function is thread-safe.** The parameters to this callback function are:

- + *errinfo* is a pointer to an error string.
- + *private\_info* and *cb* represent a pointer to binary data that is returned by the OpenCL implementation that can be used to log additional information helpful in debugging the error.
- + *user\_data* is a pointer to user supplied data.

If *pfn\_notify* is NULL, no callback function is registered.

NOTE: There are a number of cases where error notifications need to be delivered due to an error that occurs outside a context. Such notifications may not be delivered through the *pfn\_notify* callback. Where these notifications go is implementation-defined.

*user\_data* will be passed as the *user\_data* argument when *pfn\_notify* is called. *user\_data* can be NULL.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreateContext** returns a valid non-zero context and *errcode\_ret* is set to CL\_SUCCESS if the context is created successfully. Otherwise, it returns a NULL value with the following error values returned in *errcode\_ret*:

- + CL\_INVALID\_PLATFORM if *properties* is NULL and no platform could be selected or if platform value specified in *properties* is not a valid platform.
- + CL\_INVALID\_PROPERTY if context property name in *properties* is not a supported property name, if the value specified for a supported property name is not valid, or if the same property name is specified more than once.
- + CL\_INVALID\_VALUE if *devices* is NULL.
- + CL\_INVALID\_VALUE if *num\_devices* is equal to zero.
- + CL\_INVALID\_VALUE if *pfn\_notify* is NULL but *user\_data* is not NULL.
- + CL\_INVALID\_DEVICE if *devices* contains an invalid device.

---

<sup>10</sup> Duplicate devices specified in *devices* are ignored.

- ✚ CL\_DEVICE\_NOT\_AVAILABLE if a device in *devices* is currently not available even though the device was returned by **clGetDeviceIDs**.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_context
clCreateContextFromType11 (const cl_context_properties *properties,
                             cl_device_type device_type,
                             void (CL_CALLBACK *pfn_notify)(const char *errinfo,
                                                             const void *private_info, size_t cb,
                                                             void *user_data),
                             void *user_data,
                             cl_int *errcode_ret)
```

creates an OpenCL context from a device type that identifies the specific device(s) to use. Only devices that are returned by **clGetDeviceIDs** for *device\_type* are used to create the context. The context does not reference any sub-devices that may have been created from these devices.

*properties* specifies a list of context property names and their corresponding values. Each property name is immediately followed by the corresponding desired value. The list of supported properties is described in *table 4.5*. *properties* can also be NULL in which case the platform that is selected is implementation-defined.

*device\_type* is a bit-field that identifies the type of device and is described in *table 4.2* in *section 4.2*.

*pfn\_notify* and *user\_data* are described in **clCreateContext**.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreateContextFromType** returns a valid non-zero context and *errcode\_ret* is set to CL\_SUCCESS if the context is created successfully. Otherwise, it returns a NULL value with the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_PLATFORM if *properties* is NULL and no platform could be selected or if

---

<sup>11</sup> **clCreateContextFromType** may return all or a subset of the actual physical devices present in the platform and that match *device\_type*.

platform value specified in *properties* is not a valid platform.

- ✚ CL\_INVALID\_PROPERTY if context property name in *properties* is not a supported property name, if the value specified for a supported property name is not valid, or if the same property name is specified more than once.
- ✚ CL\_INVALID\_VALUE if *pfn\_notify* is NULL but *user\_data* is not NULL.
- ✚ CL\_INVALID\_DEVICE\_TYPE if *device\_type* is not a valid value.
- ✚ CL\_DEVICE\_NOT\_AVAILABLE if no devices that match *device\_type* and property values specified in *properties* are currently available.
- ✚ CL\_DEVICE\_NOT\_FOUND if no devices that match *device\_type* and property values specified in *properties* were found.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

cl\_int            **clRetainContext** (cl\_context *context*)

increments the *context* reference count. **clRetainContext** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:




- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid OpenCL context.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

**clCreateContext** and **clCreateContextFromType** perform an implicit retain. This is very helpful for 3<sup>rd</sup> party libraries, which typically get a context passed to them by the application. However, it is possible that the application may delete the context without informing the library. Allowing functions to attach to (i.e. retain) and release a context solves the problem of a context being used by a library no longer being valid.

The function

cl\_int            **clReleaseContext** (cl\_context *context*)

decrements the *context* reference count. **clReleaseContext** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_CONTEXT if *context* is not a valid OpenCL context.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

After the *context* reference count becomes zero and all the objects attached to *context* (such as memory objects, command-queues) are released, the *context* is deleted. Using this function to release a reference that was not obtained by creating the object or by calling **clRetainContext** causes undefined behavior.

The function

```
cl_int          clGetContextInfo (cl_context context,
                                     cl_context_info param_name,
                                     size_t param_value_size,
                                     void *param_value,
                                     size_t *param_value_size_ret)
```

can be used to query information about a context.

*context* specifies the OpenCL context being queried.

*param\_name* is an enumeration constant that specifies the information to query.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* specifies the size in bytes of memory pointed to by *param\_value*. This size must be greater than or equal to the size of return type as described in *table 4.6*.

*param\_value\_size\_ret* returns the actual size in bytes of data being queried by *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.


The list of supported *param\_name* values and the information returned in *param\_value* by **clGetContextInfo** is described in *table 4.6*.



<b>cl_context_info</b>	<b>Return Type</b>	<b>Information returned in param_value</b>
<b>CL_CONTEXT_REFERENCE_COUNT</b> <sup>12</sup>	cl_uint	Return the <i>context</i> reference count.
<b>CL_CONTEXT_NUM_DEVICES</b>	cl_uint	Return the number of devices in <i>context</i> .
<b>CL_CONTEXT_DEVICES</b>	cl_device_id[]	Return the list of devices and sub-devices in <i>context</i> .
<b>CL_CONTEXT_PROPERTIES</b>	cl_context_properties[]	<p>Return the <i>properties</i> argument specified in <b>clCreateContext</b> or <b>clCreateContextFromType</b>.</p> <p>If the <i>properties</i> argument specified in <b>clCreateContext</b> or <b>clCreateContextFromType</b> used to create <i>context</i> is not NULL, the implementation must return the values specified in the <i>properties</i> argument.</p> <p>If the <i>properties</i> argument specified in <b>clCreateContext</b> or <b>clCreateContextFromType</b> used to create <i>context</i> is NULL, the implementation may return either a <i>param_value_size_ret</i> of 0 i.e. there is no context property value to be returned or can return a context property value of 0 (where 0 is used to terminate the context properties list) in the memory that <i>param_value</i> points to.</p>

**Table 4.6** List of supported *param\_names* by **clGetContextInfo**

**clGetContextInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

 CL\_INVALID\_CONTEXT if *context* is not a valid context.

<sup>12</sup> The reference count returned should be considered immediately stale. It is unsuitable for general use in applications. This feature is provided for identifying memory leaks.

- ✚ CL\_INVALID\_VALUE if *param\_name* is not one of the supported values or if size in bytes specified by *param\_value\_size* is < size of return type as specified in *table 4.6* and *param\_value* is not a NULL value.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 5. The OpenCL Runtime

In this section we describe the API calls that manage OpenCL objects such as command-queues, memory objects, program objects, kernel objects for `__kernel` functions in a program and calls that allow you to enqueue commands to a command-queue such as executing a kernel, reading, or writing a memory object.

### 5.1 Command Queues

OpenCL objects such as memory, program and kernel objects are created using a context. Operations on these objects are performed using a command-queue. The command-queue can be used to queue a set of operations (referred to as commands) in order. Having multiple command-queues allows applications to queue multiple independent commands without requiring synchronization. Note that this should work as long as these objects are not being shared. **Sharing of objects across multiple command-queues will require the application to perform appropriate synchronization.** This is described in *Appendix A*.

The function

```
cl_command_queue  clCreateCommandQueueWithProperties (cl_context context,
                                                    cl_device_id device,
                                                    const cl_queue_properties *properties,
                                                    cl_int *errcode_ret)
```

creates a host or device command-queue on a specific device.

*context* must be a valid OpenCL context.

Queue Properties	Property Value	Description
CL_QUEUE_PROPERTIES	cl_command_queue_properties	<p>This is a bitfield and can be set to a combination of the following values:</p> <p>CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE – Determines whether the commands queued in the command-queue are executed in-order or out-of-order. If set, the commands in the command-queue are executed out-of-order. Otherwise, commands are executed in-order.</p> <p>CL_QUEUE_PROFILING_ENABLE – Enable or disable profiling of commands in</p>

		<p>the command-queue. If set, the profiling of commands is enabled. Otherwise profiling of commands is disabled.</p> <p>CL_QUEUE_ON_DEVICE – Indicates that this is a device queue. If CL_QUEUE_ON_DEVICE is set, CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE<sup>13</sup> must also be set.</p> <p>CL_QUEUE_ON_DEVICE_DEFAULT<sup>14</sup> – indicates that this is the default device queue. This can only be used with CL_QUEUE_ON_DEVICE.</p> <p>If CL_QUEUE_PROPERTIES is not specified an in-order host command queue is created for the specified device.</p>
CL_QUEUE_SIZE	cl_uint	<p>Specifies the size of the device queue in bytes.</p> <p>This can only be specified if CL_QUEUE_ON_DEVICE is set in CL_QUEUE_PROPERTIES. This must be a value &lt;= CL_DEVICE_QUEUE_ON_DEVICE_MAX_SIZE.</p> <p>For best performance, this should be &lt;= CL_DEVICE_QUEUE_ON_DEVICE_PREFERRED_SIZE.</p> <p>If CL_QUEUE_SIZE is not specified, the device queue is created with CL_DEVICE_QUEUE_ON_DEVICE_PREFERRED_SIZE as the size of the queue.</p>

**Table 5.1**      *List of supported cl\_queue\_properties values and description.*

*device* must be a device or sub-device associated with *context*. It can either be in the list of devices and sub-devices specified when *context* is created using **clCreateContext** or be a root

<sup>13</sup> Only out-of-order device queues are supported.







<sup>14</sup> The application **must** create the default device queue if any kernels containing calls to `get_default_queue` are enqueued. There can only be one default device queue for each device within a context. `clCreateCommandQueueWithProperties` with CL\_QUEUE\_PROPERTIES set to CL\_QUEUE\_ON\_DEVICE | CL\_QUEUE\_ON\_DEVICE\_DEFAULT will return the default device queue that has already been created and increment its retain count by 1.

device with the same device type as specified when *context* is created using **clCreateContextFromType**.

*properties* specifies a list of properties for the command-queue and their corresponding values. Each property name is immediately followed by the corresponding desired value. The list is terminated with 0. The list of supported properties is described in the table below. If a supported property and its value is not specified in *properties*, its default value will be used. *properties* can be NULL in which case the default values for supported command-queue properties will be used.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.



**clCreateCommandQueueWithProperties** returns a valid non-zero command-queue and *errcode\_ret* is set to CL\_SUCCESS if the command-queue is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

-  CL\_INVALID\_CONTEXT if *context* is not a valid context.
-  CL\_INVALID\_DEVICE if *device* is not a valid device or is not associated with *context*.
-  CL\_INVALID\_VALUE if values specified in *properties* are not valid.
-  CL\_INVALID\_QUEUE\_PROPERTIES if values specified in *properties* are valid but are not supported by the device.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

cl\_int                    **clRetainCommandQueue** (cl\_command\_queue *command\_queue*)

increments the *command\_queue* reference count. **clRetainCommandQueue** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid command-queue.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.

- CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

**clCreateCommandQueueWithProperties** performs an implicit retain. This is very helpful for 3<sup>rd</sup> party libraries, which typically get a command-queue passed to them by the application. However, it is possible that the application may delete the command-queue without informing the library. Allowing functions to attach to (i.e. retain) and release a command-queue solves the problem of a command-queue being used by a library no longer being valid.

The function

cl_int	<b>clReleaseCommandQueue</b> (cl_command_queue <i>command_queue</i> )
--------	---

decrements the *command\_queue* reference count. **clReleaseCommandQueue** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- + CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid command-queue.
- + CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- + CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

- 

- 

After the `command_queue` reference count becomes zero and all commands queued to `command_queue` have finished (eg. kernel-instances, memory object updates etc.), the command-queue is deleted.

**clReleaseCommandQueue** performs an implicit flush to issue any previously queued OpenCL commands in *command\_queue*. Using this function to release a reference that was not obtained by creating the object or by calling **clRetainCommandQueue** causes undefined behavior.

The function

[illegible]

can be used to query information about a command-queue.

*command\_queue* specifies the command-queue being queried.

*param\_name* specifies the information to query.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.2*. If *param\_value* is NULL, it is ignored.


*param\_value\_size\_ret* returns the actual size in bytes of data being queried by *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

The list of supported *param\_name* values and the information returned in *param\_value* by **clGetCommandQueueInfo** is described in *table 5.2*.

<b>cl_command_queue_info</b>	<b>Return Type</b>	<b>Information returned in <i>param_value</i></b>
<b>CL_QUEUE_CONTEXT</b>	cl_context	Return the context specified when the command-queue is created.
<b>CL_QUEUE_DEVICE</b>	cl_device_id	Return the device specified when the command-queue is created.
<b>CL_QUEUE_REFERENCE_COUNT<sup>15</sup></b>	cl_uint	Return the command-queue reference count.
<b>CL_QUEUE_PROPERTIES</b>	cl_command_queue_properties	Return the currently specified properties for the command-queue. These properties are specified by the value associated with the CL_COMMAND_QUEUE_PROPERTIES passed in <i>properties</i> argument in <b>clCreateCommandQueueWithProperties</b> .
<b>CL_QUEUE_SIZE</b>	cl_uint	Return the currently specified size for the device command-queue. This query is only supported for device command queues.

**Table 5.2**      *List of supported param\_names by clGetCommandQueueInfo*

**clGetCommandQueueInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

 **CL\_INVALID\_COMMAND\_QUEUE** if *command\_queue* is not a valid command-queue.

<sup>15</sup> The reference count returned should be considered immediately stale. It is unsuitable for general use in applications. This feature is provided for identifying memory leaks.

- ✚ CL\_INVALID\_VALUE if *param\_name* is not one of the supported values or if size in bytes specified by *param\_value\_size* is < size of return type as specified in *table 5.2* and *param\_value* is not a NULL value.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## NOTE

It is possible that a device(s) becomes unavailable after a context and command-queues that use this device(s) have been created and commands have been queued to command-queues. In this case the behavior of OpenCL API calls that use this context (and command-queues) are considered to be implementation-defined. The user callback function, if specified, when the context is created can be used to record appropriate information in the *errinfo*, *private\_info* arguments passed to the callback function when the device becomes unavailable.



## 5.2 Buffer Objects

A *buffer* object stores a one-dimensional collection of elements. Elements of a *buffer* object can be a scalar data type (such as an int, float), vector data type, or a user-defined structure.

### 5.2.1 Creating Buffer Objects

A **buffer object** is created using the following function

```
cl_mem  clCreateBuffer (cl_context context,
                        cl_mem_flags flags,
                        size_t size,
                        void *host_ptr,
                        cl_int *errcode_ret)
```

*context* is a valid OpenCL context used to create the buffer object.

*flags* is a bit-field that is used to specify allocation and usage information such as the memory arena that should be used to allocate the buffer object and how it will be used. *Table 5.3* describes the possible values for *flags*. If value specified for *flags* is 0, the default is used which is CL\_MEM\_READ\_WRITE.

cl_mem_flags	Description
CL_MEM_READ_WRITE	This flag specifies that the memory object will be read and written by a kernel. This is the default.
CL_MEM_WRITE_ONLY	This flag specifies that the memory object will be written but not read by a kernel.  Reading from a buffer or image object created with CL_MEM_WRITE_ONLY inside a kernel is undefined.  CL_MEM_READ_WRITE and CL_MEM_WRITE_ONLY are mutually exclusive.
CL_MEM_READ_ONLY	This flag specifies that the memory object is a read-only memory object when used inside a kernel.  Writing to a buffer or image object created with CL_MEM_READ_ONLY inside a kernel is undefined.  CL_MEM_READ_WRITE or CL_MEM_WRITE_ONLY and CL_MEM_READ_ONLY are mutually exclusive.
CL_MEM_USE_HOST_PTR	This flag is valid only if <i>host_ptr</i> is not NULL. If

	<p>specified, it indicates that the application wants the OpenCL implementation to use memory referenced by <i>host_ptr</i> as the storage bits for the memory object.</p> <p>OpenCL implementations are allowed to cache the buffer contents pointed to by <i>host_ptr</i> in device memory. This cached copy can be used when kernels are executed on a device.</p> <p>The result of OpenCL commands that operate on multiple buffer objects created with the same <i>host_ptr</i> or overlapping host regions is considered to be undefined.</p> <p>Also refer to <i>section C.3</i> for a description of the alignment rules for <i>host_ptr</i> for memory objects (buffer and images) created using <code>CL_MEM_USE_HOST_PTR</code>.</p>
<b>CL_MEM_ALLOC_HOST_PTR</b>	<p>This flag specifies that the application wants the OpenCL implementation to allocate memory from host accessible memory.</p> <p><code>CL_MEM_ALLOC_HOST_PTR</code> and <code>CL_MEM_USE_HOST_PTR</code> are mutually exclusive.</p>
<b>CL_MEM_COPY_HOST_PTR</b>	<p>This flag is valid only if <i>host_ptr</i> is not NULL. If specified, it indicates that the application wants the OpenCL implementation to allocate memory for the memory object and copy the data from memory referenced by <i>host_ptr</i>. The implementation will copy the memory immediately and <i>host_ptr</i> is available for reuse by the application when the <i>clCreateBuffer</i> or <i>clCreateImage</i> operation returns.</p> <p><code>CL_MEM_COPY_HOST_PTR</code> and <code>CL_MEM_USE_HOST_PTR</code> are mutually exclusive.</p> <p><code>CL_MEM_COPY_HOST_PTR</code> can be used with <code>CL_MEM_ALLOC_HOST_PTR</code> to initialize the contents of the <code>cl_mem</code> object allocated using host-accessible (e.g. PCIe) memory.</p>
<b>CL_MEM_HOST_WRITE_ONLY</b>	<p>This flag specifies that the host will only write to the memory object (using OpenCL APIs that enqueue a write or a map for write). This can be used to optimize write access from the host (e.g. enable write-</p>

	combined allocations for memory objects for devices that communicate with the host over a system bus such as PCIe).
<b>CL_MEM_HOST_READ_ONLY</b>	This flag specifies that the host will only read the memory object (using OpenCL APIs that enqueue a read or a map for read).  CL_MEM_HOST_WRITE_ONLY and CL_MEM_HOST_READ_ONLY are mutually exclusive.
<b>CL_MEM_HOST_NO_ACCESS</b>	This flag specifies that the host will not read or write the memory object.  CL_MEM_HOST_WRITE_ONLY or CL_MEM_HOST_READ_ONLY and CL_MEM_HOST_NO_ACCESS are mutually exclusive.

**Table 5.3**      *List of supported cl\_mem\_flags values*

*size* is the size in bytes of the buffer memory object to be allocated.

*host\_ptr* is a pointer to the buffer data that may already be allocated by the application. The size of the buffer that *host\_ptr* points to must be  $\geq$  *size* bytes.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreateBuffer** returns a valid non-zero buffer object and *errcode\_ret* is set to CL\_SUCCESS if the buffer object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid context.
- ✚ CL\_INVALID\_VALUE if values specified in *flags* are not valid as defined in *table 5.3*.
- ✚ CL\_INVALID\_BUFFER\_SIZE if *size* is  $0^{16}$ .
- ✚ CL\_INVALID\_HOST\_PTR if *host\_ptr* is NULL and CL\_MEM\_USE\_HOST\_PTR or CL\_MEM\_COPY\_HOST\_PTR are set in *flags* or if *host\_ptr* is not NULL but CL\_MEM\_COPY\_HOST\_PTR or CL\_MEM\_USE\_HOST\_PTR are not set in *flags*.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for

---

<sup>16</sup> Implementations may return CL\_INVALID\_BUFFER\_SIZE if *size* is greater than CL\_DEVICE\_MAX\_MEM\_ALLOC\_SIZE value specified in *table 4.3* for all *devices* in context.

buffer object.

- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_mem      clCreateSubBuffer (cl_mem buffer,  
                                cl_mem_flags flags,  
                                cl_buffer_create_type buffer_create_type,  
                                const void *buffer_create_info,  
                                cl_int *errcode_ret)
```

can be used to create a new buffer object (referred to as a sub-buffer object) from an existing buffer object.

*buffer* must be a valid buffer object and cannot be a sub-buffer object.

*flags* is a bit-field that is used to specify allocation and usage information about the sub-buffer memory object being created and is described in *table 5.3*. If the CL\_MEM\_READ\_WRITE, CL\_MEM\_READ\_ONLY or CL\_MEM\_WRITE\_ONLY values are not specified in *flags*, they are inherited from the corresponding memory access qualifiers associated with *buffer*. The CL\_MEM\_USE\_HOST\_PTR, CL\_MEM\_ALLOC\_HOST\_PTR and CL\_MEM\_COPY\_HOST\_PTR values cannot be specified in *flags* but are inherited from the corresponding memory access qualifiers associated with *buffer*. If CL\_MEM\_COPY\_HOST\_PTR is specified in the memory access qualifier values associated with *buffer* it does not imply any additional copies when the sub-buffer is created from *buffer*. If the CL\_MEM\_HOST\_WRITE\_ONLY, CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS values are not specified in *flags*, they are inherited from the corresponding memory access qualifiers associated with *buffer*.




*buffer\_create\_type* and *buffer\_create\_info* describe the type of buffer object to be created. The list of supported values for *buffer\_create\_type* and corresponding descriptor that *buffer\_create\_info* points to is described in *table 5.4*.

cl_buffer_create_type	Description
CL_BUFFER_CREATE_TYPE_REGION	Create a buffer object that represents a specific region in <i>buffer</i> .  <i>buffer_create_info</i> is a pointer to the following structure:  typedef struct cl_buffer_region {

	<pre> size_t origin; size_t size; } cl_buffer_region; </pre> <p>(<i>origin</i>, <i>size</i>) defines the offset and size in bytes in <i>buffer</i>.</p> <p>If <i>buffer</i> is created with CL_MEM_USE_HOST_PTR, the <i>host_ptr</i> associated with the buffer object returned is <i>host_ptr</i> + <i>origin</i>.</p> <p>The buffer object returned references the data store allocated for <i>buffer</i> and points to a specific region given by (<i>origin</i>, <i>size</i>) in this data store.</p> <p>CL_INVALID_VALUE is returned in <i>errcode_ret</i> if the region specified by (<i>origin</i>, <i>size</i>) is out of bounds in <i>buffer</i>.</p> <p>CL_INVALID_BUFFER_SIZE if <i>size</i> is 0.</p> <p>CL_MISALIGNED_SUB_BUFFER_OFFSET is returned in <i>errcode_ret</i> if there are no devices in context associated with <i>buffer</i> for which the <i>origin</i> value is aligned to the CL_DEVICE_MEM_BASE_ADDR_ALIGN value.</p>
--	--

**Table 5.4** List of supported names and values in *clCreateSubBuffer*.

**clCreateSubBuffer** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors in *errcode\_ret*:

-  CL\_INVALID\_MEM\_OBJECT if *buffer* is not a valid buffer object or is a sub-buffer object.
-  CL\_INVALID\_VALUE if *buffer* was created with CL\_MEM\_WRITE\_ONLY and *flags* specifies CL\_MEM\_READ\_WRITE or CL\_MEM\_READ\_ONLY, or if *buffer* was created with CL\_MEM\_READ\_ONLY and *flags* specifies CL\_MEM\_READ\_WRITE or CL\_MEM\_WRITE\_ONLY, or if *flags* specifies CL\_MEM\_USE\_HOST\_PTR or CL\_MEM\_ALLOC\_HOST\_PTR or CL\_MEM\_COPY\_HOST\_PTR.
-  CL\_INVALID\_VALUE if *buffer* was created with CL\_MEM\_HOST\_WRITE\_ONLY and *flags* specify CL\_MEM\_HOST\_READ\_ONLY, or if *buffer* was created with CL\_MEM\_HOST\_READ\_ONLY and *flags* specify CL\_MEM\_HOST\_WRITE\_ONLY, or if

*buffer* was created with CL\_MEM\_HOST\_NO\_ACCESS and *flags* specify CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_WRITE\_ONLY.

- ✚ CL\_INVALID\_VALUE if value specified in *buffer\_create\_type* is not valid.
- ✚ CL\_INVALID\_VALUE if value(s) specified in *buffer\_create\_info* (for a given *buffer\_create\_type*) is not valid or if *buffer\_create\_info* is NULL.
- ✚ CL\_INVALID\_BUFFER\_SIZE if *size* is 0.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for sub-buffer object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

NOTE:

Concurrent reading from, writing to and copying between both a buffer object and its sub-buffer object(s) is undefined. Concurrent reading from, writing to and copying between overlapping sub-buffer objects created with the same buffer object is undefined. Only reading from both a buffer object and its sub-buffer objects or reading from multiple overlapping sub-buffer objects is defined.

## 5.2.2 Reading, Writing and Copying Buffer Objects

The following functions enqueue commands to read from a buffer object to host memory or write to a buffer object from host memory.

```
cl_int  clEnqueueReadBuffer (cl_command_queue command_queue,
                               cl_mem buffer,
                               cl_bool blocking_read,
                               size_t offset,
                               size_t size,
                               void *ptr,
                               cl_uint num_events_in_wait_list,
                               const cl_event *event_wait_list,
                               cl_event *event)
```

```
cl_int  clEnqueueWriteBuffer (cl_command_queue command_queue,
                               cl_mem buffer,
                               cl_bool blocking_write,
                               size_t offset,
                               size_t size,
                               const void *ptr,
                               cl_uint num_events_in_wait_list,
                               const cl_event *event_wait_list,
                               cl_event *event)
```

*command\_queue* is a valid host command-queue in which the read / write command will be queued. *command\_queue* and *buffer* must be created with the same OpenCL context.

*buffer* refers to a valid buffer object.

*blocking\_read* and *blocking\_write* indicate if the read and write operations are *blocking* or *non-blocking*.

If *blocking\_read* is CL\_TRUE i.e. the read command is blocking, **clEnqueueReadBuffer** does not return until the buffer data has been read and copied into memory pointed to by *ptr*.

If *blocking\_read* is CL\_FALSE i.e. the read command is non-blocking, **clEnqueueReadBuffer** queues a non-blocking read command and returns. The contents of the buffer that *ptr* points to cannot be used until the read command has completed. The *event* argument returns an event object which can be used to query the execution status of the read command. When the read command has completed, the contents of the buffer that *ptr* points to can be used by the application.

If *blocking\_write* is CL\_TRUE, the OpenCL implementation copies the data referred to by *ptr* and enqueues the write operation in the command-queue. The memory pointed to by *ptr* can be reused by the application after the **clEnqueueWriteBuffer** call returns.

If *blocking\_write* is CL\_FALSE, the OpenCL implementation will use *ptr* to perform a non-blocking write. As the write is non-blocking the implementation can return immediately. The memory pointed to by *ptr* cannot be reused by the application after the call returns. The *event* argument returns an event object which can be used to query the execution status of the write command. When the write command has completed, the memory pointed to by *ptr* can then be reused by the application.

*offset* is the offset in bytes in the buffer object to read from or write to.





*size* is the size in bytes of data being read or written.

*ptr* is the pointer to buffer in host memory where data is to be read into or to be written from.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular read / write command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueReadBuffer** and **clEnqueueWriteBuffer** return CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_INVALID\_CONTEXT if the context associated with *command\_queue* and *buffer* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
-  CL\_INVALID\_MEM\_OBJECT if *buffer* is not a valid buffer object.
-  CL\_INVALID\_VALUE if the region being read or written specified by (*offset*, *size*) is out of bounds or if *ptr* is a NULL value or if *size* is 0.



- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_EXEC\_STATUS\_ERROR\_FOR\_EVENTS\_IN\_WAIT\_LIST if the read and write operations are blocking and the execution status of any of the events in *event\_wait\_list* is a negative integer value.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *buffer*.
- ✚ CL\_INVALID\_OPERATION if **clEnqueueReadBuffer** is called on *buffer* which has been created with CL\_MEM\_HOST\_WRITE\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS.
- ✚ CL\_INVALID\_OPERATION if **clEnqueueWriteBuffer** is called on *buffer* which has been created with CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The following functions enqueue commands to read a 2D or 3D rectangular region from a buffer object to host memory or write a 2D or 3D rectangular region to a buffer object from host memory.

```
cl_int  clEnqueueReadBufferRect (cl_command_queue command_queue,
                                   cl_mem buffer,
                                   cl_bool blocking_read,
                                   const size_t *buffer_origin,
                                   const size_t *host_origin,
                                   const size_t *region,
                                   size_t buffer_row_pitch,
                                   size_t buffer_slice_pitch,
                                   size_t host_row_pitch,
                                   size_t host_slice_pitch,
                                   void *ptr,
                                   cl_uint num_events_in_wait_list,
                                   const cl_event *event_wait_list,
                                   cl_event *event)
```

```

cl_int  clEnqueueWriteBufferRect (cl_command_queue command_queue,
                                   cl_mem buffer,
                                   cl_bool blocking_write,
                                   const size_t *buffer_origin,
                                   const size_t *host_origin,
                                   const size_t *region,
                                   size_t buffer_row_pitch,
                                   size_t buffer_slice_pitch,
                                   size_t host_row_pitch,
                                   size_t host_slice_pitch,
                                   const void *ptr,
                                   cl_uint num_events_in_wait_list,
                                   const cl_event *event_wait_list,
                                   cl_event *event)

```

*command\_queue* refers to a valid host command-queue in which the read / write command will be queued. *command\_queue* and *buffer* must be created with the same OpenCL context.

*buffer* refers to a valid buffer object.

*blocking\_read* and *blocking\_write* indicate if the read and write operations are *blocking* or *non-blocking*.

If *blocking\_read* is CL\_TRUE i.e. the read command is blocking, **clEnqueueReadBufferRect** does not return until the buffer data has been read and copied into memory pointed to by *ptr*.

If *blocking\_read* is CL\_FALSE i.e. the read command is non-blocking, **clEnqueueReadBufferRect** queues a non-blocking read command and returns. The contents of the buffer that *ptr* points to cannot be used until the read command has completed. The *event* argument returns an event object which can be used to query the execution status of the read command. When the read command has completed, the contents of the buffer that *ptr* points to can be used by the application.

If *blocking\_write* is CL\_TRUE, the OpenCL implementation copies the data referred to by *ptr* and enqueues the write operation in the command-queue. The memory pointed to by *ptr* can be reused by the application after the **clEnqueueWriteBufferRect** call returns.

If *blocking\_write* is CL\_FALSE, the OpenCL implementation will use *ptr* to perform a non-blocking write. As the write is non-blocking the implementation can return immediately. The memory pointed to by *ptr* cannot be reused by the application after the call returns. The *event* argument returns an event object which can be used to query the execution status of the write command. When the write command has completed, the memory pointed to by *ptr* can then be reused by the application.

*buffer\_origin* defines the (x, y, z) offset in the memory region associated with *buffer*. For a 2D rectangle region, the z value given by *buffer\_origin*[2] should be 0. The offset in bytes is computed as *buffer\_origin*[2] \* *buffer\_slice\_pitch* + *buffer\_origin*[1] \* *buffer\_row\_pitch* + *buffer\_origin*[0].

*host\_origin* defines the (x, y, z) offset in the memory region pointed to by *ptr*. For a 2D rectangle region, the z value given by *host\_origin*[2] should be 0. The offset in bytes is computed as *host\_origin*[2] \* *host\_slice\_pitch* + *host\_origin*[1] \* *host\_row\_pitch* + *host\_origin*[0].

*region* defines the (width in bytes, height in rows, depth in slices) of the 2D or 3D rectangle being read or written. For a 2D rectangle copy, the *depth* value given by *region*[2] should be 1. The values in *region* cannot be 0.

*buffer\_row\_pitch* is the length of each row in bytes to be used for the memory region associated with *buffer*. If *buffer\_row\_pitch* is 0, *buffer\_row\_pitch* is computed as *region*[0].

*buffer\_slice\_pitch* is the length of each 2D slice in bytes to be used for the memory region associated with *buffer*. If *buffer\_slice\_pitch* is 0, *buffer\_slice\_pitch* is computed as *region*[1] \* *buffer\_row\_pitch*.

*host\_row\_pitch* is the length of each row in bytes to be used for the memory region pointed to by *ptr*. If *host\_row\_pitch* is 0, *host\_row\_pitch* is computed as *region*[0].

*host\_slice\_pitch* is the length of each 2D slice in bytes to be used for the memory region pointed to by *ptr*. If *host\_slice\_pitch* is 0, *host\_slice\_pitch* is computed as *region*[1] \* *host\_row\_pitch*.

*ptr* is the pointer to buffer in host memory where data is to be read into or to be written from.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular read / write command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueReadBufferRect** and **clEnqueueWriteBufferRect** return CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue* and *buffer* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *buffer* is not a valid buffer object.
- ✚ CL\_INVALID\_VALUE if the region being read or written specified by (*buffer\_origin*, *region*, *buffer\_row\_pitch*, *buffer\_slice\_pitch*) is out of bounds.
- ✚ CL\_INVALID\_VALUE if *ptr* is a NULL value.
- ✚ CL\_INVALID\_VALUE if any *region* array element is 0.
- ✚ CL\_INVALID\_VALUE if *buffer\_row\_pitch* is not 0 and is less than *region*[0].
- ✚ CL\_INVALID\_VALUE if *host\_row\_pitch* is not 0 and is less than *region*[0].
- ✚ CL\_INVALID\_VALUE if *buffer\_slice\_pitch* is not 0 and is less than *region*[1] \* *buffer\_row\_pitch* and not a multiple of *buffer\_row\_pitch*.
- ✚ CL\_INVALID\_VALUE if *host\_slice\_pitch* is not 0 and is less than *region*[1] \* *host\_row\_pitch* and not a multiple of *host\_row\_pitch*.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_EXEC\_STATUS\_ERROR\_FOR\_EVENTS\_IN\_WAIT\_LIST if the read and write operations are blocking and the execution status of any of the events in *event\_wait\_list* is a negative integer value.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *buffer*.
- ✚ CL\_INVALID\_OPERATION if **clEnqueueReadBufferRect** is called on *buffer* which has been created with CL\_MEM\_HOST\_WRITE\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS.
- ✚ CL\_INVALID\_OPERATION if **clEnqueueWriteBufferRect** is called on *buffer* which has been created with CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS.

- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

NOTE:

Calling **clEnqueueReadBuffer** to read a region of the buffer object with the *ptr* argument value set to *host\_ptr + offset*, where *host\_ptr* is a pointer to the memory region specified when the buffer object being read is created with CL\_MEM\_USE\_HOST\_PTR, must meet the following requirements in order to avoid undefined behavior:

- All commands that use this buffer object or a memory object (buffer or image) created from this buffer object have finished execution before the read command begins execution.
- The buffer object or memory objects created from this buffer object are not mapped.
- The buffer object or memory objects created from this buffer object are not used by any command-queue until the read command has finished execution.

Calling **clEnqueueReadBufferRect** to read a region of the buffer object with the *ptr* argument value set to *host\_ptr* and *host\_origin*, *buffer\_origin* values are the same, where *host\_ptr* is a pointer to the memory region specified when the buffer object being read is created with CL\_MEM\_USE\_HOST\_PTR, must meet the same requirements given above for **clEnqueueReadBuffer**.

Calling **clEnqueueWriteBuffer** to update the latest bits in a region of the buffer object with the *ptr* argument value set to *host\_ptr + offset*, where *host\_ptr* is a pointer to the memory region specified when the buffer object being written is created with CL\_MEM\_USE\_HOST\_PTR, must meet the following requirements in order to avoid undefined behavior:

- The host memory region given by (*host\_ptr + offset*, *cb*) contains the latest bits when the enqueued write command begins execution.
- The buffer object or memory objects created from this buffer object are not mapped.
- The buffer object or memory objects created from this buffer object are not used by any command-queue until the write command has finished execution.

Calling **clEnqueueWriteBufferRect** to update the latest bits in a region of the buffer object with the *ptr* argument value set to *host\_ptr* and *host\_origin*, *buffer\_origin* values are the same, where *host\_ptr* is a pointer to the memory region specified when the buffer object being written is

created with CL\_MEM\_USE\_HOST\_PTR, must meet the following requirements in order to avoid undefined behavior:

- The host memory region given by (*buffer\_origin region*) contains the latest bits when the enqueued write command begins execution.
- The buffer object or memory objects created from this buffer object are not mapped.
- The buffer object or memory objects created from this buffer object are not used by any command-queue until the write command has finished execution.

The function

```
cl_int clEnqueueCopyBuffer (cl_command_queue command_queue,  
                             cl_mem src_buffer,  
                             cl_mem dst_buffer,  
                             size_t src_offset,  
                             size_t dst_offset,  
                             size_t size,  
                             cl_uint num_events_in_wait_list,  
                             const cl_event *event_wait_list,  
                             cl_event *event)
```

enqueues a command to copy a buffer object identified by *src\_buffer* to another buffer object identified by *dst\_buffer*.

*command\_queue* refers to a host command-queue in which the copy command will be queued. The OpenCL context associated with *command\_queue*, *src\_buffer* and *dst\_buffer* must be the same.

*src\_offset* refers to the offset where to begin copying data from *src\_buffer*.

*dst\_offset* refers to the offset where to begin copying data into *dst\_buffer*.

*size* refers to the size in bytes to copy.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular copy command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueCopyBuffer** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue*, *src\_buffer* and *dst\_buffer* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *src\_buffer* and *dst\_buffer* are not valid buffer objects.
- ✚ CL\_INVALID\_VALUE if *src\_offset*, *dst\_offset*, *size*, *src\_offset* + *size* or *dst\_offset* + *size* require accessing elements outside the *src\_buffer* and *dst\_buffer* buffer objects respectively.
- ✚ CL\_INVALID\_VALUE if *size* is 0.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *src\_buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *dst\_buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_MEM\_COPY\_OVERLAP if *src\_buffer* and *dst\_buffer* are the same buffer or sub-buffer object and the source and destination regions overlap or if *src\_buffer* and *dst\_buffer* are different sub-buffers of the same associated buffer object and they overlap. The regions overlap if  $src\_offset \leq dst\_offset \leq src\_offset + size - 1$  or if  $dst\_offset \leq src\_offset \leq dst\_offset + size - 1$ .
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *src\_buffer* or *dst\_buffer*.

- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clEnqueueCopyBufferRect (cl_command_queue command_queue,
                                cl_mem src_buffer,
                                cl_mem dst_buffer,
                                const size_t *src_origin,
                                const size_t *dst_origin,
                                const size_t *region,
                                size_t src_row_pitch,
                                size_t src_slice_pitch,
                                size_t dst_row_pitch,
                                size_t dst_slice_pitch,
                                cl_uint num_events_in_wait_list,
                                const cl_event *event_wait_list,
                                cl_event *event)
```

enqueues a command to copy a 2D or 3D rectangular region from the buffer object identified by *src\_buffer* to a 2D or 3D region in the buffer object identified by *dst\_buffer*. Copying begins at the source offset and destination offset which are computed as described below in the description for *src\_origin* and *dst\_origin*. Each byte of the region's width is copied from the source offset to the destination offset. After copying each width, the source and destination offsets are incremented by their respective source and destination row pitches. After copying each 2D rectangle, the source and destination offsets are incremented by their respective source and destination slice pitches.

NOTE: If *src\_buffer* and *dst\_buffer* are the same buffer object, *src\_row\_pitch* must equal *dst\_row\_pitch* and *src\_slice\_pitch* must equal *dst\_slice\_pitch*.

*command\_queue* refers to the host command-queue in which the copy command will be queued. The OpenCL context associated with *command\_queue*, *src\_buffer* and *dst\_buffer* must be the same.

*src\_origin* defines the (x, y, z) offset in the memory region associated with *src\_buffer*. For a 2D rectangle region, the z value given by *src\_origin*[2] should be 0. The offset in bytes is computed as *src\_origin*[2] \* *src\_slice\_pitch* + *src\_origin*[1] \* *src\_row\_pitch* + *src\_origin*[0].

*dst\_origin* defines the (x, y, z) offset in the memory region associated with *dst\_buffer*. For a 2D rectangle region, the z value given by *dst\_origin*[2] should be 0. The offset in bytes is computed as *dst\_origin*[2] \* *dst\_slice\_pitch* + *dst\_origin*[1] \* *dst\_row\_pitch* + *dst\_origin*[0].



*region* defines the (*width* in bytes, *height* in rows, *depth* in slices) of the 2D or 3D rectangle being copied. For a 2D rectangle, the *depth* value given by *region*[2] should be 1. The values in *region* cannot be 0.

*src\_row\_pitch* is the length of each row in bytes to be used for the memory region associated with *src\_buffer*. If *src\_row\_pitch* is 0, *src\_row\_pitch* is computed as *region*[0].

*src\_slice\_pitch* is the length of each 2D slice in bytes to be used for the memory region associated with *src\_buffer*. If *src\_slice\_pitch* is 0, *src\_slice\_pitch* is computed as *region*[1] \* *src\_row\_pitch*.




*dst\_row\_pitch* is the length of each row in bytes to be used for the memory region associated with *dst\_buffer*. If *dst\_row\_pitch* is 0, *dst\_row\_pitch* is computed as *region*[0].

*dst\_slice\_pitch* is the length of each 2D slice in bytes to be used for the memory region associated with *dst\_buffer*. If *dst\_slice\_pitch* is 0, *dst\_slice\_pitch* is computed as *region*[1] \* *dst\_row\_pitch*.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular copy command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueCopyBufferRect** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_INVALID\_CONTEXT if the context associated with *command\_queue*, *src\_buffer* and *dst\_buffer* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
-  CL\_INVALID\_MEM\_OBJECT if *src\_buffer* and *dst\_buffer* are not valid buffer objects.

- ✚ CL\_INVALID\_VALUE if (*src\_origin*, *region*, *src\_row\_pitch*, *src\_slice\_pitch*) or (*dst\_origin*, *region*, *dst\_row\_pitch*, *dst\_slice\_pitch*) require accessing elements outside the *src\_buffer* and *dst\_buffer* buffer objects respectively.
- ✚ CL\_INVALID\_VALUE if any *region* array element is 0.
- ✚ CL\_INVALID\_VALUE if *src\_row\_pitch* is not 0 and is less than *region*[0].
- ✚ CL\_INVALID\_VALUE if *dst\_row\_pitch* is not 0 and is less than *region*[0].
- ✚ CL\_INVALID\_VALUE if *src\_slice\_pitch* is not 0 and is less than *region*[1] \* *src\_row\_pitch* or if *src\_slice\_pitch* is not 0 and is not a multiple of *src\_row\_pitch*.
- ✚ CL\_INVALID\_VALUE if *dst\_slice\_pitch* is not 0 and is less than *region*[1] \* *dst\_row\_pitch* or if *dst\_slice\_pitch* is not 0 and is not a multiple of *dst\_row\_pitch*.
- ✚ CL\_INVALID\_VALUE if *src\_buffer* and *dst\_buffer* are the same buffer object and *src\_slice\_pitch* is not equal to *dst\_slice\_pitch* and *src\_row\_pitch* is not equal to *dst\_row\_pitch*.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MEM\_COPY\_OVERLAP if *src\_buffer* and *dst\_buffer* are the same buffer or sub-buffer object and the source and destination regions overlap or if *src\_buffer* and *dst\_buffer* are different sub-buffers of the same associated buffer object and they overlap. Refer to Appendix E for details on how to determine if source and destination regions overlap.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *src\_buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *dst\_buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *src\_buffer* or *dst\_buffer*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 5.2.3 Filling Buffer Objects

The function

```
cl_int clEnqueueFillBuffer (cl_command_queue command_queue,
                             cl_mem buffer,
                             const void *pattern,
                             size_t pattern_size,
                             size_t offset,
                             size_t size,
                             cl_uint num_events_in_wait_list,
                             const cl_event *event_wait_list,
                             cl_event *event)
```

enqueues a command to fill a buffer object with a pattern of a given pattern size. The usage information which indicates whether the memory object can be read or written by a kernel and/or the host and is given by the `cl_mem_flags` argument value specified when *buffer* is created is ignored by **clEnqueueFillBuffer**.

*command\_queue* refers to the host command-queue in which the fill command will be queued. The OpenCL context associated with *command\_queue* and *buffer* must be the same.

*buffer* is a valid buffer object.

*pattern* is a pointer to the data pattern of size *pattern\_size* in bytes. *pattern* will be used to fill a region in *buffer* starting at *offset* and is *size* bytes in size. The data pattern must be a scalar or vector integer or floating-point data type supported by OpenCL as described in sections 6.1.1 and 6.1.2. For example, if *buffer* is to be filled with a pattern of `float4` values, then *pattern* will be a pointer to a `cl_float4` value and *pattern\_size* will be `sizeof(cl_float4)`. The maximum value of *pattern\_size* is the size of the largest integer or floating-point vector data type supported by the OpenCL device. The memory associated with *pattern* can be reused or freed after the function returns.

*offset* is the location in bytes of the region being filled in *buffer* and must be a multiple of *pattern\_size*.

*size* is the size in bytes of region being filled in *buffer* and must be a multiple of *pattern\_size*.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueFillBuffer** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue* and *buffer* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *buffer* is not a valid buffer object.
- ✚ CL\_INVALID\_VALUE if *offset* or *offset + size* require accessing elements outside the *buffer* buffer object respectively.
- ✚ CL\_INVALID\_VALUE if *pattern* is NULL or if *pattern\_size* is 0 or if *pattern\_size* is not one of {1, 2, 4, 8, 16, 32, 64, 128}.
- ✚ CL\_INVALID\_VALUE if *offset* and *size* are not a multiple of *pattern\_size*.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *buffer* is a sub-buffer object and offset specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *buffer*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 5.2.4 Mapping Buffer Objects

The function

```
void * clEnqueueMapBuffer (cl_command_queue command_queue,
                           cl_mem buffer,
                           cl_bool blocking_map,
                           cl_map_flags map_flags,
                           size_t offset,
                           size_t size,
                           cl_uint num_events_in_wait_list,
                           const cl_event *event_wait_list,
                           cl_event *event,
                           cl_int *errcode_ret)
```

enqueues a command to map a region of the buffer object given by *buffer* into the host address space and returns a pointer to this mapped region.

*command\_queue* must be a valid host command-queue.

*blocking\_map* indicates if the map operation is *blocking* or *non-blocking*.

If *blocking\_map* is CL\_TRUE, **clEnqueueMapBuffer** does not return until the specified region in *buffer* is mapped into the host address space and the application can access the contents of the mapped region using the pointer returned by **clEnqueueMapBuffer**.

If *blocking\_map* is CL\_FALSE i.e. map operation is non-blocking, the pointer to the mapped region returned by **clEnqueueMapBuffer** cannot be used until the map command has completed. The *event* argument returns an event object which can be used to query the execution status of the map command. When the map command is completed, the application can access the contents of the mapped region using the pointer returned by **clEnqueueMapBuffer**.

*map\_flags* is a bit-field and is described in *table 5.5*.

*buffer* is a valid buffer object. The OpenCL context associated with *command\_queue* and *buffer* must be the same.

*offset* and *size* are the offset in bytes and the size of the region in the buffer object that is being mapped.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must

be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clEnqueueMapBuffer** will return a pointer to the mapped region. The *errcode\_ret* is set to CL\_SUCCESS.

A NULL pointer is returned otherwise with one of the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if context associated with *command\_queue* and *buffer* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *buffer* is not a valid buffer object.
- ✚ CL\_INVALID\_VALUE if region being mapped given by (*offset*, *size*) is out of bounds or if *size* is 0 or if values specified in *map\_flags* are not valid.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for the device associated with *queue*.
- ✚ CL\_MAP\_FAILURE if there is a failure to map the requested region into the host address space. This error cannot occur for buffer objects created with CL\_MEM\_USE\_HOST\_PTR or CL\_MEM\_ALLOC\_HOST\_PTR.
- ✚ CL\_EXEC\_STATUS\_ERROR\_FOR\_EVENTS\_IN\_WAIT\_LIST if the map operation is blocking and the execution status of any of the events in *event\_wait\_list* is a negative integer value.

- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *buffer*.
- ✚ CL\_INVALID\_OPERATION if *buffer* has been created with CL\_MEM\_HOST\_WRITE\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS and CL\_MAP\_READ is set in *map\_flags* or if *buffer* has been created with CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS and CL\_MAP\_WRITE or CL\_MAP\_WRITE\_INVALIDATE\_REGION is set in *map\_flags*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.
- ✚ CL\_INVALID\_OPERATION if mapping would lead to overlapping regions being mapped for writing.

The pointer returned maps a region starting at *offset* and is at least *size* bytes in size. The result of a memory access outside this region is undefined.

NOTE:

If the buffer object is created with CL\_MEM\_USE\_HOST\_PTR set in *mem\_flags*, the following will be true:

- ✚ The *host\_ptr* specified in **clCreateBuffer** is guaranteed to contain the latest bits in the region being mapped when the **clEnqueueMapBuffer** command has completed.
- ✚ The pointer value returned by **clEnqueueMapBuffer** will be derived from the *host\_ptr* specified when the buffer object is created.

Mapped buffer objects are unmapped using **clEnqueueUnmapMemObject**. This is described in *section 5.5.2*.

cl_map_flags	Description
CL_MAP_READ	<p>This flag specifies that the region being mapped in the memory object is being mapped for reading.</p> <p>The pointer returned by <b>clEnqueueMap{Buffer   Image}</b> is guaranteed to contain the latest bits in the region being mapped when the <b>clEnqueueMap{Buffer  </b></p>

	<b>Image}</b> command has completed
<b>CL_MAP_WRITE</b>	<p>This flag specifies that the region being mapped in the memory object is being mapped for writing.</p> <p>The pointer returned by <b>clEnqueueMap{Buffer   Image}</b> is guaranteed to contain the latest bits in the region being mapped when the <b>clEnqueueMap{Buffer   Image}</b> command has completed</p>
<b>CL_MAP_WRITE_INVALIDATE_REGION</b>	<p>This flag specifies that the region being mapped in the memory object is being mapped for writing.</p> <p>The contents of the region being mapped are to be discarded. This is typically the case when the region being mapped is overwritten by the host. This flag allows the implementation to no longer guarantee that the pointer returned by <b>clEnqueueMap{Buffer   Image}</b> contains the latest bits in the region being mapped which can be a significant performance enhancement.</p> <p>CL_MAP_READ or CL_MAP_WRITE and CL_MAP_WRITE_INVALIDATE_REGION are mutually exclusive.</p>

**Table 5.5**     *List of supported cl\_map\_flags values*



## 5.3 Image Objects

An *image* object is used to store a one-, two- or three- dimensional texture, frame-buffer or image. The elements of an image object are selected from a list of predefined image formats. The minimum number of elements in a memory object is one.

### 5.3.1 Creating Image Objects

A **1D image, 1D image buffer, 1D image array, 2D image, 2D image array and 3D image object** can be created using the following function

```
cl_mem  clCreateImage (cl_context context,
                      cl_mem_flags flags,
                      const cl_image_format *image_format,
                      const cl_image_desc *image_desc,
                      void *host_ptr,
                      cl_int *errcode_ret)
```

*context* is a valid OpenCL context on which the image object is to be created.

*flags* is a bit-field that is used to specify allocation and usage information about the image memory object being created and is described in *table 5.3*.

For all image types except CL\_MEM\_OBJECT\_IMAGE1D\_BUFFER, if value specified for *flags* is 0, the default is used which is CL\_MEM\_READ\_WRITE.

For CL\_MEM\_OBJECT\_IMAGE1D\_BUFFER image type, or an image created from another memory object (image or buffer), if the CL\_MEM\_READ\_WRITE, CL\_MEM\_READ\_ONLY or CL\_MEM\_WRITE\_ONLY values are not specified in *flags*, they are inherited from the corresponding memory access qualifiers associated with *mem\_object*. The CL\_MEM\_USE\_HOST\_PTR, CL\_MEM\_ALLOC\_HOST\_PTR and CL\_MEM\_COPY\_HOST\_PTR values cannot be specified in *flags* but are inherited from the corresponding memory access qualifiers associated with *mem\_object*. If CL\_MEM\_COPY\_HOST\_PTR is specified in the memory access qualifier values associated with *mem\_object* it does not imply any additional copies when the image is created from *mem\_object*. If the CL\_MEM\_HOST\_WRITE\_ONLY, CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS values are not specified in *flags*, they are inherited from the corresponding memory access qualifiers associated with *mem\_object*.

*image\_format* is a pointer to a structure that describes format properties of the image to be allocated. Refer to *section 5.3.1.1* for a detailed description of the image format descriptor.

*image\_desc* is a pointer to a structure that describes type and dimensions of the image to be allocated. Refer to *section 5.3.1.2* for a detailed description of the image descriptor.

*host\_ptr* is a pointer to the image data that may already be allocated by the application. Refer to table below for a description of how large the buffer that *host\_ptr* points to must be.

Image Type	Size of buffer that <i>host_ptr</i> points to
CL_MEM_OBJECT_IMAGE1D	$\geq \text{image\_row\_pitch}$
CL_MEM_OBJECT_IMAGE1D_BUFFER	$\geq \text{image\_row\_pitch}$
CL_MEM_OBJECT_IMAGE2D	$\geq \text{image\_row\_pitch} * \text{image\_height}$
CL_MEM_OBJECT_IMAGE3D	$\geq \text{image\_slice\_pitch} * \text{image\_depth}$
CL_MEM_OBJECT_IMAGE1D_ARRAY	$\geq \text{image\_slice\_pitch} * \text{image\_array\_size}$
CL_MEM_OBJECT_IMAGE2D_ARRAY	$\geq \text{image\_slice\_pitch} * \text{image\_array\_size}$

**clCreateImage** can be used to create a 2D image from a buffer object or a 2D image from another 2D image object.

A 2D image can be created from a buffer by specifying a buffer object in the *image\_desc->mem\_object* passed to **clCreateImage** for *image\_desc->image\_type* = CL\_MEM\_OBJECT\_IMAGE2D. If *image\_desc->mem\_object* is created with CL\_MEM\_USE\_HOST\_PTR, the *host\_ptr* specified to **clCreateBuffer** must be aligned to the minimum of the **CL\_DEVICE\_IMAGE\_BASE\_ADDRESS\_ALIGNMENT** value for all devices in the context associated with *image\_desc->mem\_object* and that support images.

A 2D image can be created from another 2D image object by specifying an image object in the *image\_desc->mem\_object* passed to **clCreateImage** for *image\_desc->image\_type* = CL\_MEM\_OBJECT\_IMAGE2D. This allows users to create a new image object that shares the image data store with *mem\_object* but views the pixels in the image with a different channel order. The restrictions are:

- ✚ all the values specified in *image\_desc* except for *mem\_object* must match the image descriptor information associated with *mem\_object*.
- ✚ The *image\_desc* used for creation of *mem\_object* may not be equivalent to image descriptor information associated with *mem\_object*. To ensure the values in *image\_desc* will match one can query *mem\_object* for associated information using **clGetImageInfo** function described in section 5.3.7.
- ✚ the channel data type specified in *image\_format* must match the channel data type associated with *mem\_object*. The channel order values<sup>17</sup> supported are:

image_channel_order	image channel order of
---------------------	------------------------

<sup>17</sup> This allows developers to create a sRGB view of the image from a linear RGB view or vice-versa i.e. the pixels stored in the image can be accessed as linear RGB or sRGB values.

specified in image_format	mem_object
CL_sBGRA	CL_BGRA
CL_BGRA	CL_sBGRA
CL_sRGBA	CL_RGBA
CL_RGBA	CL_sRGBA
CL_sRGB	CL_RGB
CL_RGB	CL_sRGB
CL_sRGBx	CL_RGBx
CL_RGBx	CL_sRGBx
CL_DEPTH	CL_R

- ✚ the channel order specified must have the same number of channels as the channel order of `mem_object`.

For a 3D image or 2D image array, the image data specified by *host\_ptr* is stored as a linear sequence of adjacent 2D image slices or 2D images respectively. Each 2D image is a linear sequence of adjacent scanlines. Each scanline is a linear sequence of image elements.

For a 2D image, the image data specified by *host\_ptr* is stored as a linear sequence of adjacent scanlines. Each scanline is a linear sequence of image elements.

For a 1D image array, the image data specified by *host\_ptr* is stored as a linear sequence of adjacent 1D images respectively. Each 1D image or 1D image buffer is a single scanline which is a linear sequence of adjacent elements.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreateImage** returns a valid non-zero image object created and the *errcode\_ret* is set to CL\_SUCCESS if the image object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid context.
- ✚ CL\_INVALID\_VALUE if values specified in *flags* are not valid.
- ✚ CL\_INVALID\_IMAGE\_FORMAT\_DESCRIPTOR if values specified in *image\_format* are not valid or if *image\_format* is NULL.
- ✚ CL\_INVALID\_IMAGE\_FORMAT\_DESCRIPTOR if a 2D image is created from a buffer and the row pitch and base address alignment does not follow the rules described for creating a 2D image from a buffer.
- ✚ CL\_INVALID\_IMAGE\_FORMAT\_DESCRIPTOR if a 2D image is created from a 2D image object and the rules described above are not followed.

- ✚ CL\_INVALID\_IMAGE\_DESCRIPTOR if values specified in *image\_desc* are not valid or if *image\_desc* is NULL.
- ✚ CL\_INVALID\_IMAGE\_SIZE if image dimensions specified in *image\_desc* exceed the maximum image dimensions described in *table 4.3* for all devices in *context*.
- ✚ CL\_INVALID\_HOST\_PTR if *host\_ptr* is NULL and CL\_MEM\_USE\_HOST\_PTR or CL\_MEM\_COPY\_HOST\_PTR are set in *flags* or if *host\_ptr* is not NULL but CL\_MEM\_COPY\_HOST\_PTR or CL\_MEM\_USE\_HOST\_PTR are not set in *flags*.
- ✚ CL\_INVALID\_VALUE if an image is being created from another memory object (buffer or image) under one of the following circumstances: 1) *mem\_object* was created with CL\_MEM\_WRITE\_ONLY and *flags* specifies CL\_MEM\_READ\_WRITE or CL\_MEM\_READ\_ONLY, 2) *mem\_object* was created with CL\_MEM\_READ\_ONLY and *flags* specifies CL\_MEM\_READ\_WRITE or CL\_MEM\_WRITE\_ONLY, 3) *flags* specifies CL\_MEM\_USE\_HOST\_PTR or CL\_MEM\_ALLOC\_HOST\_PTR or CL\_MEM\_COPY\_HOST\_PTR.
- ✚ CL\_INVALID\_VALUE if an image is being created from another memory object (buffer or image) and *mem\_object* object was created with CL\_MEM\_HOST\_WRITE\_ONLY and *flags* specifies CL\_MEM\_HOST\_READ\_ONLY, or if *mem\_object* was created with CL\_MEM\_HOST\_READ\_ONLY and *flags* specifies CL\_MEM\_HOST\_WRITE\_ONLY, or if *mem\_object* was created with CL\_MEM\_HOST\_NO\_ACCESS and *flags* specifies CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_WRITE\_ONLY.
- ✚ CL\_IMAGE\_FORMAT\_NOT\_SUPPORTED if the *image\_format* is not supported.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for image object.
- ✚ CL\_INVALID\_OPERATION if there are no devices in *context* that support images (i.e. CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_FALSE).
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

### 5.3.1.1 Image Format Descriptor

The image format descriptor structure is defined as

```
typedef struct _cl_image_format {
```

```

        cl_channel_order    image_channel_order;
        cl_channel_type     image_channel_data_type;
    } cl_image_format;

```

`image_channel_order` specifies the number of channels and the channel layout i.e. the memory layout in which channels are stored in the image. Valid values are described in *table 5.6*.

`image_channel_data_type` describes the size of the channel data type. The list of supported values is described in *table 5.7*. The number of bits per element determined by the `image_channel_data_type` and `image_channel_order` must be a power of two.

Enum values that can be specified in <code>channel_order</code>
<b>CL_R, CL_Rx or CL_A</b>
<b>CL_INTENSITY</b>
<b>CL_LUMINANCE</b>
<b>CL_DEPTH</b>
<b>CL_RG, CL_RGx or CL_RA</b>
<b>CL_RGB or CL_RGBx</b>
<b>CL_RGBA</b>
<b>CL_sRGB, CL_sRGBx, CL_sRGBA, CL_sBGRA</b>
<b>CL_ARGB, CL_BGRA, CL_ABGR</b>


**Table 5.6**      *List of supported Image Channel Order Values*

Image Channel Data Type	Description
<b>CL_SNORM_INT8</b>	Each channel component is a normalized signed 8-bit integer value
<b>CL_SNORM_INT16</b>	Each channel component is a normalized signed 16-bit integer value
<b>CL_UNORM_INT8</b>	Each channel component is a normalized unsigned 8-bit integer value
<b>CL_UNORM_INT16</b>	Each channel component is a normalized unsigned 16-bit integer value
<b>CL_UNORM_SHORT_565</b>	Represents a normalized 5-6-5 3-channel RGB image. The channel order must be CL_RGB or CL_RGBx.
<b>CL_UNORM_SHORT_555</b>	Represents a normalized x-5-5-5 4-channel xRGB image. The channel order must be CL_RGB or CL_RGBx.
<b>CL_UNORM_INT_101010</b>	Represents a normalized x-10-10-10 4-channel xRGB image. The channel order must be CL_RGB or CL_RGBx.
<b>CL_SIGNED_INT8</b>	Each channel component is an unnormalized signed 8-


	bit integer value
<b>CL_SIGNED_INT16</b>	Each channel component is an unnormalized signed 16-bit integer value
<b>CL_SIGNED_INT32</b>	Each channel component is an unnormalized signed 32-bit integer value
<b>CL_UNSIGNED_INT8</b>	Each channel component is an unnormalized unsigned 8-bit integer value
<b>CL_UNSIGNED_INT16</b>	Each channel component is an unnormalized unsigned 16-bit integer value
<b>CL_UNSIGNED_INT32</b>	Each channel component is an unnormalized unsigned 32-bit integer value
<b>CL_HALF_FLOAT</b>	Each channel component is a 16-bit half-float value
<b>CL_FLOAT</b>	Each channel component is a single precision floating-point value

**Table 5.7**      *List of supported Image Channel Data Types*

For example, to specify a normalized unsigned 8-bit / channel RGBA image, `image_channel_order = CL_RGBA`, and `image_channel_data_type = CL_UNORM_INT8`. The memory layout of this image format is described below:



Similar, if `image_channel_order = CL_RGBA` and `image_channel_data_type = CL_SIGNED_INT16`, the memory layout of this image format is described below:



`image_channel_data_type` values of `CL_UNORM_SHORT_565`, `CL_UNORM_SHORT_555` and `CL_UNORM_INT_101010` are special cases of packed image formats where the channels of each element are packed into a single unsigned short or unsigned int. For these special packed image formats, the channels are normally packed with the first channel in the most significant bits of the bitfield, and successive channels occupying progressively less significant locations. For `CL_UNORM_SHORT_565`, R is in bits 15 : 11, G is in bits 10 : 5 and B is in bits 4 : 0. For `CL_UNORM_SHORT_555`, bit 15 is undefined, R is in bits 14 : 10, G in bits 9 : 5 and B in bits 4 : 0. For `CL_UNORM_INT_101010`, bits 31 : 30 are undefined, R is in bits 29 : 20, G in bits 19 : 10 and B in bits 9 : 0.

OpenCL implementations must maintain the minimum precision specified by the number of bits

in `image_channel_data_type`. If the image format specified by `image_channel_order`, and `image_channel_data_type` cannot be supported by the OpenCL implementation, then the call to **clCreateImage** will return a NULL memory object.

### 5.3.1.2 Image Descriptor

The image descriptor structure describes the type and dimensions of the image or image array and is defined as:

```
typedef struct _cl_image_desc {
    cl_mem_object_type  image_type,
    size_t              image_width;
    size_t              image_height;
    size_t              image_depth;
    size_t              image_array_size;
    size_t              image_row_pitch;
    size_t              image_slice_pitch;
    cl_uint             num_mip_levels;
    cl_uint             num_samples;
    cl_mem              mem_object;
} cl_image_desc;
```

`image_type` describes the image type and must be either `CL_MEM_OBJECT_IMAGE1D`, `CL_MEM_OBJECT_IMAGE1D_BUFFER`, `CL_MEM_OBJECT_IMAGE1D_ARRAY`, `CL_MEM_OBJECT_IMAGE2D`, `CL_MEM_OBJECT_IMAGE2D_ARRAY` or `CL_MEM_OBJECT_IMAGE3D`.

`image_width` is the width of the image in pixels. For a 2D image and image array, the image width must be a value  $\geq 1$  and  $\leq \text{CL\_DEVICE\_IMAGE2D\_MAX\_WIDTH}$ . For a 3D image, the image width must be a value  $\geq 1$  and  $\leq \text{CL\_DEVICE\_IMAGE3D\_MAX\_WIDTH}$ . For a 1D image buffer, the image width must be a value  $\geq 1$  and  $\leq \text{CL\_DEVICE\_IMAGE\_MAX\_BUFFER\_SIZE}$ . For a 1D image and 1D image array, the image width must be a value  $\geq 1$  and  $\leq \text{CL\_DEVICE\_IMAGE2D\_MAX\_WIDTH}$ .

`image_height` is height of the image in pixels. This is only used if the image is a 2D or 3D image, or a 2D image array. For a 2D image or image array, the image height must be a value  $\geq 1$  and  $\leq \text{CL\_DEVICE\_IMAGE2D\_MAX\_HEIGHT}$ . For a 3D image, the image height must be a value  $\geq 1$  and  $\leq \text{CL\_DEVICE\_IMAGE3D\_MAX\_HEIGHT}$ .

`image_depth` is the depth of the image in pixels. This is only used if the image is a 3D image and must be a value  $\geq 1$  and  $\leq \text{CL\_DEVICE\_IMAGE3D\_MAX\_DEPTH}$ .

`image_array_size`<sup>18</sup> is the number of images in the image array. This is only used if the image is a 1D or 2D image array. The values for `image_array_size`, if specified, must be a value  $\geq 1$  and  $\leq \text{CL\_DEVICE\_IMAGE\_MAX\_ARRAY\_SIZE}$ .

`image_row_pitch` is the scan-line pitch in bytes. This must be 0 if `host_ptr` is NULL and can be either 0 or  $\geq \text{image\_width} * \text{size of element in bytes}$  if `host_ptr` is not NULL. If `host_ptr` is not NULL and `image_row_pitch = 0`, `image_row_pitch` is calculated as `image_width * size of element in bytes`. If `image_row_pitch` is not 0, it must be a multiple of the image element size in bytes. For a 2D image created from a buffer, the pitch specified (or computed if pitch specified is 0) must be a multiple of the maximum of the `CL_DEVICE_IMAGE_PITCH_ALIGNMENT` value for all devices in the context associated with `image_desc->mem_object` and that support images.

`image_slice_pitch` is the size in bytes of each 2D slice in the 3D image or the size in bytes of each image in a 1D or 2D image array. This must be 0 if `host_ptr` is NULL. If `host_ptr` is not NULL, `image_slice_pitch` can be either 0 or  $\geq \text{image\_row\_pitch} * \text{image\_height}$  for a 2D image array or 3D image and can be either 0 or  $\geq \text{image\_row\_pitch}$  for a 1D image array. If `host_ptr` is not NULL and `image_slice_pitch = 0`, `image_slice_pitch` is calculated as `image_row_pitch * image_height` for a 2D image array or 3D image and `image_row_pitch` for a 1D image array. If `image_slice_pitch` is not 0, it must be a multiple of the `image_row_pitch`.

`num_mip_levels` and `num_samples` must be 0.

`mem_object` refers to a valid buffer or image memory object. `mem_object` can be a buffer memory object if `image_type` is `CL_MEM_OBJECT_IMAGE1D_BUFFER` or `CL_MEM_OBJECT_IMAGE2D`<sup>19</sup>. `mem_object` can be an image object if `image_type` is `CL_MEM_OBJECT_IMAGE2D`<sup>20</sup>. Otherwise it must be NULL. The image pixels are taken from the memory object's data store. When the contents of the specified memory object's data store are modified, those changes are reflected in the contents of the image object and vice-versa at corresponding synchronization points. For a 1D image buffer object, the `image_width * size of element in bytes` must be  $\leq$  size of buffer object data store. For a 2D image created from a buffer, the `image_row_pitch * image_height` must be  $\leq$  size of buffer object data store. For an image object created from another image object, the values specified in the image descriptor except for `mem_object` must match the image descriptor information associated with `mem_object`.

#### NOTE:

Concurrent reading from, writing to and copying between both a buffer object and 1D image buffer or 2D image object associated with the buffer object is undefined. Only reading from both a buffer object and 1D image buffer or 2D image object associated with the buffer object is

---

<sup>18</sup> Note that reading and writing 2D image arrays from a kernel with `image_array_size = 1` may be lower performance than 2D images.

<sup>19</sup> To create a 2D image from a buffer object that share the data store between the image and buffer object.

<sup>20</sup> To create an image object from another image object that share the data store between these image objects.



defined.




Writing to an image created from a buffer and then reading from this buffer in a kernel even if appropriate synchronization operations (such as a barrier) are performed between the writes and reads is undefined. Similarly, writing to the buffer and reading from the image created from this buffer with appropriate synchronization between the writes and reads is undefined.

### 5.3.2 Querying List of Supported Image Formats

The function

```
cl_int  clGetSupportedImageFormats (cl_context context,
                                     cl_mem_flags flags,
                                     cl_mem_object_type image_type,
                                     cl_uint num_entries,
                                     cl_image_format *image_formats,
                                     cl_uint *num_image_formats)
```

can be used to get the list of image formats supported by an OpenCL implementation when the following information about an image memory object is specified:

-  Context
-  Image type – 1D, 2D, or 3D image, 1D image buffer, 1D or 2D image array.
-  Image object allocation information

**clGetSupportedImageFormats** returns a union of image formats supported by all devices in the context.

*context* is a valid OpenCL context on which the image object(s) will be created.

*flags* is a bit-field that is used to specify allocation and usage information about the image memory object being queried and is described in *table 5.3*. To get a list of supported image formats that can be read from or written to by a kernel, *flags* must be set to CL\_MEM\_READ\_WRITE (get a list of images that can be read from and written to by different kernel instances when correctly ordered by event dependencies), CL\_MEM\_READ\_ONLY (list of images that can be read from by a kernel) or CL\_MEM\_WRITE\_ONLY (list of images that can be written to by a kernel). To get a list of supported image formats that can be both read from and written to by the same kernel instance, *flags* must be set to CL\_MEM\_KERNEL\_READ\_AND\_WRITE. Please see section 5.3.2.2 for clarification.





*image\_type* describes the image type and must be either CL\_MEM\_OBJECT\_IMAGE1D, CL\_MEM\_OBJECT\_IMAGE1D\_BUFFER, CL\_MEM\_OBJECT\_IMAGE2D, CL\_MEM\_OBJECT\_IMAGE3D, CL\_MEM\_OBJECT\_IMAGE1D\_ARRAY or CL\_MEM\_OBJECT\_IMAGE2D\_ARRAY.

*num\_entries* specifies the number of entries that can be returned in the memory location given by *image\_formats*.

*image\_formats* is a pointer to a memory location where the list of supported image formats are returned. Each entry describes a *cl\_image\_format* structure supported by the OpenCL implementation. If *image\_formats* is NULL, it is ignored.

*num\_image\_formats* is the actual number of supported image formats for a specific *context* and values specified by *flags*. If *num\_image\_formats* is NULL, it is ignored.

**clGetSupportedImageFormats** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_CONTEXT if *context* is not a valid context.
-  CL\_INVALID\_VALUE if *flags* or *image\_type* are not valid, or if *num\_entries* is 0 and *image\_formats* is not NULL.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

If CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_TRUE, the values assigned to CL\_DEVICE\_MAX\_READ\_IMAGE\_ARGS, CL\_DEVICE\_MAX\_WRITE\_IMAGE\_ARGS, CL\_DEVICE\_IMAGE2D\_MAX\_WIDTH, CL\_DEVICE\_IMAGE2D\_MAX\_HEIGHT, CL\_DEVICE\_IMAGE3D\_MAX\_WIDTH, CL\_DEVICE\_IMAGE3D\_MAX\_HEIGHT, CL\_DEVICE\_IMAGE3D\_MAX\_DEPTH and CL\_DEVICE\_MAX\_SAMPLERS by the implementation must be greater than or equal to the minimum values specified in *table 4.3*.

### 5.3.2.1 Minimum List of Supported Image Formats

For 1D, 1D image from buffer, 2D, 3D image objects, 1D and 2D image array objects, the mandated minimum list of image formats that can be read from and written to by different kernel instances when correctly ordered by event dependencies and that must be supported by all devices that support images is described in *table 5.8.a*.

num_channels	channel_order	channel_data_type
--------------	---------------	-------------------

1	CL_R	CL_UNORM_INT8 CL_UNORM_INT16 CL_SNORM_INT8 CL_SNORM_INT16 CL_SIGNED_INT8 CL_SIGNED_INT16 CL_SIGNED_INT32 CL_UNSIGNED_INT8 CL_UNSIGNED_INT16 CL_UNSIGNED_INT32 CL_HALF_FLOAT CL_FLOAT
1	CL_DEPTH <sup>21</sup>	CL_UNORM_INT16 CL_FLOAT
2	CL_RG	CL_UNORM_INT8 CL_UNORM_INT16 CL_SNORM_INT8 CL_SNORM_INT16 CL_SIGNED_INT8 CL_SIGNED_INT16 CL_SIGNED_INT32 CL_UNSIGNED_INT8 CL_UNSIGNED_INT16 CL_UNSIGNED_INT32 CL_HALF_FLOAT CL_FLOAT
4	CL_RGBA	CL_UNORM_INT8 CL_UNORM_INT16 CL_SNORM_INT8 CL_SNORM_INT16 CL_SIGNED_INT8 CL_SIGNED_INT16 CL_SIGNED_INT32 CL_UNSIGNED_INT8 CL_UNSIGNED_INT16 CL_UNSIGNED_INT32 CL_HALF_FLOAT CL_FLOAT
4	CL_BGRA	CL_UNORM_INT8
4	CL_sRGBA	CL_UNORM_INT8 <sup>22</sup>

**Table 5.8.a** *Min. list of supported image formats – kernel read or write*

For 1D, 1D image from buffer, 2D, 3D image objects, 1D and 2D image array objects, the mandated minimum list of image formats that can be read from and written to by the same kernel instance and that must be supported by all devices that support images is described in *table 5.8.b*.

<sup>21</sup> CL\_DEPTH channel order is supported only for 2D image and 2D image array objects

<sup>22</sup> sRGB channel order support is not required for 1D image buffers. Writes to images with sRGB channel orders requires device support of the cl\_khr\_srgb\_image\_writes extension.

num_channels	channel_order	channel_data_type
1	CL_R	CL_UNORM_INT8 CL_SIGNED_INT8 CL_SIGNED_INT16 CL_SIGNED_INT32 CL_UNSIGNED_INT8 CL_UNSIGNED_INT16 CL_UNSIGNED_INT32 CL_HALF_FLOAT CL_FLOAT
4	CL_RGBA	CL_UNORM_INT8 CL_SIGNED_INT8 CL_SIGNED_INT16 CL_SIGNED_INT32 CL_UNSIGNED_INT8 CL_UNSIGNED_INT16 CL_UNSIGNED_INT32 CL_HALF_FLOAT CL_FLOAT

**Table 5.8.b** *Min. list of supported image formats – kernel read and write*

### 5.3.2.2 Image format mapping to OpenCL C image access qualifiers

Image arguments to kernels may have the `read_only`, `write_only` or `read_write` qualifier. Not all image formats supported by the device and platform are valid to be passed to all of these access qualifiers. For each access qualifier, only images whose format is in the list of formats returned by `clGetSupportedImageFormats` with the given flag arguments in *table 5.9* are permitted. It is not valid to pass an image supporting writing as both a `read_only` image and a `write_only` image parameter, or to a `read_write` image parameter and any other image parameter.

Access Qualifier	cl_mem_flags
<b>read_only</b>	CL_MEM_READ_ONLY, CL_MEM_READ_WRITE, CL_MEM_KERNEL_READ_AND_WRITE
<b>write_only</b>	CL_MEM_WRITE_ONLY, CL_MEM_READ_WRITE, CL_MEM_KERNEL_READ_AND_WRITE
<b>read_write</b>	CL_MEM_KERNEL_READ_AND_WRITE

**Table 5.9** *Mapping from format flags passed to `clGetSupportedImageFormats` to OpenCL C image access qualifiers*

### 5.3.3 Reading, Writing and Copying Image Objects

The following functions enqueue commands to read from an image or image array object to host memory or write to an image or image array object from host memory.

```
cl_int  clEnqueueReadImage (cl_command_queue command_queue,
                             cl_mem image,
                             cl_bool blocking_read,
                             const size_t *origin,
                             const size_t *region,
                             size_t row_pitch,
                             size_t slice_pitch,
                             void *ptr,
                             cl_uint num_events_in_wait_list,
                             const cl_event *event_wait_list,
                             cl_event *event)
```

```
cl_int  clEnqueueWriteImage (cl_command_queue command_queue,
                              cl_mem image,
                              cl_bool blocking_write,
                              const size_t *origin,
                              const size_t *region,
                              size_t input_row_pitch,
                              size_t input_slice_pitch,
                              const void * ptr,
                              cl_uint num_events_in_wait_list,
                              const cl_event *event_wait_list,
                              cl_event *event)
```

*command\_queue* refers to the host command-queue in which the read / write command will be queued. *command\_queue* and *image* must be created with the same OpenCL context.

*image* refers to a valid image or image array object.

*blocking\_read* and *blocking\_write* indicate if the read and write operations are *blocking* or *non-blocking*.

If *blocking\_read* is CL\_TRUE i.e. the read command is blocking, **clEnqueueReadImage** does not return until the buffer data has been read and copied into memory pointed to by *ptr*.

If *blocking\_read* is CL\_FALSE i.e. the read command is non-blocking, **clEnqueueReadImage** queues a non-blocking read command and returns. The contents of the buffer that *ptr* points to cannot be used until the read command has completed. The *event* argument returns an event object which can be used to query the execution status of the read command. When the read command has completed, the contents of the buffer that *ptr* points to can be used by the application.

If *blocking\_write* is CL\_TRUE, the OpenCL implementation copies the data referred to by *ptr* and enqueues the write command in the command-queue. The memory pointed to by *ptr* can be reused by the application after the **clEnqueueWriteImage** call returns.

If *blocking\_write* is CL\_FALSE, the OpenCL implementation will use *ptr* to perform a non-blocking write. As the write is non-blocking the implementation can return immediately. The memory pointed to by *ptr* cannot be reused by the application after the call returns. The *event* argument returns an event object which can be used to query the execution status of the write command. When the write command has completed, the memory pointed to by *ptr* can then be reused by the application.

*origin* defines the (*x*, *y*, *z*) offset in pixels in the 1D, 2D or 3D image, the (*x*, *y*) offset and the image index in the 2D image array or the (*x*) offset and the image index in the 1D image array. If *image* is a 2D image object, *origin*[2] must be 0. If *image* is a 1D image or 1D image buffer object, *origin*[1] and *origin*[2] must be 0. If *image* is a 1D image array object, *origin*[2] must be 0. If *image* is a 1D image array object, *origin*[1] describes the image index in the 1D image array. If *image* is a 2D image array object, *origin*[2] describes the image index in the 2D image array.

*region* defines the (*width*, *height*, *depth*) in pixels of the 1D, 2D or 3D rectangle, the (*width*, *height*) in pixels of the 2D rectangle and the number of images of a 2D image array or the (*width*) in pixels of the 1D rectangle and the number of images of a 1D image array. If *image* is a 2D image object, *region*[2] must be 1. If *image* is a 1D image or 1D image buffer object, *region*[1] and *region*[2] must be 1. If *image* is a 1D image array object, *region*[2] must be 1. The values in *region* cannot be 0.

*row\_pitch* in **clEnqueueReadImage** and *input\_row\_pitch* in **clEnqueueWriteImage** is the length of each row in bytes. This value must be greater than or equal to the element size in bytes \* *width*. If *row\_pitch* (or *input\_row\_pitch*) is set to 0, the appropriate row pitch is calculated based on the size of each element in bytes multiplied by *width*.

*slice\_pitch* in **clEnqueueReadImage** and *input\_slice\_pitch* in **clEnqueueWriteImage** is the size in bytes of the 2D slice of the 3D region of a 3D image or each image of a 1D or 2D image array being read or written respectively. This must be 0 if *image* is a 1D or 2D image. Otherwise this value must be greater than or equal to *row\_pitch* \* *height*. If *slice\_pitch* (or *input\_slice\_pitch*) is set to 0, the appropriate slice pitch is calculated based on the *row\_pitch* \* *height*.

*ptr* is the pointer to a buffer in host memory where image data is to be read from or to be written to.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in

*event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular read / write command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueReadImage** and **clEnqueueWriteImage** return CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue* and *image* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *image* is not a valid image object.
- ✚ CL\_INVALID\_VALUE if the region being read or written specified by *origin* and *region* is out of bounds or if *ptr* is a NULL value.
- ✚ CL\_INVALID\_VALUE if values in *origin* and *region* do not follow rules described in the argument description for *origin* and *region*.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_INVALID\_IMAGE\_SIZE if image dimensions (image width, height, specified or compute row and/or slice pitch) for *image* are not supported by device associated with *queue*.
- ✚ CL\_IMAGE\_FORMAT\_NOT\_SUPPORTED if image format (image channel order and data type) for *image* are not supported by device associated with *queue*.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *image*.
- ✚ CL\_INVALID\_OPERATION if the device associated with *command\_queue* does not support images (i.e. CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_FALSE).
- ✚ CL\_INVALID\_OPERATION if **clEnqueueReadImage** is called on *image* which has been created with CL\_MEM\_HOST\_WRITE\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS.

- ✚ CL\_INVALID\_OPERATION if **clEnqueueWriteImage** is called on *image* which has been created with CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS.
- ✚ CL\_EXEC\_STATUS\_ERROR\_FOR\_EVENTS\_IN\_WAIT\_LIST if the read and write operations are blocking and the execution status of any of the events in *event\_wait\_list* is a negative integer value.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

NOTE:

Calling **clEnqueueReadImage** to read a region of the *image* with the *ptr* argument value set to  $host\_ptr + (origin[2] * image\ slice\ pitch + origin[1] * image\ row\ pitch + origin[0] * bytes\ per\ pixel)$ , where *host\_ptr* is a pointer to the memory region specified when the *image* being read is created with CL\_MEM\_USE\_HOST\_PTR, must meet the following requirements in order to avoid undefined behavior:

- All commands that use this image object have finished execution before the read command begins execution.
- The *row\_pitch* and *slice\_pitch* argument values in **clEnqueueReadImage** must be set to the image row pitch and slice pitch.
- The image object is not mapped.
- The image object is not used by any command-queue until the read command has finished execution.

Calling **clEnqueueWriteImage** to update the latest bits in a region of the *image* with the *ptr* argument value set to  $host\_ptr + (origin[2] * image\ slice\ pitch + origin[1] * image\ row\ pitch + origin[0] * bytes\ per\ pixel)$ , where *host\_ptr* is a pointer to the memory region specified when the *image* being written is created with CL\_MEM\_USE\_HOST\_PTR, must meet the following requirements in order to avoid undefined behavior:

- The host memory region being written contains the latest bits when the enqueued write command begins execution.
- The *input\_row\_pitch* and *input\_slice\_pitch* argument values in **clEnqueueWriteImage** must be set to the image row pitch and slice pitch.
- The image object is not mapped.



- The image object is not used by any command-queue until the write command has finished execution.

The function

```
cl_int  clEnqueueCopyImage (cl_command_queue command_queue,
                             cl_mem src_image,
                             cl_mem dst_image,
                             const size_t *src_origin,
                             const size_t *dst_origin,
                             const size_t *region,
                             cl_uint num_events_in_wait_list,
                             const cl_event *event_wait_list,
                             cl_event *event)
```

enqueues a command to copy image objects. *src\_image* and *dst\_image* can be 1D, 2D, 3D image or a 1D, 2D image array objects. It is possible to copy subregions between any combinations of source and destination types, provided that the dimensions of the subregions are the same e.g., one can copy a rectangular region from a 2D image to a slice of a 3D image.

*command\_queue* refers to the host command-queue in which the copy command will be queued. The OpenCL context associated with *command\_queue*, *src\_image* and *dst\_image* must be the same.

*src\_origin* defines the (x, y, z) offset in pixels in the 1D, 2D or 3D image, the (x, y) offset and the image index in the 2D image array or the (x) offset and the image index in the 1D image array. If *image* is a 2D image object, *src\_origin*[2] must be 0. If *src\_image* is a 1D image object, *src\_origin*[1] and *src\_origin*[2] must be 0. If *src\_image* is a 1D image array object, *src\_origin*[2] must be 0. If *src\_image* is a 1D image array object, *src\_origin*[1] describes the image index in the 1D image array. If *src\_image* is a 2D image array object, *src\_origin*[2] describes the image index in the 2D image array.

*dst\_origin* defines the (x, y, z) offset in pixels in the 1D, 2D or 3D image, the (x, y) offset and the image index in the 2D image array or the (x) offset and the image index in the 1D image array. If *dst\_image* is a 2D image object, *dst\_origin*[2] must be 0. If *dst\_image* is a 1D image or 1D image buffer object, *dst\_origin*[1] and *dst\_origin*[2] must be 0. If *dst\_image* is a 1D image array object, *dst\_origin*[2] must be 0. If *dst\_image* is a 1D image array object, *dst\_origin*[1] describes the image index in the 1D image array. If *dst\_image* is a 2D image array object, *dst\_origin*[2] describes the image index in the 2D image array.

*region* defines the (*width*, *height*, *depth*) in pixels of the 1D, 2D or 3D rectangle, the (*width*, *height*) in pixels of the 2D rectangle and the number of images of a 2D image array or the (*width*) in pixels of the 1D rectangle and the number of images of a 1D image array. If *src\_image* or *dst\_image* is a 2D image object, *region*[2] must be 1. If *src\_image* or *dst\_image* is

a 1D image or 1D image buffer object, *region*[1] and *region*[2] must be 1. If *src\_image* or *dst\_image* is a 1D image array object, *region*[2] must be 1. The values in *region* cannot be 0.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular copy command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

It is currently a requirement that the *src\_image* and *dst\_image* image memory objects for **clEnqueueCopyImage** must have the exact same image format (i.e. the *cl\_image\_format* descriptor specified when *src\_image* and *dst\_image* are created must match).

**clEnqueueCopyImage** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue*, *src\_image* and *dst\_image* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *src\_image* and *dst\_image* are not valid image objects.
- ✚ CL\_IMAGE\_FORMAT\_MISMATCH if *src\_image* and *dst\_image* do not use the same image format.
- ✚ CL\_INVALID\_VALUE if the 2D or 3D rectangular region specified by *src\_origin* and *src\_origin* + *region* refers to a region outside *src\_image*, or if the 2D or 3D rectangular region specified by *dst\_origin* and *dst\_origin* + *region* refers to a region outside *dst\_image*.
- ✚ CL\_INVALID\_VALUE if values in *src\_origin*, *dst\_origin* and *region* do not follow rules described in the argument description for *src\_origin*, *dst\_origin* and *region*.

- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_INVALID\_IMAGE\_SIZE if image dimensions (image width, height, specified or compute row and/or slice pitch) for *src\_image* or *dst\_image* are not supported by device associated with *queue*.
- ✚ CL\_IMAGE\_FORMAT\_NOT\_SUPPORTED if image format (image channel order and data type) for *src\_image* or *dst\_image* are not supported by device associated with *queue*.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *src\_image* or *dst\_image*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.
- ✚ CL\_INVALID\_OPERATION if the device associated with *command\_queue* does not support images (i.e. CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_FALSE).
- ✚ CL\_MEM\_COPY\_OVERLAP if *src\_image* and *dst\_image* are the same image object and the source and destination regions overlap.

## 5.3.4 Filling Image Objects

The function

```
cl_int  clEnqueueFillImage (cl_command_queue command_queue,
                           cl_mem image,
                           const void *fill_color,
                           const size_t *origin,
                           const size_t *region,
                           cl_uint num_events_in_wait_list,
                           const cl_event *event_wait_list,
                           cl_event *event)
```

enqueues a command to fill an image object with a specified color. The usage information which indicates whether the memory object can be read or written by a kernel and/or the host and is given by the `cl_mem_flags` argument value specified when *image* is created is ignored by **clEnqueueFillImage**.

*command\_queue* refers to the host command-queue in which the fill command will be queued. The OpenCL context associated with *command\_queue* and *image* must be the same.

*image* is a valid image object.

*fill\_color* is the color used to fill the image. The fill color is a single floating point value if the channel order is `CL_DEPTH`. Otherwise, the fill color is a four component RGBA floating-point color value if the *image* channel data type is not an unnormalized signed or unsigned integer type, is a four component signed integer value if the *image* channel data type is an unnormalized signed integer type and is a four component unsigned integer value if the *image* channel data type is an unnormalized unsigned integer type. The fill color will be converted to the appropriate image channel format and order associated with *image* as described in *sections 6.12.14 and 8.3*.

*origin* defines the (*x*, *y*, *z*) offset in pixels in the 1D, 2D or 3D image, the (*x*, *y*) offset and the image index in the 2D image array or the (*x*) offset and the image index in the 1D image array. If *image* is a 2D image object, *origin*[2] must be 0. If *image* is a 1D image or 1D image buffer object, *origin*[1] and *origin*[2] must be 0. If *image* is a 1D image array object, *origin*[2] must be 0. If *image* is a 1D image array object, *origin*[1] describes the image index in the 1D image array. If *image* is a 2D image array object, *origin*[2] describes the image index in the 2D image array.

*region* defines the (*width*, *height*, *depth*) in pixels of the 1D, 2D or 3D rectangle, the (*width*, *height*) in pixels of the 2D rectangle and the number of images of a 2D image array or the (*width*) in pixels of the 1D rectangle and the number of images of a 1D image array. If *image* is a 2D image object, *region*[2] must be 1. If *image* is a 1D image or 1D image buffer object, *region*[1] and *region*[2] must be 1. If *image* is a 1D image array object, *region*[2] must be 1. The values in *region* cannot be 0.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueFillImage** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue* and *image* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *image* is not a valid image object.
- ✚ CL\_INVALID\_VALUE if *fill\_color* is NULL.
- ✚ CL\_INVALID\_VALUE if the region being filled as specified by *origin* and *region* is out of bounds or if *ptr* is a NULL value.
- ✚ CL\_INVALID\_VALUE if values in *origin* and *region* do not follow rules described in the argument description for *origin* and *region*.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_INVALID\_IMAGE\_SIZE if image dimensions (image width, height, specified or compute row and/or slice pitch) for *image* are not supported by device associated with *queue*.
- ✚ CL\_IMAGE\_FORMAT\_NOT\_SUPPORTED if image format (image channel order and data type) for *image* are not supported by device associated with *queue*.

- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *image*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

### 5.3.5 Copying between Image and Buffer Objects

The function

```
cl_int  clEnqueueCopyImageToBuffer (cl_command_queue command_queue,
                                     cl_mem src_image,
                                     cl_mem dst_buffer,
                                     const size_t *src_origin,
                                     const size_t *region,
                                     size_t dst_offset,
                                     cl_uint num_events_in_wait_list,
                                     const cl_event *event_wait_list,
                                     cl_event *event)
```

enqueues a command to copy an image object to a buffer object.

*command\_queue* must be a valid host command-queue. The OpenCL context associated with *command\_queue*, *src\_image* and *dst\_buffer* must be the same.

*src\_image* is a valid image object.

*dst\_buffer* is a valid buffer object.

*src\_origin* defines the (*x*, *y*, *z*) offset in pixels in the 1D, 2D or 3D image, the (*x*, *y*) offset and the image index in the 2D image array or the (*x*) offset and the image index in the 1D image array. If *src\_image* is a 2D image object, *src\_origin*[2] must be 0. If *src\_image* is a 1D image or 1D image buffer object, *src\_origin*[1] and *src\_origin*[2] must be 0. If *src\_image* is a 1D image array object, *src\_origin*[2] must be 0. If *src\_image* is a 1D image array object, *src\_origin*[1] describes the image index in the 1D image array. If *src\_image* is a 2D image array object, *src\_origin*[2] describes the image index in the 2D image array.

*region* defines the (*width*, *height*, *depth*) in pixels of the 1D, 2D or 3D rectangle, the (*width*, *height*) in pixels of the 2D rectangle and the number of images of a 2D image array or the (*width*) in pixels of the 1D rectangle and the number of images of a 1D image array. If *src\_image* is a 2D image object, *region*[2] must be 1. If *src\_image* is a 1D image or 1D image

buffer object, *region*[1] and *region*[2] must be 1. If *src\_image* is a 1D image array object, *region*[2] must be 1. The values in *region* cannot be 0.

*dst\_offset* refers to the offset where to begin copying data into *dst\_buffer*. The size in bytes of the region to be copied referred to as *dst\_cb* is computed as *width \* height \* depth \* bytes/image element* if *src\_image* is a 3D image object, is computed as *width \* height \* bytes/image element* if *src\_image* is a 2D image, is computed as *width \* height \* arraysize \* bytes/image element* if *src\_image* is a 2D image array object, is computed as *width \* bytes/image element* if *src\_image* is a 1D image or 1D image buffer object and is computed as *width \* arraysize \* bytes/image element* if *src\_image* is a 1D image array object.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular copy command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueCopyImageToBuffer** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue*, *src\_image* and *dst\_buffer* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *src\_image* is not a valid image object or *dst\_buffer* is not a valid buffer object or if *src\_image* is a 1D image buffer object created from *dst\_buffer*.
- ✚ CL\_INVALID\_VALUE if the 1D, 2D or 3D rectangular region specified by *src\_origin* and *src\_origin + region* refers to a region outside *src\_image*, or if the region specified by *dst\_offset* and *dst\_offset + dst\_cb* to a region outside *dst\_buffer*.
- ✚ CL\_INVALID\_VALUE if values in *src\_origin* and *region* do not follow rules described in the argument description for *src\_origin* and *region*.

- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *dst\_buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_INVALID\_IMAGE\_SIZE if image dimensions (image width, height, specified or compute row and/or slice pitch) for *src\_image* are not supported by device associated with *queue*.
- ✚ CL\_IMAGE\_FORMAT\_NOT\_SUPPORTED if image format (image channel order and data type) for *src\_image* are not supported by device associated with *queue*.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *src\_image* or *dst\_buffer*.
- ✚ CL\_INVALID\_OPERATION if the device associated with *command\_queue* does not support images (i.e. CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_FALSE).
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clEnqueueCopyBufferToImage (cl_command_queue command_queue,
                                     cl_mem src_buffer,
                                     cl_mem dst_image,
                                     size_t src_offset,
                                     const size_t *dst_origin,
                                     const size_t *region,
                                     cl_uint num_events_in_wait_list,
                                     const cl_event *event_wait_list,
                                     cl_event *event)
```

enqueues a command to copy a buffer object to an image object.

*command\_queue* must be a valid host command-queue. The OpenCL context associated with *command\_queue*, *src\_buffer* and *dst\_image* must be the same.

*src\_buffer* is a valid buffer object.



*dst\_image* is a valid image object.

*src\_offset* refers to the offset where to begin copying data from *src\_buffer*.

*dst\_origin* defines the (x, y, z) offset in pixels in the 1D, 2D or 3D image, the (x, y) offset and the image index in the 2D image array or the (x) offset and the image index in the 1D image array. If *dst\_image* is a 2D image object, *dst\_origin*[2] must be 0. If *dst\_image* is a 1D image or 1D image buffer object, *dst\_origin*[1] and *dst\_origin*[2] must be 0. If *dst\_image* is a 1D image array object, *dst\_origin*[2] must be 0. If *dst\_image* is a 1D image array object, *dst\_origin*[1] describes the image index in the 1D image array. If *dst\_image* is a 2D image array object, *dst\_origin*[2] describes the image index in the 2D image array.


*region* defines the (width, height, depth) in pixels of the 1D, 2D or 3D rectangle, the (width, height) in pixels of the 2D rectangle and the number of images of a 2D image array or the (width) in pixels of the 1D rectangle and the number of images of a 1D image array. If *dst\_image* is a 2D image object, *region*[2] must be 1. If *dst\_image* is a 1D image or 1D image buffer object, *region*[1] and *region*[2] must be 1. If *dst\_image* is a 1D image array object, *region*[2] must be 1. The values in *region* cannot be 0.

The size in bytes of the region to be copied from *src\_buffer* referred to as *src\_cb* is computed as  $width * height * depth * bytes/image\ element$  if *dst\_image* is a 3D image object, is computed as  $width * height * bytes/image\ element$  if *dst\_image* is a 2D image, is computed as  $width * height * arraysize * bytes/image\ element$  if *dst\_image* is a 2D image array object, is computed as  $width * bytes/image\ element$  if *dst\_image* is a 1D image or 1D image buffer object and is computed as  $width * arraysize * bytes/image\ element$  if *dst\_image* is a 1D image array object.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular copy command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueCopyBufferToImage** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.

- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue*, *src\_buffer* and *dst\_image* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if *src\_buffer* is not a valid buffer object or *dst\_image* is not a valid image object or if *dst\_image* is a 1D image buffer object created from *src\_buffer*.
- ✚ CL\_INVALID\_VALUE if the 1D, 2D or 3D rectangular region specified by *dst\_origin* and *dst\_origin + region* refer to a region outside *dst\_image*, or if the region specified by *src\_offset* and *src\_offset + src\_cb* refer to a region outside *src\_buffer*.
- ✚ CL\_INVALID\_VALUE if values in *dst\_origin* and *region* do not follow rules described in the argument description for *dst\_origin* and *region*.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if *src\_buffer* is a sub-buffer object and *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_INVALID\_IMAGE\_SIZE if image dimensions (image width, height, specified or compute row and/or slice pitch) for *dst\_image* are not supported by device associated with *queue*.
- ✚ CL\_IMAGE\_FORMAT\_NOT\_SUPPORTED if image format (image channel order and data type) for *dst\_image* are not supported by device associated with *queue*.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *src\_buffer* or *dst\_image*.
- ✚ CL\_INVALID\_OPERATION if the device associated with *command\_queue* does not support images (i.e. CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_FALSE).
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

### 5.3.6 Mapping Image Objects

The function

```
void * clEnqueueMapImage (cl_command_queue command_queue,  
                          cl_mem image,  
                          cl_bool blocking_map,  
                          cl_map_flags map_flags,  
                          const size_t *origin,  
                          const size_t *region,  
                          size_t *image_row_pitch,  
                          size_t *image_slice_pitch,  
                          cl_uint num_events_in_wait_list,  
                          const cl_event *event_wait_list,  
                          cl_event *event,  
                          cl_int *errcode_ret)
```

enqueues a command to map a region in the image object given by *image* into the host address space and returns a pointer to this mapped region.

*command\_queue* must be a valid host command-queue.

*image* is a valid image object. The OpenCL context associated with *command\_queue* and *image* must be the same.

*blocking\_map* indicates if the map operation is *blocking* or *non-blocking*.

If *blocking\_map* is CL\_TRUE, **clEnqueueMapImage** does not return until the specified region in *image* is mapped into the host address space and the application can access the contents of the mapped region using the pointer returned by **clEnqueueMapImage**.

If *blocking\_map* is CL\_FALSE i.e. map operation is non-blocking, the pointer to the mapped region returned by **clEnqueueMapImage** cannot be used until the map command has completed. The *event* argument returns an event object which can be used to query the execution status of the map command. When the map command is completed, the application can access the contents of the mapped region using the pointer returned by **clEnqueueMapImage**.

*map\_flags* is a bit-field and is described in *table 5.5*.

*origin* defines the (*x*, *y*, *z*) offset in pixels in the 1D, 2D or 3D image, the (*x*, *y*) offset and the image index in the 2D image array or the (*x*) offset and the image index in the 1D image array. If *image* is a 2D image object, *origin*[2] must be 0. If *image* is a 1D image or 1D image buffer object, *origin*[1] and *origin*[2] must be 0. If *image* is a 1D image array object, *origin*[2] must be 0. If *image* is a 1D image array object, *origin*[1] describes the image index in the 1D image

array. If *image* is a 2D image array object, *origin*[2] describes the image index in the 2D image array.

*region* defines the (*width*, *height*, *depth*) in pixels of the 1D, 2D or 3D rectangle, the (*width*, *height*) in pixels of the 2D rectangle and the number of images of a 2D image array or the (*width*) in pixels of the 1D rectangle and the number of images of a 1D image array. If *image* is a 2D image object, *region*[2] must be 1. If *image* is a 1D image or 1D image buffer object, *region*[1] and *region*[2] must be 1. If *image* is a 1D image array object, *region*[2] must be 1. The values in *region* cannot be 0.

*image\_row\_pitch* returns the scan-line pitch in bytes for the mapped region. This must be a non-NULL value.

*image\_slice\_pitch* returns the size in bytes of each 2D slice of a 3D image or the size of each 1D or 2D image in a 1D or 2D image array for the mapped region. For a 1D and 2D image, zero is returned if this argument is not NULL. For a 3D image, 1D and 2D image array, *image\_slice\_pitch* must be a non-NULL value.



*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before **clEnqueueMapImage** can be executed. If *event\_wait\_list* is NULL, then **clEnqueueMapImage** does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clEnqueueMapImage** will return a pointer to the mapped region. The *errcode\_ret* is set to CL\_SUCCESS.

A NULL pointer is returned otherwise with one of the following error values returned in *errcode\_ret*:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_INVALID\_CONTEXT if context associated with *command\_queue* and *image* are not

the same or if context associated with *command\_queue* and events in *event\_wait\_list* are not the same.

- ✚ CL\_INVALID\_MEM\_OBJECT if *image* is not a valid image object.
- ✚ CL\_INVALID\_VALUE if region being mapped given by (*origin*, *origin+region*) is out of bounds or if values specified in *map\_flags* are not valid.
- ✚ CL\_INVALID\_VALUE if values in *origin* and *region* do not follow rules described in the argument description for *origin* and *region*.
- ✚ CL\_INVALID\_VALUE if *image\_row\_pitch* is NULL.
- ✚ CL\_INVALID\_VALUE if *image* is a 3D image, 1D or 2D image array object and *image\_slice\_pitch* is NULL.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_INVALID\_IMAGE\_SIZE if image dimensions (image width, height, specified or compute row and/or slice pitch) for *image* are not supported by device associated with *queue*.
- ✚ CL\_IMAGE\_FORMAT\_NOT\_SUPPORTED if image format (image channel order and data type) for *image* are not supported by device associated with *queue*.
- ✚ CL\_MAP\_FAILURE if there is a failure to map the requested region into the host address space. This error cannot occur for image objects created with CL\_MEM\_USE\_HOST\_PTR or CL\_MEM\_ALLOC\_HOST\_PTR.
- ✚ CL\_EXEC\_STATUS\_ERROR\_FOR\_EVENTS\_IN\_WAIT\_LIST if the map operation is blocking and the execution status of any of the events in *event\_wait\_list* is a negative integer value.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with *image*.
- ✚ CL\_INVALID\_OPERATION if the device associated with *command\_queue* does not support images (i.e. CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_FALSE).
- ✚ CL\_INVALID\_OPERATION if *image* has been created with CL\_MEM\_HOST\_WRITE\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS and CL\_MAP\_READ is set in *map\_flags* or if *image* has been created with CL\_MEM\_HOST\_READ\_ONLY or CL\_MEM\_HOST\_NO\_ACCESS and CL\_MAP\_WRITE or CL\_MAP\_WRITE\_INVALIDATE\_REGION is set in *map\_flags*.

- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.
- ✚ CL\_INVALID\_OPERATION if mapping would lead to overlapping regions being mapped for writing.

The pointer returned maps a 1D, 2D or 3D region starting at *origin* and is at least *region[0]* pixels in size for a 1D image, 1D image buffer or 1D image array, (*image\_row\_pitch* \* *region[1]*) pixels in size for a 2D image or 2D image array, and (*image\_slice\_pitch* \* *region[2]*) pixels in size for a 3D image. The result of a memory access outside this region is undefined.

If the image object is created with CL\_MEM\_USE\_HOST\_PTR set in *mem\_flags*, the following will be true:

- ✚ The *host\_ptr* specified in **clCreateImage** is guaranteed to contain the latest bits in the region being mapped when the **clEnqueueMapImage** command has completed.
- ✚ The pointer value returned by **clEnqueueMapImage** will be derived from the *host\_ptr* specified when the image object is created.

Mapped image objects are unmapped using **clEnqueueUnmapMemObject**. This is described in *section 5.5.2*.

### 5.3.7 Image Object Queries

To get information that is common to all memory objects, use the **clGetMemObjectInfo** function described in *section 5.5.5*.

To get information specific to an image object created with **clCreateImage**, use the following function

```
cl_int          clGetImageInfo (cl_mem image,
                                   cl_image_info param_name,
                                   size_t param_value_size,
                                   void *param_value,
                                   size_t *param_value_size_ret)
```

*image* specifies the image object being queried.

*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetImageInfo** is described in *table 5.10*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.10*.

*param\_value\_size\_ret* returns the actual size in bytes of data being queried by *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

**clGetImageInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is  $<$  size of return type as described in *table 5.10* and *param\_value* is not NULL.
- ✚ CL\_INVALID\_MEM\_OBJECT if *image* is a not a valid image object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

cl_image_info	Return type	Info. returned in <i>param_value</i>
CL_IMAGE_FORMAT	cl_image_format	Return image format descriptor specified when <i>image</i> is created with <b>clCreateImage</b> .
CL_IMAGE_ELEMENT_SIZE	size_t	Return size of each element of the image memory object given by <i>image</i> in bytes. An element is made up of <i>n</i> channels. The value of <i>n</i> is given in <i>cl_image_format</i> descriptor.
CL_IMAGE_ROW_PITCH	size_t	Return calculated row pitch in bytes of a row of elements of the image object given by <i>image</i> .
CL_IMAGE_SLICE_PITCH	size_t	Return calculated slice pitch in bytes of a 2D slice for the 3D image object or size of each image in a 1D or 2D image array given by <i>image</i> . For a 1D image, 1D image buffer and 2D image object return 0.
CL_IMAGE_WIDTH	size_t	Return width of the image in pixels.
CL_IMAGE_HEIGHT	size_t	Return height of the image in pixels. For a 1D image, 1D image buffer and 1D image

		array object, height = 0.
<b>CL_IMAGE_DEPTH</b>	size_t	Return depth of the image in pixels. For a 1D image, 1D image buffer, 2D image or 1D and 2D image array object, depth = 0.
<b>CL_IMAGE_ARRAY_SIZE</b>	size_t	Return number of images in the image array. If <i>image</i> is not an image array, 0 is returned.
<b>CL_IMAGE_NUM_MIP_LEVELS</b>	cl_uint	Return num_mip_levels associated with <i>image</i> .
<b>CL_IMAGE_NUM_SAMPLES</b>	cl_uint	Return num_samples associated with <i>image</i> .

**Table 5.10**    *List of supported param\_names by clGetImageInfo*



## 5.4 Pipes

A *pipe* is a memory object that stores data organized as a FIFO. Pipe objects can only be accessed using built-in functions that read from and write to a pipe. Pipe objects are not accessible from the host. A pipe object encapsulates the following information:

- ✚ Packet size in bytes
- ✚ Maximum capacity in packets
- ✚ Information about the number of packets currently in the pipe
- ✚ Data packets

### 5.4.1 Creating Pipe Objects

A **pipe object** is created using the following function

```
cl_mem  clCreatePipe (cl_context context,
                    cl_mem_flags flags,
                    cl_uint pipe_packet_size,
                    cl_uint pipe_max_packets,
                    const cl_pipe_properties *properties,
                    cl_int *errcode_ret)
```

*context* is a valid OpenCL context used to create the pipe object.

*flags* is a bit-field that is used to specify allocation and usage information such as the memory arena that should be used to allocate the pipe object and how it will be used. *Table 5.3* describes the possible values for *flags*. Only CL\_MEM\_READ\_WRITE and CL\_MEM\_HOST\_NO\_ACCESS can be specified when creating a pipe object. If the value specified for *flags* is 0, the default is used which is CL\_MEM\_READ\_WRITE | CL\_MEM\_HOST\_NO\_ACCESS.

*pipe\_packet\_size* is the size in bytes of a pipe packet.

*pipe\_max\_packets* specifies the pipe capacity by specifying the maximum number of packets the pipe can hold.

*properties* specifies a list of properties for the pipe and their corresponding values. Each property name is immediately followed by the corresponding desired value. The list is terminated with 0. In OpenCL 2.0, *properties* must be NULL.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreatePipe** returns a valid non-zero pipe object and *errcode\_ret* is set to CL\_SUCCESS if the pipe object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid context.
- ✚ CL\_INVALID\_VALUE if values specified in *flags* are not as defined above.
- ✚ CL\_INVALID\_VALUE if *properties* is not NULL.
- ✚ CL\_INVALID\_PIPE\_SIZE if *pipe\_packet\_size* is 0 or the *pipe\_packet\_size* exceeds CL\_DEVICE\_PIPE\_MAX\_PACKET\_SIZE value specified in *table 4.3* for all devices in *context* or if *pipe\_max\_packets* is 0.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for the pipe object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

Pipes follow the same memory consistency model as defined for buffer and image objects. The pipe state i.e. contents of the pipe across kernel-instances (on the same or different devices) is enforced at a synchronization point.

## 5.4.2 Pipe Object Queries

To get information that is common to all memory objects, use the **clGetMemObjectInfo** function described in *section 5.5.5*.

To get information specific to a pipe object created with **clCreatePipe**, use the following function

```
cl_int          clGetPipeInfo (cl_mem pipe,
                               cl_pipe_info param_name,
                               size_t param_value_size,
                               void *param_value,
                               size_t *param_value_size_ret)
```

*pipe* specifies the pipe object being queried.





*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetPipeInfo** is described in *table 5.11*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.11*.

*param\_value\_size\_ret* returns the actual size in bytes of data being queried by *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

**clGetPipeInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is  $<$  size of return type as described in *table 5.11* and *param\_value* is not NULL.
-  CL\_INVALID\_MEM\_OBJECT if *pipe* is a not a valid pipe object.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

cl_pipe_info	Return type	Info. returned in <i>param_value</i>
CL_PIPE_PACKET_SIZE	cl_uint	Return pipe packet size specified when <i>pipe</i> is created with <b>clCreatePipe</b> .
CL_PIPE_MAX_PACKETS	cl_uint	Return max. number of packets specified when <i>pipe</i> is created with <b>clCreatePipe</b> .

**Table 5.11**    *List of supported param\_names by clGetPipeInfo*




## 5.5 Querying, Unmapping, Migrating, Retaining and Releasing Memory Objects

### 5.5.1 Retaining and Releasing Memory Objects

The function

`cl_int`                    **clRetainMemObject** (`cl_mem memobj`)

increments the *memobj* reference count. **clRetainMemObject** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:




-  CL\_INVALID\_MEM\_OBJECT if *memobj* is not a valid memory object (buffer or image object).
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

**clCreateBuffer**, **clCreateSubBuffer**, **clCreateImage** and **clCreatePipe** perform an implicit retain.

The function

`cl_int`                    **clReleaseMemObject** (`cl_mem memobj`)

decrements the *memobj* reference count. **clReleaseMemObject** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_MEM\_OBJECT if *memobj* is not a valid memory object.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

After the *memobj* reference count becomes zero and commands queued for execution on a command-queue(s) that use *memobj* have finished, the memory object is deleted. If *memobj* is a buffer object, *memobj* cannot be deleted until all sub-buffer objects associated with *memobj* are

deleted. Using this function to release a reference that was not obtained by creating the object or by calling **clRetainMemObject** causes undefined behavior.

The function

```
cl_int          clSetMemObjectDestructorCallback (cl_mem memobj,
                                                    void (CL_CALLBACK *pfn_notify)(cl_mem memobj,
                                                                    void *user_data),
                                                    void *user_data)
```

registers a user callback function with a memory object. Each call to **clSetMemObjectDestructorCallback** registers the specified user callback function on a callback stack associated with *memobj*. The registered user callback functions are called in the reverse order in which they were registered. The user callback functions are called and then the memory object's resources are freed and the memory object is deleted. This provides a mechanism for the application (and libraries) using *memobj* to be notified when the memory referenced by *host\_ptr*, specified when the memory object is created and used as the storage bits for the memory object, can be reused or freed.

*memobj* is a valid memory object.

*pfn\_notify* is the callback function that can be registered by the application. This callback function may be called asynchronously by the OpenCL implementation. It is the application's responsibility to ensure that the callback function is thread-safe. The parameters to this callback function are:

- ✚ *memobj* is the memory object being deleted. When the user callback is called by the implementation, this memory object is not longer valid. *memobj* is only provided for reference purposes.
- ✚ *user\_data* is a pointer to user supplied data.

*user\_data* will be passed as the *user\_data* argument when *pfn\_notify* is called. *user\_data* can be NULL.

**clSetMemObjectDestructorCallback** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_MEM\_OBJECT if *memobj* is not a valid memory object.
- ✚ CL\_INVALID\_VALUE if *pfn\_notify* is NULL.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

NOTE: When the user callback function is called by the implementation, the contents of the memory region pointed to by *host\_ptr* (if the memory object is created with CL\_MEM\_USE\_HOST\_PTR) are undefined. The callback function is typically used by the application to either free or reuse the memory region pointed to by *host\_ptr*.

The behavior of calling expensive system routines, OpenCL API calls to create contexts or command-queues, or blocking OpenCL operations from the following list below, in a callback is undefined.

**clFinish,**  
**clWaitForEvents,**  
blocking calls to **clEnqueueReadBuffer, clEnqueueReadBufferRect,**  
**clEnqueueWriteBuffer, clEnqueueWriteBufferRect,**  
blocking calls to **clEnqueueReadImage** and **clEnqueueWriteImage,**  
blocking calls to **clEnqueueMapBuffer,**  
**clEnqueueMapImage,**  
blocking calls to **clBuildProgram, clCompileProgram** or **clLinkProgram**

If an application needs to wait for completion of a routine from the above list in a callback, please use the non-blocking form of the function, and assign a completion callback to it to do the remainder of your work. Note that when a callback (or other code) enqueues commands to a command-queue, the commands are not required to begin execution until the queue is flushed. In standard usage, blocking enqueue calls serve this role by implicitly flushing the queue. Since blocking calls are not permitted in callbacks, those callbacks that enqueue commands on a command queue should either call **clFlush** on the queue before returning or arrange for **clFlush** to be called later on another thread.

The user callback function may not call OpenCL APIs with the memory object for which the callback function is invoked and for such cases the behavior of OpenCL APIs is considered to be undefined.

## 5.5.2 Unmapping Mapped Memory Objects

The function

```
cl_int  clEnqueueUnmapMemObject (cl_command_queue command_queue,
                                cl_mem memobj,
                                void *mapped_ptr,
                                cl_uint num_events_in_wait_list,
                                const cl_event *event_wait_list,
                                cl_event *event)
```

enqueues a command to unmap a previously mapped region of a memory object. Reads or writes from the host using the pointer returned by **clEnqueueMapBuffer** or **clEnqueueMapImage** are considered to be complete.

*command\_queue* must be a valid host command-queue.






*memobj* is a valid memory (buffer or image) object. The OpenCL context associated with *command\_queue* and *memobj* must be the same.

*mapped\_ptr* is the host address returned by a previous call to **clEnqueueMapBuffer**, or **clEnqueueMapImage** for *memobj*.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before **clEnqueueUnmapMemObject** can be executed. If *event\_wait\_list* is NULL, then **clEnqueueUnmapMemObject** does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueUnmapMemObject** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_INVALID\_MEM\_OBJECT if *memobj* is not a valid memory object or is a pipe object.
-  CL\_INVALID\_VALUE if *mapped\_ptr* is not a valid pointer returned by **clEnqueueMapBuffer**, or **clEnqueueMapImage** for *memobj*.
-  CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or if *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.

- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.
- ✚ CL\_INVALID\_CONTEXT if context associated with *command\_queue* and *memobj* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.

**clEnqueueMapBuffer**, and **clEnqueueMapImage** increments the mapped count of the memory object. The initial mapped count value of the memory object is zero. Multiple calls to **clEnqueueMapBuffer**, or **clEnqueueMapImage** on the same memory object will increment this mapped count by appropriate number of calls. **clEnqueueUnmapMemObject** decrements the mapped count of the memory object.

**clEnqueueMapBuffer**, and **clEnqueueMapImage** act as synchronization points for a region of the buffer object being mapped.

### 5.5.3 Accessing mapped regions of a memory object

This section describes the behavior of OpenCL commands that access mapped regions of a memory object.

The contents of the region of a memory object and associated memory objects (sub-buffer objects or 1D image buffer objects that overlap this region) mapped for writing (i.e. CL\_MAP\_WRITE or CL\_MAP\_WRITE\_INVALIDATE\_REGION is set in *map\_flags* argument to **clEnqueueMapBuffer**, or **clEnqueueMapImage**) are considered to be undefined until this region is unmapped.

Multiple commands in command-queues can map a region or overlapping regions of a memory object and associated memory objects (sub-buffer objects or 1D image buffer objects that overlap this region) for reading (i.e. *map\_flags* = CL\_MAP\_READ). The contents of the regions of a memory object mapped for reading can also be read by kernels and other OpenCL commands (such as **clEnqueueCopyBuffer**) executing on a device(s).

Mapping (and unmapping) overlapped regions in a memory object and/or associated memory objects (sub-buffer objects or 1D image buffer objects that overlap this region) for writing is an error and will result in CL\_INVALID\_OPERATION error returned by **clEnqueueMapBuffer**, or **clEnqueueMapImage**.

If a memory object is currently mapped for writing, the application must ensure that the memory object is unmapped before any enqueued kernels or commands that read from or write to this memory object or any of its associated memory objects (sub-buffer or 1D image buffer objects) or its parent object (if the memory object is a sub-buffer or 1D image buffer object) begin execution; otherwise the behavior is undefined.

If a memory object is currently mapped for reading, the application must ensure that the memory



object is unmapped before any enqueued kernels or commands that write to this memory object or any of its associated memory objects (sub-buffer or 1D image buffer objects) or its parent object (if the memory object is a sub-buffer or 1D image buffer object) begin execution; otherwise the behavior is undefined.

A memory object is considered as mapped if there are one or more active mappings for the memory object irrespective of whether the mapped regions span the entire memory object.

Accessing the contents of the memory region referred to by the mapped pointer that has been unmapped is undefined.

The mapped pointer returned by **clEnqueueMapBuffer** or **clEnqueueMapImage** can be used as *ptr* argument value to **clEnqueue{Read | Write}Buffer**, **clEnqueue{Read | Write}BufferRect**, **clEnqueue{Read | Write}Image** provided the rules described above are adhered to.

## 5.5.4 Migrating Memory Objects

This section describes a mechanism for assigning which device an OpenCL memory object resides. A user may wish to have more explicit control over the location of their memory objects on creation. This could be used to:

- ✚ Ensure that an object is allocated on a specific device prior to usage.
- ✚ Preemptively migrate an object from one device to another.

The function

```
cl_int  clEnqueueMigrateMemObjects (cl_command_queue command_queue,
                                     cl_uint num_mem_objects,
                                     const cl_mem *mem_objects,
                                     cl_mem_migration_flags flags,
                                     cl_uint num_events_in_wait_list,
                                     const cl_event *event_wait_list,
                                     cl_event *event)
```

enqueues a command to indicate which device a set of memory objects should be associated with. Typically, memory objects are implicitly migrated to a device for which enqueued commands, using the memory object, are targeted. **clEnqueueMigrateMemObjects** allows this migration to be explicitly performed ahead of the dependent commands. This allows a user to preemptively change the association of a memory object, through regular command queue scheduling, in order to prepare for another upcoming command. This also permits an application to overlap the placement of memory objects with other unrelated operations before these memory objects are needed potentially hiding transfer latencies. Once the event, returned from **clEnqueueMigrateMemObjects**, has been marked CL\_COMPLETE the memory objects specified in *mem\_objects* have been successfully migrated to the device associated with

*command\_queue*. The migrated memory object shall remain resident on the device until another command is enqueued that either implicitly or explicitly migrates it away.

**clEnqueueMigrateMemObjects** can also be used to direct the initial placement of a memory object, after creation, possibly avoiding the initial overhead of instantiating the object on the first enqueued command to use it.

The user is responsible for managing the event dependencies, associated with this command, in order to avoid overlapping access to memory objects. Improperly specified event dependencies passed to **clEnqueueMigrateMemObjects** could result in undefined results.

*command\_queue* is a valid host command-queue. The specified set of memory objects in *mem\_objects* will be migrated to the OpenCL device associated with *command\_queue* or to the host if the CL\_MIGRATE\_MEM\_OBJECT\_HOST has been specified.

*num\_mem\_objects* is the number of memory objects specified in *mem\_objects*.

*mem\_objects* is a pointer to a list of memory objects.

*flags* is a bit-field that is used to specify migration options. The *table 5.12* describes the possible values for flags.

cl_mem_migration flags	Description
CL_MIGRATE_MEM_OBJECT_HOST	This flag indicates that the specified set of memory objects are to be migrated to the host, regardless of the target command-queue.
CL_MIGRATE_MEM_OBJECT_CONTENT_UNDEFINED	This flag indicates that the contents of the set of memory objects are undefined after migration. The specified set of memory objects are migrated to the device associated with <i>command_queue</i> without incurring the overhead of migrating their contents.

**Table 5.12** Supported *cl\_mem\_migration flags*.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will

not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueMigrateMemObjects** return CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue* and memory objects in *mem\_objects* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_MEM\_OBJECT if any of the memory objects in *mem\_objects* is not a valid memory object.
- ✚ CL\_INVALID\_VALUE if *num\_mem\_objects* is zero or if *mem\_objects* is NULL.
- ✚ CL\_INVALID\_VALUE if *flags* is not 0 or is not any of the values described in the table above.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for the specified set of memory objects in *mem\_objects*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 5.5.5 Memory Object Queries

To get information that is common to all memory objects (buffer and image objects), use the following function

```
cl_int          clGetMemObjectInfo (cl_mem memobj,
                                     cl_mem_info param_name,
                                     size_t param_value_size,
                                     void *param_value,
```

size\_t \*param\_value\_size\_ret)

*memobj* specifies the memory object being queried.

*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetMemObjectInfo** is described in *table 5.13*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be >= size of return type as described in *table 5.13*.

*param\_value\_size\_ret* returns the actual size in bytes of data being queried by *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

**clGetMemObjectInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is < size of return type as described in *table 5.13* and *param\_value* is not NULL.
- ✚ CL\_INVALID\_MEM\_OBJECT if *memobj* is a not a valid memory object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

cl_mem_info	Return type	Info. returned in <i>param_value</i>
CL_MEM_TYPE	cl_mem_object_type	Returns one of the following values:  CL_MEM_OBJECT_BUFFER if <i>memobj</i> is created with <b>clCreateBuffer</b> or <b>clCreateSubBuffer</b> .  cl_image_desc.image_type argument value if <i>memobj</i> is created with <b>clCreateImage</b> .  CL_MEM_OBJECT_PIPE if <i>memobj</i> is created with <b>clCreatePipe</b> .
CL_MEM_FLAGS	cl_mem_flags	Return the <i>flags</i> argument value specified

		<p>when <i>memobj</i> is created with <b>clCreateBuffer</b>, <b>clCreateSubBuffer</b>, <b>clCreateImage</b> or <b>clCreatePipe</b>.</p> <p>If <i>memobj</i> is a sub-buffer the memory access qualifiers inherited from parent buffer is also returned.</p>
<b>CL_MEM_SIZE</b>	size_t	Return actual size of the data store associated with <i>memobj</i> in bytes.
<b>CL_MEM_HOST_PTR</b>	void *	<p>If <i>memobj</i> is created with <b>clCreateBuffer</b> or <b>clCreateImage</b> and CL_MEM_USE_HOST_PTR is specified in <i>mem_flags</i>, return the <i>host_ptr</i> argument value specified when <i>memobj</i> is created. Otherwise a NULL value is returned.</p> <p>If <i>memobj</i> is created with <b>clCreateSubBuffer</b>, return the <i>host_ptr</i> + <i>origin</i> value specified when <i>memobj</i> is created. <i>host_ptr</i> is the argument value specified to <b>clCreateBuffer</b> and CL_MEM_USE_HOST_PTR is specified in <i>mem_flags</i> for memory object from which <i>memobj</i> is created. Otherwise a NULL value is returned.</p>
<b>CL_MEM_MAP_COUNT</b> <sup>23</sup>	cl_uint	Map count.
<b>CL_MEM_REFERENCE_COUNT</b> <sup>24</sup>	cl_uint	Return <i>memobj</i> reference count.
<b>CL_MEM_CONTEXT</b>	cl_context	Return context specified when memory object is created. If <i>memobj</i> is created using <b>clCreateSubBuffer</b> , the context associated with the memory object specified as the <i>buffer</i> argument to <b>clCreateSubBuffer</b> is returned.
<b>CL_MEM_ASSOCIATED_MEMOBJECT</b>	cl_mem	<p>Return memory object from which <i>memobj</i> is created.</p> <p>This returns the memory object specified</p>

<sup>23</sup> The map count returned should be considered immediately stale. It is unsuitable for general use in applications. This feature is provided for debugging.

<sup>24</sup> The reference count returned should be considered immediately stale. It is unsuitable for general use in applications. This feature is provided for identifying memory leaks.

		<p>as <i>buffer</i> argument to <b>clCreateSubBuffer</b> if <i>memobj</i> is a sub-buffer object created using <b>clCreateSubBuffer</b>.</p> <p>This returns the <i>mem_object</i> specified in <i>cl_image_desc</i> if <i>memobj</i> is an image object.</p> <p>Otherwise a NULL value is returned.</p>
<b>CL_MEM_OFFSET</b>	size_t	<p>Return offset if <i>memobj</i> is a sub-buffer object created using <b>clCreateSubBuffer</b>.</p> <p>This return 0 if <i>memobj</i> is not a sub-buffer object.</p>
<b>CL_MEM_USES_SVM_POINTER</b>	cl_bool	<p>Return CL_TRUE if <i>memobj</i> is a buffer object that was created with CL_MEM_USE_HOST_PTR or is a sub-buffer object of a buffer object that was created with CL_MEM_USE_HOST_PTR and the <i>host_ptr</i> specified when the buffer object was created is a SVM pointer; otherwise returns CL_FALSE.</p>



**Table 5.13**    *List of supported param\_names by clGetMemObjectInfo*

## 5.6 Shared Virtual Memory

OpenCL 2.0 adds support for shared virtual memory (a.k.a. SVM). SVM allows the host and kernels executing on devices to directly share complex, pointer-containing data structures such as trees and linked lists. It also eliminates the need to marshal data between the host and devices. As a result, SVM substantially simplifies OpenCL programming and may improve performance.

### 5.6.1 SVM sharing granularity: coarse- and fine- grained sharing

OpenCL maintains memory consistency in a coarse-grained fashion in regions of buffers. We call this coarse-grained sharing. Many platforms such as those with integrated CPU-GPU processors and ones using the SVM-related PCI-SIG IOMMU services can do better, and can support sharing at a granularity smaller than a buffer. We call this fine-grained sharing. OpenCL 2.0 requires that the host and all OpenCL 2.0 devices support coarse-grained sharing at a minimum.

-  Coarse-grained sharing: Coarse-grain sharing may be used for memory and virtual pointer sharing between multiple devices as well as between the host and one or more devices. The shared memory region is a memory buffer allocated using **clSVMAlloc**. Memory consistency is guaranteed at synchronization points and the host can use calls to **clEnqueueSVMMMap** and **clEnqueueSVMUnmap** or create a `cl_mem` buffer object using the SVM pointer and use OpenCL's existing host API functions **clEnqueueMapBuffer** and **clEnqueueUnmapMemObject** to update regions of the buffer. What coarse-grain buffer SVM adds to OpenCL's earlier buffer support are the ability to share virtual memory pointers and a guarantee that concurrent access to the same memory allocation from multiple kernels on a single device is valid. The coarse-grain buffer SVM provides a memory consistency model similar to the global memory consistency model described in *sections 3.3.1 and 3.4.3* of the OpenCL 1.2 specification. This memory consistency applies to the regions of buffers being shared in a coarse-grained fashion. It is enforced at the synchronization points between commands enqueued to command queues in a single context with the additional consideration that multiple kernels concurrently running on the same device may safely share the data.
-  Fine-grained sharing: Shared virtual memory where memory consistency is maintained at a granularity smaller than a buffer. How fine-grained SVM is used depends on whether the device supports SVM atomic operations.
  - If SVM atomic operations are supported, they provide memory consistency for loads and stores by the host and kernels executing on devices supporting SVM. This means that the host and devices can concurrently read and update the same memory. The consistency provided by SVM atomics is in addition to the consistency provided at synchronization points. There is no need for explicit calls

to **clEnqueueSVMMMap** and **clEnqueueSVMUnmap** or **clEnqueueMapBuffer** and **clEnqueueUnmapMemObject** on a `cl_mem` buffer object created using the SVM pointer.

- If SVM atomic operations are not supported, the host and devices can concurrently read the same memory locations and can concurrently update non-overlapping memory regions, but attempts to update the same memory locations are undefined. Memory consistency is guaranteed at synchronization points without the need for explicit calls to **clEnqueueSVMMMap** and **clEnqueueSVMUnmap** or **clEnqueueMapBuffer** and **clEnqueueUnmapMemObject** on a `cl_mem` buffer object created using the SVM pointer.

There are two kinds of fine-grain sharing support. Devices may support either fine-grain buffer sharing or fine-grain system sharing.

- Fine-grain buffer sharing provides fine-grain SVM only within buffers and is an extension of coarse-grain sharing. To support fine-grain buffer sharing in an OpenCL context, all devices in the context must support `CL_DEVICE_SVM_FINE_GRAIN_BUFFER`.
- Fine-grain system sharing enables fine-grain sharing of the host's entire virtual memory, including memory regions allocated by the system **malloc** API. OpenCL buffer objects are unnecessary and programmers can pass pointers allocated using **malloc** to OpenCL kernels.

As an illustration of fine-grain SVM using SVM atomic operations to maintain memory consistency, consider the following example. The host and a set of devices can simultaneously access and update a shared work-queue data structure holding work-items to be done. The host can use atomic operations to insert new work-items into the queue at the same time as the devices using similar atomic operations to remove work-items for processing.

It is the programmer's responsibility to ensure that no host code or executing kernels attempt to access a shared memory region after that memory is freed. We require the SVM implementation to work with either 32- or 64- bit host applications subject to the following requirement: the address space size must be the same for the host and all OpenCL devices in the context.

The function

```
void * clSVMAlloc (cl_context context,  
                  cl_svm_mem_flags flags,  
                  size_t size,  
                  cl_uint alignment)
```

allocates a shared virtual memory buffer (referred to as a SVM buffer) that can be shared by the host and all devices in an OpenCL context that support shared virtual memory.



*context* is a valid OpenCL context used to create the SVM buffer.

*flags* is a bit-field that is used to specify allocation and usage information. *Table 5.14* describes the possible values for *flags*.

<b>cl_svm_mem_flags</b>	<b>Description</b>
<b>CL_MEM_READ_WRITE</b>	This flag specifies that the SVM buffer will be read and written by a kernel. This is the default.
<b>CL_MEM_WRITE_ONLY</b>	This flag specifies that the SVM buffer will be written but not read by a kernel.  Reading from a SVM buffer created with CL_MEM_WRITE_ONLY inside a kernel is undefined.  CL_MEM_READ_WRITE and CL_MEM_WRITE_ONLY are mutually exclusive.
<b>CL_MEM_READ_ONLY</b>	This flag specifies that the SVM buffer object is a read-only memory object when used inside a kernel.  Writing to a SVM buffer created with CL_MEM_READ_ONLY inside a kernel is undefined.  CL_MEM_READ_WRITE or CL_MEM_WRITE_ONLY and CL_MEM_READ_ONLY are mutually exclusive.
<b>CL_MEM_SVM_FINE_GRAIN_BUFFER</b>	This specifies that the application wants the OpenCL implementation to do a fine-grained allocation.
<b>CL_MEM_SVM_ATOMICS</b>	This flag is valid only if CL_MEM_SVM_FINE_GRAIN_BUFFER is specified in flags. It is used to indicate that SVM atomic operations can control visibility of memory accesses in this SVM buffer.

**Table 5.14** *List of supported cl\_svm\_mem\_flags values*

If CL\_MEM\_SVM\_FINE\_GRAIN\_BUFFER is not specified, the buffer can be created as a coarse grained SVM allocation. Similarly, if CL\_MEM\_SVM\_ATOMICS is not specified, the buffer can be created without support for the OpenCL 2.0 SVM atomic operations (refer to *section 6.13.11* of the OpenCL C 2.0 specification).

*size* is the size in bytes of the SVM buffer to be allocated.

*alignment* is the minimum alignment in bytes that is required for the newly created buffer's memory region. It must be a power of two up to the largest data type supported by the OpenCL device. For the full profile, the largest data type is long16. For the embedded profile, it is

`long16` if the device supports 64-bit integers; otherwise it is `int16`. If alignment is 0, a default alignment will be used that is equal to the size of largest data type supported by the OpenCL implementation.

**clSVMAlloc** returns a valid non-NULL shared virtual memory address if the SVM buffer is successfully allocated. Otherwise, like **malloc**, it returns a NULL pointer value. **clSVMAlloc** will fail if

- ✚ *context* is not a valid context.
- ✚ *flags* does not contain `CL_MEM_SVM_FINE_GRAIN_BUFFER` but does contain `CL_MEM_SVM_ATOMICS`.
- ✚ Values specified in *flags* do not follow rules described for supported values in *table 5.14*.
- ✚ `CL_MEM_SVM_FINE_GRAIN_BUFFER` or `CL_MEM_SVM_ATOMICS` is specified in *flags* and these are not supported by at least one device in *context*.
- ✚ The values specified in *flags* are not valid i.e. don't match those defined in *table 5.14*.
- ✚ *size* is 0 or > `CL_DEVICE_MAX_MEM_ALLOC_SIZE` value for any device in *context*.
- ✚ *alignment* is not a power of two or the OpenCL implementation cannot support the specified alignment for at least one device in *context*.
- ✚ There was a failure to allocate resources.

Calling **clSVMAlloc** does not itself provide consistency for the shared memory region. When the host can't use the SVM atomic operations, it must rely on OpenCL's guaranteed memory consistency at synchronization points. To initialize a buffer to be shared with a kernel, the host can create the buffer and use the resulting virtual memory pointer to initialize the buffer's contents.

For SVM to be used efficiently, the host and any devices sharing a buffer containing virtual memory pointers should have the same endianness. If the context passed to **clSVMAlloc** has devices with mixed endianness and the OpenCL implementation is unable to implement SVM because of that mixed endianness, **clSVMAlloc** will fail and return NULL.

Although SVM is generally not supported for image objects, **clCreateImage** may create an image from a buffer (a 1D image from a buffer or a 2D image from buffer) if the buffer specified in its image description parameter is a SVM buffer. Such images have a linear memory representation so their memory can be shared using SVM. However, fine grained sharing and atomics are not supported for image reads and writes in a kernel.

If **clCreateBuffer** is called with a pointer returned by **clSVMAlloc** as its *host\_ptr* argument, and

CL\_MEM\_USE\_HOST\_PTR is set in its *flags* argument, **clCreateBuffer** will succeed and return a valid non-zero buffer object as long as the *size* argument to **clCreateBuffer** is no larger than the *size* argument passed in the original **clSVMAlloc** call. The new buffer object returned has the shared memory as the underlying storage. Locations in the buffer's underlying shared memory can be operated on using, e.g., atomic operations if the device supports them.

The function

```
void    clSVMFree (cl_context context,  
                  void * svm_pointer)
```

frees a shared virtual memory buffer allocated using **clSVMAlloc**.

*context* is a valid OpenCL context used to create the SVM buffer.

*svm\_pointer* must be the value returned by a call to **clSVMAlloc**. If a NULL pointer is passed in *svm\_pointer*, no action occurs.

Note that **clSVMFree** does not wait for previously enqueued commands that may be using *svm\_pointer* to finish before freeing *svm\_pointer*. It is the responsibility of the application to make sure that enqueued commands that use *svm\_pointer* have finished before freeing *svm\_pointer*. This can be done by enqueueing a blocking operation such as **clFinish**, **clWaitForEvents**, **clEnqueueReadBuffer** or by registering a callback with the events associated with enqueued commands and when the last enqueued command has finished freeing *svm\_pointer*.

The behavior of using *svm\_pointer* after it has been freed is undefined. In addition, if a buffer object is created using **clCreateBuffer** with *svm\_pointer*, the buffer object must first be released before the *svm\_pointer* is freed.

The **clEnqueueSVMFree** API can also be used to enqueue a callback to free the shared virtual memory buffer allocated using **clSVMAlloc** or a shared system memory pointer.

The function

```
cl_int  clEnqueueSVMFree (cl_command_queue command_queue,
                          cl_uint num_svm_pointers,
                          void *svm_pointers[],
                          void (CL_CALLBACK *pfn_free_func)(
                              cl_command_queue queue,
                              cl_uint num_svm_pointers,
                              void *svm_pointers[],
                              void *user_data),
                          void *user_data,
                          cl_uint num_events_in_wait_list,
                          const cl_event *event_wait_list,
                          cl_event *event)
```

enqueues a command to free the shared virtual memory allocated using **clSVMAlloc** or a shared system memory pointer.

*command\_queue* is a valid host command-queue.

*svm\_pointers* and *num\_svm\_pointers* specify shared virtual memory pointers to be freed. Each pointer in *svm\_pointers* that was allocated using **clSVMAlloc** must have been allocated from the same context from which *command\_queue* was created. The memory associated with *svm\_pointers* can be reused or freed after the function returns.

*pfn\_free\_func* specifies the callback function to be called to free the SVM pointers. *pfn\_free\_func* takes four arguments: *queue* which is the command queue in which **clEnqueueSVMFree** was enqueued, the count and list of SVM pointers to free and *user\_data* which is a pointer to user specified data. If *pfn\_free\_func* is NULL, all pointers specified in *svm\_pointers* must be allocated using **clSVMAlloc** and the OpenCL implementation will free these SVM pointers. *pfn\_free\_func* must be a valid callback function if any SVM pointer to be freed is a shared system memory pointer i.e. not allocated using **clSVMAlloc**. If *pfn\_free\_func* is a valid callback function, the OpenCL implementation will call *pfn\_free\_func* to free all the SVM pointers specified in *svm\_pointers*.

*user\_data* will be passed as the *user\_data* argument when *pfn\_free\_func* is called. *user\_data* can be NULL.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before **clEnqueueSVMFree** can be executed. If *event\_wait\_list* is NULL, then **clEnqueueSVMFree** does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in

*event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueSVMFree** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_VALUE if *num\_svm\_pointers* is 0 or if *svm\_pointers* is NULL or if any of the pointers specified in *svm\_pointers* array is NULL.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The following function enqueues a command to do a memcpy operation.

```
cl_int  clEnqueueSVMMemcpy (cl_command_queue command_queue,
                           cl_bool blocking_copy,
                           void *dst_ptr,
                           const void *src_ptr,
                           size_t size,
                           cl_uint num_events_in_wait_list,
                           const cl_event *event_wait_list,
                           cl_event *event)
```

*command\_queue* refers to the host command-queue in which the read / write command will be queued. If either *dst\_ptr* or *src\_ptr* is allocated using clSVMAlloc then the OpenCL context allocated against must match that of *command\_queue*.

*blocking\_copy* indicates if the copy operation is *blocking* or *non-blocking*.

If *blocking\_copy* is CL\_TRUE i.e. the copy command is blocking, **clEnqueueSVMMemcpy** does not return until the buffer data has been copied into memory pointed to by *dst\_ptr*.

If *blocking\_copy* is CL\_FALSE i.e. the copy command is non-blocking, **clEnqueueSVMMemcpy** queues a non-blocking copy command and returns. The contents of the buffer that *dst\_ptr* point to cannot be used until the copy command has completed. The *event* argument returns an event object which can be used to query the execution status of the read command. When the copy command has completed, the contents of the buffer that *dst\_ptr* points to can be used by the application.

*size* is the size in bytes of data being copied.

*dst\_ptr* is the pointer to a memory region where data is copied to.




*src\_ptr* is the pointer to a memory region where data is copied from.

If *dst\_ptr* and/or *src\_ptr* are allocated using **clSVMAlloc** and either is not allocated from the same context from which *command\_queue* was created the behavior is undefined.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular read / write command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueSVMMemcpy** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_INVALID\_CONTEXT if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
-  CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.

- ✚ CL\_EXEC\_STATUS\_ERROR\_FOR\_EVENTS\_IN\_WAIT\_LIST if the copy operation is blocking and the execution status of any of the events in *event\_wait\_list* is a negative integer value.
- ✚ CL\_INVALID\_VALUE if *dst\_ptr* or *src\_ptr* are NULL.
- ✚ CL\_INVALID\_VALUE if *size* is 0.
- ✚ CL\_MEM\_COPY\_OVERLAP if the values specified for *dst\_ptr*, *src\_ptr* and *size* result in an overlapping copy.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int clEnqueueSVMMemFill (cl_command_queue command_queue,
                             void *svm_ptr,
                             const void *pattern,
                             size_t pattern_size,
                             size_t size,
                             cl_uint num_events_in_wait_list,
                             const cl_event *event_wait_list,
                             cl_event *event)
```

enqueues a command to fill a region in memory with a pattern of a given pattern size.

*command\_queue* refers to the host command-queue in which the fill command will be queued. The OpenCL context associated with *command\_queue* and SVM pointer referred to by *svm\_ptr* must be the same.

*svm\_ptr* is a pointer to a memory region that will be filled with *pattern*. It must be aligned to *pattern\_size* bytes. If *svm\_ptr* is allocated using **clSVMAlloc** then it must be allocated from the same context from which *command\_queue* was created. Otherwise the behavior is undefined.

*pattern* is a pointer to the data pattern of size *pattern\_size* in bytes. *pattern* will be used to fill a region in *buffer* starting at *svm\_ptr* and is *size* bytes in size. The data pattern must be a scalar or vector integer or floating-point data type supported by OpenCL as described in sections 6.1.1 and 6.1.2. For example, if region pointed to by *svm\_ptr* is to be filled with a pattern of `float4` values, then *pattern* will be a pointer to a `cl_float4` value and *pattern\_size* will be `sizeof(cl_float4)`. The maximum value of *pattern\_size* is the size of the largest integer

or floating-point vector data type supported by the OpenCL device. The memory associated with *pattern* can be reused or freed after the function returns.

*size* is the size in bytes of region being filled starting with *svm\_ptr* and must be a multiple of *pattern\_size*.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueSVMMemFill** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_VALUE if *svm\_ptr* is NULL.
- ✚ CL\_INVALID\_VALUE if *svm\_ptr* is not aligned to *pattern\_size* bytes.
- ✚ CL\_INVALID\_VALUE if *pattern* is NULL or if *pattern\_size* is 0 or if *pattern\_size* is not one of {1, 2, 4, 8, 16, 32, 64, 128}.
- ✚ CL\_INVALID\_VALUE if *size* is 0 or is not a multiple of *pattern\_size*.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.



- CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clEnqueueSVMMap (cl_command_queue command_queue,
                        cl_bool blocking_map,
                        cl_map_flags map_flags,
                        void *svm_ptr,
                        size_t size,
                        cl_uint num_events_in_wait_list,
                        const cl_event *event_wait_list,
                        cl_event *event)
```

enqueues a command that will allow the host to update a region of a SVM buffer. Note that since we are enqueueing a command with a SVM buffer, the region is already mapped in the host address space.

*command\_queue* must be a valid host command-queue.

*blocking\_map* indicates if the map operation is *blocking* or *non-blocking*.

If *blocking\_map* is CL\_TRUE, **clEnqueueSVMMap** does not return until the application can access the contents of the SVM region specified by *svm\_ptr* and *size* on the host.

If *blocking\_map* is CL\_FALSE i.e. map operation is non-blocking, the region specified by *svm\_ptr* and *size* cannot be used until the map command has completed. The *event* argument returns an event object which can be used to query the execution status of the map command. When the map command is completed, the application can access the contents of the region specified by *svm\_ptr* and *size*.

*map\_flags* is a bit-field and is described in *table 5.5*.

*svm\_ptr* and *size* are a pointer to a memory region and size in bytes that will be updated by the host. If *svm\_ptr* is allocated using **clSVMAlloc** then it must be allocated from the same context from which *command\_queue* was created. Otherwise the behavior is undefined.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in

*event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueSVMMap** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_VALUE if *svm\_ptr* is NULL.
- ✚ CL\_INVALID\_VALUE if *size* is 0 or if values specified in *map\_flags* are not valid.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_EXEC\_STATUS\_ERROR\_FOR\_EVENTS\_IN\_WAIT\_LIST if the map operation is blocking and the execution status of any of the events in *event\_wait\_list* is a negative integer value.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clEnqueueSVMUnmap (cl_command_queue command_queue,
                           void *svm_ptr,
                           cl_uint num_events_in_wait_list,
                           const cl_event *event_wait_list,
                           cl_event *event)
```

enqueues a command to indicate that the host has completed updating the region given by *svm\_ptr* and which was specified in a previous call to **clEnqueueSVMMap**.







*command\_queue* must be a valid host command-queue.

*svm\_ptr* is a pointer that was specified in a previous call to **clEnqueueSVMMap**. If *svm\_ptr* is allocated using **clSVMAlloc** then it must be allocated from the same context from which *command\_queue* was created. Otherwise the behavior is undefined.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before **clEnqueueSVMUnmap** can be executed. If *event\_wait\_list* is NULL, then **clEnqueueSVMUnmap** does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular command and can be used to query or queue a wait for this particular command to complete. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. **clEnqueueBarrierWithWaitList** can be used instead. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueSVMUnmap** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_INVALID\_CONTEXT if context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
-  CL\_INVALID\_VALUE if *svm\_ptr* is NULL.
-  CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or if *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

**clEnqueueSVMMMap**, and **clEnqueueSVMUnmap** act as synchronization points for the region of the SVM buffer specified in these calls.

NOTE:

If a coarse-grained SVM buffer is currently mapped for writing, the application must ensure that the SVM buffer is unmapped before any enqueued kernels or commands that read from or write to this SVM buffer or any of its associated `cl_mem` buffer objects begin execution; otherwise the behavior is undefined.

If a coarse-grained SVM buffer is currently mapped for reading, the application must ensure that the SVM buffer is unmapped before any enqueued kernels or commands that write to this memory object or any of its associated `cl_mem` buffer objects begin execution; otherwise the behavior is undefined.

A SVM buffer is considered as mapped if there are one or more active mappings for the SVM buffer irrespective of whether the mapped regions span the entire SVM buffer.

The above note does not apply to fine-grained SVM buffers (fine-grained buffers allocated using **clSVMAlloc** or fine-grained system allocations).

## 5.6.2 Memory consistency for SVM allocations

To ensure memory consistency in SVM allocations, the program can rely on the guaranteed memory consistency at synchronization points. This consistency support already exists in OpenCL 1.x and can be used for coarse-grained SVM allocations or for fine-grained buffer SVM allocations; what SVM adds is the ability to share pointers between the host and all SVM devices.

In addition, sub-buffers can also be used to ensure that each device gets a consistent view of a SVM buffer's memory when it is shared by multiple devices. For example, assume that two devices share a SVM pointer. The host can create a `cl_mem` buffer object using **clCreateBuffer** with `CL_MEM_USE_HOST_PTR` and *host\_ptr* set to the SVM pointer and then create two disjoint sub-buffers with starting virtual addresses *sb1\_ptr* and *sb2\_ptr*. These pointers (*sb1\_ptr* and *sb2\_ptr*) can be passed to kernels executing on the two devices. **clEnqueueMapBuffer** and **clEnqueueUnmapMemObject** and the existing access rules for memory objects (in *section 5.5.3*) can be used to ensure consistency for buffer regions (*sb1\_ptr* and *sb2\_ptr*) read and written by these kernels.

When the host and devices are able to use SVM atomic operations (i.e. `CL_DEVICE_SVM_ATOMICS` is set in `CL_DEVICE_SVM_CAPABILITIES`), these atomic operations can be used to provide memory consistency at a fine grain in a shared memory region. The effect of these operations is visible to the host and all devices with which that memory is shared.



## 5.7 Sampler Objects

A sampler object describes how to sample an image when the image is read in the kernel. The built-in functions to read from an image in a kernel take a sampler as an argument. The sampler arguments to the image read function can be sampler objects created using OpenCL functions and passed as argument values to the kernel or can be samplers declared inside a kernel. In this section we discuss how sampler objects are created using OpenCL functions.

### 5.7.1 Creating Sampler Objects

The function

```
cl_sampler    clCreateSamplerWithProperties (cl_context context,
                                           const cl_sampler_properties *sampler_properties,
                                           cl_int *errcode_ret)
```

creates a sampler object.

*context* must be a valid OpenCL context.

*sampler\_properties* specifies a list of sampler property names and their corresponding values. Each sampler property name is immediately followed by the corresponding desired value. The list is terminated with 0. The list of supported properties is described in *table 5.15*. If a supported property and its value is not specified in *sampler\_properties*, its default value will be used. *sampler\_properties* can be NULL in which case the default values for supported sampler properties will be used.

cl_sampler_properties enum	Property Value	Description
CL_SAMPLER_NORMALIZED_COORDS	cl_bool	A boolean value that specifies whether the image coordinates specified are normalized or not.  The default value (i.e. the value used if this property is not specified in <i>sampler_properties</i> ) is CL_TRUE.
CL_SAMPLER_ADDRESSING_MODE	cl_addressing_mode	Specifies how out-of-range image coordinates are handled when reading from an image.  Valid values are:  CL_ADDRESS_MIRRORED_REPEAT

		CL_ADDRESS_REPEAT CL_ADDRESS_CLAMP_TO_EDGE CL_ADDRESS_CLAMP CL_ADDRESS_NONE  The default is CL_ADDRESS_CLAMP.
<b>CL_SAMPLER_FILTER_MODE</b>	cl_filter_mode	Specifies the type of filter that must be applied when reading an image. Valid values are:  CL_FILTER_NEAREST CL_FILTER_LINEAR  The default value is CL_FILTER_NEAREST.

**Table 5.15** *List of supported cl\_sampler\_properties values and description.*

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.




**clCreateSamplerWithProperties** returns a valid non-zero sampler object and *errcode\_ret* is set to CL\_SUCCESS if the sampler object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid context.
- ✚ CL\_INVALID\_VALUE if the property name in *sampler\_properties* is not a supported property name, if the value specified for a supported property name is not valid, or if the same property name is specified more than once.
- ✚ CL\_INVALID\_OPERATION if images are not supported by any device associated with *context* (i.e. CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_FALSE).
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

cl\_int            **clRetainSampler** (cl\_sampler *sampler*)




increments the *sampler* reference count. **clCreateSamplerWithProperties** performs an implicit retain. **clRetainSampler** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_SAMPLER if *sampler* is not a valid sampler object.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

cl\_int                   **clReleaseSampler** (cl\_sampler *sampler*)

decrements the *sampler* reference count. The sampler object is deleted after the reference count becomes zero and commands queued for execution on a command-queue(s) that use *sampler* have finished. **clReleaseSampler** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_SAMPLER if *sampler* is not a valid sampler object.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

Using this function to release a reference that was not obtained by creating the object or by calling **clRetainSampler** causes undefined behavior.

## 5.7.2 Sampler Object Queries

The function

cl\_int                   **clGetSamplerInfo** (cl\_sampler *sampler*,  
  cl\_sampler\_info *param\_name*,  
  size\_t *param\_value\_size*,  
  void \**param\_value*,  
  size\_t \**param\_value\_size\_ret*)

returns information about the sampler object.



*sampler* specifies the sampler being queried.

*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetSamplerInfo** is described in *table 5.16*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.





*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.16*.

*param\_value\_size\_ret* returns the actual size in bytes of data copied to *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

cl_sampler_info	Return Type	Info. returned in <i>param_value</i>
CL_SAMPLER_REFERENCE_COUNT <sup>25</sup>	cl_uint	Return the <i>sampler</i> reference count.
CL_SAMPLER_CONTEXT	cl_context	Return the context specified when the sampler is created.
CL_SAMPLER_NORMALIZED_COORDS	cl_bool	Return the normalized coords value associated with <i>sampler</i> .
CL_SAMPLER_ADDRESSING_MODE	cl_addressing_mode	Return the addressing mode value associated with <i>sampler</i> .
CL_SAMPLER_FILTER_MODE	cl_filter_mode	Return the filter mode value associated with <i>sampler</i> .

**Table 5.16** *clGetSamplerInfo* parameter queries.

**clGetSamplerInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is  $<$  size of return type as described in *table 5.16* and *param\_value* is not NULL.
-  CL\_INVALID\_SAMPLER if *sampler* is a not a valid sampler object.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.





---

<sup>25</sup> The reference count returned should be considered immediately stale. It is unsuitable for general use in applications. This feature is provided for identifying memory leaks.

## 5.8 Program Objects

An OpenCL program consists of a set of kernels that are identified as functions declared with the `__kernel` qualifier in the program source. OpenCL programs may also contain auxiliary functions and constant data that can be used by `__kernel` functions. The program executable can be generated *online* or *offline* by the OpenCL compiler for the appropriate target device(s).

A program object encapsulates the following information:

-  An associated context.
-  A program source or binary.
-  The latest successfully built program executable, library or compiled binary, the list of devices for which the program executable, library or compiled binary is built, the build options used and a build log.
-  The number of kernel objects currently attached.

### 5.8.1 Creating Program Objects

The function

```
cl_program    clCreateProgramWithSource (cl_context context,  
                                         cl_uint count,  
                                         const char **strings,  
                                         const size_t *lengths,  
                                         cl_int *errcode_ret)
```

creates a program object for a context, and loads the source code specified by the text strings in the *strings* array into the program object. The devices associated with the program object are the devices associated with *context*. The source code specified by *strings* is either an OpenCL C program source, header or implementation-defined source for custom devices that support an online compiler.

*context* must be a valid OpenCL context.

*strings* is an array of *count* pointers to optionally null-terminated character strings that make up the source code.

The *lengths* argument is an array with the number of chars in each string (the string length). If an element in *lengths* is zero, its accompanying string is null-terminated. If *lengths* is NULL, all strings in the *strings* argument are considered null-terminated. Any length value passed in that is greater than zero excludes the null terminator in its count.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreateProgramWithSource** returns a valid non-zero program object and *errcode\_ret* is set to CL\_SUCCESS if the program object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid context.
- ✚ CL\_INVALID\_VALUE if *count* is zero or if *strings* or any entry in *strings* is NULL.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_program    clCreateProgramWithBinary (cl_context context,
                                           cl_uint num_devices,
                                           const cl_device_id *device_list,
                                           const size_t *lengths,
                                           const unsigned char **binaries,
                                           cl_int *binary_status,
                                           cl_int *errcode_ret)
```

creates a program object for a context, and loads the binary bits specified by *binary* into the program object.

*context* must be a valid OpenCL context.

*device\_list* is a pointer to a list of devices that are in *context*. *device\_list* must be a non-NULL value. The binaries are loaded for devices specified in this list.

*num\_devices* is the number of devices listed in *device\_list*.

The devices associated with the program object will be the list of devices specified by *device\_list*. The list of devices specified by *device\_list* must be devices associated with *context*.

*lengths* is an array of the size in bytes of the program binaries to be loaded for devices specified by *device\_list*.

*binaries* is an array of pointers to program binaries to be loaded for devices specified by *device\_list*. For each device given by *device\_list*[*i*], the pointer to the program binary for that

device is given by *binaries[i]* and the length of this corresponding binary is given by *lengths[i]*. *lengths[i]* cannot be zero and *binaries[i]* cannot be a NULL pointer.

The program binaries specified by *binaries* contain the bits that describe one of the following:

- ✚ a program executable to be run on the device(s) associated with *context*,
- ✚ a compiled program for device(s) associated with *context*, or
- ✚ a library of compiled programs for device(s) associated with *context*.

The program binary can consist of either or both:

- ✚ Device-specific code and/or,
- ✚ Implementation-specific intermediate representation (IR) which will be converted to the device-specific code.

*binary\_status* returns whether the program binary for each device specified in *device\_list* was loaded successfully or not. It is an array of *num\_devices* entries and returns CL\_SUCCESS in *binary\_status[i]* if binary was successfully loaded for device specified by *device\_list[i]*; otherwise returns CL\_INVALID\_VALUE if *lengths[i]* is zero or if *binaries[i]* is a NULL value or CL\_INVALID\_BINARY in *binary\_status[i]* if program binary is not a valid binary for the specified device. If *binary\_status* is NULL, it is ignored.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreateProgramWithBinary** returns a valid non-zero program object and *errcode\_ret* is set to CL\_SUCCESS if the program object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid context.
- ✚ CL\_INVALID\_VALUE if *device\_list* is NULL or *num\_devices* is zero.
- ✚ CL\_INVALID\_DEVICE if OpenCL devices listed in *device\_list* are not in the list of devices associated with *context*.
- ✚ CL\_INVALID\_VALUE if *lengths* or *binaries* are NULL or if any entry in *lengths[i]* is zero or *binaries[i]* is NULL.
- ✚ CL\_INVALID\_BINARY if an invalid program binary was encountered for any device. *binary\_status* will return specific status for each device.

- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

OpenCL allows applications to create a program object using the program source or binary and build appropriate program executables. This can be very useful as it allows applications to load program source and then compile and link to generate a program executable online on its first instance for appropriate OpenCL devices in the system. These executables can now be queried and cached by the application. Future instances of the application launching will no longer need to compile and link the program executables. The cached executables can be read and loaded by the application, which can help significantly reduce the application initialization time.

The function

```
cl_program  clCreateProgramWithBuiltInKernels (cl_context context,
                                              cl_uint num_devices,
                                              const cl_device_id *device_list,
                                              const char *kernel_names,
                                              cl_int *errcode_ret)
```

creates a program object for a context, and loads the information related to the built-in kernels into a program object.

*context* must be a valid OpenCL context.

*num\_devices* is the number of devices listed in *device\_list*.

*device\_list* is a pointer to a list of devices that are in *context*. *device\_list* must be a non-NULL value. The built-in kernels are loaded for devices specified in this list.

The devices associated with the program object will be the list of devices specified by *device\_list*. The list of devices specified by *device\_list* must be devices associated with *context*.

*kernel\_names* is a semi-colon separated list of built-in kernel names.

**clCreateProgramWithBuiltInKernels** returns a valid non-zero program object and *errcode\_ret* is set to CL\_SUCCESS if the program object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid context.
- ✚ CL\_INVALID\_VALUE if *device\_list* is NULL or *num\_devices* is zero.

- ✚ CL\_INVALID\_VALUE if *kernel\_names* is NULL or *kernel\_names* contains a kernel name that is not supported by any of the devices in *device\_list*.
- ✚ CL\_INVALID\_DEVICE if devices listed in *device\_list* are not in the list of devices associated with *context*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

cl\_int            **clRetainProgram** (cl\_program *program*)

increments the *program* reference count. **clCreateProgram** does an implicit retain. **clRetainProgram** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_PROGRAM if *program* is not a valid program object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

cl\_int            **clReleaseProgram** (cl\_program *program*)

decrements the *program* reference count. The program object is deleted after all kernel objects associated with *program* have been deleted and the *program* reference count becomes zero. **clReleaseProgram** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_PROGRAM if *program* is not a valid program object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

Using this function to release a reference that was not obtained by creating the object or by calling **clRetainProgram** causes undefined behavior.

## 5.8.2 Building Program Executables

The function

```
cl_int      clBuildProgram (cl_program program,
                             cl_uint num_devices,
                             const cl_device_id *device_list,
                             const char *options,
                             void (CL_CALLBACK *pfn_notify)(cl_program program,
                                                             void *user_data),
                             void *user_data)
```

builds (compiles & links) a program executable from the program source or binary for all the devices or a specific device(s) in the OpenCL context associated with *program*. OpenCL allows program executables to be built using the source or the binary. **clBuildProgram** must be called for *program* created using either **clCreateProgramWithSource** or **clCreateProgramWithBinary** to build the program executable for one or more devices associated with *program*. If *program* is created with **clCreateProgramWithBinary**, then the program binary must be an executable binary (not a compiled binary or library).

The executable binary can be queried using **clGetProgramInfo**(*program*, CL\_PROGRAM\_BINARIES, ...) and can be specified to **clCreateProgramWithBinary** to create a new program object.

*program* is the program object.

*device\_list* is a pointer to a list of devices associated with *program*. If *device\_list* is a NULL value, the program executable is built for all devices associated with *program* for which a source or binary has been loaded. If *device\_list* is a non-NULL value, the program executable is built for devices specified in this list for which a source or binary has been loaded.

*num\_devices* is the number of devices listed in *device\_list*.

*options* is a pointer to a null-terminated string of characters that describes the build options to be used for building the program executable. The list of supported options is described in section 5.8.4.

*pfn\_notify* is a function pointer to a notification routine. The notification routine is a callback function that an application can register and which will be called when the program executable has been built (successfully or unsuccessfully). If *pfn\_notify* is not NULL, **clBuildProgram** does not need to wait for the build to complete and can return immediately once the build operation can begin. The build operation can begin if the context, program whose sources are being compiled and linked, list of devices and build options specified are all valid and appropriate host and device resources needed to perform the build are available. If *pfn\_notify* is NULL, **clBuildProgram** does not return until the build has completed. This callback function may be



called asynchronously by the OpenCL implementation. It is the application's responsibility to ensure that the callback function is thread-safe.

*user\_data* will be passed as an argument when *pfn\_notify* is called. *user\_data* can be NULL.

**clBuildProgram** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_PROGRAM if *program* is not a valid program object.
- ✚ CL\_INVALID\_VALUE if *device\_list* is NULL and *num\_devices* is greater than zero, or if *device\_list* is not NULL and *num\_devices* is zero.
- ✚ CL\_INVALID\_VALUE if *pfn\_notify* is NULL but *user\_data* is not NULL.
- ✚ CL\_INVALID\_DEVICE if OpenCL devices listed in *device\_list* are not in the list of devices associated with *program*
- ✚ CL\_INVALID\_BINARY if *program* is created with **clCreateProgramWithBinary** and devices listed in *device\_list* do not have a valid program binary loaded.
- ✚ CL\_INVALID\_BUILD\_OPTIONS if the build options specified by *options* are invalid.
- ✚ CL\_INVALID\_OPERATION if the build of a program executable for any of the devices listed in *device\_list* by a previous call to **clBuildProgram** for *program* has not completed.
- ✚ CL\_COMPILER\_NOT\_AVAILABLE if *program* is created with **clCreateProgramWithSource** and a compiler is not available i.e. CL\_DEVICE\_COMPILER\_AVAILABLE specified in *table 4.3* is set to CL\_FALSE.
- ✚ CL\_BUILD\_PROGRAM\_FAILURE if there is a failure to build the program executable. This error will be returned if **clBuildProgram** does not return until the build has completed.
- ✚ CL\_INVALID\_OPERATION if there are kernel objects attached to *program*.
- ✚ CL\_INVALID\_OPERATION if *program* was not created with **clCreateProgramWithSource** or **clCreateProgramWithBinary**.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 5.8.3 Separate Compilation and Linking of Programs

OpenCL programs are compiled and linked to support the following:

- ✚ Separate compilation and link stages. Program sources can be compiled to generate a compiled binary object and linked in a separate stage with other compiled program objects to the program executable.
- ✚ Embedded headers. In OpenCL 1.0 and 1.1, the `-I` build option could be used to specify the list of directories to be searched for header files that are included by a program source(s). OpenCL 1.2 extends this by allowing the header sources to come from program objects instead of just header files.
- ✚ Libraries. The linker can be used to link compiled objects and libraries into a program executable or to create a library of compiled binaries.

The function

```
cl_int      clCompileProgram (cl_program program,
                                cl_uint num_devices,
                                const cl_device_id *device_list,
                                const char *options,
                                cl_uint num_input_headers,
                                const cl_program *input_headers,
                                const char **header_include_names,
                                void (CL_CALLBACK *pfn_notify)(cl_program program,
                                                                void *user_data),
                                void *user_data)
```

compiles a program's source for all the devices or a specific device(s) in the OpenCL context associated with *program*. The pre-processor runs before the program sources are compiled. The compiled binary is built for all devices associated with *program* or the list of devices specified. The compiled binary can be queried using **clGetProgramInfo**(*program*, CL\_PROGRAM\_BINARIES, ...) and can be specified to **clCreateProgramWithBinary** to create a new program object.

*program* is the program object that is the compilation target.

*device\_list* is a pointer to a list of devices associated with *program*. If *device\_list* is a NULL value, the compile is performed for all devices associated with *program*. If *device\_list* is a non-NULL value, the compile is performed for devices specified in this list.

*num\_devices* is the number of devices listed in *device\_list*.

*options* is a pointer to a null-terminated string of characters that describes the compilation options to be used for building the program executable. The list of supported options is as described in *section 5.8.4*.

*num\_input\_headers* specifies the number of programs that describe headers in the array referenced by *input\_headers*.

*input\_headers* is an array of program embedded headers created with **clCreateProgramWithSource**.

*header\_include\_names* is an array that has a one to one correspondence with *input\_headers*. Each entry in *header\_include\_names* specifies the include name used by source in *program* that comes from an embedded header. The corresponding entry in *input\_headers* identifies the program object which contains the header source to be used. The embedded headers are first searched before the headers in the list of directories specified by the `-I` compile option (as described in *section 5.8.4.1*). If multiple entries in *header\_include\_names* refer to the same header name, the first one encountered will be used.

For example, consider the following program source:

```
#include <foo.h>
#include <mydir/myinc.h>

__kernel void
image_filter (int n, int m,
              __constant float *filter_weights,
              __read_only image2d_t src_image,
              __write_only image2d_t dst_image)
{
    ...
}
```

This kernel includes two headers `foo.h` and `mydir/myinc.h`. The following describes how these headers can be passed as embedded headers in program objects:

```
cl_program foo_pg = clCreateProgramWithSource(context,
                                              1, &foo_header_src, NULL, &err);
cl_program myinc_pg = clCreateProgramWithSource(context,
                                                1, &myinc_header_src, NULL, &err);

// let's assume the program source described above is given
// by program_A and is loaded via clCreateProgramWithSource

cl_program input_headers[2] = { foo_pg, myinc_pg };
char * input_header_names[2] = { "foo.h", "mydir/myinc.h" };
clCompileProgram(program_A,
```

```

0, NULL,    // num_devices & device_list
NULL,      // compile_options
2,         // num_input_headers
input_headers,
input_header_names,
NULL, NULL); // pfn_notify & user_data

```

*pfn\_notify* is a function pointer to a notification routine. The notification routine is a callback function that an application can register and which will be called when the program executable has been built (successfully or unsuccessfully). If *pfn\_notify* is not NULL, **clCompileProgram** does not need to wait for the compiler to complete and can return immediately once the compilation can begin. The compilation can begin if the context, program whose sources are being compiled, list of devices, input headers, programs that describe input headers and compiler options specified are all valid and appropriate host and device resources needed to perform the compile are available. If *pfn\_notify* is NULL, **clCompileProgram** does not return until the compiler has completed. This callback function may be called asynchronously by the OpenCL implementation. It is the application's responsibility to ensure that the callback function is thread-safe.

*user\_data* will be passed as an argument when *pfn\_notify* is called. *user\_data* can be NULL.

**clCompileProgram** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_PROGRAM if *program* is not a valid program object.
- ✚ CL\_INVALID\_VALUE if *device\_list* is NULL and *num\_devices* is greater than zero, or if *device\_list* is not NULL and *num\_devices* is zero.
- ✚ CL\_INVALID\_VALUE if *num\_input\_headers* is zero and *header\_include\_names* or *input\_headers* are not NULL or if *num\_input\_headers* is not zero and *header\_include\_names* or *input\_headers* are NULL.
- ✚ CL\_INVALID\_VALUE if *pfn\_notify* is NULL but *user\_data* is not NULL.
- ✚ CL\_INVALID\_DEVICE if OpenCL devices listed in *device\_list* are not in the list of devices associated with *program*
- ✚ CL\_INVALID\_COMPILER\_OPTIONS if the compiler options specified by *options* are invalid.
- ✚ CL\_INVALID\_OPERATION if the compilation or build of a program executable for any of the devices listed in *device\_list* by a previous call to **clCompileProgram** or **clBuildProgram** for *program* has not completed.
- ✚ CL\_COMPILER\_NOT\_AVAILABLE if a compiler is not available i.e.

CL\_DEVICE\_COMPILER\_AVAILABLE specified in *table 4.3* is set to CL\_FALSE.

- ✚ CL\_COMPILE\_PROGRAM\_FAILURE if there is a failure to compile the program source. This error will be returned if **clCompileProgram** does not return until the compile has completed.
- ✚ CL\_INVALID\_OPERATION if there are kernel objects attached to *program*.
- ✚ CL\_INVALID\_OPERATION if *program* has no source i.e. it has not been created with **clCreateProgramWithSource**.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_program    clLinkProgram (cl_context context,
                               cl_uint num_devices,
                               const cl_device_id *device_list,
                               const char *options,
                               cl_uint num_input_programs,
                               const cl_program *input_programs,
                               void (CL_CALLBACK *pfn_notify)(cl_program program,
                                                                void *user_data),
                               void *user_data,
                               cl_int *errcode_ret)
```

links a set of compiled program objects and libraries for all the devices or a specific device(s) in the OpenCL context and creates a library or executable. **clLinkProgram** creates a new program object which contains the library or executable. The library or executable binary can be queried using **clGetProgramInfo**(*program*, CL\_PROGRAM\_BINARIES, ...) and can be specified to **clCreateProgramWithBinary** to create a new program object.

The devices associated with the returned program object will be the list of devices specified by *device\_list* or if *device\_list* is NULL it will be the list of devices associated with *context*.

*context* must be a valid OpenCL context.

*device\_list* is a pointer to a list of devices that are in *context*. If *device\_list* is a NULL value, the link is performed for all devices associated with *context* for which a compiled object is available. If *device\_list* is a non-NULL value, the link is performed for devices specified in this list for which a compiled object is available.

*num\_devices* is the number of devices listed in *device\_list*.

*options* is a pointer to a null-terminated string of characters that describes the link options to be used for building the program executable. The list of supported options is as described in section 5.8.5.

*num\_input\_programs* specifies the number of programs in array referenced by *input\_programs*.

*input\_programs* is an array of program objects that are compiled binaries or libraries that are to be linked to create the program executable. For each device in *device\_list* or if *device\_list* is NULL the list of devices associated with context, the following cases occur:

- ✚ All programs specified by *input\_programs* contain a compiled binary or library for the device. In this case, a link is performed to generate a program executable for this device.
- ✚ None of the programs contain a compiled binary or library for that device. In this case, no link is performed and there will be no program executable generated for this device.
- ✚ All other cases will return a CL\_INVALID\_OPERATION error.

*pfn\_notify* is a function pointer to a notification routine. The notification routine is a callback function that an application can register and which will be called when the program executable has been built (successfully or unsuccessfully).

If *pfn\_notify* is not NULL, **clLinkProgram** does not need to wait for the linker to complete and can return immediately once the linking operation can begin. Once the linker has completed, the *pfn\_notify* callback function is called which returns the program object returned by **clLinkProgram**. The application can query the link status and log for this program object. This callback function may be called asynchronously by the OpenCL implementation. It is the application's responsibility to ensure that the callback function is thread-safe.

If *pfn\_notify* is NULL, **clLinkProgram** does not return until the linker has completed.

*user\_data* will be passed as an argument when *pfn\_notify* is called. *user\_data* can be NULL.

The linking operation can begin if the context, list of devices, input programs and linker options specified are all valid and appropriate host and device resources needed to perform the link are available. If the linking operation can begin, **clLinkProgram** returns a valid non-zero program object.

If *pfn\_notify* is NULL, the *errcode\_ret* will be set to CL\_SUCCESS if the link operation was successful and CL\_LINK\_FAILURE if there is a failure to link the compiled binaries and/or libraries.

If *pfn\_notify* is not NULL, **clLinkProgram** does not have to wait until the linker to complete and

can return CL\_SUCCESS in *errcode\_ret* if the linking operation can begin. The *pfn\_notify* callback function will return a CL\_SUCCESS or CL\_LINK\_FAILURE if the linking operation was successful or not.

Otherwise **clLinkProgram** returns a NULL program object with an appropriate error in *errcode\_ret*. The application should query the linker status of this program object to check if the link was successful or not. The list of errors that can be returned are:

- ✚ CL\_INVALID\_CONTEXT if *context* is not a valid context.
- ✚ CL\_INVALID\_VALUE if *device\_list* is NULL and *num\_devices* is greater than zero, or if *device\_list* is not NULL and *num\_devices* is zero.
- ✚ CL\_INVALID\_VALUE if *num\_input\_programs* is zero and *input\_programs* is NULL or if *num\_input\_programs* is zero and *input\_programs* is not NULL or if *num\_input\_programs* is not zero and *input\_programs* is NULL.
- ✚ CL\_INVALID\_PROGRAM if programs specified in *input\_programs* are not valid program objects.
- ✚ CL\_INVALID\_VALUE if *pfn\_notify* is NULL but *user\_data* is not NULL.
- ✚ CL\_INVALID\_DEVICE if OpenCL devices listed in *device\_list* are not in the list of devices associated with *context*
- ✚ CL\_INVALID\_LINKER\_OPTIONS if the linker options specified by *options* are invalid.
- ✚ CL\_INVALID\_OPERATION if the compilation or build of a program executable for any of the devices listed in *device\_list* by a previous call to **clCompileProgram** or **clBuildProgram** for *program* has not completed.
- ✚ CL\_INVALID\_OPERATION if the rules for devices containing compiled binaries or libraries as described in *input\_programs* argument above are not followed.
- ✚ CL\_LINKER\_NOT\_AVAILABLE if a linker is not available i.e. CL\_DEVICE\_LINKER\_AVAILABLE specified in *table 4.3* is set to CL\_FALSE.
- ✚ CL\_LINK\_PROGRAM\_FAILURE if there is a failure to link the compiled binaries and/or libraries.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 5.8.4 Compiler Options

The compiler options are categorized as pre-processor options, options for math intrinsics, options that control optimization and miscellaneous options. This specification defines a standard set of options that must be supported by the OpenCL C compiler when building program executables online or offline. These may be extended by a set of vendor- or platform-specific options.

### 5.8.4.1 Preprocessor options

These options control the OpenCL C preprocessor which is run on each program source before actual compilation.

**-D *name***

Predefine *name* as a macro, with definition 1.

**-D *name=definition***

The contents of *definition* are tokenized and processed as if they appeared during translation phase three in a `#define` directive. In particular, the definition will be truncated by embedded newline characters.

**-D** options are processed in the order they are given in the *options* argument to **clBuildProgram** or **clCompileProgram**.

**-I *dir***

Add the directory *dir* to the list of directories to be searched for header files.

### 5.8.4.2 Math Intrinsics Options

These options control compiler behavior regarding floating-point arithmetic. These options trade off between speed and correctness.

**-cl-single-precision-constant**

Treat double precision floating-point constant as single precision constant.

**-cl-denorms-are-zero**

This option controls how single precision and double precision denormalized numbers are handled. If specified as a build option, the single precision denormalized numbers may be flushed to zero; double precision denormalized numbers may also be flushed to zero if the optional extension for double precision is supported. This is intended to be a performance hint and the OpenCL compiler can choose not to flush denorms to zero if the device supports single precision (or double precision) denormalized numbers.



This option is ignored for single precision numbers if the device does not support single precision denormalized numbers i.e. `CL_FP_DENORM` bit is not set in `CL_DEVICE_SINGLE_FP_CONFIG`.

This option is ignored for double precision numbers if the device does not support double precision or if it does support double precision but not double precision denormalized numbers i.e. `CL_FP_DENORM` bit is not set in `CL_DEVICE_DOUBLE_FP_CONFIG`.

This flag only applies for scalar and vector single precision floating-point variables and computations on these floating-point variables inside a program. It does not apply to reading from or writing to image objects.

#### `-cl-fp32-correctly-rounded-divide-sqrt`

The `-cl-fp32-correctly-rounded-divide-sqrt` build option to **clBuildProgram** or **clCompileProgram** allows an application to specify that single precision floating-point divide ( $x/y$  and  $1/x$ ) and `sqrt` used in the program source are correctly rounded. If this build option is not specified, the minimum numerical accuracy of single precision floating-point divide and sqrt are as defined in *section 7.4* of the OpenCL specification.

This build option can only be specified if the `CL_FP_CORRECTLY_ROUNDED_DIVIDE_SQRT` is set in `CL_DEVICE_SINGLE_FP_CONFIG` (as defined in *table 4.3*) for devices that the program is being build. **clBuildProgram** or **clCompileProgram** will fail to compile the program for a device if the `-cl-fp32-correctly-rounded-divide-sqrt` option is specified and `CL_FP_CORRECTLY_ROUNDED_DIVIDE_SQRT` is not set for the device.

### 5.8.4.3 Optimization Options

These options control various sorts of optimizations. Turning on optimization flags makes the compiler attempt to improve the performance and/or code size at the expense of compilation time and possibly the ability to debug the program.

#### `-cl-opt-disable`

This option disables all optimizations. The default is optimizations are enabled.

The following options control compiler behavior regarding floating-point arithmetic. These options trade off between performance and correctness and must be specifically enabled. These options are not turned on by default since it can result in incorrect output for programs which depend on an exact implementation of IEEE 754 rules/specifications for math functions.

#### `-cl-mad-enable`

Allow  $a * b + c$  to be replaced by a `mad`. The `mad` computes  $a * b + c$  with reduced accuracy. For example, some OpenCL devices implement `mad` as truncate the result of  $a * b$  before adding it to  $c$ .

#### **-cl-no-signed-zeros**

Allow optimizations for floating-point arithmetic that ignore the signedness of zero. IEEE 754 arithmetic specifies the distinct behavior of  $+0.0$  and  $-0.0$  values, which then prohibits simplification of expressions such as  $x+0.0$  or  $0.0*x$  (even with **-cl-finite-math-only**). This option implies that the sign of a zero result isn't significant.

#### **-cl-unsafe-math-optimizations**

Allow optimizations for floating-point arithmetic that (a) assume that arguments and results are valid, (b) may violate IEEE 754 standard and (c) may violate the OpenCL numerical compliance requirements as defined in *section 7.4* for single precision and double precision floating-point, and edge case behavior in *section 7.5*. This option includes the **-cl-no-signed-zeros** and **-cl-mad-enable** options.

#### **-cl-finite-math-only**

Allow optimizations for floating-point arithmetic that assume that arguments and results are not NaNs or  $\pm\infty$ . This option may violate the OpenCL numerical compliance requirements defined in *section 7.4* for single precision and double precision floating-point, and edge case behavior in *section 7.5*.

#### **-cl-fast-relaxed-math**

Sets the optimization options **-cl-finite-math-only** and **-cl-unsafe-math-optimizations**. This allows optimizations for floating-point arithmetic that may violate the IEEE 754 standard and the OpenCL numerical compliance requirements defined in *section 7.4* for single precision and double precision floating-point, and edge case behavior in *section 7.5*. This option also relaxes the precision of commonly used math functions (refer to *table 7.2* defined in *section 7.4*). This option causes the preprocessor macro `__FAST_RELAXED_MATH__` to be defined in the OpenCL program.

#### **-cl-uniform-work-group-size**

This requires that the global work-size be a multiple of the work-group size specified to **clEnqueueNDRangeKernel**. Allow optimizations that are made possible by this restriction.

### **5.8.4.4 Options to Request or Suppress Warnings**

Warnings are diagnostic messages that report constructions which are not inherently erroneous but which are risky or suggest there may have been an error. The following language-independent options do not enable specific warnings but control the kinds of diagnostics produced by the OpenCL compiler.

#### **-w**

Inhibit all warning messages.

#### **-Werror**

Make all warnings into errors.

### 5.8.4.5 Options Controlling the OpenCL C version

The following option controls the version of OpenCL C that the compiler accepts.

**-cl-std=**

Determine the OpenCL C language version to use. A value for this option must be provided. Valid values are:

CL1.1 – Support all OpenCL C programs that use the OpenCL C language features defined in *section 6* of the OpenCL 1.1 specification.

CL1.2 – Support all OpenCL C programs that use the OpenCL C language features defined in *section 6* of the OpenCL 1.2 specification.

CL2.0 – Support all OpenCL C programs that use the OpenCL C language features defined in the OpenCL C 2.0 specification.

Calls to **clBuildProgram** or **clCompileProgram** with the **-cl-std=CL1.1** option **will fail** to compile the program for any devices with `CL_DEVICE_OPENCL_C_VERSION = OpenCL C 1.0`. Calls to **clBuildProgram** or **clCompileProgram** with the **-cl-std=CL1.2** option **will fail** to compile the program for any devices with `CL_DEVICE_OPENCL_C_VERSION = OpenCL C 1.0` or OpenCL C 1.1.

Calls to **clBuildProgram** or **clCompileProgram** with the **-cl-std=CL2.0** option **will fail** to compile the program for any devices with `CL_DEVICE_OPENCL_C_VERSION = OpenCL C 1.0`, OpenCL C 1.1 or OpenCL C 1.2.

If the **-cl-std** build option is not specified, the highest OpenCL C 1.x language version supported by each device is used when compiling the program for each device. Applications are required to specify the **-cl-std=CL2.0** option if they want to compile or build their programs with OpenCL C 2.0.

### 5.8.4.6 Options for Querying Kernel Argument Information

**-cl-kernel-arg-info**

This option allows the compiler to store information about the arguments of a kernel(s) in the program executable. The argument information stored includes the argument name, its type, the address space and access qualifiers used. Refer to description of **clGetKernelArgInfo** on how to query this information.

### 5.8.4.7 Options for debugging your program

The following option is available.

-g

This option can currently be used to generate additional errors for the built-in functions that allow you to enqueue commands on a device (refer to *section 6.13.17*).

## 5.8.5 Linker Options

This specification defines a standard set of linker options that must be supported by the OpenCL C compiler when linking compiled programs online or offline. These linker options are categorized as library linking options and program linking options. These may be extended by a set of vendor- or platform-specific options.

### 5.8.5.1 Library Linking Options

The following options can be specified when creating a library of compiled binaries.

-create-library

Create a library of compiled binaries specified in *input\_programs* argument to **clLinkProgram**.

-enable-link-options

Allows the linker to modify the library behavior based on one or more link options (described in *section 5.8.5.2*) when this library is linked with a program executable. This option must be specified with the `-create-library` option.

### 5.8.5.2 Program Linking Options

The following options can be specified when linking a program executable.

-cl-denorms-are-zero

-cl-no-signed-zeroes

-cl-unsafe-math-optimizations

-cl-finite-math-only

-cl-fast-relaxed-math

The options are described in *section 5.8.4.2* and *section 5.8.4.3*. The linker may apply these options to all compiled program objects specified to **clLinkProgram**. The linker may apply these options only to libraries which were created with the `-enable-link-option`.


## 5.8.6 Unloading the OpenCL Compiler

The function

```
cl_int  clUnloadPlatformCompiler (cl_platform_id platform)
```

allows the implementation to release the resources allocated by the OpenCL compiler for *platform*. This is a hint from the application and does not guarantee that the compiler will not be used in the future or that the compiler will actually be unloaded by the implementation. Calls to **clBuildProgram**, **clCompileProgram** or **clLinkProgram** after **clUnloadPlatformCompiler** will reload the compiler, if necessary, to build the appropriate program executable.

**clUnloadPlatformCompiler** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

 CL\_INVALID\_PLATFORM if *platform* is not a valid platform.

## 5.8.7 Program Object Queries

The function

```
cl_int          clGetProgramInfo (cl_program program,
                                cl_program_info param_name,
                                size_t param_value_size,
                                void *param_value,
                                size_t *param_value_size_ret)
```

returns information about the program object.

*program* specifies the program object being queried.

*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetProgramInfo** is described in *table 5.17*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.17*.

*param\_value\_size\_ret* returns the actual size in bytes of data copied to *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

cl_program_info	Return Type	Info. returned in <i>param_value</i>
CL_PROGRAM_REFERENCE_COUNT <sup>26</sup>	cl_uint	Return the <i>program</i> reference count.
CL_PROGRAM_CONTEXT	cl_context	Return the context specified when the program object is created
CL_PROGRAM_NUM_DEVICES	cl_uint	Return the number of devices associated with <i>program</i> .
CL_PROGRAM_DEVICES	cl_device_id[]	Return the list of devices associated with the program object. This can be the devices associated with context on which the program object has been created or can be a subset of devices that are specified when a program object is created using <b>clCreateProgramWithBinary</b> .

<sup>26</sup> The reference count returned should be considered immediately stale. It is unsuitable for general use in applications. This feature is provided for identifying memory leaks.

<b>CL_PROGRAM_SOURCE</b>	char[]	<p>Return the program source code specified by <b>clCreateProgramWithSource</b>. The source string returned is a concatenation of all source strings specified to <b>clCreateProgramWithSource</b> with a null terminator. The concatenation strips any nulls in the original source strings.</p> <p>If <i>program</i> is created using <b>clCreateProgramWithBinary</b> or <b>clCreateProgramWithBuiltinKernels</b>, a null string or the appropriate program source code is returned depending on whether or not the program source code is stored in the binary.</p> <p>The actual number of characters that represents the program source code including the null terminator is returned in <i>param_value_size_ret</i>.</p>
<b>CL_PROGRAM_BINARY_SIZES</b>	size_t[]	<p>Returns an array that contains the size in bytes of the program binary (could be an executable binary, compiled binary or library binary) for each device associated with <i>program</i>. The size of the array is the number of devices associated with <i>program</i>. If a binary is not available for a device(s), a size of zero is returned.</p> <p>If <i>program</i> is created using <b>clCreateProgramWithBuiltinKernels</b>, the implementation may return zero in any entries of the returned array.</p>
<b>CL_PROGRAM_BINARIES</b>	unsigned char *[]	<p>Return the program binaries (could be an executable binary, compiled binary or library binary) for all devices associated with <i>program</i>. For each device in <i>program</i>, the binary returned can be the binary specified for the device when <i>program</i> is created with <b>clCreateProgramWithBinary</b> or it can be the executable binary generated by <b>clBuildProgram</b> or <b>clLinkProgram</b>. If <i>program</i> is created with</p>

		<p><b>clCreateProgramWithSource</b>, the binary returned is the binary generated by <b>clBuildProgram</b>, <b>clCompileProgram</b> or <b>clLinkProgram</b>. The bits returned can be an implementation-specific intermediate representation (a.k.a. IR) or device specific executable bits or both. The decision on which information is returned in the binary is up to the OpenCL implementation.</p> <p><i>param_value</i> points to an array of <i>n</i> pointers allocated by the caller, where <i>n</i> is the number of devices associated with program. The buffer sizes needed to allocate the memory that these <i>n</i> pointers refer to can be queried using the CL_PROGRAM_BINARY_SIZES query as described in this table.</p> <p>Each entry in this array is used by the implementation as the location in memory where to copy the program binary for a specific device, if there is a binary available. To find out which device the program binary in the array refers to, use the CL_PROGRAM_DEVICES query to get the list of devices. There is a one-to-one correspondence between the array of <i>n</i> pointers returned by CL_PROGRAM_BINARIES and array of devices returned by CL_PROGRAM_DEVICES.</p> <p>If an entry value in the array is NULL, the implementation skips copying the program binary for the specific device identified by the array index.</p>
<b>CL_PROGRAM_NUM_KERNELS</b>	size_t	Returns the number of kernels declared in <i>program</i> that can be created with <b>clCreateKernel</b> . This information is only available after a successful program executable has been built for at least one device in the list of devices associated with <i>program</i> .
<b>CL_PROGRAM_KERNEL_</b>	char[]	Returns a semi-colon separated list of



NAMES		kernel names in <i>program</i> that can be created with <b>clCreateKernel</b> . This information is only available after a successful program executable has been built for at least one device in the list of devices associated with <i>program</i> .
-------	--	---

**Table 5.17** *clGetProgramInfo* parameter queries.

**clGetProgramInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is < size of return type as described in *table 5.17* and *param\_value* is not NULL.
- ✚ CL\_INVALID\_PROGRAM if *program* is a not a valid program object.
- ✚ CL\_INVALID\_PROGRAM\_EXECUTABLE if *param\_name* is CL\_PROGRAM\_NUM\_KERNELS or CL\_PROGRAM\_KERNEL\_NAMES and a successful program executable has not been built for at least one device in the list of devices associated with *program*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int          clGetProgramBuildInfo (cl_program program,
                                         cl_device_id device,
                                         cl_program_build_info param_name,
                                         size_t param_value_size,
                                         void *param_value,
                                         size_t *param_value_size_ret)
```

returns build information for each device in the program object.

*program* specifies the program object being queried.

*device* specifies the device for which build information is being queried. *device* must be a valid device associated with *program*.

*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetProgramBuildInfo** is described in *table 5.18*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.18*.

*param\_value\_size\_ret* returns the actual size in bytes of data copied to *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

<b>cl_program_build_info</b>	<b>Return Type</b>	<b>Info. returned in <i>param_value</i></b>
<b>CL_PROGRAM_BUILD_STATUS</b>	cl_build_status	<p>Returns the build, compile or link status, whichever was performed last on <i>program</i> for <i>device</i>.</p> <p>This can be one of the following:</p> <p>CL_BUILD_NONE. The build status returned if no <b>clBuildProgram</b>, <b>clCompileProgram</b> or <b>clLinkProgram</b> has been performed on the specified program object for <i>device</i>.</p> <p>CL_BUILD_ERROR. The build status returned if <b>clBuildProgram</b>, <b>clCompileProgram</b> or <b>clLinkProgram</b> whichever was performed last on the specified program object for <i>device</i> generated an error.</p> <p>CL_BUILD_SUCCESS. The build status returned if <b>clBuildProgram</b>, <b>clCompileProgram</b> or <b>clLinkProgram</b> whichever was performed last on the specified program object for <i>device</i> was successful.</p> <p>CL_BUILD_IN_PROGRESS. The build status returned if <b>clBuildProgram</b>, <b>clCompileProgram</b> or <b>clLinkProgram</b> whichever was performed last on the specified program object for <i>device</i> has not finished.</p>
<b>CL_PROGRAM_BUILD_OPTIONS</b>	char[]	<p>Return the build, compile or link options specified by the <i>options</i> argument in <b>clBuildProgram</b>, <b>clCompileProgram</b> or <b>clLinkProgram</b>, whichever was performed last</p>

		<p>on <i>program</i> for <i>device</i>.</p> <p>If build status of <i>program</i> for <i>device</i> is CL_BUILD_NONE, an empty string is returned.</p>
<b>CL_PROGRAM_BUILD_LOG</b>	char[]	<p>Return the build or compile log for <b>clBuildProgram</b> or <b>clCompileProgram</b> whichever was performed last on <i>program</i> for <i>device</i>.</p> <p>If build status of <i>program</i> for <i>device</i> is CL_BUILD_NONE, an empty string is returned.</p>
<b>CL_PROGRAM_BINARY_TYPE</b>	cl_program_binary_type	<p>Return the program binary type for <i>device</i>. This can be one of the following values:</p> <p>CL_PROGRAM_BINARY_TYPE_NONE – There is no binary associated with <i>device</i>.</p> <p>CL_PROGRAM_BINARY_TYPE_COMPILED_OBJECT – A compiled binary is associated with <i>device</i>. This is the case if <i>program</i> was created using <b>clCreateProgramWithSource</b> and compiled using <b>clCompileProgram</b> or a compiled binary is loaded using <b>clCreateProgramWithBinary</b>.</p> <p>CL_PROGRAM_BINARY_TYPE_LIBRARY – A library binary is associated with <i>device</i>. This is the case if <i>program</i> was created by <b>clLinkProgram</b> which is called with the –create-library link option or if a library binary is loaded using <b>clCreateProgramWithBinary</b>.</p> <p>CL_PROGRAM_BINARY_TYPE_EXECUTABLE – An executable binary is associated with <i>device</i>. This is the case if <i>program</i> was created by <b>clLinkProgram</b> without the –create-library link option or <i>program</i> was created by <b>clBuildProgram</b> or an executable binary is loaded using <b>clCreateProgramWithBinary</b>.</p>
<b>CL_PROGRAM_BUILD_GLOBAL_VARIABLE_TOTAL_SIZE</b>	size_t	<p>The total amount of storage, in bytes, used by program variables in the global address space.</p>

**Table 5.18** *clGetProgramBuildInfo* parameter queries.

**clGetProgramBuildInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_DEVICE if *device* is not in the list of devices associated with *program*.
- ✚ CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is < size of return type as described in *table 5.18* and *param\_value* is not NULL.
- ✚ CL\_INVALID\_PROGRAM if *program* is not a valid program object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

NOTE:

A program binary (compiled binary, library binary or executable binary) built for a parent device can be used by all its sub-devices. If a program binary has not been built for a sub-device, the program binary associated with the parent device will be used.

A program binary for a device specified with **clCreateProgramWithBinary** or queried using **clGetProgramInfo** can be used as the binary for the associated root device, and all sub-devices created from the root-level device or sub-devices thereof.

## 5.9 Kernel Objects

A kernel is a function declared in a program. A kernel is identified by the `__kernel` qualifier applied to any function in a program. A kernel object encapsulates the specific `__kernel` function declared in a program and the argument values to be used when executing this `__kernel` function.

### 5.9.1 Creating Kernel Objects

To create a kernel object, use the function








```
cl_kernel      clCreateKernel (cl_program program,  
                                const char *kernel_name,  
                                cl_int *errcode_ret)
```

*program* is a program object with a successfully built executable.

*kernel\_name* is a function name in the program declared with the `__kernel` qualifier.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreateKernel** returns a valid non-zero kernel object and *errcode\_ret* is set to CL\_SUCCESS if the kernel object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

-  CL\_INVALID\_PROGRAM if *program* is not a valid program object.
-  CL\_INVALID\_PROGRAM\_EXECUTABLE if there is no successfully built executable for *program*.
-  CL\_INVALID\_KERNEL\_NAME if *kernel\_name* is not found in *program*.
-  CL\_INVALID\_KERNEL\_DEFINITION if the function definition for `__kernel` function given by *kernel\_name* such as the number of arguments, the argument types are not the same for all devices for which the *program* executable has been built.
-  CL\_INVALID\_VALUE if *kernel\_name* is NULL.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the

OpenCL implementation on the host.

The function

```
cl_int      clCreateKernelsInProgram (cl_program program,
                                       cl_uint num_kernels,
                                       cl_kernel *kernels,
                                       cl_uint *num_kernels_ret)
```

creates kernel objects for all kernel functions in *program*. Kernel objects are not created for any `__kernel` functions in *program* that do not have the same function definition across all devices for which a program executable has been successfully built.






*program* is a program object with a successfully built executable.

*num\_kernels* is the size of memory pointed to by *kernels* specified as the number of `cl_kernel` entries.

*kernels* is the buffer where the kernel objects for kernels in *program* will be returned. If *kernels* is NULL, it is ignored. If *kernels* is not NULL, *num\_kernels* must be greater than or equal to the number of kernels in *program*.

*num\_kernels\_ret* is the number of kernels in *program*. If *num\_kernels\_ret* is NULL, it is ignored.

**clCreateKernelsInProgram** will return `CL_SUCCESS` if the kernel objects were successfully allocated. Otherwise, it returns one of the following errors:

-  `CL_INVALID_PROGRAM` if *program* is not a valid program object.
-  `CL_INVALID_PROGRAM_EXECUTABLE` if there is no successfully built executable for any device in *program*.
-  `CL_INVALID_VALUE` if *kernels* is not NULL and *num\_kernels* is less than the number of kernels in *program*.
-  `CL_OUT_OF_RESOURCES` if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  `CL_OUT_OF_HOST_MEMORY` if there is a failure to allocate resources required by the OpenCL implementation on the host.




Kernel objects can only be created once you have a program object with a valid program source or binary loaded into the program object and the program executable has been successfully built for one or more devices associated with program. No changes to the program executable are allowed while there are kernel objects associated with a program object. This means that calls to

**clBuildProgram** and **clCompileProgram** return CL\_INVALID\_OPERATION if there are kernel objects attached to a program object. The OpenCL context associated with *program* will be the context associated with *kernel*. The list of devices associated with *program* are the devices associated with *kernel*. Devices associated with a program object for which a valid program executable has been built can be used to execute kernels declared in the program object.

The function

cl\_int            **clRetainKernel** (cl\_kernel *kernel*)

increments the *kernel* reference count. **clRetainKernel** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:




-  CL\_INVALID\_KERNEL if *kernel* is not a valid kernel object.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

**clCreateKernel** or **clCreateKernelsInProgram** do an implicit retain.

The function

cl\_int            **clReleaseKernel** (cl\_kernel *kernel*)

decrements the *kernel* reference count. **clReleaseKernel** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_KERNEL if *kernel* is not a valid kernel object.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The kernel object is deleted once the number of instances that are retained to *kernel* become zero and the kernel object is no longer needed by any enqueued commands that use *kernel*. Using this function to release a reference that was not obtained by creating the object or by calling **clRetainKernel** causes undefined behavior.

## 5.9.2 Setting Kernel Arguments

To execute a kernel, the kernel arguments must be set.

The function

```
cl_int  clSetKernelArg (cl_kernel kernel,
                        cl_uint arg_index,
                        size_t arg_size,
                        const void *arg_value)
```

is used to set the argument value for a specific argument of a kernel.

*kernel* is a valid kernel object.

*arg\_index* is the argument index. Arguments to the kernel are referred by indices that go from 0 for the leftmost argument to  $n - 1$ , where  $n$  is the total number of arguments declared by a kernel.

For example, consider the following kernel:

```
kernel void
image_filter (int n, int m,
              constant float *filter_weights,
              read_only image2d_t src_image,
              write_only image2d_t dst_image)
{
    ...
}
```

Argument index values for *image\_filter* will be 0 for *n*, 1 for *m*, 2 for *filter\_weights*, 3 for *src\_image* and 4 for *dst\_image*.

*arg\_value* is a pointer to data that should be used as the argument value for argument specified by *arg\_index*. The argument data pointed to by *arg\_value* is copied and the *arg\_value* pointer can therefore be reused by the application after **clSetKernelArg** returns. The argument value specified is the value used by all API calls that enqueue *kernel* (**clEnqueueNDRangeKernel**) until the argument value is changed by a call to **clSetKernelArg** for *kernel*.

If the argument is a memory object (buffer, pipe, image or image array), the *arg\_value* entry will be a pointer to the appropriate buffer, pipe, image or image array object. The memory object must be created with the context associated with the kernel object. If the argument is a buffer object, the *arg\_value* pointer can be NULL or point to a NULL value in which case a NULL value will be used as the value for the argument declared as a pointer to global or constant memory in the kernel. If the argument is declared with the `local` qualifier, the *arg\_value* entry



must be NULL. If the argument is of type *sampler\_t*, the *arg\_value* entry must be a pointer to the sampler object. If the argument is of type *queue\_t*, the *arg\_value* entry must be a pointer to the device queue object.

If the argument is declared to be a pointer of a built-in scalar or vector type, or a user defined structure type in the global or constant address space, the memory object specified as argument value must be a buffer object (or NULL). If the argument is declared with the `constant` qualifier, the size in bytes of the memory object cannot exceed `CL_DEVICE_MAX_CONSTANT_BUFFER_SIZE` and the number of arguments declared as pointers to *constant* memory cannot exceed `CL_DEVICE_MAX_CONSTANT_ARGS`.

The memory object specified as argument value must be a pipe object if the argument is declared with the *pipe* qualifier.

The memory object specified as argument value must be a 2D image object if the argument is declared to be of type *image2d\_t*. The memory object specified as argument value must be a 2D image object with image channel order = `CL_DEPTH` if the argument is declared to be of type *image2d\_depth\_t*. The memory object specified as argument value must be a 3D image object if argument is declared to be of type *image3d\_t*. The memory object specified as argument value must be a 1D image object if the argument is declared to be of type *image1d\_t*. The memory object specified as argument value must be a 1D image buffer object if the argument is declared to be of type *image1d\_buffer\_t*. The memory object specified as argument value must be a 1D image array object if argument is declared to be of type *image1d\_array\_t*. The memory object specified as argument value must be a 2D image array object if argument is declared to be of type *image2d\_array\_t*. The memory object specified as argument value must be a 2D image array object with image channel order = `CL_DEPTH` if argument is declared to be of type *image2d\_array\_depth\_t*.

For all other kernel arguments, the *arg\_value* entry must be a pointer to the actual data to be used as argument value.

*arg\_size* specifies the size of the argument value. If the argument is a memory object, the size is the size of the memory object. For arguments declared with the `local` qualifier, the size specified will be the size in bytes of the buffer that must be allocated for the `local` argument. If the argument is of type *sampler\_t*, the *arg\_size* value must be equal to `sizeof(cl_sampler)`. If the argument is of type *queue\_t*, the *arg\_size* value must be equal to `sizeof(cl_command_queue)`. For all other arguments, the size will be the size of argument type.

NOTE: A kernel object does not update the reference count for objects such as memory, sampler objects specified as argument values by **clSetKernelArg**. Users may not rely on a kernel object to retain objects specified as argument values to the kernel<sup>27</sup>.

---

<sup>27</sup> Implementations shall not allow `cl_kernel` objects to hold reference counts to `cl_kernel` arguments, because no mechanism is provided for the user to tell the kernel to release that ownership right. If the kernel holds ownership rights on kernel args, that would make it impossible for the user to tell with certainty when he may safely release

**clSetKernelArg** returns CL\_SUCCESS if the function was executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_KERNEL if *kernel* is not a valid kernel object.
- ✚ CL\_INVALID\_ARG\_INDEX if *arg\_index* is not a valid argument index.
- ✚ CL\_INVALID\_ARG\_VALUE if *arg\_value* specified is not a valid value.
- ✚ CL\_INVALID\_MEM\_OBJECT for an argument declared to be a memory object when the specified *arg\_value* is not a valid memory object.
- ✚ CL\_INVALID\_SAMPLER for an argument declared to be of type *sampler\_t* when the specified *arg\_value* is not a valid sampler object.
- ✚ CL\_INVALID\_DEVICE\_QUEUE for an argument declared to be of type *queue\_t* when the specified *arg\_value* is not a valid device queue object.
- ✚ CL\_INVALID\_ARG\_SIZE if *arg\_size* does not match the size of the data type for an argument that is not a memory object or if the argument is a memory object and *arg\_size*  $\neq \text{sizeof}(\text{cl\_mem})$  or if *arg\_size* is zero and the argument is declared with the `local` qualifier or if the argument is a sampler and *arg\_size*  $\neq \text{sizeof}(\text{cl\_sampler})$ .
- ✚ CL\_INVALID\_ARG\_VALUE if the argument is an image declared with the `read_only` qualifier and *arg\_value* refers to an image object created with *cl\_mem\_flags* of CL\_MEM\_WRITE or if the image argument is declared with the `write_only` qualifier and *arg\_value* refers to an image object created with *cl\_mem\_flags* of CL\_MEM\_READ.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clSetKernelArgSVMPointer (cl_kernel kernel,
                                cl_uint arg_index,
                                const void *arg_value)
```

---

user allocated resources associated with OpenCL objects such as the `cl_mem` backing store used with CL\_MEM\_USE\_HOST\_PTR.






is used to set a SVM pointer as the argument value for a specific argument of a kernel.

*kernel* is a valid kernel object.

*arg\_index* is the argument index. Arguments to the kernel are referred by indices that go from 0 for the leftmost argument to  $n - 1$ , where  $n$  is the total number of arguments declared by a kernel.

*arg\_value* is the SVM pointer that should be used as the argument value for argument specified by *arg\_index*. The SVM pointer specified is the value used by all API calls that enqueue *kernel* (**clEnqueueNDRangeKernel**) until the argument value is changed by a call to **clSetKernelArgSVMPointer** for *kernel*. The SVM pointer can only be used for arguments that are declared to be a pointer to global or constant memory. The SVM pointer value must be aligned according to the argument's type. For example, if the argument is declared to be `global float4 *p`, the SVM pointer value passed for *p* must be at a minimum aligned to a `float4`. The SVM pointer value specified as the argument value can be the pointer returned by **clSVMAlloc** or can be a pointer + offset into the SVM region.

**clSetKernelArgSVMPointer** returns `CL_SUCCESS` if the function was executed successfully. Otherwise, it returns one of the following errors:

-  `CL_INVALID_KERNEL` if *kernel* is not a valid kernel object.
-  `CL_INVALID_ARG_INDEX` if *arg\_index* is not a valid argument index.
-  `CL_INVALID_ARG_VALUE` if *arg\_value* specified is not a valid value.
-  `CL_OUT_OF_RESOURCES` if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  `CL_OUT_OF_HOST_MEMORY` if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clSetKernelExecInfo (cl_kernel kernel,
                             cl_kernel_exec_info param_name,
                             size_t param_value_size,
                             const void *param_value)
```

can be used to pass additional information other than argument values to a kernel.

*kernel* specifies the kernel object being queried.

*param\_name* specifies the information to be passed to kernel. The list of supported *param\_name* types and the corresponding values passed in *param\_value* is described in *table 5.19*.





*param\_value\_size* specifies the size in bytes of the memory pointed to by *param\_value*.

*param\_value* is a pointer to memory where the appropriate values determined by *param\_name* are specified.

cl_kernel_exec_info	Type	Description
CL_KERNEL_EXEC_INFO_SVM_PTRS	void *[]	SVM pointers must reference locations contained entirely within buffers that are passed to <i>kernel</i> as arguments, or that are passed through the execution information.  Non-argument SVM buffers must be specified by passing pointers to those buffers via <b>clSetKernelExecInfo</b> for coarse-grain and fine-grain buffer SVM allocations but not for fine-grain system SVM allocations.
CL_KERNEL_EXEC_INFO_SVM_FINE_GRAIN_SYSTEM	cl_bool	This flag indicates whether the kernel uses pointers that are fine grain system SVM allocations. These fine grain system SVM pointers may be passed as arguments or defined in SVM buffers that are passed as arguments to <i>kernel</i> .

**Table 5.19** *clSetKernelExecInfo* parameter values.

**clSetKernelExecInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_KERNEL if *kernel* is a not a valid kernel object.
-  CL\_INVALID\_VALUE if *param\_name* is not valid, if *param\_value* is NULL or if the size specified by *param\_value\_size* is not valid.
-  CL\_INVALID\_OPERATION if *param\_name* = CL\_KERNEL\_EXEC\_INFO\_SVM\_FINE\_GRAIN\_SYSTEM and *param\_value* = CL\_TRUE but no devices in context associated with *kernel* support fine-grain system SVM allocations.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.

- CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## NOTES

1. Coarse-grain or fine-grain buffer SVM pointers used by a kernel which are not passed as a kernel arguments must be specified using **clSetKernelExecInfo** with **CL\_KERNEL\_EXEC\_INFO\_SVM\_PTRS**. For example, if SVM buffer A contains a pointer to another SVM buffer B, and the kernel dereferences that pointer, then a pointer to B must either be passed as an argument in the call to that kernel or it must be made available to the kernel using **clSetKernelExecInfo**. For example, we might pass extra SVM pointers as follows:

```
clSetKernelExecInfo(kernel,
                    CL_KERNEL_EXEC_INFO_SVM_PTRS,
                    num_ptrs * sizeof(void *),
                    extra_svm_ptr_list);
```

Here `num_ptrs` specifies the number of additional SVM pointers while `extra_svm_ptr_list` specifies a pointer to memory containing those SVM pointers.

When calling **clSetKernelExecInfo** with **CL\_KERNEL\_EXEC\_INFO\_SVM\_PTRS** to specify pointers to non-argument SVM buffers as extra arguments to a kernel, each of these pointers can be the SVM pointer returned by **clSVMAlloc** or can be a pointer + offset into the SVM region. It is sufficient to provide one pointer for each SVM buffer used.

2. **CL\_KERNEL\_EXEC\_INFO\_SVM\_FINE\_GRAIN\_SYSTEM** is used to indicate whether SVM pointers used by a kernel will refer to system allocations or not.

**CL\_KERNEL\_EXEC\_INFO\_SVM\_FINE\_GRAIN\_SYSTEM = CL\_FALSE** indicates that the OpenCL implementation may assume that system pointers are not passed as kernel arguments and are not stored inside SVM allocations passed as kernel arguments.

**CL\_KERNEL\_EXEC\_INFO\_SVM\_FINE\_GRAIN\_SYSTEM = CL\_TRUE** indicates that the OpenCL implementation must assume that system pointers might be passed as kernel arguments and/or stored inside SVM allocations passed as kernel arguments. In this case, if the device to which the kernel is enqueued does not support system SVM pointers, **clEnqueueNDRangeKernel** will return a **CL\_INVALID\_OPERATION** error. If none of the devices in the context associated with kernel support fine-grain system SVM allocations, **clSetKernelExecInfo** will return a **CL\_INVALID\_OPERATION** error.

If **clSetKernelExecInfo** has not been called with a value for **CL\_KERNEL\_EXEC\_INFO\_SVM\_FINE\_GRAIN\_SYSTEM**, the default value is used for this kernel attribute. The default value depends on whether the device on which the kernel is enqueued supports fine-grain system SVM allocations. If so, the default value used is **CL\_TRUE** (system pointers might be passed); otherwise, the default is **CL\_FALSE**.

3. A call to **clSetKernelExecInfo** for a given value of *param\_name* replaces any prior value passed for that value of *param\_name*. Only one *param\_value* will be stored for each value of *param\_name*.

### 5.9.3 Kernel Object Queries

The function

```
cl_int          clGetKernelInfo (cl_kernel kernel,
                                cl_kernel_info param_name,
                                size_t param_value_size,
                                void *param_value,
                                size_t *param_value_size_ret)
```

returns information about the kernel object.

*kernel* specifies the kernel object being queried.

*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetKernelInfo** is described in *table 5.20*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.20*.

*param\_value\_size\_ret* returns the actual size in bytes of data copied to *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

cl_kernel_info	Return Type	Info. returned in <i>param_value</i>
CL_KERNEL_FUNCTION_NAME	char[]	Return the kernel function name.
CL_KERNEL_NUM_ARGS	cl_uint	Return the number of arguments to <i>kernel</i> .
CL_KERNEL_REFERENCE_COUNT <sup>28</sup>	cl_uint	Return the <i>kernel</i> reference count.
CL_KERNEL_CONTEXT	cl_context	Return the context associated with

<sup>28</sup> The reference count returned should be considered immediately stale. It is unsuitable for general use in applications. This feature is provided for identifying memory leaks.

		<i>kernel</i> .
<b>CL_KERNEL_PROGRAM</b>	cl_program	Return the program object associated with <i>kernel</i> .
<b>CL_KERNEL_ATTRIBUTES</b>	char[]	<p>Returns any attributes specified using the <code>__attribute__</code> qualifier with the kernel function declaration in the program source. These attributes include attributes described in <i>section 6.11.2</i> and other attributes supported by an implementation.</p> <p>Attributes are returned as they were declared inside <code>__attribute__((...))</code>, with any surrounding whitespace and embedded newlines removed. When multiple attributes are present, they are returned as a single, space delimited string.</p>

**Table 5.20** *clGetKernelInfo parameter queries.*

**clGetKernelInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is < size of return type as described in *table 5.20* and *param\_value* is not NULL.
- ✚ CL\_INVALID\_KERNEL if *kernel* is a not a valid kernel object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int clGetKernelWorkGroupInfo (cl_kernel kernel,
                                cl_device_id device,
                                cl_kernel_work_group_info param_name,
                                size_t param_value_size,
                                void *param_value,
                                size_t *param_value_size_ret)
```

returns information about the kernel object that may be specific to a device.

*kernel* specifies the kernel object being queried.

*device* identifies a specific device in the list of devices associated with *kernel*. The list of devices is the list of devices in the OpenCL context that is associated with *kernel*. If the list of devices associated with *kernel* is a single device, *device* can be a NULL value.

*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetKernelWorkGroupInfo** is described in *table 5.21*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.21*.

*param\_value\_size\_ret* returns the actual size in bytes of data copied to *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

<b>cl_kernel_work_group_info</b>	<b>Return Type</b>	<b>Info. returned in <i>param_value</i></b>
<b>CL_KERNEL_GLOBAL_WORK_SIZE</b>	size_t[3]	<p>This provides a mechanism for the application to query the maximum global size that can be used to execute a kernel (i.e. <i>global_work_size</i> argument to <b>clEnqueueNDRangeKernel</b>) on a custom device given by <i>device</i> or a built-in kernel on an OpenCL device given by <i>device</i>.</p> <p>If <i>device</i> is not a custom device and <i>kernel</i> is not a built-in kernel, <b>clGetKernelWorkGroupInfo</b> returns the error CL_INVALID_VALUE.</p>
<b>CL_KERNEL_WORK_GROUP_SIZE</b>	size_t	<p>This provides a mechanism for the application to query the maximum work-group size that can be used to execute the kernel on a specific device given by <i>device</i>. The OpenCL implementation uses the resource requirements of the kernel (register usage etc.) to determine what this work-group size should be.</p> <p>As a result and unlike</p>



		<p><b>CL_DEVICE_MAX_WORK_GROUP_SIZE</b> this value may vary from one kernel to another as well as one device to another.</p> <p><b>CL_KERNEL_WORK_GROUP_SIZE</b> will be less than or equal to <b>CL_DEVICE_MAX_WORK_GROUP_SIZE</b> for a given kernel object.</p>
<b>CL_KERNEL_COMPILE_WORK_GROUP_SIZE</b>	size_t[3]	<p>Returns the work-group size specified by the <code>__attribute__((reqd_work_group_size(X, Y, Z)))</code> qualifier. Refer to <i>section 6.7.2</i>.</p> <p>If the work-group size is not specified using the above attribute qualifier (0, 0, 0) is returned.</p>
<b>CL_KERNEL_LOCAL_MEM_SIZE</b>	cl_ulong	<p>Returns the amount of local memory in bytes being used by a kernel. This includes local memory that may be needed by an implementation to execute the kernel, variables declared inside the kernel with the <code>__local</code> address qualifier and local memory to be allocated for arguments to the kernel declared as pointers with the <code>__local</code> address qualifier and whose size is specified with <b>clSetKernelArg</b>.</p> <p>If the local memory size, for any pointer argument to the kernel declared with the <code>__local</code> address qualifier, is not specified, its size is assumed to be 0.</p>
<b>CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE</b>	size_t	<p>Returns the preferred multiple of work-group size for launch. This is a performance hint. Specifying a work-group size that is not a multiple of the value returned by this query as the value of the local work size argument to <b>clEnqueueNDRangeKernel</b> will not fail to enqueue the kernel for execution unless the work-group size specified is larger than the device maximum.</p>
<b>CL_KERNEL_PRIVATE_MEM_SIZE</b>	cl_ulong	<p>Returns the minimum amount of private</p>

<b>SIZE</b>		memory, in bytes, used by each work-item in the kernel. This value may include any private memory needed by an implementation to execute the kernel, including that used by the language built-ins and variable declared inside the kernel with the <code>private</code> qualifier.
-------------	--	---

**Table 5.21** *clGetKernelWorkGroupInfo* parameter queries.

**clGetKernelWorkGroupInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_DEVICE if *device* is not in the list of devices associated with *kernel* or if *device* is NULL but there is more than one device associated with *kernel*.
- ✚ CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is < size of return type as described in *table 5.21* and *param\_value* is not NULL.
- ✚ CL\_INVALID\_VALUE if *param\_name* is CL\_KERNEL\_GLOBAL\_WORK\_SIZE and *device* is not a custom device and *kernel* is not a built-in kernel.
- ✚ CL\_INVALID\_KERNEL if *kernel* is a not a valid kernel object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clGetKernelArgInfo (cl_kernel kernel,
                           cl_uint arg_indx,
                           cl_kernel_arg_info param_name,
                           size_t param_value_size,
                           void *param_value,
                           size_t *param_value_size_ret)
```

returns information about the arguments of a kernel. Kernel argument information is only available if the program object associated with *kernel* is created with **clCreateProgramWithSource** and the program executable is built with the `-cl-kernel-arg-info` option specified in *options* argument to **clBuildProgram** or **clCompileProgram**.

*kernel* specifies the kernel object being queried.

*arg\_indx* is the argument index. Arguments to the kernel are referred by indices that go from 0 for the leftmost argument to  $n - 1$ , where  $n$  is the total number of arguments declared by a kernel.

*param\_name* specifies the argument information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetKernelArgInfo** is described in [table 5.22](#).

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $>$  size of return type as described in [table 5.22](#). *param\_value\_size\_ret* returns the actual size in bytes of data copied to *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

<b>cl_kernel_arg_info</b>	<b>Return Type</b>	<b>Info. returned in <i>param_value</i></b>
<b>CL_KERNEL_ARG_ADDRESS_QUALIFIER</b>	cl_kernel_arg_address_qualifier	<p>Returns the address qualifier specified for the argument given by <i>arg_indx</i>. This can be one of the following values:</p> <p>CL_KERNEL_ARG_ADDRESS_GLOBAL  CL_KERNEL_ARG_ADDRESS_LOCAL  CL_KERNEL_ARG_ADDRESS_CONSTANT  CL_KERNEL_ARG_ADDRESS_PRIVATE</p> <p>If no address qualifier is specified, the default address qualifier which is CL_KERNEL_ARG_ADDRESS_PRIVATE is returned.</p>
<b>CL_KERNEL_ARG_ACCESS_QUALIFIER</b>	cl_kernel_arg_access_qualifier	<p>Returns the access qualifier specified for the argument given by <i>arg_indx</i>. This can be one of the following values:</p> <p>CL_KERNEL_ARG_ACCESS_READ_ONLY  CL_KERNEL_ARG_ACCESS_WRITE_ONLY  CL_KERNEL_ARG_ACCESS_READ_WRITE  CL_KERNEL_ARG_ACCESS_NONE</p> <p>If argument is not an image type and is not declared with the pipe qualifier, CL_KERNEL_ARG_ACCESS_NONE is returned. If argument is an image type, the access qualifier specified or the default access qualifier is returned.</p>
<b>CL_KERNEL_ARG_TYPE_NAME</b>	char[]	Returns the type name specified for the argument given by <i>arg_indx</i> . The type name returned will be

		the argument type name as it was declared with any whitespace removed. If argument type name is an unsigned scalar type (i.e. unsigned char, unsigned short, unsigned int, unsigned long), uchar, ushort, uint and ulong will be returned. The argument type name returned does not include any type qualifiers.
<b>CL_KERNEL_ARG_TYPE_QUALIFIER</b>	cl_kernel_arg_type_qualifier	<p>Returns the type qualifier specified for the argument given by <i>arg_indx</i>. The returned value can be:</p> <p>CL_KERNEL_ARG_TYPE_CONST  CL_KERNEL_ARG_TYPE_RESTRICT  CL_KERNEL_ARG_TYPE_VOLATILE, a combination of the above enums,  CL_KERNEL_ARG_TYPE_PIPE or  CL_KERNEL_ARG_TYPE_NONE</p> <p>NOTE: CL_KERNEL_ARG_TYPE_VOLATILE is returned if the argument is a pointer and the referenced type is declared with the volatile qualifier. For example, a kernel argument declared as <code>global int volatile *x</code> returns CL_KERNEL_ARG_TYPE_VOLATILE but a kernel argument declared as <code>global int *volatile x</code> does not. Similarly, CL_KERNEL_ARG_TYPE_CONST is returned if the argument is a pointer and the referenced type is declared with the restrict or const qualifier. For example, a kernel argument declared as <code>global int const *x</code> returns CL_KERNEL_ARG_TYPE_CONST but a kernel argument declared as <code>global int *const x</code> does not. CL_KERNEL_ARG_TYPE_RESTRICT will be returned if the pointer type is marked restrict. For example, <code>global int * restrict x</code> returns CL_KERNEL_ARG_TYPE_RESTRICT.</p> <p>If the argument is declared with the constant address space qualifier, the CL_KERNEL_ARG_TYPE_CONST type qualifier will be set.</p>
<b>CL_KERNEL_ARG_NAME</b>	char[]	Returns the name specified for the argument given by <i>arg_indx</i> .

**Table 5.22** *clGetKernelArgInfo parameter queries.*

**clGetKernelArgInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_ARG\_INDEX if *arg\_indx* is not a valid argument index.
- ✚ CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value* size is < size of return type as described in *table 5.22* and *param\_value* is not NULL.
- ✚ CL\_KERNEL\_ARG\_INFO\_NOT\_AVAILABLE if the argument information is not available for kernel.
- ✚ CL\_INVALID\_KERNEL if *kernel* is a not a valid kernel object.

## 5.10 Executing Kernels

The function

```
cl_int      clEnqueueNDRangeKernel (cl_command_queue command_queue,
                                     cl_kernel kernel,
                                     cl_uint work_dim,
                                     const size_t *global_work_offset,
                                     const size_t *global_work_size,
                                     const size_t *local_work_size,
                                     cl_uint num_events_in_wait_list,
                                     const cl_event *event_wait_list,
                                     cl_event *event)
```

enqueues a command to execute a kernel on a device.

*command\_queue* is a valid host command-queue. The kernel will be queued for execution on the device associated with *command\_queue*.

*kernel* is a valid kernel object. The OpenCL context associated with *kernel* and *command\_queue* must be the same.

*work\_dim* is the number of dimensions used to specify the global work-items and work-items in the work-group. *work\_dim* must be greater than zero and less than or equal to CL\_DEVICE\_MAX\_WORK\_ITEM\_DIMENSIONS.

*global\_work\_offset* can be used to specify an array of *work\_dim* unsigned values that describe the offset used to calculate the global ID of a work-item. If *global\_work\_offset* is NULL, the global IDs start at offset (0, 0, ... 0).

*global\_work\_size* points to an array of *work\_dim* unsigned values that describe the number of global work-items in *work\_dim* dimensions that will execute the kernel function. The total number of global work-items is computed as *global\_work\_size*[0] \* ... \* *global\_work\_size*[*work\_dim* - 1].

*local\_work\_size* points to an array of *work\_dim* unsigned values that describe the number of work-items that make up a work-group (also referred to as the size of the work-group) that will execute the kernel specified by *kernel*. The total number of work-items in a work-group is computed as *local\_work\_size*[0] \* ... \* *local\_work\_size*[*work\_dim* - 1]. The total number of work-items in the work-group must be less than or equal to the CL\_KERNEL\_WORK\_GROUP\_SIZE value specified in *table 5.21* and the number of work-items specified in *local\_work\_size*[0], ... *local\_work\_size*[*work\_dim* - 1] must be less than or equal to the corresponding values specified by CL\_DEVICE\_MAX\_WORK\_ITEM\_SIZES[0], ..., CL\_DEVICE\_MAX\_WORK\_ITEM\_SIZES[*work\_dim* - 1]. The explicitly specified *local\_work\_size* will be used to determine how to break the global work-items specified by

*global\_work\_size* into appropriate work-group instances.

Enabling non-uniform work-groups requires a kernel to be compiled with the `-cl-std=CL2.0` flag and without the `-cl-uniform-work-group-size` flag. If the program was created using `clLinkProgram` and any of the linked programs were compiled in a way that only supports uniform work-group sizes, the linked program only supports uniform work group sizes. If *local\_work\_size* is specified and the OpenCL kernel is compiled without non-uniform work-groups enabled, the values specified in *global\_work\_size*[0], ... *global\_work\_size*[work\_dim - 1] must be evenly divisible by the corresponding values specified in *local\_work\_size*[0], ... *local\_work\_size*[work\_dim - 1].

If non-uniform work-groups are enabled for the kernel, any single dimension for which the global size is not divisible by the local size will be partitioned into two regions. One region will have work-groups that have the same number of work items as was specified by the local size parameter in that dimension. The other region will have work-groups with less than the number of work items specified by the local size parameter in that dimension. The global IDs and group IDs of the work items in the first region will be numerically lower than those in the second, and the second region will be at most one work-group wide in that dimension. Workgroup sizes could be non-uniform in multiple dimensions, potentially producing work-groups of up to 4 different sizes in a 2D range and 8 different sizes in a 3D range.

If *local\_work\_size* is NULL and the kernel is compiled without support for non-uniform work-groups, the OpenCL runtime will implement the ND-range with uniform work-group sizes. If *local\_work\_size* is NULL and non-uniform-work-groups are enabled, the OpenCL runtime is free to implement the ND-range using uniform or non-uniform work-group sizes, regardless of the divisibility of the global work size. If the ND-range is implemented using non-uniform work-group sizes, the work-group sizes, global IDs and group IDs will follow the same pattern as described in above paragraph.

The work-group size to be used for *kernel* can also be specified in the program source using the `__attribute__((reqd_work_group_size(X, Y, Z)))` qualifier (refer to *section 6.7.2*). In this case the size of work group specified by *local\_work\_size* must match the value specified by the `reqd_work_group_size` attribute qualifier.

These work-group instances are executed in parallel across multiple compute units or concurrently on the same compute unit.

Each work-item is uniquely identified by a global identifier. The global ID, which can be read inside the kernel, is computed using the value given by *global\_work\_size* and *global\_work\_offset*. In addition, a work-item is also identified within a work-group by a unique local ID. The local ID, which can also be read by the kernel, is computed using the value given by *local\_work\_size*. The starting local ID is always (0, 0, ... 0).

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed. If *event\_wait\_list* is NULL, then this particular command does not wait on any event to complete. If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list*

must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

*event* returns an event object that identifies this particular kernel-instance. Event objects are unique and can be used to identify a particular kernel-instance later on. If *event* is NULL, no event will be created for this kernel-instance and therefore it will not be possible for the application to query or queue a wait for this particular kernel-instance. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueNDRangeKernel** returns CL\_SUCCESS if the kernel-instance was successfully queued. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_PROGRAM\_EXECUTABLE if there is no successfully built program executable available for device associated with *command\_queue*.
- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_KERNEL if *kernel* is not a valid kernel object.
- ✚ CL\_INVALID\_CONTEXT if context associated with *command\_queue* and *kernel* are not the same or if the context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_KERNEL\_ARGS if the kernel argument values have not been specified or if a kernel argument declared to be a pointer to a type does not point to a named address space.
- ✚ CL\_INVALID\_WORK\_DIMENSION if *work\_dim* is not a valid value (i.e. a value between 1 and 3).
- ✚ CL\_INVALID\_GLOBAL\_WORK\_SIZE if *global\_work\_size* is NULL, or if any of the values specified in *global\_work\_size*[0], ... *global\_work\_size*[*work\_dim* - 1] are 0 or exceed the range given by the `sizeof(size_t)` for the device on which the kernel-instance will be enqueued.
- ✚ CL\_INVALID\_GLOBAL\_OFFSET if the value specified in *global\_work\_size* + the corresponding values in *global\_work\_offset* for any dimensions is greater than the `sizeof(size_t)` for the device on which the kernel-instance will be enqueued.



- ✚ CL\_INVALID\_WORK\_GROUP\_SIZE if *local\_work\_size* is specified and does not match the work-group size for *kernel* in the program source given by the `__attribute__((reqd_work_group_size(X, Y, Z)))` qualifier.
- ✚ CL\_INVALID\_WORK\_GROUP\_SIZE if *local\_work\_size* is specified and the total number of work-items in the work-group computed as *local\_work\_size*[0] \* ... *local\_work\_size*[*work\_dim* - 1] is greater than the value specified by CL\_KERNEL\_WORK\_GROUP\_SIZE in *table 5.21*.
- ✚ CL\_INVALID\_WORK\_GROUP\_SIZE if the program was compiled with `-cl-uniform-work-group-size` and the number of work-items specified by *global\_work\_size* is not evenly divisible by size of work-group given by *local\_work\_size* or if the size of work-group is specified using `__attribute__((reqd_work_group_size(X, Y, Z)))` qualifier and *local\_work\_size* is NULL.
- ✚ CL\_INVALID\_WORK\_ITEM\_SIZE if the number of work-items specified in any of *local\_work\_size*[0], ... *local\_work\_size*[*work\_dim* - 1] is greater than the corresponding values specified by CL\_DEVICE\_MAX\_WORK\_ITEM\_SIZES[0], .... CL\_DEVICE\_MAX\_WORK\_ITEM\_SIZES[*work\_dim* - 1].
- ✚ CL\_MISALIGNED\_SUB\_BUFFER\_OFFSET if a sub-buffer object is specified as the value for an argument that is a buffer object and the *offset* specified when the sub-buffer object is created is not aligned to CL\_DEVICE\_MEM\_BASE\_ADDR\_ALIGN value for device associated with *queue*.
- ✚ CL\_INVALID\_IMAGE\_SIZE if an image object is specified as an argument value and the image dimensions (image width, height, specified or compute row and/or slice pitch) are not supported by device associated with *queue*.
- ✚ CL\_IMAGE\_FORMAT\_NOT\_SUPPORTED if an image object is specified as an argument value and the image format (image channel order and data type) is not supported by device associated with *queue*.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to queue the execution instance of *kernel* on the command-queue because of insufficient resources needed to execute the kernel. For example, the explicitly specified *local\_work\_size* causes a failure to execute the kernel because of insufficient resources such as registers or local memory. Another example would be the number of read-only image args used in *kernel* exceed the CL\_DEVICE\_MAX\_READ\_IMAGE\_ARGS value for device or the number of write-only and read-write image args used in *kernel* exceed the CL\_DEVICE\_MAX\_READ\_WRITE\_IMAGE\_ARGS value for device or the number of samplers used in *kernel* exceed CL\_DEVICE\_MAX\_SAMPLERS for device.
- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with image or buffer objects specified as arguments to *kernel*.

- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_INVALID\_OPERATION if SVM pointers are passed as arguments to a kernel and the device does not support SVM or if system pointers are passed as arguments to a kernel and/or stored inside SVM allocations passed as kernel arguments and the device does not support fine grain system SVM allocations.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int  clEnqueueNativeKernel (cl_command_queue command_queue,
                               void (CL_CALLBACK *user_func)(void *)
                               void *args,
                               size_t cb_args,
                               cl_uint num_mem_objects,
                               const cl_mem *mem_list,
                               const void **args_mem_loc,
                               cl_uint num_events_in_wait_list,
                               const cl_event *event_wait_list,
                               cl_event *event)
```

enqueues a command to execute a native C/C++ function not compiled using the OpenCL compiler.

*command\_queue* is a valid host command-queue. A native user function can only be executed on a command-queue created on a device that has CL\_EXEC\_NATIVE\_KERNEL capability set in CL\_DEVICE\_EXECUTION\_CAPABILITIES as specified in *table 4.3*.

*user\_func* is a pointer to a host-callable user function.

*args* is a pointer to the args list that *user\_func* should be called with.

*cb\_args* is the size in bytes of the args list that *args* points to.

The data pointed to by *args* and *cb\_args* bytes in size will be copied and a pointer to this copied region will be passed to *user\_func*. The copy needs to be done because the memory objects (cl\_mem values) that *args* may contain need to be modified and replaced by appropriate

pointers to global memory. When **clEnqueueNativeKernel** returns, the memory region pointed to by *args* can be reused by the application.

*num\_mem\_objects* is the number of buffer objects that are passed in *args*.

*mem\_list* is a list of valid buffer objects, if *num\_mem\_objects* > 0. The buffer object values specified in *mem\_list* are memory object handles (*cl\_mem* values) returned by **clCreateBuffer** or NULL.

*args\_mem\_loc* is a pointer to appropriate locations that *args* points to where memory object handles (*cl\_mem* values) are stored. Before the user function is executed, the memory object handles are replaced by pointers to global memory.

*event\_wait\_list*, *num\_events\_in\_wait\_list* and *event* are as described in **clEnqueueNDRangeKernel**.

**clEnqueueNativeKernel** returns CL\_SUCCESS if the user function execution instance was successfully queued. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_VALUE if *user\_func* is NULL.
- ✚ CL\_INVALID\_VALUE if *args* is a NULL value and *cb\_args* > 0, or if *args* is a NULL value and *num\_mem\_objects* > 0.
- ✚ CL\_INVALID\_VALUE if *args* is not NULL and *cb\_args* is 0.
- ✚ CL\_INVALID\_VALUE if *num\_mem\_objects* > 0 and *mem\_list* or *args\_mem\_loc* are NULL.
- ✚ CL\_INVALID\_VALUE if *num\_mem\_objects* = 0 and *mem\_list* or *args\_mem\_loc* are not NULL.
- ✚ CL\_INVALID\_OPERATION if the device associated with *command\_queue* cannot execute the native kernel.
- ✚ CL\_INVALID\_MEM\_OBJECT if one or more memory objects specified in *mem\_list* are not valid or are not buffer objects.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to queue the execution instance of *kernel* on the command-queue because of insufficient resources needed to execute the kernel.

- ✚ CL\_MEM\_OBJECT\_ALLOCATION\_FAILURE if there is a failure to allocate memory for data store associated with buffer objects specified as arguments to *kernel*.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_INVALID\_OPERATION if SVM pointers are passed as arguments to a kernel and the device does not support SVM or if system pointers are passed as arguments to a kernel and/or stored inside SVM allocations passed as kernel arguments and the device does not support fine grain system SVM allocations.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

NOTE:

The total number of read-only images specified as arguments to a kernel cannot exceed CL\_DEVICE\_MAX\_READ\_IMAGE\_ARGS. Each image array argument to a kernel declared with the `read_only` qualifier counts as one image.

The total number of write-only images specified as arguments to a kernel cannot exceed CL\_DEVICE\_MAX\_WRITE\_IMAGE\_ARGS. Each image array argument to a kernel declared with the `write_only` qualifier counts as one image.






The total number of read-write images specified as arguments to a kernel cannot exceed CL\_DEVICE\_MAX\_READ\_WRITE\_IMAGE\_ARGS. Each image array argument to a kernel declared with the `read_write` qualifier counts as one image.

## 5.11 Event Objects

Event objects can be used to refer to a kernel-instance command (**clEnqueueNDRangeKernel**, **clEnqueueNativeKernel**), read, write, map and copy commands on memory objects (**clEnqueue{Read|Write|Map}Buffer**, **clEnqueueUnmapMemObject**, **clEnqueue{Read|Write}BufferRect**, **clEnqueue{Read|Write|Map}Image**, **clEnqueueCopy{Buffer|Image}**, **clEnqueueCopyBufferRect**, **clEnqueueCopyBufferToImage**, **clEnqueueCopyImageToBuffer**), **clEnqueueSVMMemcpy**, **clEnqueueSVMMemFill**, **clEnqueueSVMMap**, **clEnqueueSVMUnmap**, **clEnqueueSVMFree**, **clEnqueueMarkerWithWaitList**, **clEnqueueBarrierWithWaitList** (refer to *section 5.12*) or user events.

An event object can be used to track the execution status of a command. The API calls that enqueue commands to a command-queue create a new event object that is returned in the *event* argument. In case of an error enqueueing the command in the command-queue the event argument does not return an event object.

The execution status of an enqueued command at any given point in time can be one of the following:

-  **CL\_QUEUED** – This indicates that the command has been enqueued in a command-queue. This is the initial state of all events except user events.
-  **CL\_SUBMITTED** – This is the initial state for all user events. For all other events, this indicates that the command has been submitted by the host to the device.
-  **CL\_RUNNING** – This indicates that the device has started executing this command. In order for the execution status of an enqueued command to change from **CL\_SUBMITTED** to **CL\_RUNNING**, all events that this command is waiting on must have completed successfully i.e. their execution status must be **CL\_COMPLETE**.
-  **CL\_COMPLETE** – This indicates that the command has successfully completed.
-  **Error code** – The error code is a negative integer value and indicates that the command was abnormally terminated. Abnormal termination may occur for a number of reasons such as a bad memory access.

**NOTE:** A command is considered to be complete if its execution status is **CL\_COMPLETE** or is a negative integer value.

If the execution of a command is terminated, the command-queue associated with this terminated command, and the associated context (and all other command-queues in this context) may no longer be available. The behavior of OpenCL API calls that use this context (and command-queues associated with this context) are now considered to be implementation-

defined. The user registered callback function specified when context is created can be used to report appropriate error information.

The function




`cl_event`            **clCreateUserEvent** (`cl_context context`, `cl_int *errcode_ret`)

creates a user event object. User events allow applications to enqueue commands that wait on a user event to finish before the command is executed by the device.

*context* must be a valid OpenCL context.

*errcode\_ret* will return an appropriate error code. If *errcode\_ret* is NULL, no error code is returned.

**clCreateUserEvent** returns a valid non-zero event object and *errcode\_ret* is set to CL\_SUCCESS if the user event object is created successfully. Otherwise, it returns a NULL value with one of the following error values returned in *errcode\_ret*:

-  CL\_INVALID\_CONTEXT if *context* is not a valid context.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The execution status of the user event object created is set to CL\_SUBMITTED.

The function

`cl_int`            **clSetUserEventStatus** (`cl_event event`, `cl_int execution_status`)

sets the execution status of a user event object.

*event* is a user event object created using **clCreateUserEvent**.

*execution\_status* specifies the new execution status to be set and can be CL\_COMPLETE or a negative integer value to indicate an error. A negative integer value causes all enqueued commands that wait on this user event to be terminated. **clSetUserEventStatus** can only be called once to change the execution status of *event*.

**clSetUserEventStatus** returns CL\_SUCCESS if the function was executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_EVENT if *event* is not a valid user event object.

- ✚ CL\_INVALID\_VALUE if the *execution\_status* is not CL\_COMPLETE or a negative integer value.
- ✚ CL\_INVALID\_OPERATION if the *execution\_status* for *event* has already been changed by a previous call to **clSetUserEventStatus**.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

NOTE: If there are enqueued commands with user events in the *event\_wait\_list* argument of **clEnqueue\*\*\*** commands, the user must ensure that the status of these user events being waited on are set using **clSetUserEventStatus** before any OpenCL APIs that release OpenCL objects except for event objects are called; otherwise the behavior is undefined.

For example, the following code sequence will result in undefined behavior of **clReleaseMemObject**.

```
ev1 = clCreateUserEvent(ctx, NULL);
clEnqueueWriteBuffer(cq, buf1, CL_FALSE, ...,
                    1, &ev1, NULL);
clEnqueueWriteBuffer(cq, buf2, CL_FALSE, ...);
clReleaseMemObject(buf2);
clSetUserEventStatus(ev1, CL_COMPLETE);
```

The following code sequence, however, works correctly.

```
ev1 = clCreateUserEvent(ctx, NULL);
clEnqueueWriteBuffer(cq, buf1, CL_FALSE, ...,
                    1, &ev1, NULL);
clEnqueueWriteBuffer(cq, buf2, CL_FALSE, ...);
clSetUserEventStatus(ev1, CL_COMPLETE);
clReleaseMemObject(buf2);
```

The function

```
cl_int      clWaitForEvents (cl_uint num_events, const cl_event *event_list)
```

waits on the host thread for commands identified by event objects in *event\_list* to complete. A command is considered complete if its execution status is CL\_COMPLETE or a negative value. The events specified in *event\_list* act as synchronization points.

**clWaitForEvents** returns CL\_SUCCESS if the execution status of all events in *event\_list* is CL\_COMPLETE. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_VALUE if *num\_events* is zero or *event\_list* is NULL.
- ✚ CL\_INVALID\_CONTEXT if events specified in *event\_list* do not belong to the same context.
- ✚ CL\_INVALID\_EVENT if event objects specified in *event\_list* are not valid event objects.
- ✚ CL\_EXEC\_STATUS\_ERROR\_FOR\_EVENTS\_IN\_WAIT\_LIST if the execution status of any of the events in *event\_list* is a negative integer value.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int          clGetEventInfo (cl_event event,
                                cl_event_info param_name,
                                size_t param_value_size,
                                void *param_value,
                                size_t *param_value_size_ret)
```

returns information about the event object.

*event* specifies the event object being queried.

*param\_name* specifies the information to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetEventInfo** is described in *table 5.23*.

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.23*.

*param\_value\_size\_ret* returns the actual size in bytes of data copied to *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.

cl_event_info	Return	Info. returned in <i>param_value</i>
---------------	--------	--------------------------------------



	Type	
<b>CL_EVENT_COMMAND_QUEUE</b>	cl_command_queue	Return the command-queue associated with <i>event</i> . For user event objects, a NULL value is returned.
<b>CL_EVENT_CONTEXT</b>	cl_context	Return the context associated with <i>event</i> .
<b>CL_EVENT_COMMAND_TYPE</b>	cl_command_type	<p>Return the command associated with <i>event</i>. Can be one of the following values:</p> <p>CL_COMMAND_NDRANGE_KERNEL  CL_COMMAND_NATIVE_KERNEL  CL_COMMAND_READ_BUFFER  CL_COMMAND_WRITE_BUFFER  CL_COMMAND_COPY_BUFFER  CL_COMMAND_READ_IMAGE  CL_COMMAND_WRITE_IMAGE  CL_COMMAND_COPY_IMAGE  CL_COMMAND_COPY_BUFFER_TO_IMAGE  CL_COMMAND_COPY_IMAGE_TO_BUFFER  CL_COMMAND_MAP_BUFFER  CL_COMMAND_MAP_IMAGE  CL_COMMAND_UNMAP_MEM_OBJECT  CL_COMMAND_MARKER  CL_COMMAND_ACQUIRE_GL_OBJECTS  CL_COMMAND_RELEASE_GL_OBJECTS</p> <p>CL_COMMAND_READ_BUFFER_RECT  CL_COMMAND_WRITE_BUFFER_RECT  CL_COMMAND_COPY_BUFFER_RECT  CL_COMMAND_USER  CL_COMMAND_BARRIER  CL_COMMAND_MIGRATE_MEM_OBJECTS  CL_COMMAND_FILL_BUFFER  CL_COMMAND_FILL_IMAGE</p> <p>CL_COMMAND_SVM_FREE  CL_COMMAND_SVM_MEMCPY  CL_COMMAND_SVM_MEMFILL  CL_COMMAND_SVM_MAP  CL_COMMAND_SVM_UNMAP</p>
<b>CL_EVENT_COMMAND_EXECUTION_STATUS<sup>29</sup></b>	cl_int	<p>Return the execution status of the command identified by <i>event</i>.</p> <p>Valid values are:</p> <p>CL_QUEUED (command has been enqueued in the command-queue),</p>

<sup>29</sup> The error code values are negative, and event state values are positive. The event state values are ordered from the largest value (CL\_QUEUED) for the first or initial state to the smallest value (CL\_COMPLETE or negative integer value) for the last or complete state. The value of CL\_COMPLETE and CL\_SUCCESS are the same.

		<p>CL_SUBMITTED (enqueued command has been submitted by the host to the device associated with the command-queue),</p> <p>CL_RUNNING (device is currently executing this command),</p> <p>CL_COMPLETE (the command has completed), or</p> <p>Error code given by a negative integer value. (command was abnormally terminated – this may be caused by a bad memory access etc.). These error codes come from the same set of error codes that are returned from the platform or runtime API calls as return values or <i>errcode_ret</i> values.</p>
<b>CL_EVENT_REFERENCE_COUNT</b> <sup>30</sup>	cl_uint	Return the <i>event</i> reference count.

**Table 5.23** *clGetEventInfo parameter queries.*

Using **clGetEventInfo** to determine if a command identified by *event* has finished execution (i.e. CL\_EVENT\_COMMAND\_EXECUTION\_STATUS returns CL\_COMPLETE) is not a synchronization point. There are no guarantees that the memory objects being modified by command associated with *event* will be visible to other enqueued commands.

**clGetEventInfo** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_VALUE if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is < size of return type as described in *table 5.23* and *param\_value* is not NULL.
- ✚ CL\_INVALID\_VALUE if information to query given in *param\_name* cannot be queried for *event*.
- ✚ CL\_INVALID\_EVENT if *event* is not a valid event object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.

<sup>30</sup> The reference count returned should be considered immediately stale. It is unsuitable for general use in applications. This feature is provided for identifying memory leaks.

- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int      clSetEventCallback (cl_event event,  
                                cl_int command_exec_callback_type,  
                                void (CL_CALLBACK *pfn_event_notify)(cl_event event,  
                                cl_int event_command_exec_status,  
                                void *user_data),  
                                void *user_data)
```

registers a user callback function for a specific command execution status. The registered callback function will be called when the execution status of command associated with *event* changes to an execution status equal to or past the status specified by *command\_exec\_status*.

Each call to **clSetEventCallback** registers the specified user callback function on a callback stack associated with *event*. The order in which the registered user callback functions are called is undefined.

*event* is a valid event object.

*command\_exec\_callback\_type* specifies the command execution status for which the callback is registered. The command execution callback values for which a callback can be registered are: CL\_SUBMITTED, CL\_RUNNING or CL\_COMPLETE<sup>31</sup>. There is no guarantee that the callback functions registered for various execution status values for an event will be called in the exact order that the execution status of a command changes. Furthermore, it should be noted that receiving a call back for an event with a status other than CL\_COMPLETE, in no way implies that the memory model or execution model as defined by the OpenCL specification has changed. For example, it is not valid to assume that a corresponding memory transfer has completed unless the event is in a state CL\_COMPLETE.

*pfn\_event\_notify* is the event callback function that can be registered by the application. This callback function may be called asynchronously by the OpenCL implementation. It is the application's responsibility to ensure that the callback function is thread-safe. The parameters to this callback function are:

- ✚ *event* is the event object for which the callback function is invoked.
- ✚ *event\_command\_exec\_status* represents the execution status of command for which this callback function is invoked. Refer to *table 5.23* for the command execution status values. If the callback is called as the result of the command associated with

---

<sup>31</sup>The callback function registered for a *command\_exec\_callback\_type* value of CL\_COMPLETE will be called when the command has completed successfully or is abnormally terminated.

event being abnormally terminated, an appropriate error code for the error that caused the termination will be passed to *event\_command\_exec\_status* instead.

 *user\_data* is a pointer to user supplied data.

*user\_data* will be passed as the *user\_data* argument when *pfn\_notify* is called. *user\_data* can be NULL.

All callbacks registered for an event object must be called. All enqueued callbacks shall be called before the event object is destroyed. Callbacks must return promptly. The behavior of calling expensive system routines, OpenCL API calls to create contexts or command-queues, or blocking OpenCL operations from the following list below, in a callback is undefined.

**clFinish,**

**clWaitForEvents,**

blocking calls to **clEnqueueReadBuffer**, **clEnqueueReadBufferRect**,

**clEnqueueWriteBuffer**, **clEnqueueWriteBufferRect**,

blocking calls to **clEnqueueReadImage** and **clEnqueueWriteImage**,


blocking calls to **clEnqueueMapBuffer** and **clEnqueueMapImage**,


blocking calls to **clBuildProgram**, **clCompileProgram** or **clLinkProgram**,


blocking calls to **clEnqueueSVMMemcpy** or **clEnqueueSVMMap**


If an application needs to wait for completion of a routine from the above list in a callback, please use the non-blocking form of the function, and assign a completion callback to it to do the remainder of your work. Note that when a callback (or other code) enqueues commands to a command-queue, the commands are not required to begin execution until the queue is flushed. In standard usage, blocking enqueue calls serve this role by implicitly flushing the queue. Since blocking calls are not permitted in callbacks, those callbacks that enqueue commands on a command queue should either call **clFlush** on the queue before returning or arrange for **clFlush** to be called later on another thread.

**clSetEventCallback** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

 CL\_INVALID\_EVENT if *event* is not a valid event object.

 CL\_INVALID\_VALUE if *pfn\_event\_notify* is NULL or if *command\_exec\_callback\_type* is not CL\_SUBMITTED, CL\_RUNNING or CL\_COMPLETE.

 CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.

 CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

cl\_int            **clRetainEvent** (cl\_event *event*)

increments the *event* reference count. The OpenCL commands that return an event perform an implicit retain.

**clRetainEvent** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_EVENT if *event* is not a valid event object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

To release an event, use the following function

cl\_int            **clReleaseEvent** (cl\_event *event*)

decrements the *event* reference count.

**clReleaseEvent** returns CL\_SUCCESS if the function is executed successfully. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_EVENT if *event* is not a valid event object.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The event object is deleted once the reference count becomes zero, the specific command identified by this event has completed (or terminated) and there are no commands in the command-queues of a context that require a wait for this event to complete. Using this function to release a reference that was not obtained by creating the object or by calling **clRetainEvent** causes undefined behavior.

NOTE: Developers should be careful when releasing their last reference count on events created by **clCreateUserEvent** that have not yet been set to status of CL\_COMPLETE or an error. If the user event was used in the event\_wait\_list argument passed to a clEnqueue\*\*\* API or another application host thread is waiting for it in **clWaitForEvents**, those commands and host threads will continue to wait for the event status to reach CL\_COMPLETE or error, even after the user has released the object. Since in this scenario the developer has released his last reference count to

the user event, it would be in principle no longer valid for him to change the status of the event to unblock all the other machinery. As a result the waiting tasks will wait forever, and associated events, cl\_mem objects, command queues and contexts are likely to leak. In-order command queues caught up in this deadlock may cease to do any work.

## 5.12 Markers, Barriers and Waiting for Events

The function

```
cl_int          clEnqueueMarkerWithWaitList (cl_command_queue command_queue,
                                             cl_uint num_events_in_wait_list,
                                             const cl_event *event_wait_list,
                                             cl_event *event)
```

enqueues a marker command which waits for either a list of events to complete, or if the list is empty it waits for all commands previously enqueued in *command\_queue* to complete before it completes. This command returns an *event* which can be waited on, i.e. this event can be waited on to insure that all events either in the *event\_wait\_list* or all previously enqueued commands, queued before this command to *command\_queue*, have completed.

*command\_queue* is a valid host command-queue.



*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed.

If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

If *event\_wait\_list* is NULL, then this particular command waits until all previous enqueued commands to *command\_queue* have completed.

*event* returns an event object that identifies this particular command. Event objects are unique and can be used to identify this marker command later on. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.

**clEnqueueMarkerWithWaitList** returns CL\_SUCCESS if the function is successfully executed. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_INVALID\_CONTEXT if context associated with *command\_queue* and events in *event\_wait\_list* are not the same.

- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

The function

```
cl_int          clEnqueueBarrierWithWaitList (cl_command_queue command_queue,
                                                cl_uint num_events_in_wait_list,
                                                const cl_event *event_wait_list,
                                                cl_event *event)
```

enqueues a barrier command which waits for either a list of events to complete, or if the list is empty it waits for all commands previously enqueued in *command\_queue* to complete before it completes. This command blocks command execution, that is, any following commands enqueued after it do not execute until it completes. This command returns an *event* which can be waited on, i.e. this event can be waited on to insure that all events either in the *event\_wait\_list* or all previously enqueued commands, queued before this command to *command\_queue*, have completed

*command\_queue* is a valid host command-queue.

*event\_wait\_list* and *num\_events\_in\_wait\_list* specify events that need to complete before this particular command can be executed.

If *event\_wait\_list* is NULL, *num\_events\_in\_wait\_list* must be 0. If *event\_wait\_list* is not NULL, the list of events pointed to by *event\_wait\_list* must be valid and *num\_events\_in\_wait\_list* must be greater than 0. The events specified in *event\_wait\_list* act as synchronization points. The context associated with events in *event\_wait\_list* and *command\_queue* must be the same. The memory associated with *event\_wait\_list* can be reused or freed after the function returns.

If *event\_wait\_list* is NULL, then this particular command waits until all previous enqueued commands to *command\_queue* have completed.

*event* returns an event object that identifies this particular command. Event objects are unique and can be used to identify this barrier command later on. *event* can be NULL in which case it will not be possible for the application to query the status of this command or queue a wait for this command to complete. If the *event\_wait\_list* and the *event* arguments are not NULL, the *event* argument should not refer to an element of the *event\_wait\_list* array.



**clEnqueueBarrierWithWaitList** returns CL\_SUCCESS if the function is successfully executed. Otherwise, it returns one of the following errors:

- ✚ CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
- ✚ CL\_INVALID\_CONTEXT if context associated with *command\_queue* and events in *event\_wait\_list* are not the same.
- ✚ CL\_INVALID\_EVENT\_WAIT\_LIST if *event\_wait\_list* is NULL and *num\_events\_in\_wait\_list* > 0, or *event\_wait\_list* is not NULL and *num\_events\_in\_wait\_list* is 0, or if event objects in *event\_wait\_list* are not valid events.
- ✚ CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
- ✚ CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## 5.13 Out-of-order Execution of Kernels and Memory Object Commands

The OpenCL functions that are submitted to a command-queue are enqueued in the order the calls are made but can be configured to execute in-order or out-of-order. The *properties* argument in **clCreateCommandQueueWithProperties** can be used to specify the execution order.

If the `CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE` property of a command-queue is not set, the commands enqueued to a command-queue execute in order. For example, if an application calls **clEnqueueNDRangeKernel** to execute kernel A followed by a **clEnqueueNDRangeKernel** to execute kernel B, the application can assume that kernel A finishes first and then kernel B is executed. If the memory objects output by kernel A are inputs to kernel B then kernel B will see the correct data in memory objects produced by execution of kernel A. If the `CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE` property of a command-queue is set, then there is no guarantee that kernel A will finish before kernel B starts execution.

Applications can configure the commands enqueued to a command-queue to execute out-of-order by setting the `CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE` property of the command-queue. This can be specified when the command-queue is created. In out-of-order execution mode there is no guarantee that the enqueued commands will finish execution in the order they were queued. As there is no guarantee that kernels will be executed in order, i.e. based on when the **clEnqueueNDRangeKernel** calls are made within a command-queue, it is therefore possible that an earlier **clEnqueueNDRangeKernel** call to execute kernel A identified by event A may execute and/or finish later than a **clEnqueueNDRangeKernel** call to execute kernel B which was called by the application at a later point in time. To guarantee a specific order of execution of kernels, a wait on a particular event (in this case event A) can be used. The wait for event A can be specified in the *event\_wait\_list* argument to **clEnqueueNDRangeKernel** for kernel B.

In addition, a marker (**clEnqueueMarkerWithWaitList**) or a barrier (**clEnqueueBarrierWithWaitList**) command can be enqueued to the command-queue. The marker command ensures that previously enqueued commands identified by the list of events to wait for (or all previous commands) have finished. A barrier command is similar to a marker command, but additionally guarantees that no later-enqueued commands will execute until the waited-for commands have executed.

Similarly, commands to read, write, copy or map memory objects that are enqueued after **clEnqueueNDRangeKernel** or **clEnqueueNativeKernel** commands are not guaranteed to wait for kernels scheduled for execution to have completed (if the `CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE` property is set). To ensure correct ordering of commands, the event object returned by **clEnqueueNDRangeKernel** or **clEnqueueNativeKernel** can be used to enqueue a wait for event or a barrier command can be enqueued that must complete before reads or writes to the memory object(s) occur.

## 5.14 Profiling Operations on Memory Objects and Kernels

This section describes profiling of OpenCL functions that are enqueued as commands to a command-queue. The specific functions<sup>32</sup> being referred to are:

**clEnqueue{Read|Write|Map}Buffer**, **clEnqueue{Read|Write}BufferRect**,  
**clEnqueue{Read|Write|Map}Image**, **clEnqueueUnmapMemObject**,  
**clEnqueueSVMMemcpy**, **clEnqueueSVMMemFill**, **clEnqueueSVMMap**,  
**clEnqueueSVMUnmap**, **clEnqueueSVMFree**, **clEnqueueCopyBuffer**,  
**clEnqueueCopyBufferRect**, **clEnqueueCopyImage**, **clEnqueueCopyImageToBuffer**,  
**clEnqueueCopyBufferToImage**, **clEnqueueNDRangeKernel** and **clEnqueueNativeKernel**.  
These enqueued commands are identified by unique event objects.

Event objects can be used to capture profiling information that measure execution time of a command. Profiling of OpenCL commands can be enabled either by using a command-queue created with `CL_QUEUE_PROFILING_ENABLE` flag set in *properties* argument to **clCreateCommandQueueWithProperties**.

If profiling is enabled, the function

```
cl_int          clGetEventProfilingInfo (cl_event event,  
                                         cl_profiling_info param_name,  
                                         size_t param_value_size,  
                                         void *param_value,  
                                         size_t *param_value_size_ret)
```

returns profiling information for the command associated with event.

*event* specifies the event object.

*param\_name* specifies the profiling data to query. The list of supported *param\_name* types and the information returned in *param\_value* by **clGetEventProfilingInfo** is described in *table 5.24*

*param\_value* is a pointer to memory where the appropriate result being queried is returned. If *param\_value* is NULL, it is ignored.

*param\_value\_size* is used to specify the size in bytes of memory pointed to by *param\_value*. This size must be  $\geq$  size of return type as described in *table 5.24*.

---

<sup>32</sup> **clEnqueueAcquireGLObjects** and **clEnqueueReleaseGLObjects** defined in *section 9.6.6* of the OpenCL 2.0 Extension Specification are also included.

*param\_value\_size\_ret* returns the actual size in bytes of data copied to *param\_value*. If *param\_value\_size\_ret* is NULL, it is ignored.






<b>cl_profiling_info</b>	<b>Return Type</b>	<b>Info. returned in <i>param_value</i></b>
<b>CL_PROFILING_COMMAND_QUEUED</b>	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> is enqueued in a command-queue by the host.
<b>CL_PROFILING_COMMAND_SUBMIT</b>	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> that has been enqueued is submitted by the host to the device associated with the command-queue.
<b>CL_PROFILING_COMMAND_START</b>	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> starts execution on the device.
<b>CL_PROFILING_COMMAND_END</b>	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> has finished execution on the device.
<b>CL_PROFILING_COMMAND_COMPLETE</b>	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> and any child commands enqueued by this command on the device have finished execution.

**Table 5.24** *clGetEventProfilingInfo* parameter queries.

The unsigned 64-bit values returned can be used to measure the time in nano-seconds consumed by OpenCL commands.

OpenCL devices are required to correctly track time across changes in device frequency and power states. The `CL_DEVICE_PROFILING_TIMER_RESOLUTION` specifies the resolution of the timer i.e. the number of nanoseconds elapsed before the timer is incremented.

**clGetEventProfilingInfo** returns `CL_SUCCESS` if the function is executed successfully and the profiling information has been recorded. Otherwise, it returns one of the following errors:

-  `CL_PROFILING_INFO_NOT_AVAILABLE` if the `CL_QUEUE_PROFILING_ENABLE` flag is not set for the command-queue, if the execution status of the command identified by *event* is not `CL_COMPLETE` or if *event* refers to the **clEnqueueSVMFree** command or is a user event object.
-  `CL_INVALID_VALUE` if *param\_name* is not valid, or if size in bytes specified by *param\_value\_size* is < size of return type as described in *table 5.24* and *param\_value* is not NULL.
-  `CL_INVALID_EVENT` if *event* is not a valid event object.
-  `CL_OUT_OF_RESOURCES` if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  `CL_OUT_OF_HOST_MEMORY` if there is a failure to allocate resources required by the OpenCL implementation on the host.




## 5.15 Flush and Finish

The function

cl\_int            **clFlush** (cl\_command\_queue *command\_queue*)

issues all previously queued OpenCL commands in *command\_queue* to the device associated with *command\_queue*. **clFlush** only guarantees that all queued commands to *command\_queue* will eventually be submitted to the appropriate device. There is no guarantee that they will be complete after **clFlush** returns.

**clFlush** returns CL\_SUCCESS if the function call was executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

Any blocking commands queued in a command-queue and **clReleaseCommandQueue** perform an implicit flush of the command-queue. These blocking commands are **clEnqueueReadBuffer**, **clEnqueueReadBufferRect**, **clEnqueueReadImage**, with *blocking\_read* set to CL\_TRUE; **clEnqueueWriteBuffer**, **clEnqueueWriteBufferRect**, **clEnqueueWriteImage** with *blocking\_write* set to CL\_TRUE; **clEnqueueMapBuffer**, **clEnqueueMapImage** with *blocking\_map* set to CL\_TRUE; **clEnqueueSVMMemcpy** with *blocking\_copy* set to CL\_TRUE; **clEnqueueSVMMap** with *blocking\_map* set to CL\_TRUE or **clWaitForEvents**.




To use event objects that refer to commands enqueued in a command-queue as event objects to wait on by commands enqueued in a different command-queue, the application must call a **clFlush** or any blocking commands that perform an implicit flush of the command-queue where the commands that refer to these event objects are enqueued.

The function

cl\_int            **clFinish** (cl\_command\_queue *command\_queue*)

blocks until all previously queued OpenCL commands in *command\_queue* are issued to the associated device and have completed. **clFinish** does not return until all previously queued commands in *command\_queue* have been processed and completed. **clFinish** is also a synchronization point.

**clFinish** returns CL\_SUCCESS if the function call was executed successfully. Otherwise, it returns one of the following errors:

-  CL\_INVALID\_COMMAND\_QUEUE if *command\_queue* is not a valid host command-queue.
-  CL\_OUT\_OF\_RESOURCES if there is a failure to allocate resources required by the OpenCL implementation on the device.
-  CL\_OUT\_OF\_HOST\_MEMORY if there is a failure to allocate resources required by the OpenCL implementation on the host.

## **6. The OpenCL C Programming Language**

## **7. OpenCL Numerical Compliance**

## **8. Image Addressing and Filtering**

Sections 6 – 8 are described in the OpenCL C 2.0 Specification document.

## **9. Optional Extensions**

The list of optional features supported by OpenCL 2.0 is described in the OpenCL 2.0 Extension Specification document.



## 10. OpenCL Embedded Profile

The OpenCL 2.0 specification describes the feature requirements for desktop platforms. This section describes the OpenCL 2.0 embedded profile that allows us to target a subset of the OpenCL 2.0 specification for handheld and embedded platforms. The optional extensions defined in the OpenCL 2.0 Extension Specification apply to both profiles.

The OpenCL 2.0 embedded profile has the following restrictions:

1. 64 bit integers i.e. long, along including the appropriate vector data types and operations on 64-bit integers are optional. The **cles\_khr\_int64**<sup>33</sup> extension string will be reported if the embedded profile implementation supports 64-bit integers.

- ~~2. Support for 3D images is optional.~~

~~If CL\_DEVICE\_IMAGE3D\_MAX\_WIDTH, CL\_DEVICE\_IMAGE3D\_MAX\_HEIGHT and CL\_DEVICE\_IMAGE3D\_MAX\_DEPTH are zero, the call to **clCreateImage** in the embedded profile will fail to create the 3D image. The *errcode\_ret* argument in **clCreateImage** returns CL\_INVALID\_OPERATION. Declaring arguments of type `image3d_t` in a kernel will result in a compilation error.~~

~~If CL\_DEVICE\_IMAGE3D\_MAX\_WIDTH, CL\_DEVICE\_IMAGE3D\_HEIGHT and CL\_DEVICE\_IMAGE3D\_MAX\_DEPTH > 0, 3D images are supported by the OpenCL embedded profile implementation. **clCreateImage** will work as defined by the OpenCL specification. The `image3d_t` data type can be used in a kernel(s).~~

- ~~3. Support for 2D image array writes is optional. If the **cles\_khr\_2d\_image\_array\_writes** extension is supported by the embedded profile, writes to 2D image arrays are supported.~~
- ~~4. Image and image arrays created with an `image_channel_data_type` value of CL\_FLOAT or CL\_HALF\_FLOAT can only be used with samplers that use a filter mode of CL\_FILTER\_NEAREST. The values returned by **read\_imagef** and **read\_imageh**<sup>34</sup> for 2D and 3D images if `image_channel_data_type` value is CL\_FLOAT or CL\_HALF\_FLOAT and sampler with `filter_mode = CL_FILTER_LINEAR` are undefined.~~
5. The mandated minimum single precision floating-point capability given by CL\_DEVICE\_SINGLE\_FP\_CONFIG is CL\_FP\_ROUND\_TO\_ZERO or CL\_FP\_ROUND\_TO\_NEAREST. If CL\_FP\_ROUND\_TO\_NEAREST is supported, the default rounding mode will be round to nearest even; otherwise the default rounding

---

<sup>33</sup> Note that the performance of 64-bit integer arithmetic can vary significantly between embedded devices.

<sup>34</sup> If **cl\_khr\_fp16** extension is supported.

mode will be round to zero.

6. The single precision floating-point operations (addition, subtraction and multiplication) shall be correctly rounded. Zero results may always be positive 0.0. The accuracy of division and sqrt are given by *table 10.1*.

If CL\_FP\_INF\_NAN is not set in CL\_DEVICE\_SINGLE\_FP\_CONFIG, and one of the operands or the result of addition, subtraction, multiplication or division would signal the overflow or invalid exception (see IEEE 754 specification), the value of the result is implementation-defined. Likewise, single precision comparison operators (<, >, <=, >=, ==, !=) return implementation-defined values when one or more operands is a NaN.

In all cases, conversions (*section 6.2* and *6.13.7*) shall be correctly rounded as described for the FULL\_PROFILE, including those that consume or produce an INF or NaN. The built-in math functions (*section 6.13.2*) shall behave as described for the FULL\_PROFILE, including edge case behavior described in *section 7.5.1* but with accuracy as described by *table 10.1*.

**Note:** If addition, subtraction and multiplication have default round to zero rounding mode, then **fract**, **fma** and **fdim** shall produce the correctly rounded result for round to zero rounding mode.

This relaxation of the requirement to adhere to IEEE 754 requirements for basic floating-point operations, though extremely undesirable, is to provide flexibility for embedded devices that have lot stricter requirements on hardware area budgets.

7. Denormalized numbers for the half data type which may be generated when converting a float to a half using variants of the **vstore\_half** function or when converting from a half to a float using variants of the **vload\_half** function can be flushed to zero. Refer to *section 6.1.1.1*.
8. The precision of conversions from CL\_UNORM\_INT8, CL\_SNORM\_INT8, CL\_UNORM\_INT16 and CL\_SNORM\_INT16 to float is <= 2 ulp for the embedded profile instead of <= 1.5 ulp as defined in *section 8.3.1.1*. The exception cases described in *section 8.3.1.1* and given below apply to the embedded profile.

For CL\_UNORM\_INT8

```
0 must convert to 0.0f and
255 must convert to 1.0f
```

For CL\_UNORM\_INT16

```
0 must convert to 0.0f and
65535 must convert to 1.0f
```

For CL\_SNORM\_INT8

-128 and -127 must convert to -1.0f,  
 0 must convert to 0.0f and  
 127 must convert to 1.0f

For CL\_SNORM\_INT16

-32768 and -32767 must convert to -1.0f,  
 0 must convert to 0.0f and  
 32767 must convert to 1.0f

For CL\_UNORM\_INT\_101010

0 must convert to 0.0f and  
 1023 must convert to 1.0f

~~9. Built-in atomic functions as defined in section 6.12.11 are optional.~~

The following optional extensions defined in the OpenCL 2.0 Extension Specification are available to the embedded profile:

- ✚ **cl\_khr\_int64\_base\_atomics**
- ✚ **cl\_khr\_int64\_extended\_atomics**
- ✚ **cl\_khr\_fp16**
- ✚ **cles\_khr\_int64.** If double precision is supported i.e. CL\_DEVICE\_DOUBLE\_FP\_CONFIG is not zero, then cles\_khr\_int64 must also be supported.

Table 10.1 describes the minimum accuracy of single precision floating-point arithmetic operations given as ULP values for the embedded profile. The reference value used to compute the ULP value of an arithmetic operation is the infinitely precise result.

Function	Min Accuracy - ULP values <sup>35</sup>
$x + y$	Correctly rounded
$x - y$	Correctly rounded
$x * y$	Correctly rounded
$1.0 / x$	$\leq 3$ ulp
$x / y$	$\leq 3$ ulp
<b>acos</b>	$\leq 4$ ulp
<b>acospi</b>	$\leq 5$ ulp
<b>asin</b>	$\leq 4$ ulp

<sup>35</sup> 0 ulp is used for math functions that do not require rounding.

<b>asinpi</b>	<= 5 ulp
<b>atan</b>	<= 5 ulp
<b>atan2</b>	<= 6 ulp
<b>atanpi</b>	<= 5 ulp
<b>atan2pi</b>	<= 6 ulp
<b>acosh</b>	<= 4 ulp
<b>asinh</b>	<= 4 ulp
<b>atanh</b>	<= 5 ulp
<b>cbrt</b>	<= 4 ulp
<b>ceil</b>	Correctly rounded
<b>copysign</b>	0 ulp
<b>cos</b>	<= 4 ulp
<b>cosh</b>	<= 4 ulp
<b>cospi</b>	<= 4 ulp
<b>erfc</b>	<= 16 ulp
<b>erf</b>	<= 16 ulp
<b>exp</b>	<= 4 ulp
<b>exp2</b>	<= 4 ulp
<b>exp10</b>	<= 4 ulp
<b>expm1</b>	<= 4 ulp
<b>fabs</b>	0 ulp
<b>fdim</b>	Correctly rounded
<b>floor</b>	Correctly rounded
<b>fma</b>	Correctly rounded
<b>fmax</b>	0 ulp
<b>fmin</b>	0 ulp
<b>fmod</b>	0 ulp
<b>fract</b>	Correctly rounded
<b>frexp</b>	0 ulp
<b>hypot</b>	<= 4 ulp
<b>ilogb</b>	0 ulp
<b>ldexp</b>	Correctly rounded
<b>log</b>	<= 4 ulp
<b>log2</b>	<= 4 ulp
<b>log10</b>	<= 4 ulp
<b>log1p</b>	<= 4 ulp
<b>logb</b>	0 ulp
<b>mad</b>	Any value allowed (infinite ulp)
<b>maxmag</b>	0 ulp
<b>minmag</b>	0 ulp
<b>modf</b>	0 ulp
<b>nan</b>	0 ulp
<b>nextafter</b>	0 ulp
<b>pow(x, y)</b>	<= 16 ulp
<b>pown(x, y)</b>	<= 16 ulp

<b>powr(x, y)</b>	<= 16 ulp
<b>remainder</b>	0 ulp
<b>remquo</b>	0 ulp
<b>rint</b>	Correctly rounded
<b>rootn</b>	<= 16 ulp
<b>round</b>	Correctly rounded
<b>rsqrt</b>	<= 4 ulp
<b>sin</b>	<= 4 ulp
<b>sincos</b>	<= 4 ulp for sine and cosine values
<b>sinh</b>	<= 4 ulp
<b>sinpi</b>	<= 4 ulp
<b>sqrt</b>	<= 4 ulp
<b>tan</b>	<= 5 ulp
<b>tanh</b>	<= 5 ulp
<b>tanpi</b>	<= 6 ulp
<b>tgamma</b>	<= 16 ulp
<b>trunc</b>	Correctly rounded
<b>half_cos</b>	<= 8192 ulp
<b>half_divide</b>	<= 8192 ulp
<b>half_exp</b>	<= 8192 ulp
<b>half_exp2</b>	<= 8192 ulp
<b>half_exp10</b>	<= 8192 ulp
<b>half_log</b>	<= 8192 ulp
<b>half_log2</b>	<= 8192 ulp
<b>half_log10</b>	<= 8192 ulp
<b>half_powr</b>	<= 8192 ulp
<b>half_recip</b>	<= 8192 ulp
<b>half_rsqrt</b>	<= 8192 ulp
<b>half_sin</b>	<= 8192 ulp
<b>half_sqrt</b>	<= 8192 ulp
<b>half_tan</b>	<= 8192 ulp
<b>native_cos</b>	Implementation-defined
<b>native_divide</b>	Implementation-defined
<b>native_exp</b>	Implementation-defined
<b>native_exp2</b>	Implementation-defined
<b>native_exp10</b>	Implementation-defined
<b>native_log</b>	Implementation-defined
<b>native_log2</b>	Implementation-defined
<b>native_log10</b>	Implementation-defined
<b>native_powr</b>	Implementation-defined
<b>native_recip</b>	Implementation-defined
<b>native_rsqrt</b>	Implementation-defined
<b>native_sin</b>	Implementation-defined

<b>native_sqrt</b>	Implementation-defined
<b>native_tan</b>	Implementation-defined

**Table 10.1** *ULP values for built-in math functions*

The `__EMBEDDED_PROFILE__` macro is added to the language (refer to *section 6.10*). It will be the integer constant 1 for OpenCL devices that implement the embedded profile and is undefined otherwise.

`CL_PLATFORM_PROFILE` defined in *table 4.1* will return the string `EMBEDDED_PROFILE` if the OpenCL implementation supports the embedded profile only.

The minimum maximum values specified in *table 4.3* that have been modified for the OpenCL embedded profile are listed below:

<b>cl_device_info</b>	<b>Return Type</b>	<b>Description</b>
<b>CL_DEVICE_MAX_READ_IMAGE_ARGS</b>	cl_uint	Max number of image objects arguments of a kernel declared with the <code>read_only</code> qualifier. The minimum value is 8 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_MAX_WRITE_IMAGE_ARGS</b>	cl_uint	Max number of image objects arguments of a kernel declared with the <code>write_only</code> qualifier. The minimum value is 8 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_MAX_READ_WRITE_IMAGE_ARGS</b>	cl_uint	Max number of image objects arguments of a kernel declared with the <code>write_only</code> or <code>read_write</code> qualifier. The minimum value is 8 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE2D_MAX_WIDTH</b>	size_t	Max width of 2D image in pixels. The minimum value is 2048 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE2D_MAX_HEIGHT</b>	size_t	Max height of 2D image in pixels. The minimum value is 2048 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .
<b>CL_DEVICE_IMAGE3D_MAX_WIDTH</b>	size_t	Max width of 3D image in pixels. The minimum value is 2048 if <code>CL_DEVICE_IMAGE_SUPPORT</code> is <code>CL_TRUE</code> .

<b>CL_DEVICE_IMAGE3D_MAX_HEIGHT</b>	size_t	Max height of 3D image in pixels. The minimum value is 2048.
<b>CL_DEVICE_IMAGE3D_MAX_DEPTH</b>	size_t	Max depth of 3D image in pixels. The minimum value is 2048.
<b>CL_DEVICE_IMAGE_MAX_BUFFER_SIZE</b>	size_t	Max number of pixels for a 1D image created from a buffer object.  The minimum value is 2048 if CL_DEVICE_IMAGE_SUPPORT is CL_TRUE.
<b>CL_DEVICE_IMAGE_MAX_ARRAY_SIZE</b>	size_t	Max number of images in a 1D or 2D image array.  The minimum value is 256 if CL_DEVICE_IMAGE_SUPPORT is CL_TRUE.
<b>CL_DEVICE_MAX_SAMPLERS</b>	unsigned int	Maximum number of samplers that can be used in a kernel.  The minimum value is 8 if CL_DEVICE_IMAGE_SUPPORT is CL_TRUE.
<b>CL_DEVICE_MAX_PARAMETER_SIZE</b>	size_t	Max size in bytes of all arguments that can be passed to a kernel. The minimum value is 256 bytes for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_SINGLE_FP_CONFIG</b>	cl_device_fp_config	Describes single precision floating-point capability of the device. This is a bit-field that describes one or more of the following values:  CL_FP_DENORM – denorms are supported  CL_FP_INF_NAN – INF and quiet NaNs are supported.  CL_FP_ROUND_TO_NEAREST – round to nearest even rounding mode supported  CL_FP_ROUND_TO_ZERO – round to zero rounding mode supported  CL_FP_ROUND_TO_INF – round to positive and negative infinity rounding modes supported

		<p>CL_FP_FMA – IEEE754-2008 fused multiply-add is supported.</p> <p>CL_FP_CORRECTLY_ROUNDED_DIVIDE_SQRT – divide and sqrt are correctly rounded as defined by the IEEE754 specification.</p> <p>CL_FP_SOFT_FLOAT – Basic floating-point operations (such as addition, subtraction, multiplication) are implemented in software.</p> <p>The mandated minimum floating-point capability is:  CL_FP_ROUND_TO_ZERO or  CL_FP_ROUND_TO_NEAREST  for devices that are not of type  CL_DEVICE_TYPE_CUSTOM.</p>
<b>CL_DEVICE_MAX_CONSTANT_BUFFER_SIZE</b>	unsigned long long	Max size in bytes of a constant buffer allocation. The minimum value is 1 KB for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_MAX_CONSTANT_ARGS</b>	unsigned int	Max number of arguments declared with the __constant qualifier in a kernel. The minimum value is 4 for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_LOCAL_MEM_SIZE</b>	cl_ulong	Size of local memory arena in bytes. The minimum value is 1 KB for devices that are not of type CL_DEVICE_TYPE_CUSTOM.
<b>CL_DEVICE_COMPILER_AVAILABLE</b>	cl_bool	<p>Is CL_FALSE if the implementation does not have a compiler available to compile the program source.  Is CL_TRUE if the compiler is available.</p> <p>This can be CL_FALSE for the embedded platform profile only.</p>
<b>CL_DEVICE_LINKER_AVAILABLE</b>	cl_bool	<p>Is CL_FALSE if the implementation does not have a linker available.  Is CL_TRUE if the linker is available.</p> <p>This can be CL_FALSE for the embedded platform profile only.</p>



		This must be CL_TRUE if <b>CL_DEVICE_COMPILER_AVAILABLE</b> is CL_TRUE.
<b>CL_DEVICE_QUEUE_ON_DEVICE_MAX_SIZE</b>	cl_uint	The max. size of the device queue in bytes. The minimum value is 64 KB for the embedded profile
<b>CL_DEVICE_PRINTF_BUFFER_SIZE</b>	size_t	Maximum size in bytes of the internal buffer that holds the output of printf calls from a kernel. The minimum value for the EMBEDDED profile is 1 KB.

If CL\_DEVICE\_IMAGE\_SUPPORT specified in *table 4.3* is CL\_TRUE, the values assigned to CL\_DEVICE\_MAX\_READ\_IMAGE\_ARGS, CL\_DEVICE\_MAX\_WRITE\_IMAGE\_ARGS, CL\_DEVICE\_IMAGE2D\_MAX\_WIDTH, CL\_DEVICE\_IMAGE2D\_MAX\_HEIGHT, CL\_DEVICE\_IMAGE3D\_MAX\_WIDTH, CL\_DEVICE\_IMAGE3D\_MAX\_HEIGHT, CL\_DEVICE\_IMAGE3D\_MAX\_DEPTH and CL\_DEVICE\_MAX\_SAMPLERS by the implementation must be greater than or equal to the minimum values specified in the embedded profile version of *table 4.3* given above.

# Appendix A

## A.1 Shared OpenCL Objects



This section describes which objects can be shared across multiple command-queues created within a host process.

OpenCL memory objects, program objects and kernel objects are created using a context and can be shared across multiple command-queues created using the same context. Event objects can be created when a command is queued to a command-queue. These event objects can be shared across multiple command-queues created using the same context.

The application needs to implement appropriate synchronization across threads on the host processor to ensure that the changes to the state of a shared object (such as a command-queue object, memory object, program or kernel object) happen in the correct order (deemed correct by the application) when multiple command-queues in multiple threads are making changes to the state of a shared object.

A command-queue can cache changes to the state of a memory object on the device associated with the command-queue. To synchronize changes to a memory object across command-queues, the application must do the following:

In the command-queue that includes commands that modify the state of a memory object, the application must do the following:

-  Get appropriate event objects for commands that modify the state of the shared memory object.
-  Call the **clFlush** (or **clFinish**) API to issue any outstanding commands from this command-queue.

In the command-queue that wants to synchronize to the latest state of a memory object, commands queued by the application must use the appropriate event objects that represent commands that modify the state of the shared memory object as event objects to wait on. This is to ensure that commands that use this shared memory object complete in the previous command-queue before the memory objects are used by commands executing in this command-queue.

The results of modifying a shared resource in one command-queue while it is being used by another command-queue are undefined.

## A.2 Multiple Host Threads

All OpenCL API calls are thread-safe<sup>36</sup> except those that modify the state of `cl_kernel` objects: **`clSetKernelArg`**, **`clSetKernelArgSVMPointer`** and **`clSetKernelExecInfo`**. **`clSetKernelArg`**, **`clSetKernelArgSVMPointer`** and **`clSetKernelExecInfo`** are safe to call from any host thread, and safe to call re-entrantly so long as concurrent calls to any combination of these API calls operate on different `cl_kernel` objects. The state of the `cl_kernel` object is undefined if **`clSetKernelArg`**, **`clSetKernelArgSVMPointer`** or **`clSetKernelExecInfo`** are called from multiple host threads on the same `cl_kernel` object at the same time<sup>37</sup>. Please note that there are additional limitations as to which OpenCL APIs may be called from OpenCL callback functions -- please see *section 5.11*.

The behavior of OpenCL APIs called from an interrupt or signal handler is implementation-defined

The OpenCL implementation should be able to create multiple command-queues for a given OpenCL context and multiple OpenCL contexts in an application running on the host processor.

---

<sup>36</sup> Please refer to the OpenCL glossary for the OpenCL definition of thread-safe. This definition may be different from usage of the term in other contexts.

<sup>37</sup> There is an inherent race condition in the design of OpenCL that occurs between setting a kernel argument and using the kernel with `clEnqueueNDRangeKernel`. Another host thread might change the kernel arguments between when a host thread sets the kernel arguments and then enqueues the kernel, causing the wrong kernel arguments to be enqueued. Rather than attempt to share `cl_kernel` objects among multiple host threads, applications are strongly encouraged to make additional `cl_kernel` objects for kernel functions for each host thread.

## Appendix B — Portability

OpenCL is designed to be portable to other architectures and hardware designs. OpenCL uses at its core a C99 based programming language. Floating-point arithmetic is based on the **IEEE-754** and **IEEE-754-2008** standards. The memory objects, pointer qualifiers and weakly ordered memory are designed to provide maximum compatibility with discrete memory architectures implemented by OpenCL devices. Command-queues and barriers allow for synchronization between the host and OpenCL devices. The design, capabilities and limitations of OpenCL are very much a reflection of the capabilities of underlying hardware.

Unfortunately, there are a number of areas where idiosyncrasies of one hardware platform may allow it to do some things that do not work on another. By virtue of the rich operating system resident on the CPU, on some implementations the kernels executing on a CPU may be able to call out to system services whereas the same calls on the GPU will likely fail for now. (Please see *section 6.9*). Since there is some advantage to having these services available for debugging purposes, implementations can use the OpenCL extension mechanism to implement these services.

Likewise, the heterogeneity of computing architectures might mean that a particular loop construct might execute at an acceptable speed on the CPU but very poorly on a GPU, for example. CPUs are designed in general to work well on latency sensitive algorithms on single threaded tasks, whereas common GPUs may encounter extremely long latencies, potentially orders of magnitude worse. A developer interested in writing portable code may find that it is necessary to test his design on a diversity of hardware designs to make sure that key algorithms are structured in a way that works well on a diversity of hardware. We suggest favoring more work-items over fewer. It is anticipated that over the coming months and years experience will produce a set of best practices that will help foster a uniformly favorable experience on a diversity of computing devices.

Of somewhat more concern is the topic of endianness. Since a majority of devices supported by the initial implementation of OpenCL are little-endian, developers need to make sure that their kernels are tested on both big-endian and little-endian devices to ensure source compatibility with OpenCL devices now and in the future. The endian attribute qualifier is supported by the OpenCL C programming language to allow developers to specify whether the data uses the endianness of the host or the OpenCL device. This allows the OpenCL compiler to do appropriate endian-conversion on load and store operations from or to this data.

We also describe how endianness can leak into an implementation causing kernels to produce unintended results:

When a big-endian vector machine (e.g. AltiVec, CELL SPE) loads a vector, the order of the data is retained. That is both the order of the bytes within each element and the order of the elements in the vector are the same as in memory. When a little-endian vector machine (e.g. SSE) loads a vector, the order of the data in register (where all the work is done) is reversed. **Both** the order of the bytes within each element and the order of the elements with respect to one

another in the vector are reversed.

Memory:

uint4 a =	<b>0x00010203</b>	0x04050607	0x08090A0B	0x0C0D0E0F
-----------	-------------------	------------	------------	------------

In register (big-endian):

uint4 a =	<b>0x00010203</b>	0x04050607	0x08090A0B	0x0C0D0E0F
-----------	-------------------	------------	------------	------------

In register (little-endian):

uint4 a =	0x0F0E0D0C	0x0B0A0908	0x07060504	<b>0x03020100</b>
-----------	------------	------------	------------	-------------------

This allows little-endian machines to use a single vector load to load little-endian data, regardless of how large each piece of data is in the vector. That is the transformation is equally valid whether that vector was a `uchar16` or a `ulong2`. Of course, as is well known, little-endian machines actually<sup>38</sup> store their data in reverse byte order to compensate for the little-endian storage format of the array elements:

Memory (big-endian):

uint4 a =	<b>0x00010203</b>	0x04050607	0x08090A0B	0x0C0D0E0F
-----------	-------------------	------------	------------	------------

Memory (little-endian):

uint4 a =	<b>0x03020100</b>	0x07060504	0x0B0A0908	0x0F0E0D0C
-----------	-------------------	------------	------------	------------

Once that data is loaded into a vector, we end up with this:

In register (big-endian):

uint4 a =	<b>0x00010203</b>	0x04050607	0x08090A0B	0x0C0D0E0F
-----------	-------------------	------------	------------	------------

In register (little-endian):

uint4 a =	0x0C0D0E0F	0x08090A0B	0x04050607	<b>0x00010203</b>
-----------	------------	------------	------------	-------------------

---

<sup>38</sup> Note that we are talking about the programming model here. In reality, little endian systems might choose to simply address their bytes from "the right" or reverse the "order" of the bits in the byte. Either of these choices would mean that no big swap would need to occur in hardware.

That is, in the process of correcting the endianness of the bytes within each element, the machine ends up reversing the order that the elements appear in the vector with respect to each other within the vector. `0x00010203` appears at the left of the big-endian vector and at the right of the little-endian vector.

When the host and device have different endianness, the developer must ensure that kernel argument values are processed correctly. The implementation may or may not automatically convert endianness of kernel arguments. Developers should consult vendor documentation for guidance on how to handle kernel arguments in these situations.

OpenCL provides a consistent programming model across architectures by numbering elements according to their order in memory. Concepts such as `even/odd` and `high/low` follow accordingly. Once the data is loaded into registers, we find that element 0 is at the left of the big-endian vector and element 0 is at the right of the little-endian vector:

```
float x[4];
float4 v = vload4( 0, x );
```

Big-endian:

```
v contains { x[0], x[1], x[2], x[3] }
```

Little-endian:

```
v contains { x[3], x[2], x[1], x[0] }
```

The compiler is aware that this swap occurs and references elements accordingly. So long as we refer to them by a numeric index such as `.s0123456789abcdef` or by descriptors such as `.xyzw`, `.hi`, `.lo`, `.even` and `.odd`, everything works transparently. Any ordering reversal is undone when the data is stored back to memory. The developer should be able to work with a big endian programming model and ignore the element ordering problem in the vector ... for most problems. This mechanism relies on the fact that we can rely on a consistent element numbering. Once we change numbering system, for example by conversion-free casting (using `as_type_n`) a vector to another vector of the same size but a different number of elements, then we get different results on different implementations depending on whether the system is big- endian, or little-endian or indeed has no vector unit at all. (Thus, the behavior of bitcasts to vectors of different numbers of elements is implementation-defined, see *section 6.2.4*)

An example follows:

```
float x[4] = { 0.0f, 1.0f, 2.0f, 3.0f };
float4 v = vload4( 0, x );
uint4 y = as_uint4(v);           // legal, portable
ushort8 z = as_ushort8(v);       // legal, not portable
                                   // element size changed
```

Big-endian:

```

v contains { 0.0f, 1.0f, 2.0f, 3.0f }
y contains { 0x00000000, 0x3f800000,
             0x40000000, 0x40400000 }
z contains { 0x0000, 0x0000, 0x3f80, 0x0000,
             0x4000, 0x0000, 0x4040, 0x0000 }
z.z is 0x3f80

```

Little-endian:

```

v contains { 3.0f, 2.0f, 1.0f, 0.0f }
y contains { 0x40400000, 0x40000000,
             0x3f800000, 0x00000000 }
z contains { 0x4040, 0x0000, 0x4000,
             0x0000, 0x3f80, 0x0000, 0x0000, 0x0000
             }
z.z is 0



```

Here, the value in `z.z` is not the same between big- and little-endian vector machines

OpenCL could have made it illegal to do a conversion free cast that changes the number of elements in the name of portability. However, while OpenCL provides a common set of operators drawing from the set that are typically found on vector machines, it can not provide access to everything every ISA may offer in a consistent uniform portable manner. Many vector ISAs provide special purpose instructions that greatly accelerate specific operations such as DCT, SAD, or 3D geometry. It is not intended for OpenCL to be so heavy handed that time-critical performance sensitive algorithms can not be written by knowledgeable developers to perform at near peak performance. Developers willing to throw away portability should be able to use the platform-specific instructions in their code. For this reason, OpenCL is designed to allow traditional vector C language programming extensions, such as the AltiVec C Programming Interface or the Intel C programming interfaces (such as those found in `emmintrin.h`) to be used directly in OpenCL with OpenCL data types as an extension to OpenCL. As these interfaces rely on the ability to do conversion-free casts that change the number of elements in the vector to function properly, OpenCL allows them too.

As a general rule, any operation that operates on vector types in segments that are not the same size as the vector element size may break on other hardware with different endianness or different vector architecture.

Examples might include:

-  Combining two `uchar8`'s containing high and low bytes of a `ushort`, to make a `ushort8` using `.even` and `.odd` operators (please use **`upsample()`** for this, see *section 6.13.3*)
-  Any bitcast that changes the number of elements in the vector. (Operations on the new type are non-portable.)

- ✚ Swizzle operations that change the order of data using chunk sizes that are not the same as the element size

Examples of operations that are portable:

- ✚ Combining two `uint8`'s to make a `uchar16` using `.even` and `.odd` operators. For example to interleave left and right audio streams.
- ✚ Any bitcast that does not change the number of elements (e.g. `(float4) uint4 --` we define the storage format for floating-point types)
- ✚ Swizzle operations that swizzle elements of the same size as the elements of the vector.

OpenCL has made some additions to C to make application behavior more dependable than C. Most notably in a few cases OpenCL defines the behavior of some operations that are undefined in C99:

- ✚ OpenCL provides `convert_` operators for conversion between all types. C99 does not define what happens when a floating-point type is converted to integer type and the floating-point value lies outside the representable range of the integer type after rounding. When the `_sat` variant of the conversion is used, the float shall be converted to the nearest representable integer value. Similarly, OpenCL also makes recommendations about what should happen with NaN. Hardware manufacturers that provide the saturated conversion in hardware may use the saturated conversion hardware for both the saturated and non-saturated versions of the OpenCL `convert_` operator. OpenCL does not define what happens for the non-saturated conversions when floating-point operands are outside the range representable integers after rounding.
- ✚ The format of `half`, `float`, and `double` types is defined to be the `binary16`, `binary32` and `binary64` formats in the draft IEEE-754 standard. (The latter two are identical to the existing IEEE-754 standard.) You may depend on the positioning and meaning of the bits in these types.
- ✚ OpenCL defines behavior for oversized shift values. Shift operations that shift greater than or equal to the number of bits in the first operand reduce the shift value modulo the number of bits in the element. For example, if we shift an `int4` left by 33 bits, OpenCL treats this as shift left by  $33\%32 = 1$  bit.
- ✚ A number of edge cases for math library functions are more rigorously defined than in C99. Please see *section 7.5*.



# Appendix C — Application Data Types

This section documents the provided host application types and constant definitions. The documented material describes the commonly defined data structures, types and constant values available to all platforms and architectures. The addition of these details demonstrates our commitment to maintaining a portable programming environment and potentially deters changes to the supplied headers.

## C.1 Shared Application Scalar Data Types

The following application scalar types are provided for application convenience.

```
cl_char
cl_uchar
cl_short
cl_ushort
cl_int
cl_uint
cl_long
cl_ulong
cl_half
cl_float
cl_double
```

## C.2 Supported Application Vector Data Types

Application vector types are unions used to create vectors of the above application scalar types. The following application vector types are provided for application convenience.

```
cl_charn
cl_ucharn
cl_shortn
cl_ushortn
cl_intn
cl_uintn
cl_longn
cl_ulongn
cl_halfn
cl_floatn
cl_doublen
```

*n* can be 2, 3, 4, 8 or 16.

The application scalar and vector data types are defined in the **cl\_platform.h** header file.

## C.3 Alignment of Application Data Types

The user is responsible for ensuring that data passed into and out of OpenCL buffers are natively aligned relative to the start of the buffer per requirements in *section 6.1.5*. This implies that OpenCL buffers created with CL\_MEM\_USE\_HOST\_PTR need to provide an appropriately aligned host memory pointer that is aligned to the data types used to access these buffers in a kernel(s). As well, the user is responsible to ensure that data passed into and out of OpenCL images are properly aligned to the granularity of the data representing a single pixel (e.g. `image_num_channels * sizeof(image_channel_data_type)`) except for CL\_RGB and CL\_RGBx images where the data must be aligned to the granularity of a single channel in a pixel (i.e. `sizeof(image_channel_data_type)`).

OpenCL makes no requirement about the alignment of OpenCL application defined data types outside of buffers and images, except that the underlying vector primitives (e.g. `__cl_float4`) where defined shall be directly accessible as such using appropriate named fields in the `cl_type` union (see *section C.5*). Nevertheless, it is recommended that the **cl\_platform.h** header should attempt to naturally align OpenCL defined application data types (e.g. `cl_float4`) according to their type.

## C.4 Vector Literals

Application vector literals may be used in assignments of individual vector components. Literal usage follows the convention of the underlying application compiler.

```
cl_float2 foo = { .s[1] = 2.0f };
cl_int8 bar = {{ 2, 4, 6, 8, 10, 12, 14, 16 }};
```

## C.5 Vector Components

The components of application vector types can be addressed using the `<vector_name>.s[<index>]` notation.

For example:

```
foo.s[0] = 1.0f; // Sets the 1st vector component of foo
pos.s[6] = 2;   // Sets the 7th vector component of bar
```

In some cases vector components may also be accessed using the following notations. These notations are not guaranteed to be supported on all implementations, so their use should be accompanied by a check of the corresponding preprocessor symbol.

### C.5.1 Named vector components notation

Vector data type components may be accessed using the `.sN`, `.sn` or `.xyzw` field naming convention, similar to how they are used within the OpenCL language. Use of the `.xyzw` field naming convention only allows accessing of the first 4 component fields. Support of these notations is identified by the `CL_HAS_NAMED_VECTOR_FIELDS` preprocessor symbol. For example:

```
#ifdef CL_HAS_NAMED_VECTOR_FIELDS
    cl_float4 foo;
    cl_int16 bar;
    foo.x = 1.0f; // Set first component
    foo.s0 = 1.0f; // Same as above
    bar.z = 3;    // Set third component
    bar.se = 11;  // Same as bar.s[0xe]
    bar.sD = 12;  // Same as bar.s[0xd]
#endif
```

Unlike the OpenCL language type usage of named vector fields, only one component field may be accessed at a time. This restriction prevents the ability to swizzle or replicate components as is possible with the OpenCL language types. Attempting to access beyond the number of components for a type also results in a failure.

```
foo.xy    // illegal - illegal field name combination
bar.s1234 // illegal - illegal field name combination
foo.s7    // illegal - no component s7
```

### C.5.2 High/Low vector component notation

Vector data type components may be accessed using the `.hi` and `.lo` notation similar to that supported within the language types. Support of this notation is identified by the `CL_HAS_HI_LO_VECTOR_FIELDS` preprocessor symbol. For example:

```
#ifdef CL_HAS_HI_LO_VECTOR_FIELDS
    cl_float4 foo;
    cl_float2 new_hi = 2.0f, new_lo = 4.0f;
    foo.hi = new_hi;
    foo.lo = new_lo;
#endif
```

### C.5.3 Native vector type notation

Certain native vector types are defined for providing a mapping of vector types to architecturally builtin vector types. Unlike the above described application vector types, these native types are supported on a limited basis depending on the supporting architecture and compiler.

These types are not unions, but rather convenience mappings to the underlying architectures' builtin vector types. The native types share the name of their application counterparts but are preceded by a double underscore "\_\_".

For example, `__cl_float4` is the native builtin vector type equivalent of the `cl_float4` application vector type. The `__cl_float4` type may provide direct access to the architectural builtin `__m128` or vector float type, whereas the `cl_float4` is treated as a union.

In addition, the above described application data types may have native vector data type members for access convenience. The native components are accessed using the `.vN` sub-vector notation, where N is the number of elements in the sub-vector. In cases where the native type is a subset of a larger type (more components), the notation becomes an index based array of the sub-vector type.

Support of the native vector types is identified by a `__CL_TYPEN__` preprocessor symbol matching the native type name. For example:

```
#ifdef __CL_FLOAT4__ // Check for native cl_float4 type
    cl_float8 foo;
    __cl_float4 bar; // Use of native type
    bar = foo.v4[1]; // Access the second native float4
    vector
#endif
```

## C.6 Implicit Conversions

Implicit conversions between application vector types are not supported.

## C.7 Explicit Casts

Explicit casting of application vector types (`cl_typen`) is not supported. Explicit casting of native vector types (`__cl_typen`) is defined by the external compiler.

## C.8 Other operators and functions

The behavior of standard operators and function on both application vector types (`cl_typed`) and native vector types (`__cl_typed`) is defined by the external compiler.

## C.9 Application constant definitions

In addition to the above application type definitions, the following literal definitions are also available.

<b>CL_CHAR_BIT</b>	Bit width of a character
<b>CL_SCHAR_MAX</b>	Maximum value of a type <code>cl_char</code>
<b>CL_SCHAR_MIN</b>	Minimum value of a type <code>cl_char</code>
<b>CL_CHAR_MAX</b>	Maximum value of a type <code>cl_char</code>
<b>CL_CHAR_MIN</b>	Minimum value of a type <code>cl_char</code>
<b>CL_UCHAR_MAX</b>	Maximum value of a type <code>cl_uchar</code>
<b>CL_SHORT_MAX</b>	Maximum value of a type <code>cl_short</code>
<b>CL_SHORT_MIN</b>	Minimum value of a type <code>cl_short</code>
<b>CL_USHORT_MAX</b>	Maximum value of a type <code>cl_ushort</code>
<b>CL_INT_MAX</b>	Maximum value of a type <code>cl_int</code>
<b>CL_INT_MIN</b>	Minimum value of a type <code>cl_int</code>
<b>CL_UINT_MAX</b>	Maximum value of a type <code>cl_uint</code>
<b>CL_LONG_MAX</b>	Maximum value of a type <code>cl_long</code>
<b>CL_LONG_MIN</b>	Minimum value of a type <code>cl_long</code>
<b>CL_ULONG_MAX</b>	Maximum value of a type <code>cl_ulong</code>
<b>CL_FLT_DIAG</b>	Number of decimal digits of precision for the type <code>cl_float</code>
<b>CL_FLT_MANT_DIG</b>	Number of digits in the mantissa of type <code>cl_float</code>
<b>CL_FLT_MAX_10_EXP</b>	Maximum positive integer such that 10 raised to this power minus one can be represented as a normalized floating-point number of type <code>cl_float</code>
<b>CL_FLT_MAX_EXP</b>	Maximum exponent value of type <code>cl_float</code>
<b>CL_FLT_MIN_10_EXP</b>	Minimum negative integer such that 10 raised to this power minus one can be represented as a normalized floating-point number of type <code>cl_float</code>
<b>CL_FLT_MIN_EXP</b>	Minimum exponent value of type <code>cl_float</code>
<b>CL_FLT_RADIX</b>	Base value of type <code>cl_float</code>
<b>CL_FLT_MAX</b>	Maximum value of type <code>cl_float</code>
<b>CL_FLT_MIN</b>	Minimum value of type <code>cl_float</code>
<b>CL_FLT_EPSILON</b>	Minimum positive floating-point number of type

	<code>cl_float</code> such that <code>1.0 + CL_FLT_EPSILON != 1</code> is true.
<b>CL_DBL_DIG</b>	Number of decimal digits of precision for the type <code>cl_double</code>
<b>CL_DBL_MANT_DIG</b>	Number of digits in the mantissa of type <code>cl_double</code>
<b>CL_DBL_MAX_10_EXP</b>	Maximum positive integer such that 10 raised to this power minus one can be represented as a normalized floating-point number of type <code>cl_double</code>
<b>CL_DBL_MAX_EXP</b>	Maximum exponent value of type <code>cl_double</code>
<b>CL_DBL_MIN_10_EXP</b>	Minimum negative integer such that 10 raised to this power minus one can be represented as a normalized floating-point number of type <code>cl_double</code>
<b>CL_DBL_MIN_EXP</b>	Minimum exponent value of type <code>cl_double</code>
<b>CL_DBL_RADIX</b>	Base value of type <code>cl_double</code>
<b>CL_DBL_MAX</b>	Maximum value of type <code>cl_double</code>
<b>CL_DBL_MIN</b>	Minimum value of type <code>cl_double</code>
<b>CL_DBL_EPSILON</b>	Minimum positive floating-point number of type <code>cl_double</code> such that <code>1.0 + CL_DBL_EPSILON != 1</code> is true.
<b>CL_NAN</b>	Macro expanding to a value representing NaN
<b>CL_HUGE_VALF</b>	Largest representative value of type <code>cl_float</code>
<b>CL_HUGE_VAL</b>	Largest representative value of type <code>cl_double</code>
<b>CL_MAXFLOAT</b>	Maximum value of type <code>cl_float</code>
<b>CL_INFINITY</b>	Macro expanding to a value representing infinity

These literal definitions are defined in the **cl\_platform.h** header.

## Appendix D — OpenCL C++ Wrapper API

The OpenCL C++ wrapper API provides a C++ interface to the platform and runtime API. The C++ wrapper is built on top of the OpenCL 2.0 C API (platform and runtime) and is not a replacement. It is **required** that any implementation of the C++ wrapper API will make calls to the underlying C API and it is assumed that the C API is a compliant implementation of the OpenCL 2.0 specification.

Refer to the OpenCL C++ Wrapper API specification for details. The OpenCL C++ Wrapper API specification can be found at <http://www.khronos.org/registry/cl/>.

## Appendix E — CL\_MEM\_COPY\_OVERLAP

The following code describes how to determine if there is overlap between the source and destination rectangles specified to **clEnqueueCopyBufferRect** provided the source and destination buffers refer to the same buffer object.

Copyright (c) 2011 The Khronos Group Inc.

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and /or associated documentation files (the "Materials "), to deal in the Materials without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Materials, and to permit persons to whom the Materials are furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Materials.

THE MATERIALS ARE PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE MATERIALS OR THE USE OR OTHER DEALINGS IN THE MATERIALS.

```
unsigned int
check_copy_overlap( const size_t src_origin[],
                   const size_t dst_origin[],
                   const size_t region[],
                   const size_t row_pitch,
                   const size_t slice_pitch )
{
    const size_t slice_size = (region[1] - 1) * row_pitch + region[0];
    const size_t block_size = (region[2] - 1) * slice_pitch + slice_size;

    const size_t src_start  = src_origin[2] * slice_pitch +
                              src_origin[1] * row_pitch +
                              src_origin[0];
    const size_t src_end    = src_start + block_size;

    const size_t dst_start  = dst_origin[2] * slice_pitch +
                              dst_origin[1] * row_pitch +
                              dst_origin[0];
    const size_t dst_end    = dst_start + block_size;

    /* No overlap if dst ends before src starts or if src ends
     * before dst starts.
     */
    if( (dst_end <= src_start) || (src_end <= dst_start) )
```



```

{
    return 0;
}

/* No overlap if region[0] for dst or src fits in the gap
 * between region[0] and row_pitch.
 */
{
    const size_t src_dx = src_origin[0] % row_pitch;
    const size_t dst_dx = dst_origin[0] % row_pitch;

    if( ((dst_dx >= src_dx + region[0]) &&
        (dst_dx + region[0] <= src_dx + row_pitch)) ||
        ((src_dx >= dst_dx + region[0]) &&
        (src_dx + region[0] <= dst_dx + row_pitch)) )
    {
        return 0;
    }
}

/* No overlap if region[1] for dst or src fits in the gap
 * between region[1] and slice_pitch.
 */
{
    const size_t src_dy =
        (src_origin[1] * row_pitch + src_origin[0]) % slice_pitch;
    const size_t dst_dy =
        (dst_origin[1] * row_pitch + dst_origin[0]) % slice_pitch;

    if( ((dst_dy >= src_dy + slice_size) &&
        (dst_dy + slice_size <= src_dy + slice_pitch)) ||
        ((src_dy >= dst_dy + slice_size) &&
        (src_dy + slice_size <= dst_dy + slice_pitch)) )
    {
        return 0;
    }
}

/* Otherwise src and dst overlap. */
return 1;
}

```

# Appendix F – Changes

## F.1 Summary of changes from OpenCL 1.0

The following features are added to the OpenCL 1.1 platform layer and runtime (*sections 4 and 5*):

- ✚ Following queries to *table 4.3*
  - CL\_DEVICE\_NATIVE\_VECTOR\_WIDTH\_{CHAR | SHORT | INT | LONG | FLOAT | DOUBLE | HALF}
  - CL\_DEVICE\_HOST\_UNIFIED\_MEMORY
  - CL\_DEVICE\_OPENCL\_C\_VERSION
- ✚ CL\_CONTEXT\_NUM\_DEVICES to the list of queries specified to **clGetContextInfo**.
- ✚ Optional image formats: CL\_Rx, CL\_RGx and CL\_RGBx.
- ✚ Support for sub-buffer objects – ability to create a buffer object that refers to a specific region in another buffer object using **clCreateSubBuffer**.
- ✚ **clEnqueueReadBufferRect**, **clEnqueueWriteBufferRect** and **clEnqueueCopyBufferRect** APIs to read from, write to and copy a rectangular region of a buffer object respectively.
- ✚ **clSetMemObjectDestructorCallback** API to allow a user to register a callback function that will be called when the memory object is deleted and its resources freed.
- ✚ Options that control the OpenCL C version used when building a program executable. These are described in *section 5.8.4.5*.
- ✚ CL\_KERNEL\_PREFERRED\_WORK\_GROUP\_SIZE\_MULTIPLE to the list of queries specified to **clGetKernelWorkGroupInfo**.
- ✚ Support for user events. User events allow applications to enqueue commands that wait on a user event to finish before the command is executed by the device. Following new APIs are added - **clCreateUserEvent** and **clSetUserEventStatus**.
- ✚ **clSetEventCallback** API to register a callback function for a specific command execution status.

The following modifications are made to the OpenCL 1.1 platform layer and runtime (*sections 4 and 5*):

- ✚ Following queries in *table 4.3*

- CL\_DEVICE\_MAX\_PARAMETER\_SIZE from 256 to 1024 bytes
- CL\_DEVICE\_LOCAL\_MEM\_SIZE from 16 KB to 32 KB.

✚ The *global\_work\_offset* argument in **clEnqueueNDRangeKernel** can be a non-NULL value.

✚ All API calls except **clSetKernelArg** are thread-safe.

The following features are added to the OpenCL C programming language (*section 6*) in OpenCL 1.1:

✚ 3-component vector data types.

✚ New built-in functions

- **get\_global\_offset** work-item function defined in *section 6.12.1*.
- **minmag**, **maxmag** math functions defined in *section 6.12.2*.
- **clamp** integer function defined in *section 6.12.3*.
- (vector, scalar) variant of integer functions **min** and **max** in *section 6.12.3*.
- **async\_work\_group\_strided\_copy** defined in *section 6.12.10*.
- **vec\_step**, **shuffle** and **shuffle2** defined in *section 6.12.12*.

✚ **cl\_khr\_byte\_addressable\_store** extension is a core feature.

✚ **cl\_khr\_global\_int32\_base\_atomics**, **cl\_khr\_global\_int32\_extended\_atomics**, **cl\_khr\_local\_int32\_base\_atomics** and **cl\_khr\_local\_int32\_extended\_atomics** extensions are core features. The built-in atomic function names are changed to use the **atomic\_** prefix instead of **atom\_**.

✚ Macros CL\_VERSION\_1\_0 and CL\_VERSION\_1\_1.

The following features in OpenCL 1.0 are deprecated in OpenCL 1.1:

- ✚ The **clSetCommandQueueProperty** API is no longer supported in OpenCL 1.1.
- ✚ The `__ROUNDING_MODE__` macro is no longer supported in OpenCL C 1.1.
- ✚ The `-cl-strict-aliasing` option that can be specified in *options* argument to **clBuildProgram** is no longer supported in OpenCL 1.1.

The following new extensions are added to *section 9* in OpenCL 1.1:

- ✚ **cl\_khr\_gl\_event** – Creating a CL event object from a GL sync object.
- ✚ **cl\_khr\_d3d10\_sharing** – Sharing memory objects with Direct3D 10.

The following modifications are made to the OpenCL ES Profile described in *section 10* in OpenCL 1.1:

✚ 64-bit integer support is optional.

## F.2 Summary of changes from OpenCL 1.1

The following features are added to the OpenCL 1.2 platform layer and runtime (*sections 4 and 5*):

- ✚ Custom devices and built-in kernels are supported.
- ✚ Device partitioning that allows a device to be partitioned based on a number of partitioning schemes supported by the device.
- ✚ Extend *cl\_mem\_flags* to describe how the host accesses the data in a *cl\_mem* object.
- ✚ **clEnqueueFillBuffer** and **clEnqueueFillImage** to support filling a buffer with a pattern or an image with a color.
- ✚ Add *CL\_MAP\_WRITE\_INVALIDATE\_REGION* to *cl\_map\_flags*. Appropriate clarification to the behavior of *CL\_MAP\_WRITE* has been added to the spec.
- ✚ New image types: 1D image, 1D image from a buffer object, 1D image array and 2D image arrays.
- ✚ **clCreateImage** to create an image object.
- ✚ **clEnqueueMigrateMemObjects** API that allows a developer to have explicit control over the location of memory objects or to migrate a memory object from one device to another.
- ✚ Support separate compilation and linking of programs.
- ✚ Additional queries to get the number of kernels and kernel names in a program have been added to **clGetProgramInfo**.
- ✚ Additional queries to get the compile and link status and options have been added to **clGetProgramBuildInfo**.
- ✚ **clGetKernelArgInfo** API that returns information about the arguments of a kernel.
- ✚ **clEnqueueMarkerWithWaitList** and **clEnqueueBarrierWithWaitList** APIs.

The following features are added to the OpenCL C programming language (*section 6*) in OpenCL 1.2:

- ✚ Double-precision is now an optional core feature instead of an extension.
- ✚ New built in image types: **image1d\_t**, **image1d\_array\_t** and **image2d\_array\_t**.

- ✚ New built-in functions
  - Functions to read from and write to a 1D image, 1D and 2D image arrays described in *sections 6.12.14.2, 6.12.14.3 and 6.12.14.4.*
  - Sampler-less image read functions described in *section 6.12.14.3.*
  - **popcount** integer function described in *section 6.12.3.*
  - **printf** function described in *section 6.12.13.*
- ✚ Storage class specifiers `extern` and `static` as described in *section 6.8.*
- ✚ Macros `CL_VERSION_1_2` and `__OPENCL_C_VERSION__`.

The following APIs in OpenCL 1.1 are deprecated in OpenCL 1.2:

- ✚ **clEnqueueMarker**, **clEnqueueBarrier** and **clEnqueueWaitForEvents**
- ✚ **clCreateImage2D** and **clCreateImage3D**
- ✚ **clUnloadCompiler** and **clGetExtensionFunctionAddress**
- ✚ **clCreateFromGLTexture2D** and **clCreateFromGLTexture3D**

The following queries are deprecated in OpenCL 1.2:

- ✚ **CL\_DEVICE\_MIN\_DATA\_TYPE\_ALIGN\_SIZE** in *table 4.3* queried using **clGetDeviceInfo**.

## F.3 Summary of changes from OpenCL 1.2

The following features are added to the OpenCL 2.0 platform layer and runtime (*sections 4 and 5*):

- ✚ Shared virtual memory.
- ✚ Device queues used to enqueue kernels on the device.
- ✚ Pipes.
- ✚ Images – support for 2D image from buffer, depth images and sRGB images.

The following modifications are made to the OpenCL 2.0 platform layer and runtime (*sections 4 and 5*):

- ✚ All API calls except **clSetKernelArg**, **clSetKernelArgSVMPointer** and **clSetKernelExecInfo** are thread-safe.

The following features are added to the OpenCL C programming language (*section 6*) in OpenCL 2.0:

- ✚ Clang Blocks.
- ✚ Kernels enqueueing kernels to a device queue.
- ✚ Program scope variables in global address space.
- ✚ Generic address space.
- ✚ C1x atomics.
- ✚ New built-in functions (sections 6.13.9, 6.13.11, 6.13.15 and 6.14).
- ✚ Support images with the `read_write` qualifier.
- ✚ 3D image writes are a core feature.
- ✚ The `CL_VERSION_2_0` macro.

The following APIs are deprecated in OpenCL 2.0:

- ✚ **`clCreateCommandQueue`, `clCreateSampler` and `clEnqueueTask`**

The following queries are deprecated in OpenCL 2.0:

- ✚ **`CL_DEVICE_HOST_UNIFIED_MEMORY` in *table 4.3* queried using `clGetDeviceInfo`.**
- ✚ **`CL_IMAGE_BUFFER` in *table 5.10* is deprecated.**
- ✚ **`CL_DEVICE_QUEUE_PROPERTIES` is replaced by `CL_DEVICE_QUEUE_ON_HOST_PROPERTIES`.**
- ✚ The explicit memory fence functions defined in section 6.12.9 of the OpenCL 1.2 specification.
- ✚ The OpenCL 1.2 atomic built-in functions for 32-bit integer and floating-point data types defined in section 6.12.11 of the OpenCL 1.2 specification.

# Index - APIs

**clBuildProgram**, 192  
**clCompileProgram**, 194  
**clCreateBuffer**, 97  
**clCreateCommandQueueWithProperties**, 91  
**clCreateContext**, 84  
**clCreateContextFromType**, 86  
**clCreateImage**, 121  
**clCreateKernel**, 213  
**clCreateKernelsInProgram**, 214  
**clCreatePipe**, 153  
**clCreateProgramWithBinary**, 187  
**clCreateProgramWithBuiltInKernels**, 189  
**clCreateProgramWithSource**, 186  
**clCreateSamplerWithProperties**, 182  
**clCreateSubBuffer**, 100  
**clCreateSubDevices**, 79  
**clCreateUserEvent**, 238  
**clEnqueueBarrierWithWaitList**, 248  
**clEnqueueCopyBuffer**, 110  
**clEnqueueCopyBufferRect**, 112  
**clEnqueueCopyBufferToImage**, 144  
**clEnqueueCopyImage**, 137  
**clEnqueueCopyImageToBuffer**, 142  
**clEnqueueFillBuffer**, 115  
**clEnqueueFillImage**, 140  
**clEnqueueMapBuffer**, 117  
**clEnqueueMapImage**, 147  
**clEnqueueMarkerWithWaitList**, 247  
**clEnqueueMigrateMemObjects**, 161  
**clEnqueueNativeKernel**, 234  
**clEnqueueNDRangeKernel**, 230  
**clEnqueueReadBuffer**, 103  
**clEnqueueReadBufferRect**, 105  
**clEnqueueReadImage**, 133  
**clEnqueueSVMFree**, 172  
**clEnqueueSVMMMap**, 177  
**clEnqueueSVMMemcpy**, 173  
**clEnqueueSVMMemFill**, 175  
**clEnqueueSVMUnmap**, 178  
**clEnqueueUnmapMemObject**, 158  
**clEnqueueWriteBuffer**, 103  
**clEnqueueWriteBufferRect**, 106  
**clEnqueueWriteImage**, 133  
**clFinish**, 254  
**clFlush**, 254  
**clGetCommandQueueInfo**, 94  
**clGetContextInfo**, 88  
**clGetDeviceIDs**, 60  
**clGetDeviceInfo**, 62  
**clGetEventInfo**, 240  
**clGetEventProfilingInfo**, 251  
**clGetImageInfo**, 150  
**clGetKernelArgInfo**, 226  
**clGetKernelInfo**, 222  
**clGetKernelWorkGroupInfo**, 223  
**clGetMemObjectInfo**, 163  
**clGetPipeInfo**, 154  
**clGetPlatformIDs**, 58  
**clGetPlatformInfo**, 58  
**clGetProgramBuildInfo**, 209  
**clGetProgramInfo**, 206  
**clGetSamplerInfo**, 184  
**clGetSupportedImageFormats**, 129  
**clLinkProgram**, 197  
**clReleaseCommandQueue**, 94  
**clReleaseContext**, 87  
**clReleaseDevice**, 82  
**clReleaseEvent**, 245  
**clReleaseKernel**, 215  
**clReleaseMemObject**, 156  
**clReleaseProgram**, 190, 191  
**clReleaseSampler**, 184  
**clRetainCommandQueue**, 93  
**clRetainContext**, 87  
**clRetainDevice**, 82  
**clRetainEvent**, 245  
**clRetainKernel**, 215  
**clRetainMemObject**, 156  
**clRetainProgram**, 190  
**clRetainSampler**, 183  
**clSetEventCallback**, 243  
**clSetKernelArg**, 216, 218, 219  
**clSetKernelArgSVMPointer**, 218  
**clSetKernelExecInfo**, 219  
**clSetMemObjectDestructorCallback**, 157

**clSetUserEventStatus**, 238  
**clSVMAlloc**, 168  
**clSVMFree**, 171

**clUnloadPlatformCompiler**, 205  
**clWaitForEvents**, 239